

# Meet Christine, Our Heroine



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# Meet Christine, Our Heroine



**Noreen Hasan**  
Training Architect

Christine's Surprising Day

Lesson Summary



**Meet Christine, Our Heroine**

## **Christine's Surprising Day**

**— Placeholder —**

**A three minute animation will be added here**



# Lesson Summary



Meet **Christine**, Our Heroine

- Data engineer working on a **startup**
- **Unpredictable spike** in demand
- **Cliffhanger:** how will Christine respond to the spike?



# Introduction to Unpacking Lambda



**Noreen Hasan**

Training Architect



## SECTION BREAKDOWN

# Introduction to Unpacking Lambda



**Noreen Hasan**  
Training Architect

## Section Overview

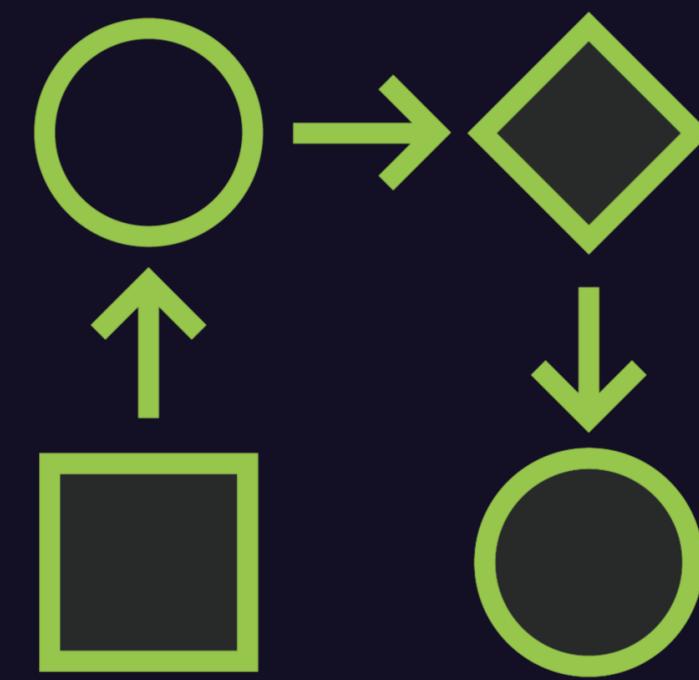


# Introduction to Unpacking Lambda

## Section Overview



What Is Lambda?



How Lambda Works

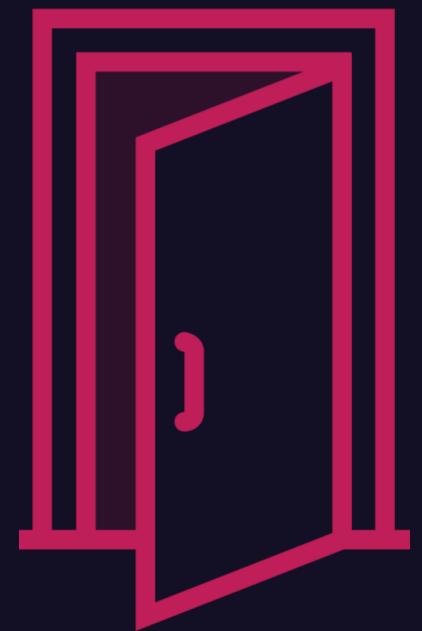


Lambda Use Cases



### Section Overview

Lambda is no



#### When to Use Lambda

Compare it against EC2, Elastic  
Beanstalk, and ECS

#### Accessing Lambda

CLI, console, SDK  
+ demo



# Let's get started!



# Unpacking Lambda: Recap



**Noreen Hasan**  
Training Architect



## SECTION BREAKDOWN

# Unpacking Lambda: Recap

## Lessons Summary



**Noreen Hasan**  
Training Architect



### Lessons Summary



#### What Is Lambda?

Event-driven serverless AWS service

#### Usages

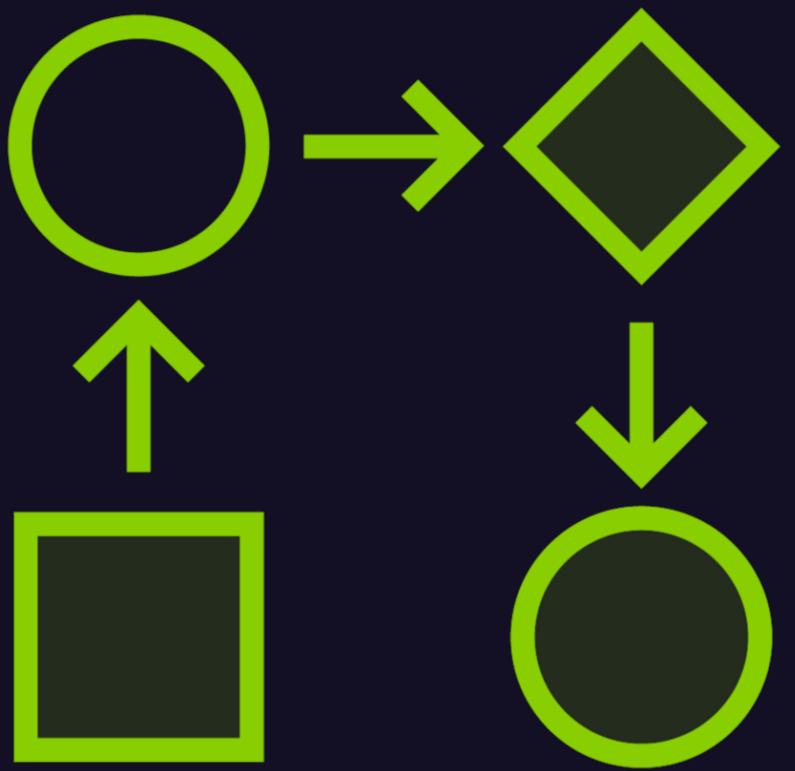
- IoT devices
- Serverless websites

#### High-Level Workflow

1. Create a **function**
2. Write or **upload code** to the function
3. Specify the **triggering event**



### Lessons Summary



### How Does Lambda Work?

Lambda gets invoked to run your code only when needed.

### Main Building Blocks

1. Function
2. Trigger
3. Event
4. Execution Environment

### Under the Hood

Load balancer and a fleet of EC2 instances



### Lessons Summary



#### Lambda Use Cases

- Processing data
  - Processing S3 objects
  - Monitoring and analyzing logs
  - Processing streaming data
- Orchestrating workflows
- Operating serverless applications
- Processing messages



### Lessons Summary

Lambda is no



#### Alternative Services

- AWS Elastic Beanstalk
- EC2
- ECS

#### When to Use Lambda

- If you want to focus more on the code and less on servers
- If your jobs take < 15 mins
- Never pay \$ for idle time



# Unpacking Lambda: Recap

## Lessons Summary



### Accessing Lambda

- AWS CLI
- AWS SDK
- AWS Lambda console

### Lambda Demo

- Walkthrough of the console
- Create a Hello World  application



**See you in the  
next section!**



# What Is Lambda?



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# What Is Lambda?



**Noreen Hasan**  
Training Architect

**Lambda in the Life of a Data Professional**

**Lambda and You**

**High-Level Workflow**

**Lesson Summary**



## What Is Lambda?

# Lambda in the Life of a Data Professional

— Placeholder —

A three minute animation will be added here



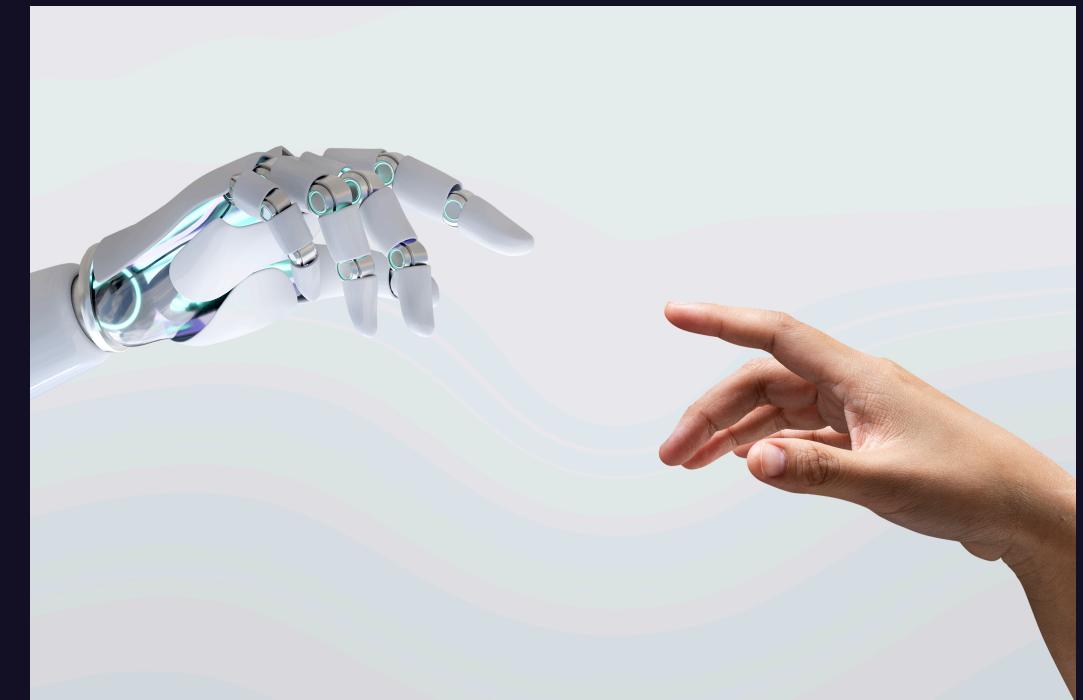
# What Is Lambda?

## Lambda and You



IoT Devices

Serverless Websites



Small Event Driven



## What Is Lambda?

### Lambda and You

You're not in charge  
of provisioning the  
servers.

You're in charge of  
writing event-driven  
code.

You're in charge of  
persisting results.



## What Is Lambda?

### Lambda and You

**Amazon is in charge  
of everything  
hardware related.**

**Amazon is in charge  
of operating systems  
and security.**

**You pay for running  
code measured in  
milliseconds (ms).**



## What Is Lambda?

# High-Level Workflow

1

**Create a function.**

2

**Write or upload code to the function.**

3

**Specify the triggering event.**



## What Is Lambda?

# High-Level Workflow

**What?**

Write data to DynamoDB.

**When?**

Users upload to S3.

**What information to process?**

The data inside the files uploaded to S3.

**Where will the code run?**

In an isolated runtime environment.



## What Is Lambda?

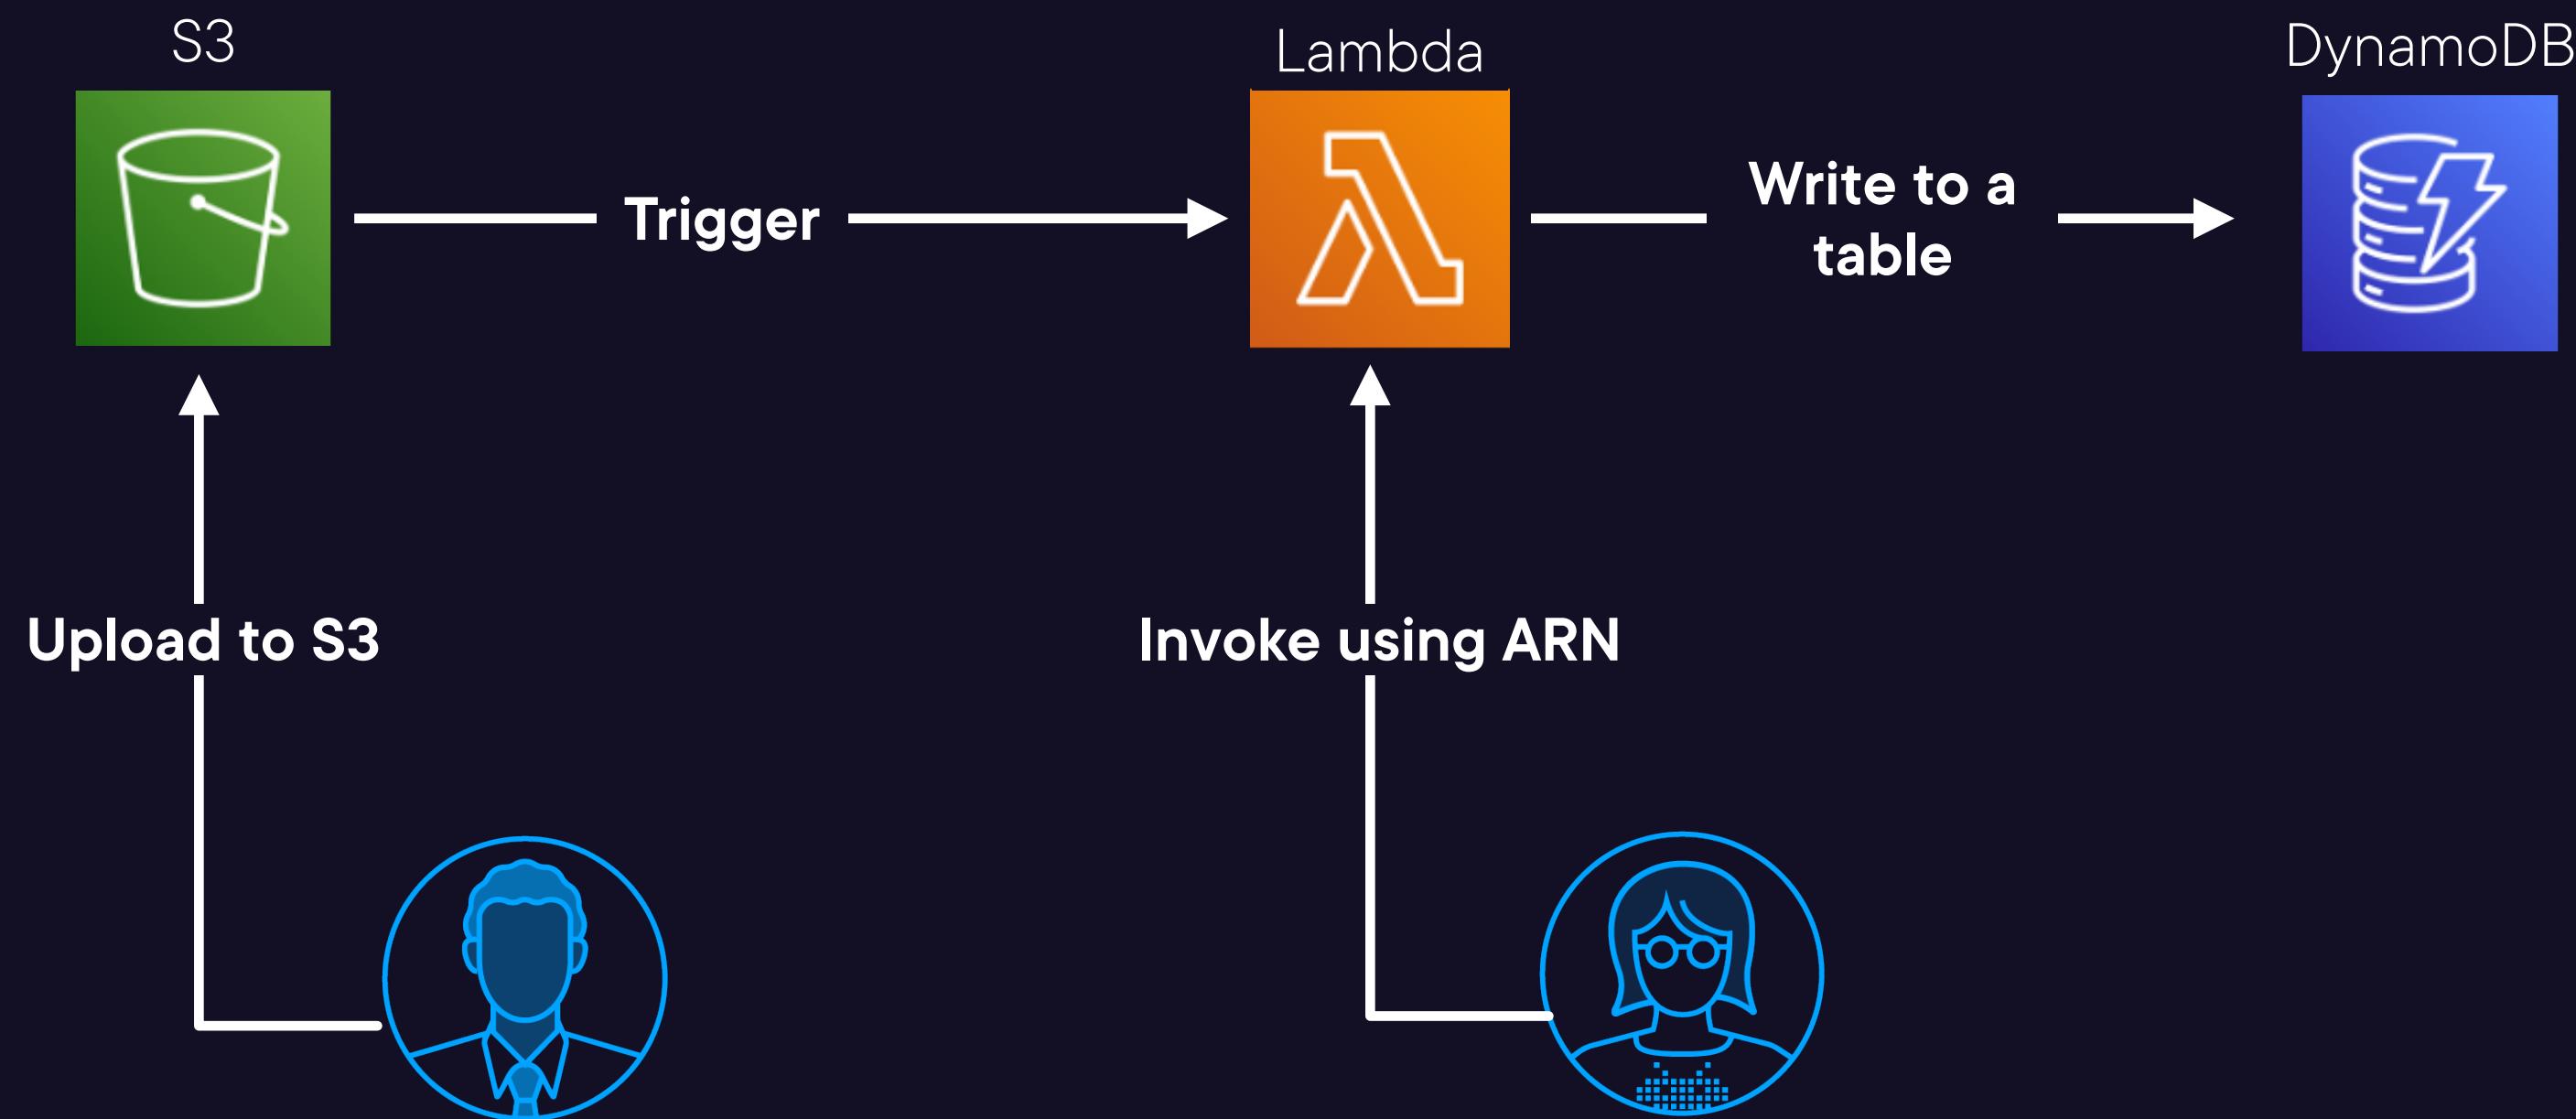
### High-Level Workflow



## Lambda Requirements

- The Lambda Function
- IAM Execution Role
- Memory Size Specification
- Execution Timeout
  - The limit is 15 mins
- Event Source Mapping





## Lesson Summary



### Lambda in the Life of a Data Professional

- Explored how Lambda serves Christine's startup

### Lambda and You

- Everyday examples
- Your responsibilities in Lambda

### High-Level Workflow

- High-level steps for working with Lambda
- Lambda requirements



# How Lambda Works



**Noreen Hasan**

Training Architect



# How Lambda Works

## LESSON BREAKDOWN

Overview of Lambda's Building Blocks

First Block: Function

Second Block: Trigger

Third Block: Event

Fourth Block: Execution Environment

Under the Hood

Lesson Summary



**Noreen Hasan**  
Training Architect



## How Lambda Works

# Overview of Lambda's Building Blocks

Function

Trigger

Event

Execution environment



## How Lambda Works

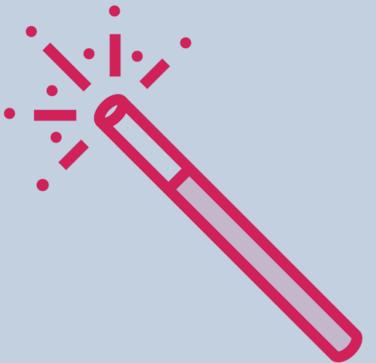
# Overview of Lambda's Building Blocks

## Function

A standalone unit of code that processes events

## Trigger

A resource that invokes the function into a running state



## How Lambda Works

# Overview of Lambda's Building Blocks

## Event

Relevant information that gets passed to functions to process the data

## Execution Environment

An isolated and secure environment to run the function code



## How Lambda Works

# Overview of Lambda's Building Blocks

Function Code

What?

Write data to DynamoDB.

Trigger

When?

Users upload to S3.

Event

What information to process?

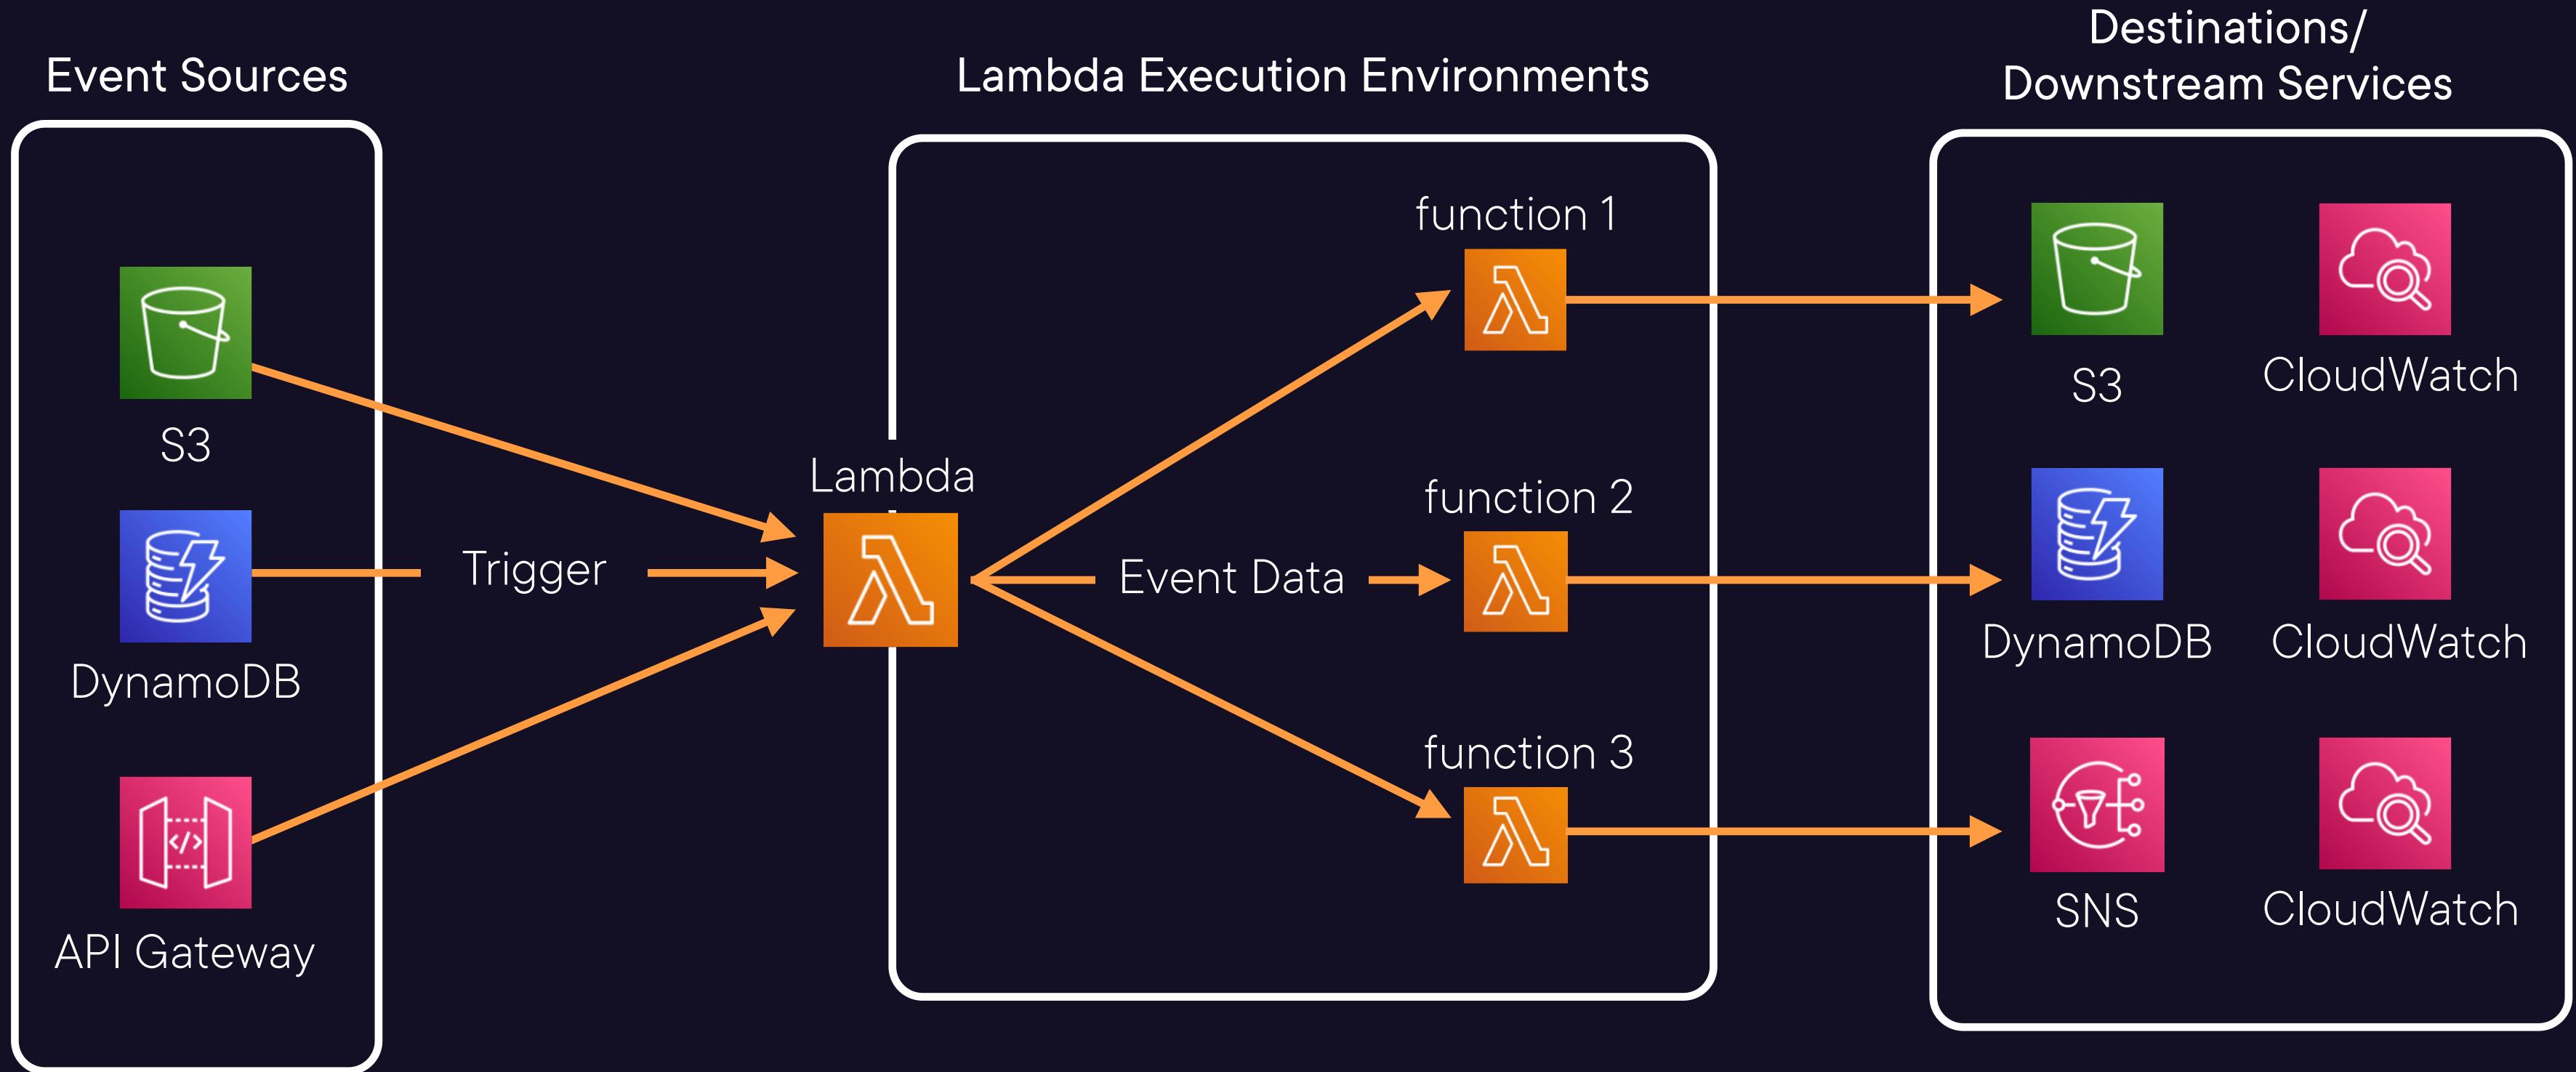
The data inside the files uploaded to S3.

Execution Environment

Where will the code run?

In an isolated runtime environment.

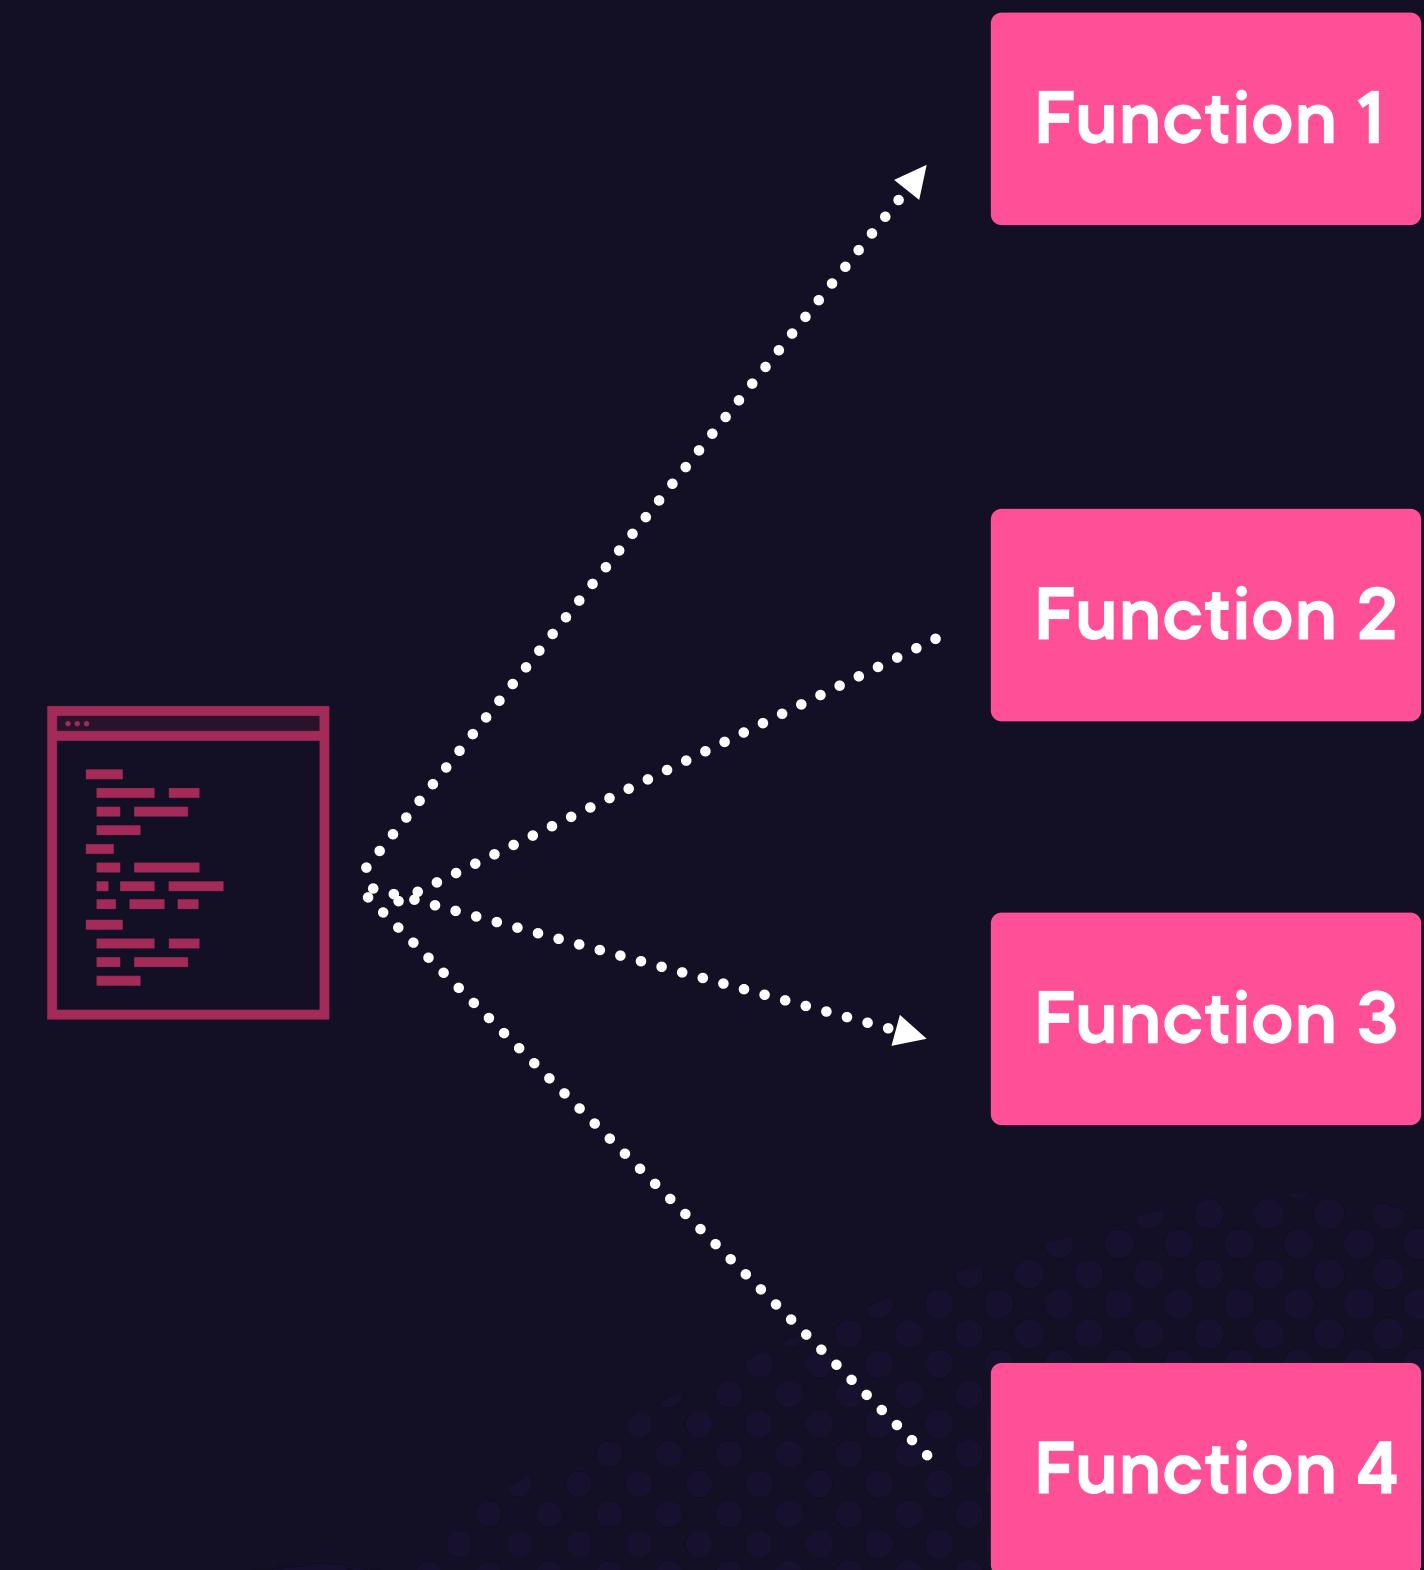




# How Lambda Works

## First Block: Function

In Lambda, code is organized and split into functions.



## How Lambda Works

### First Block: Function

You can configure your code to time out after a specific amount of time.



< 1 sec



< 3 secs



< 15 mins



## How Lambda Works

### First Block: Function

```
def handler_name(event, context)  
    . . .  
return some_value
```

#### Lambda Function Handler

- A method to respond to events

Event handlers take two parameters:

- Event
- Context



## How Lambda Works

### First Block: Function

## (Python)

### lambda\_function.py

```
import logging } libraries
import json

handler name      parameters
                ↓   ↘
def lambda_handler(event, context)
    message = 'You might like the course {}'.format(event['courseName'])
    return {
        'message': message
    }
```

### Parameters:

1. Event: has the details (usually type = `dict`) of the event that triggered the function.
2. Context: has the runtime information (type = `context`).



### Second Block: Trigger



## Trigger

In the classic tale from *The Arabian Nights*, the thieves hid their treasures in a cave behind an invisible door.

Uttering the magical words “Open, Sesame” was the **event** used to **trigger** the door to open.



## How Lambda Works

### Second Block: Trigger



Anything configured to invoke a Lambda function.

#### Other AWS Services

- S3 Upload
- Alexa Skill

#### Lambda Resources

- The Lambda API
- The Lambda Console
- Event Source Mapping (to process queued items)

Functions can have multiple triggers.



## How Lambda Works

### Third Block: Event

```
{  
  "FirstName": "Christine",  
  "LastName": "Hernandez"  
}
```

A diagram illustrating the flow of data. At the top, a white box contains a JSON event object. A green arrow points from this box down to a purple rectangular area representing a Lambda function. Inside this area, a Python-style code snippet defines a handler named `handler_name` that takes `event` and `context` as parameters and returns a `processed_value`.

```
def handler_name(event, context)  
    . . .  
    return processed_value
```

```
{  
  "key": "CourseThumbnail.jpg",  
  "size": 1024,  
  "eTag": "d41d8cd9f0b204e980ec",  
  "versionId": "09fKTRTt9fV.nflo"  
}
```

**Events are JSON-formatted documents.**

- They get passed into functions to pass in the relevant data.
- The runtime converts the document into an object.
- The structure and content depends on the invoking service.



## How Lambda Works

### Third Block: Event

# How is an event object created?

lambda\_function.py

```
import pandas

def lambda_handler(event, context)
    message = 'You might like the course - {} -'.format(event['courseName'])
    return {
        'message': message
    }
```

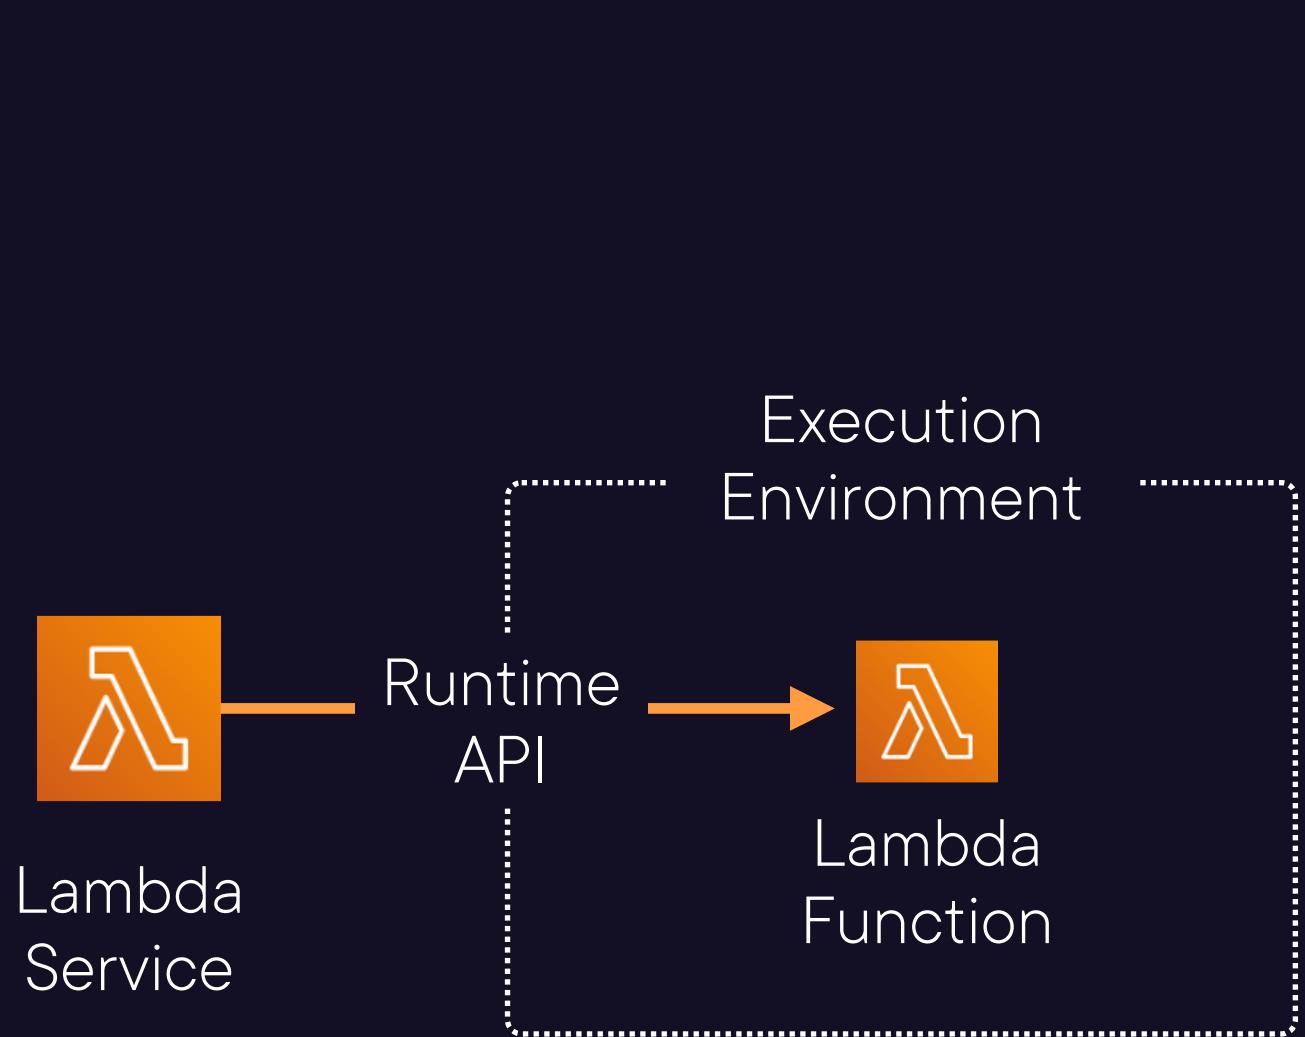
```
{  
    "Course": "Processing Serverless Data Using Lambda"  
}
```



```
{
    "message": "You might like the course – Processing Serverless Data Using Lambda"
}
```



### Fourth Block: Execution Environment



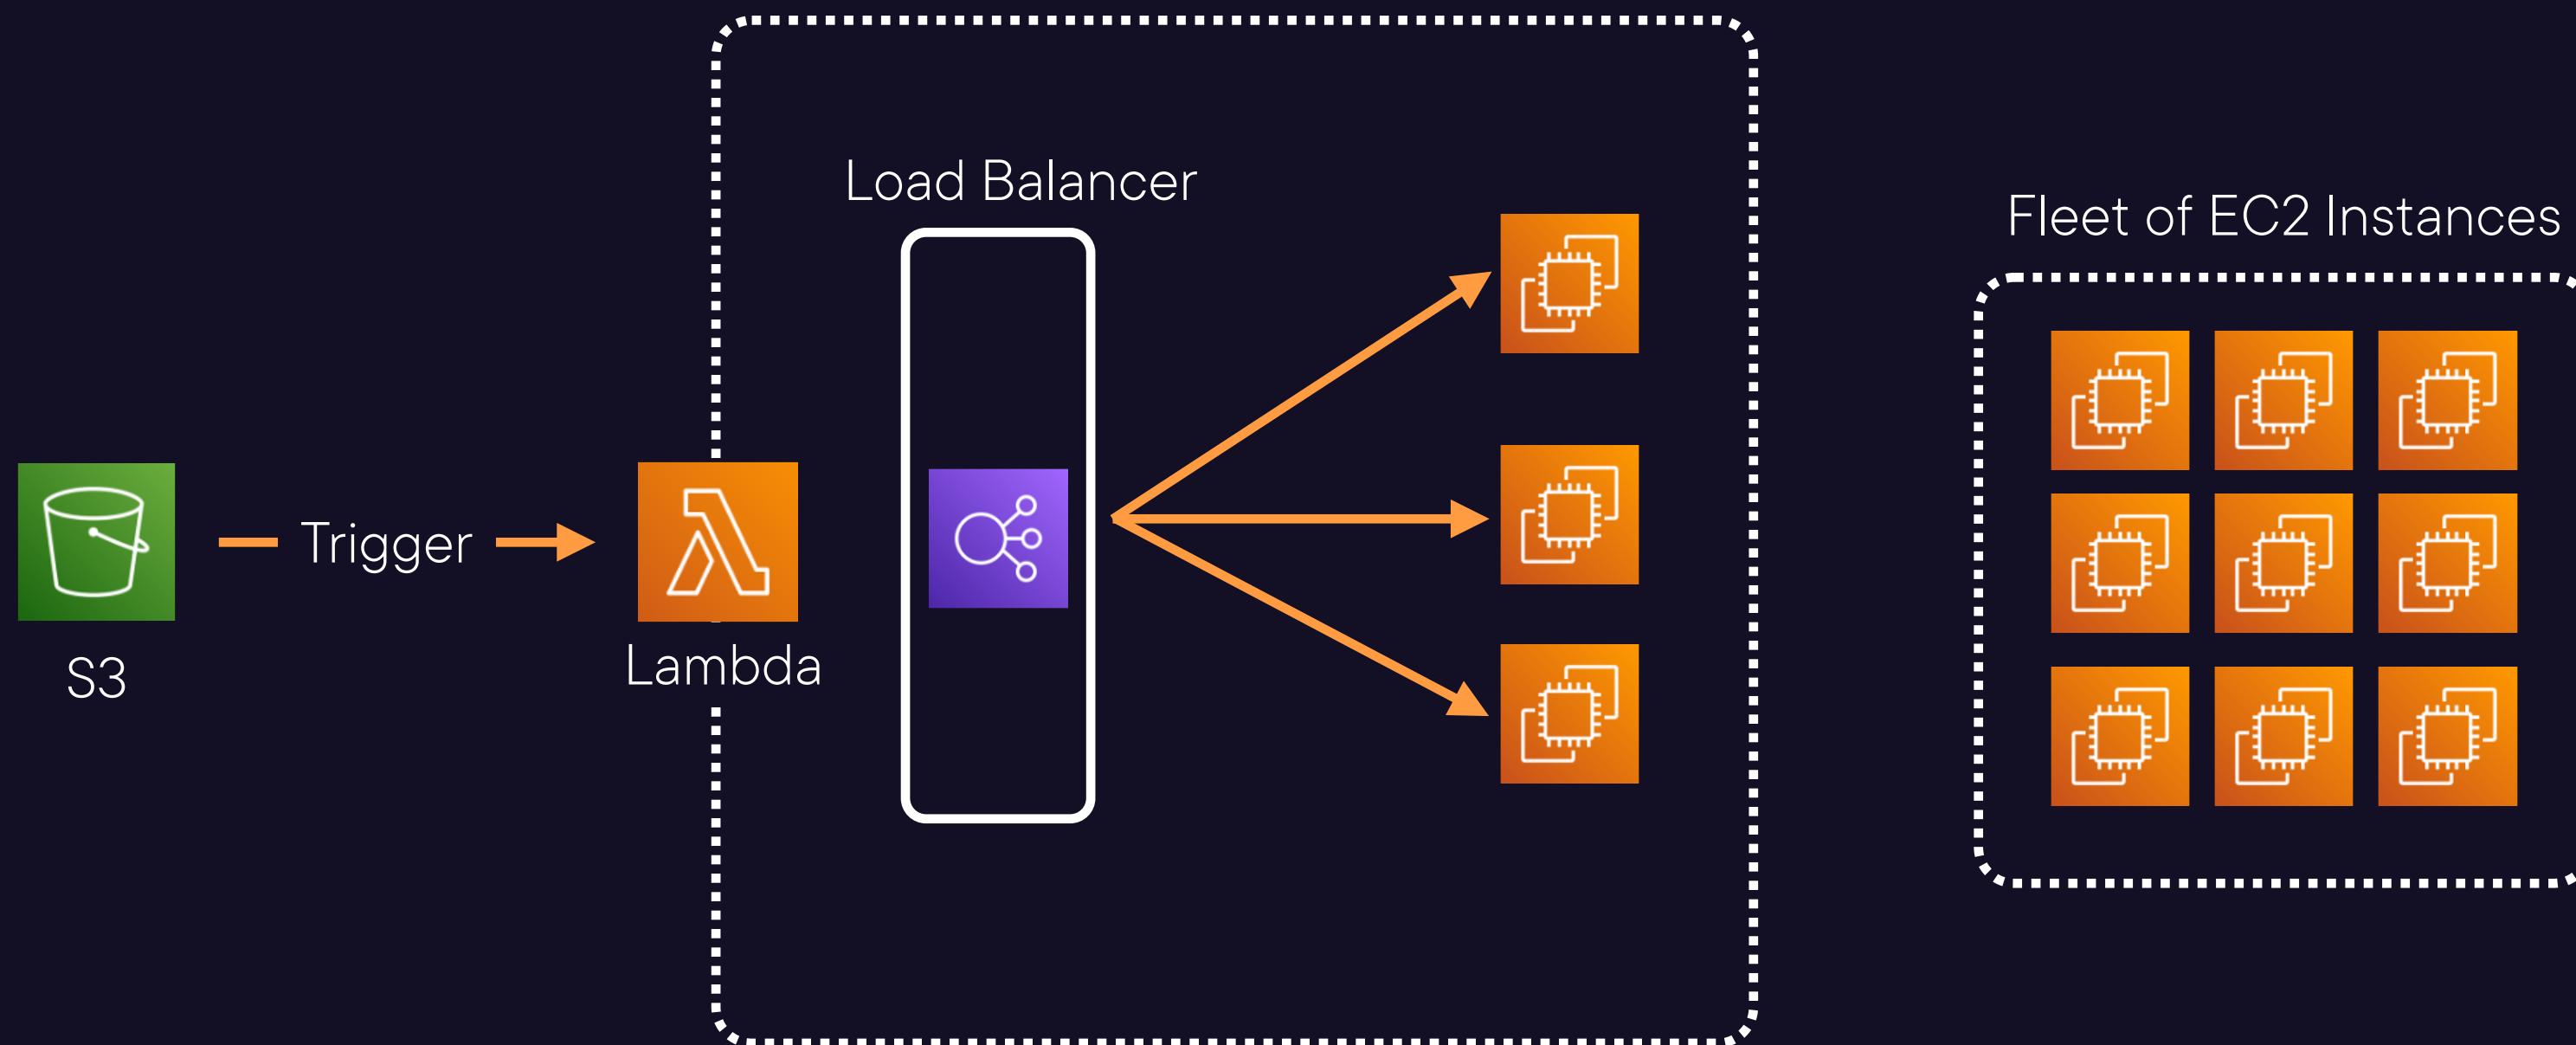
Each Lambda function runs in an isolated execution environment.

#### Execution environments:

- Manage the required resources
- Provide lifecycle support
- Language-specific



# How Lambda Works Under the Hood



# Lesson Summary



## Lambda's Building Blocks

- **Function:** A standalone unit of code
- **Trigger:** Resources that invoke Lambda
- **Event:** The data that is passed into functions
- **Execution Environment:** The runtime environment that manages the resources

## Under the Hood

- Load balancer and a fleet of EC2 instances



# Lambda Use Cases



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

---

# Lambda Use Cases

Processing Data

Orchestrating Workflows

Operating Serverless Applications

Processing Messages

Other Use Cases

Lesson Summary

---



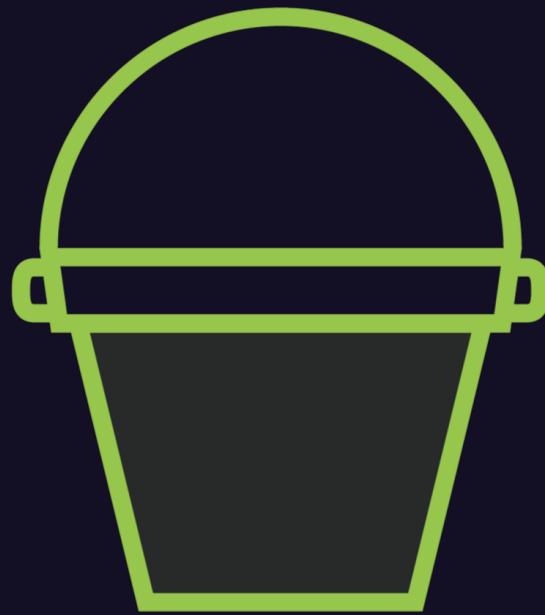
**Noreen Hasan**  
Training Architect



### Processing Data



Converting  
documents rapidly



Processing S3  
objects



Monitoring and  
analyzing logs

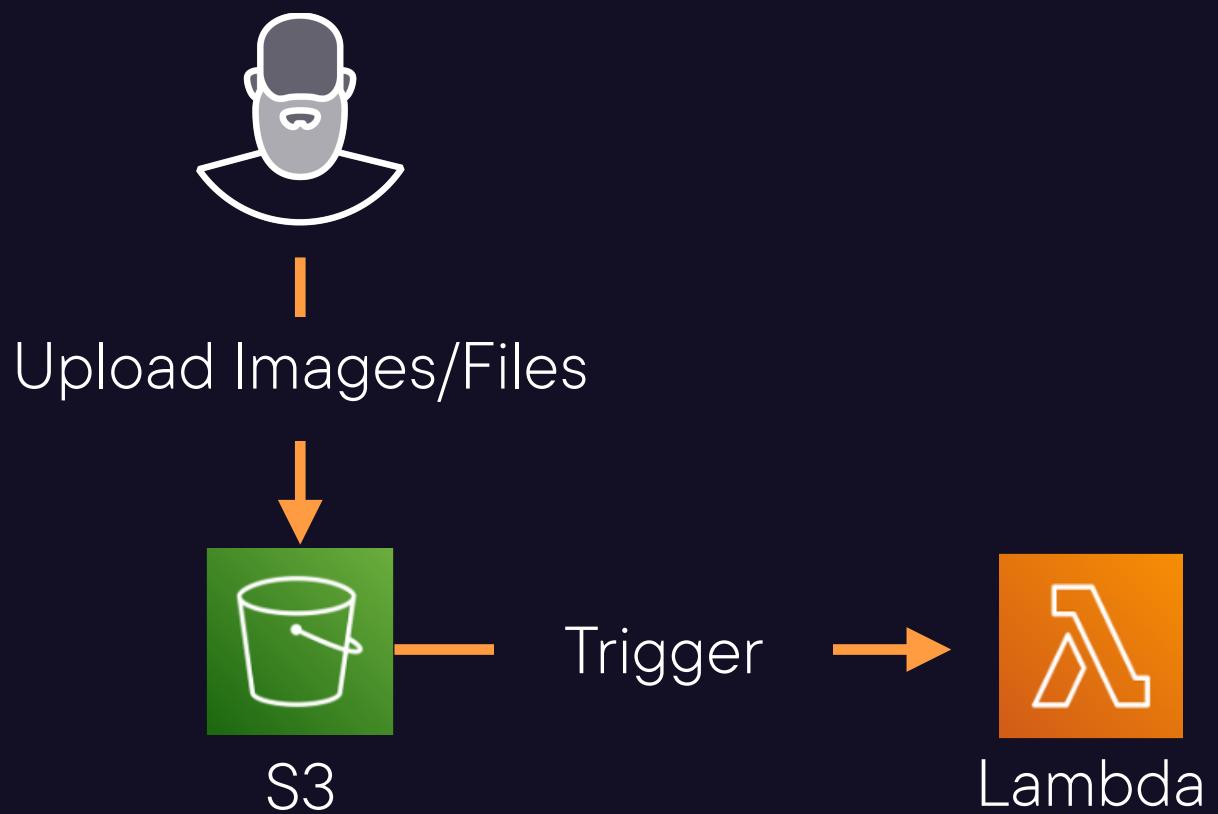


Processing  
streaming data



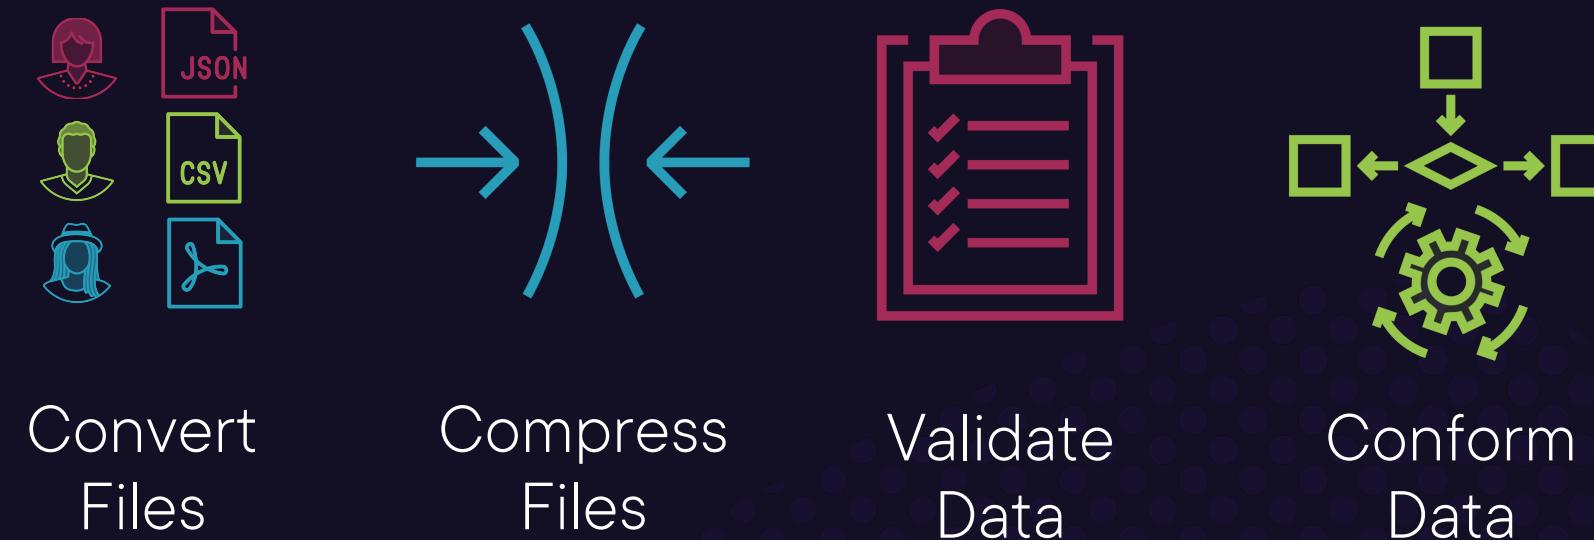
## Lambda Use Cases

### Processing Data

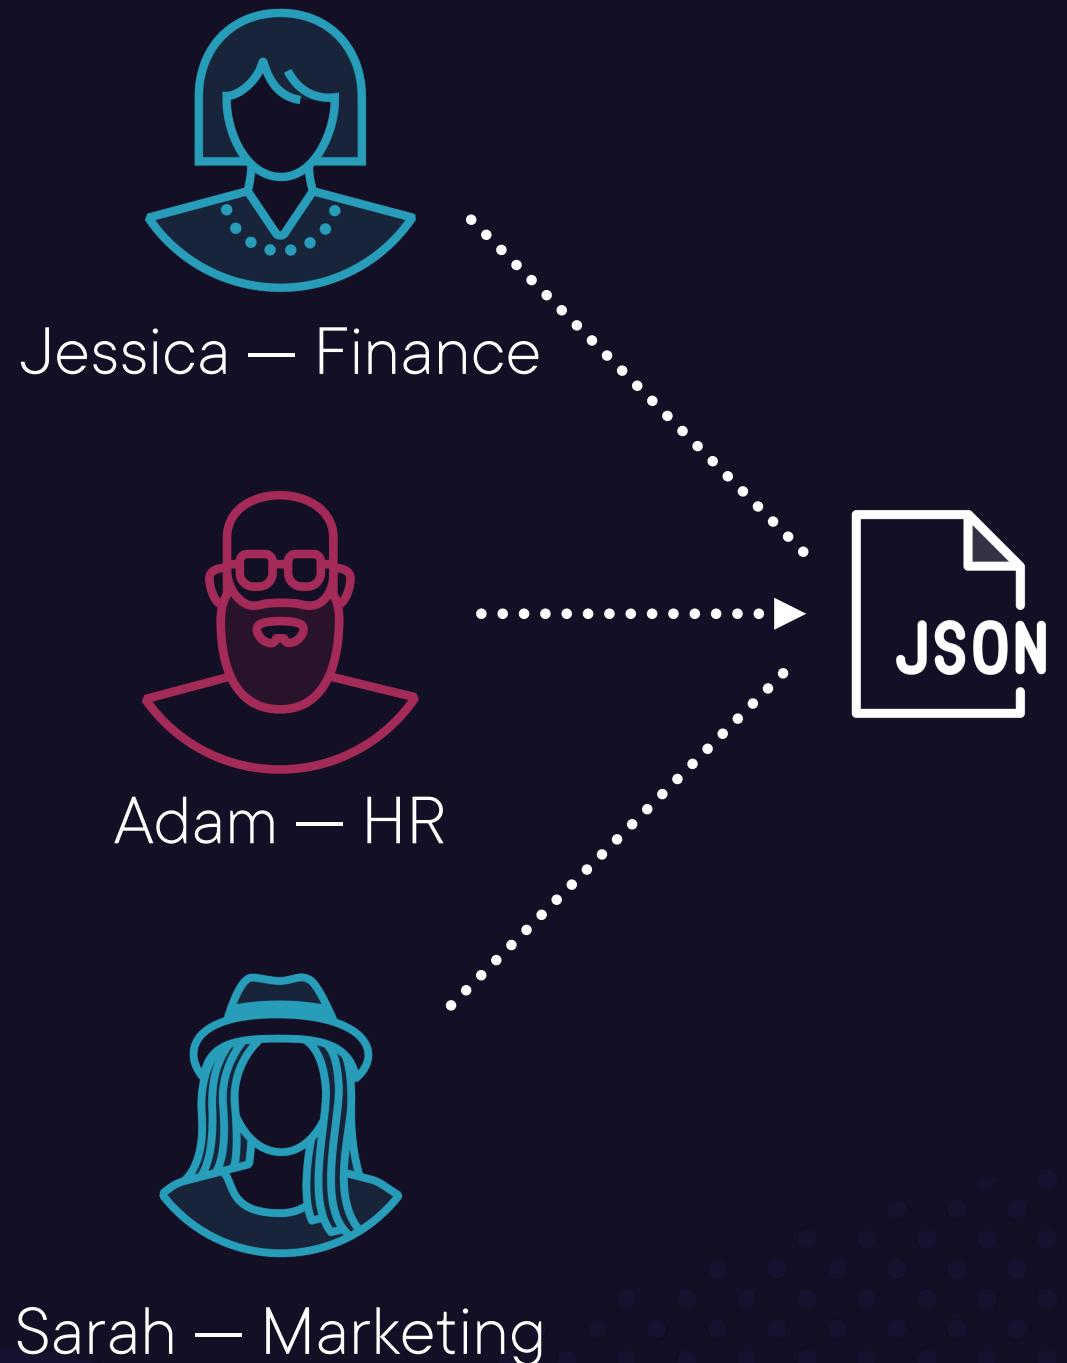


#### Processing S3 Objects

Trigger Lambda to modify files in S3:



# Processing Data



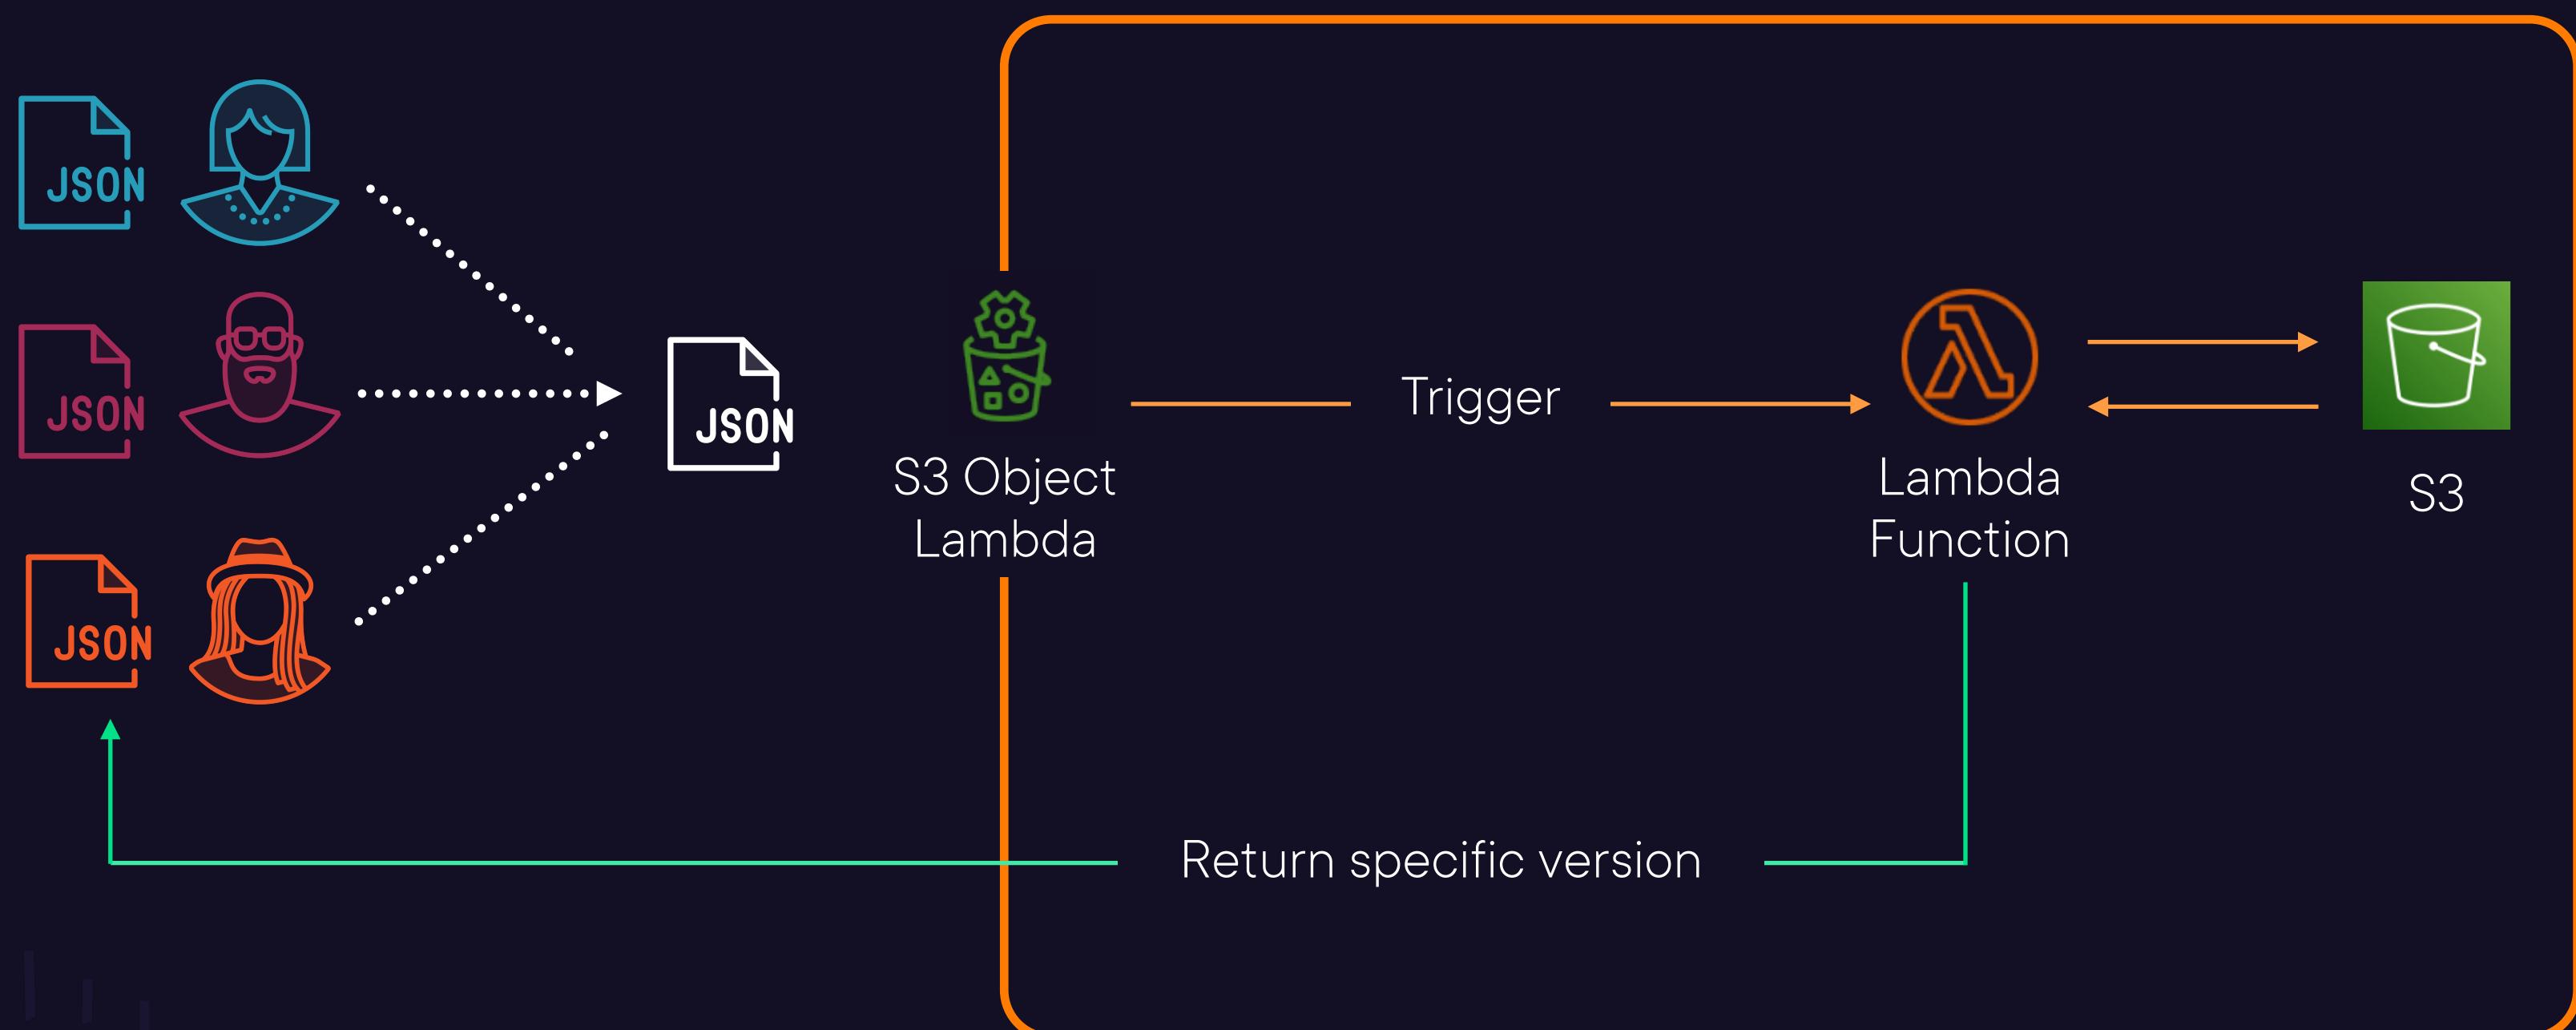
## S3 Object Lambda

- Allows you to create different views of the data for different users
- Automatically transform your data as it is being retrieved from S3



## Lambda Use Cases

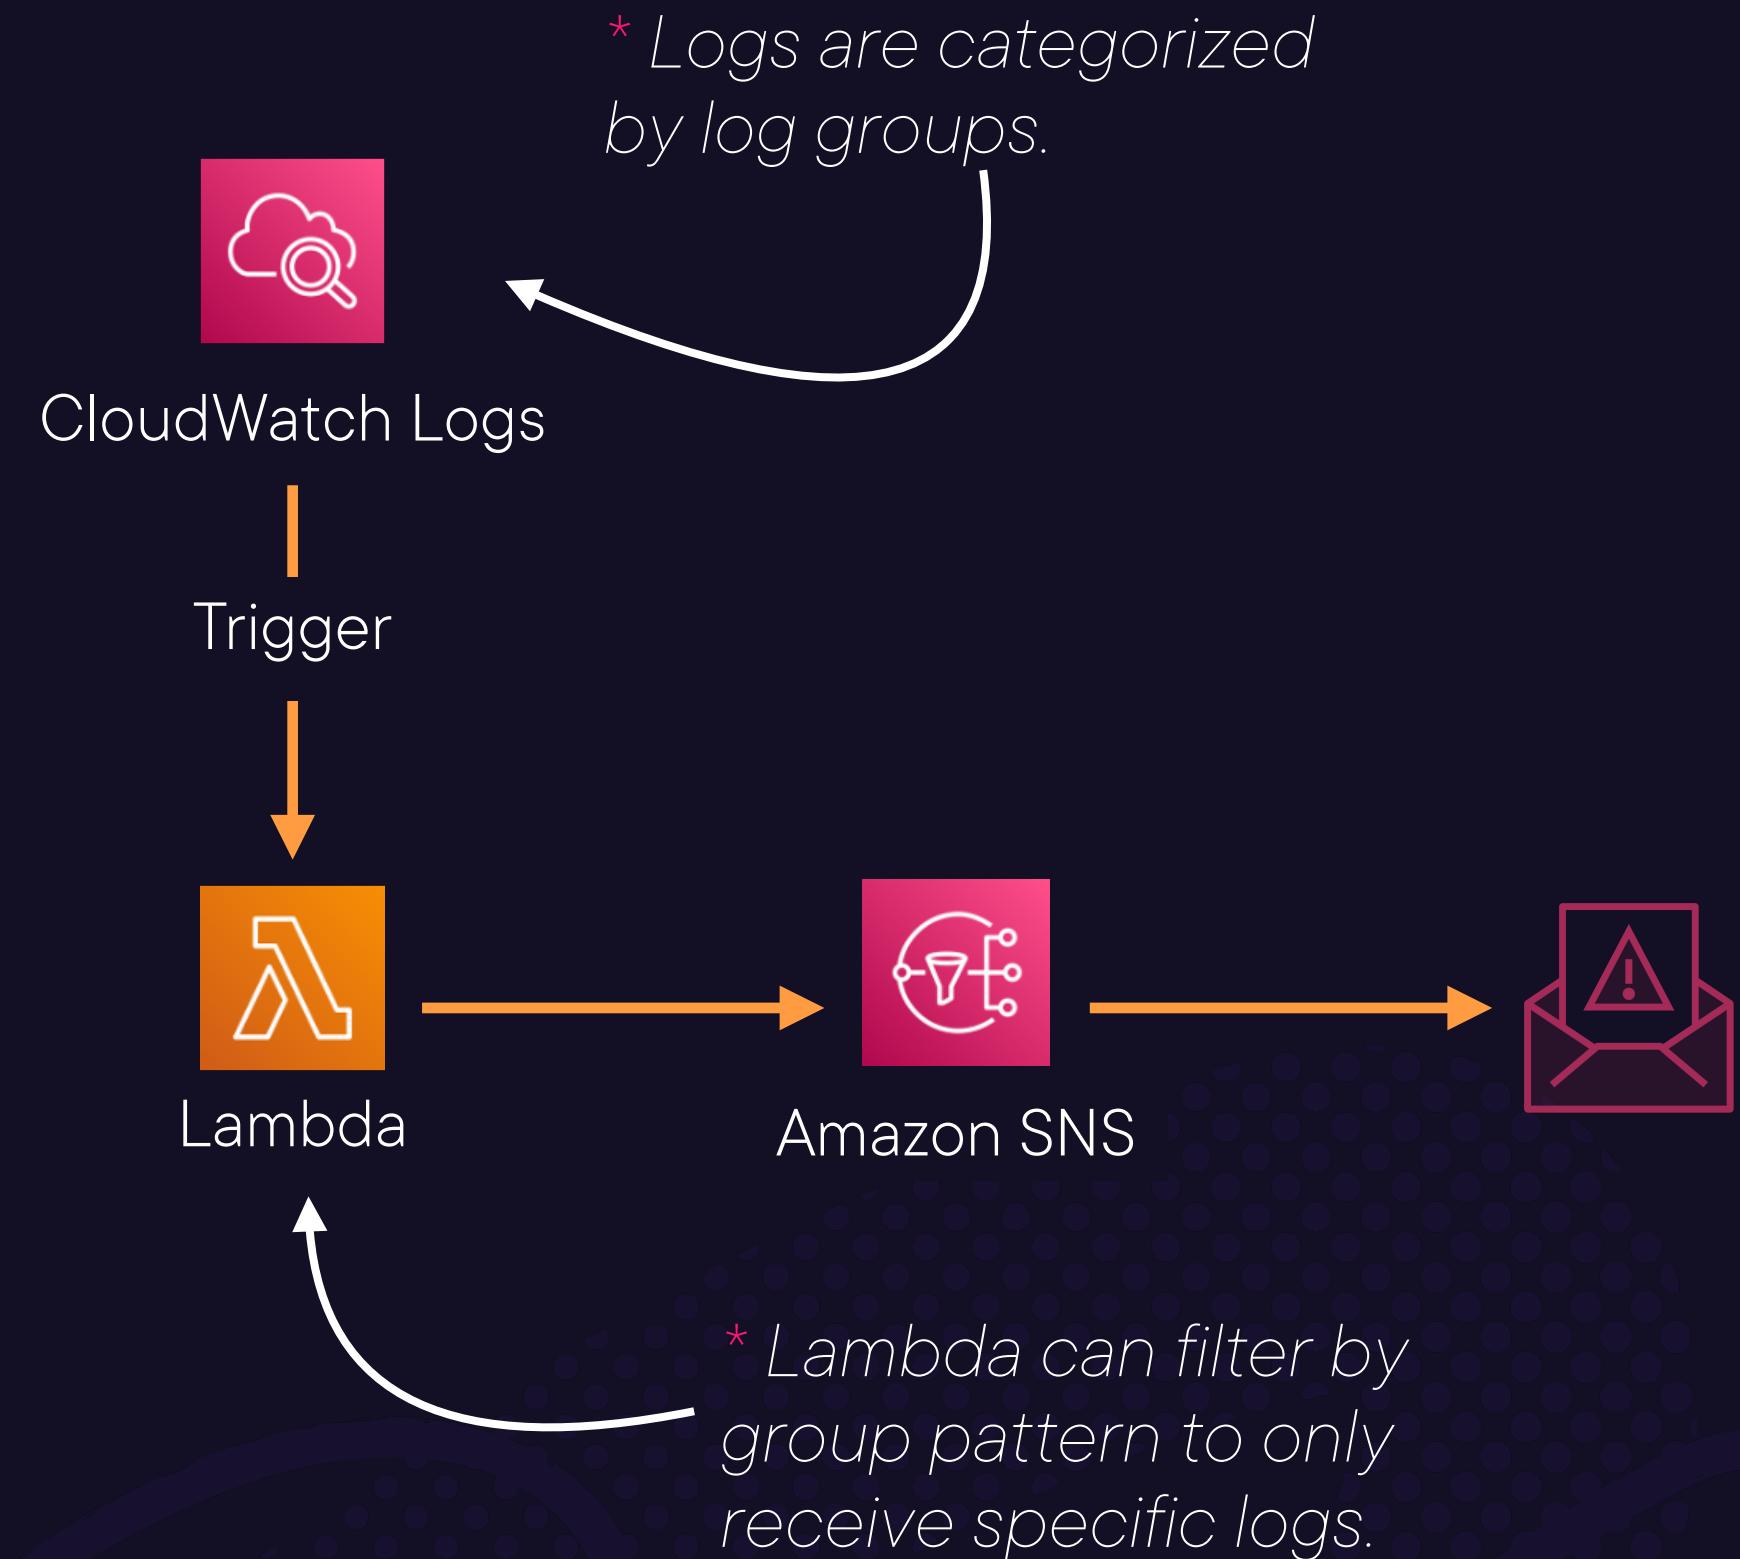
### Processing Data



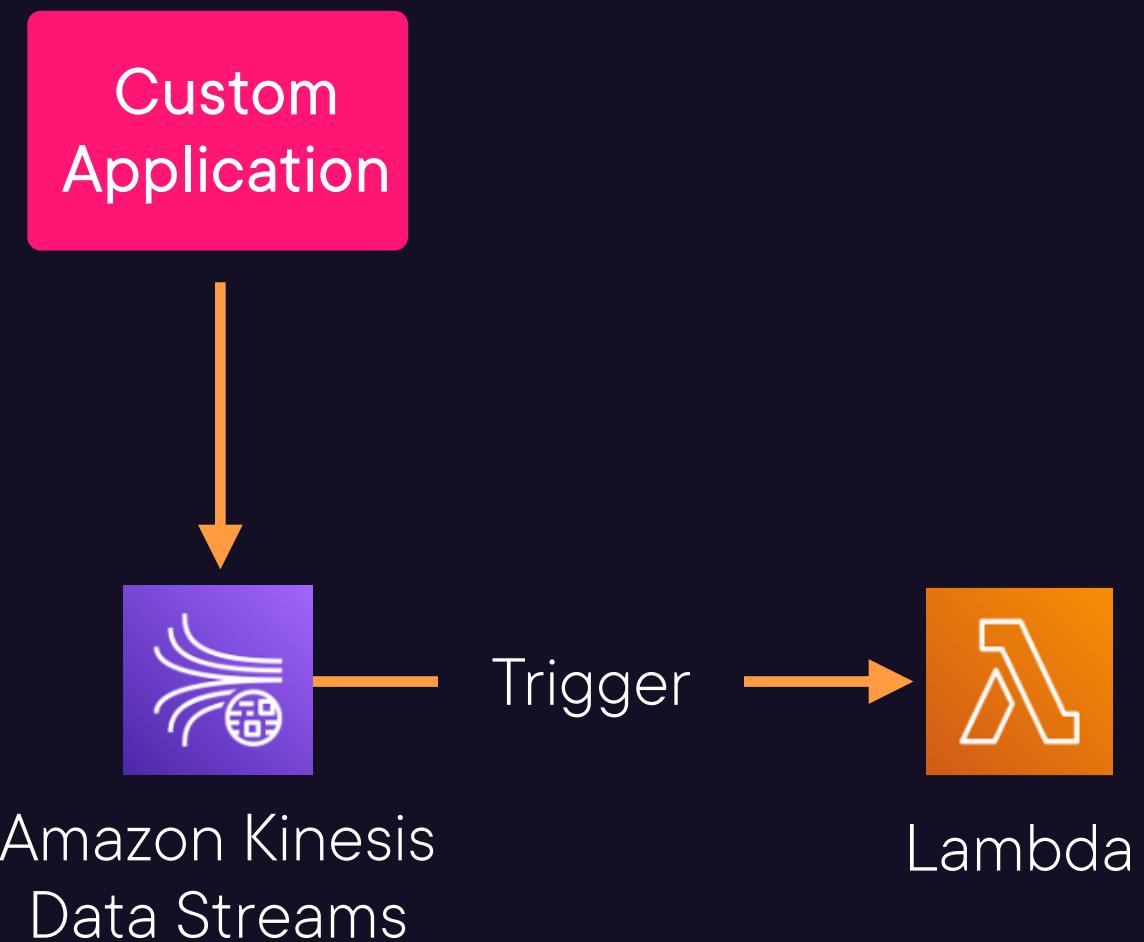
### Processing Data

Monitoring and analyzing logs.

Subscribe to Lambda to receive log events from CloudWatch Logs.



### Processing Data



#### Real-Time Stream Processing

Trigger Lambda to process data in real time one record at a time.

#### AWS Services

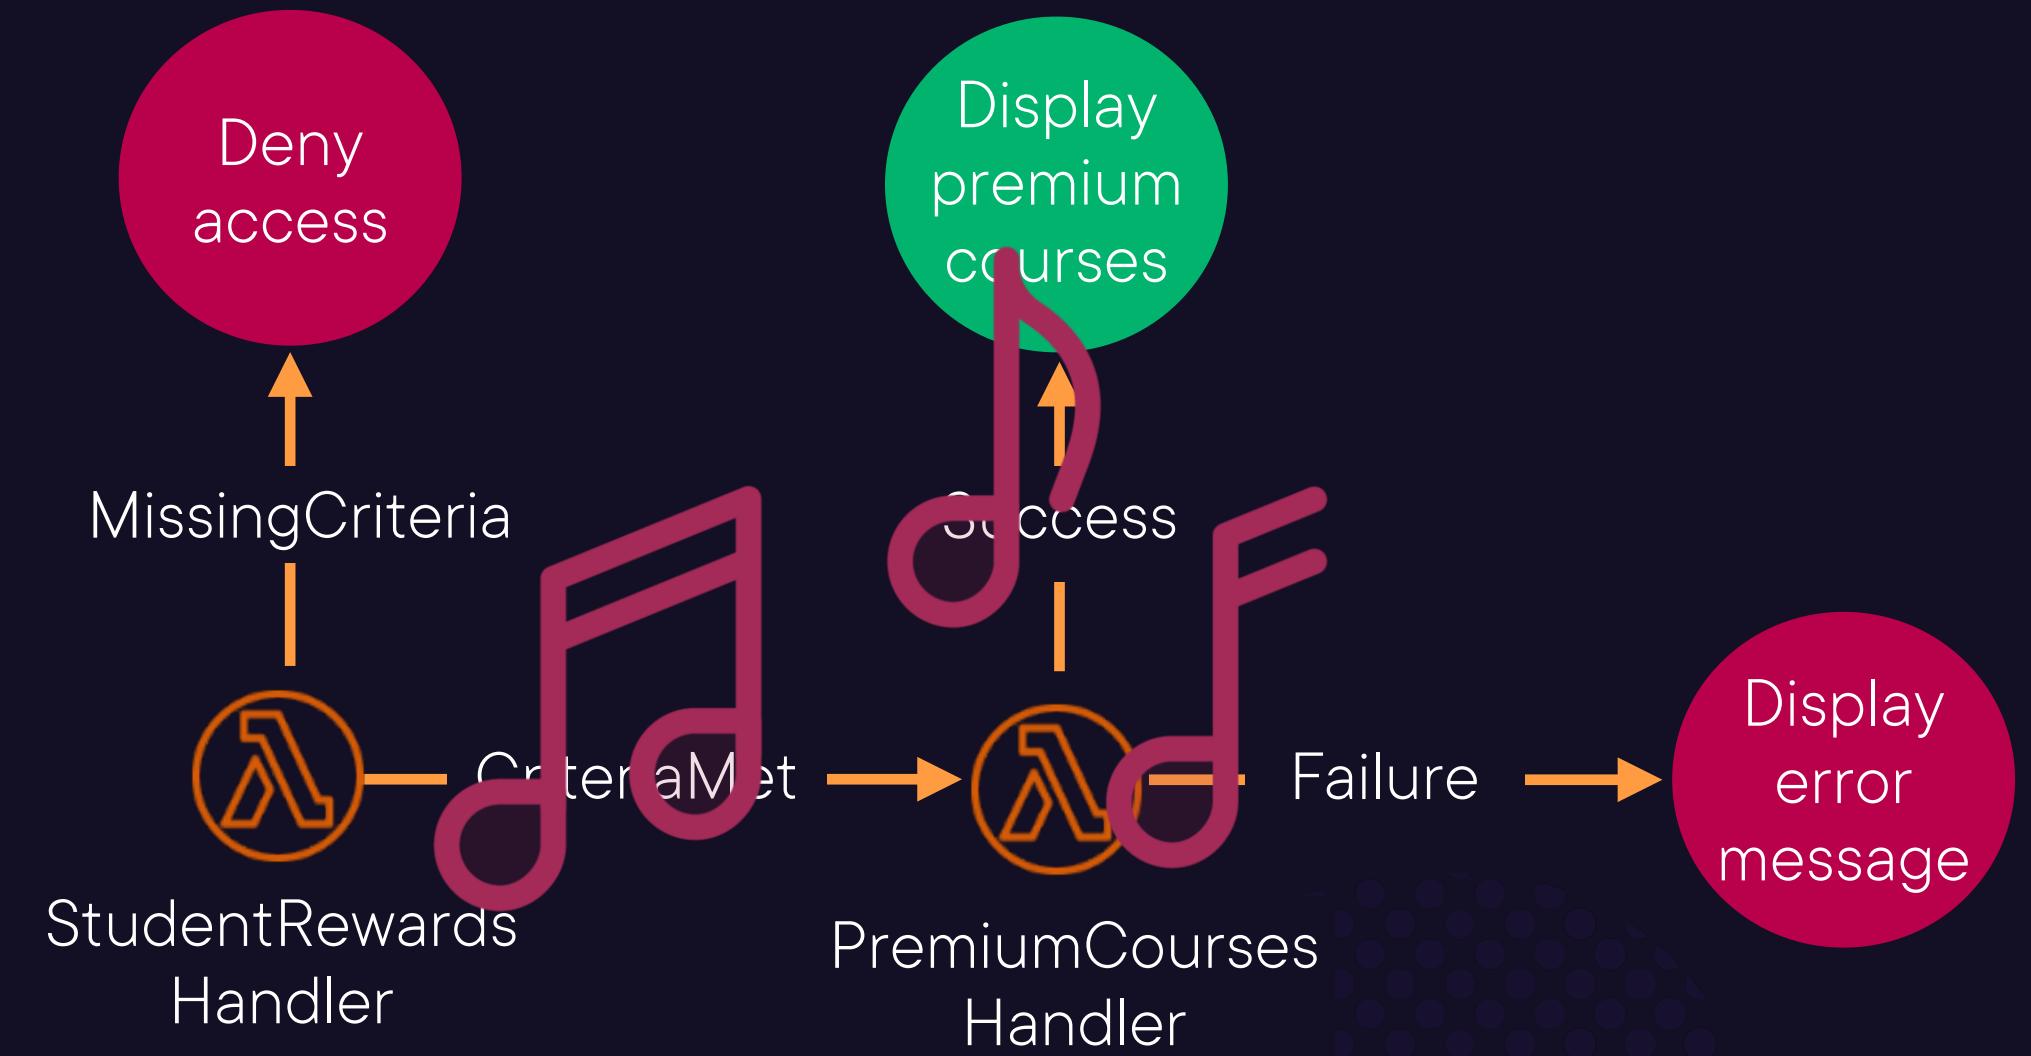
- Processing data in IoT
- Processing data from Kinesis Data Streams



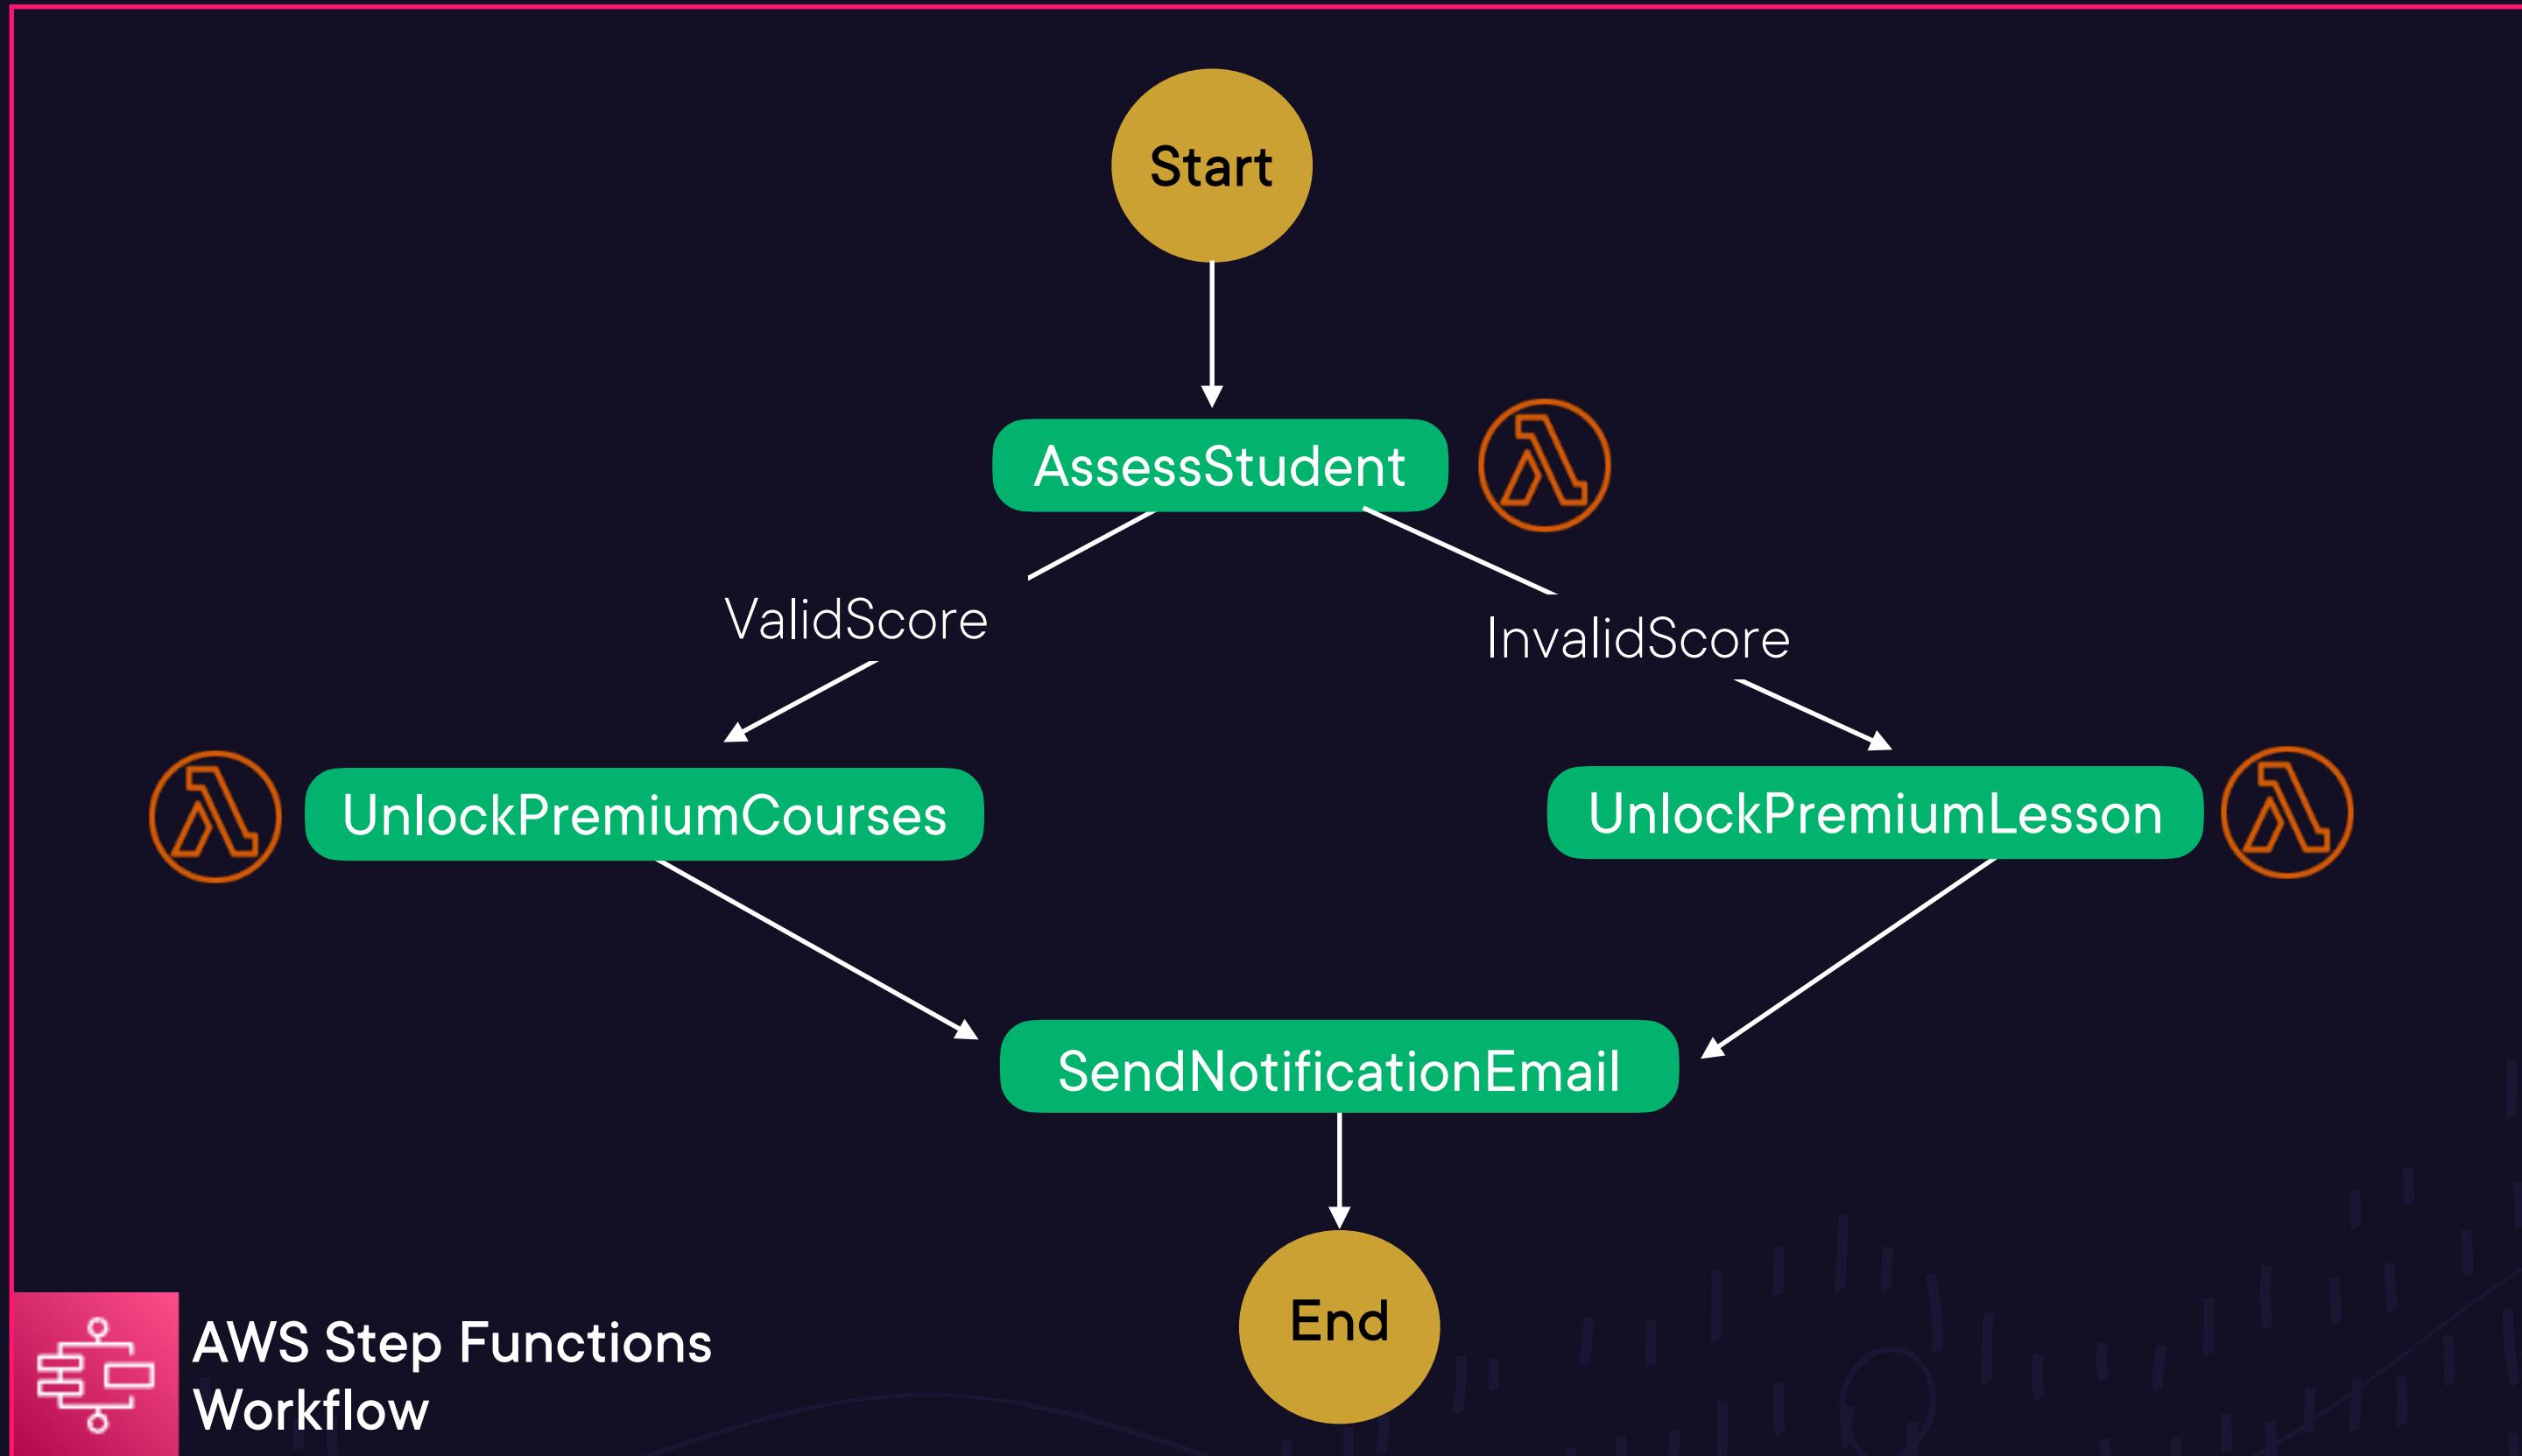
# Orchestrating Workflows

## Orchestration

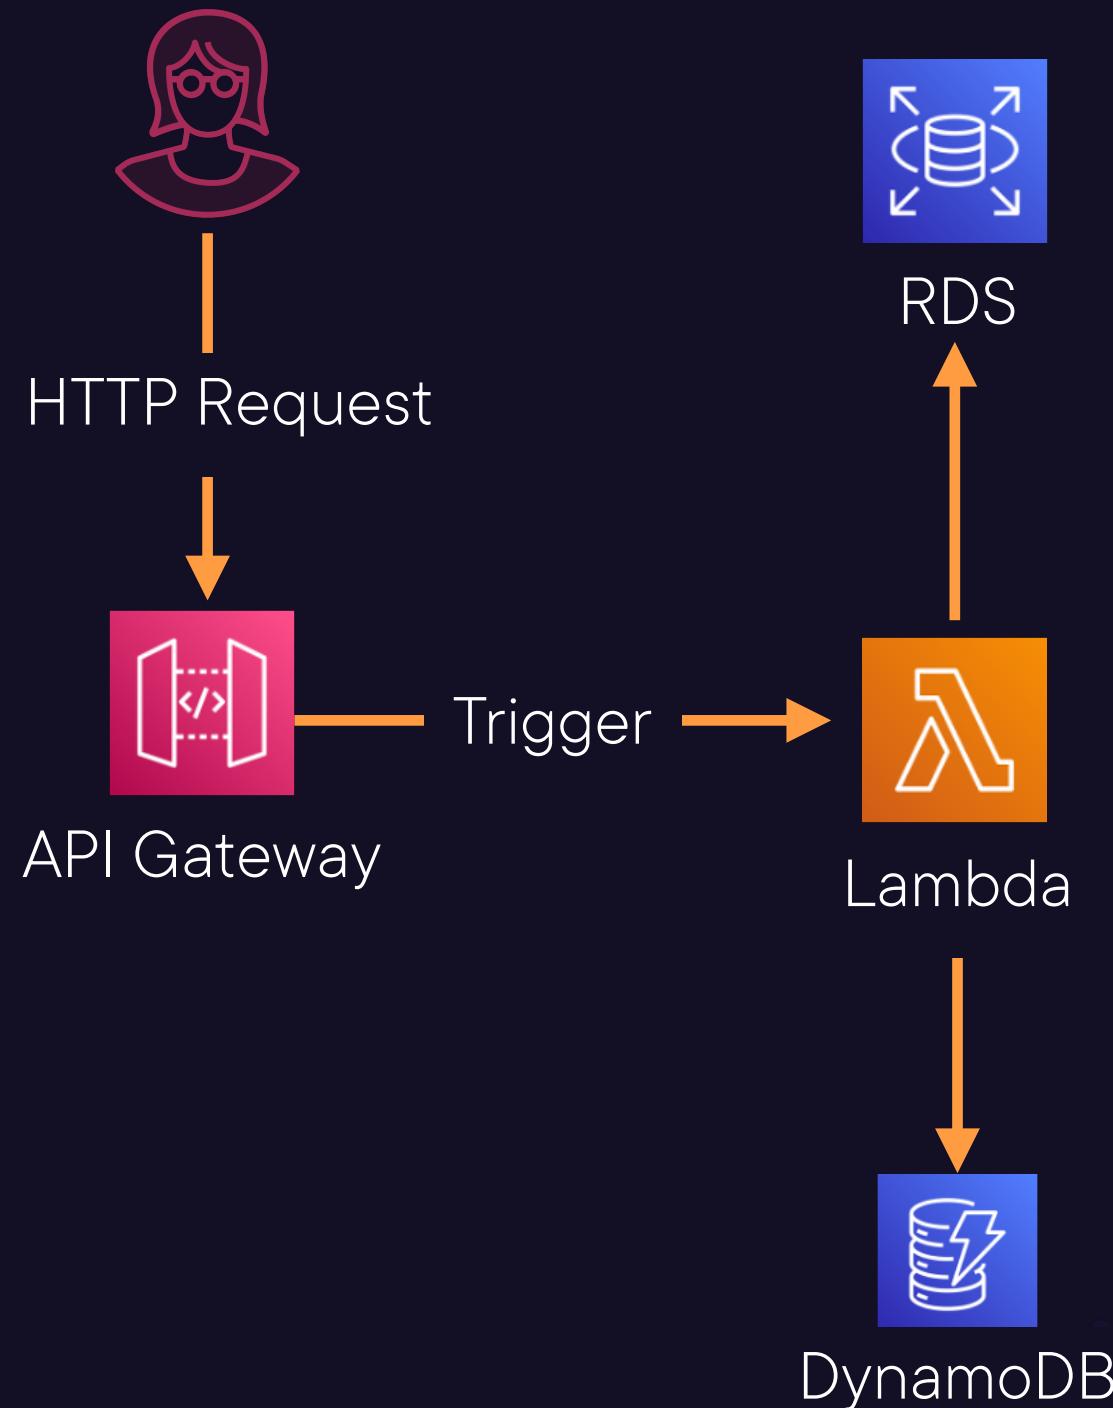
- Workflows that involve branching logic
- Trigger Lambda to determine the next step in an orchestration workflow
- Integrates with AWS Step Functions



# Orchestrating Workflows



# Operating Serverless Applications

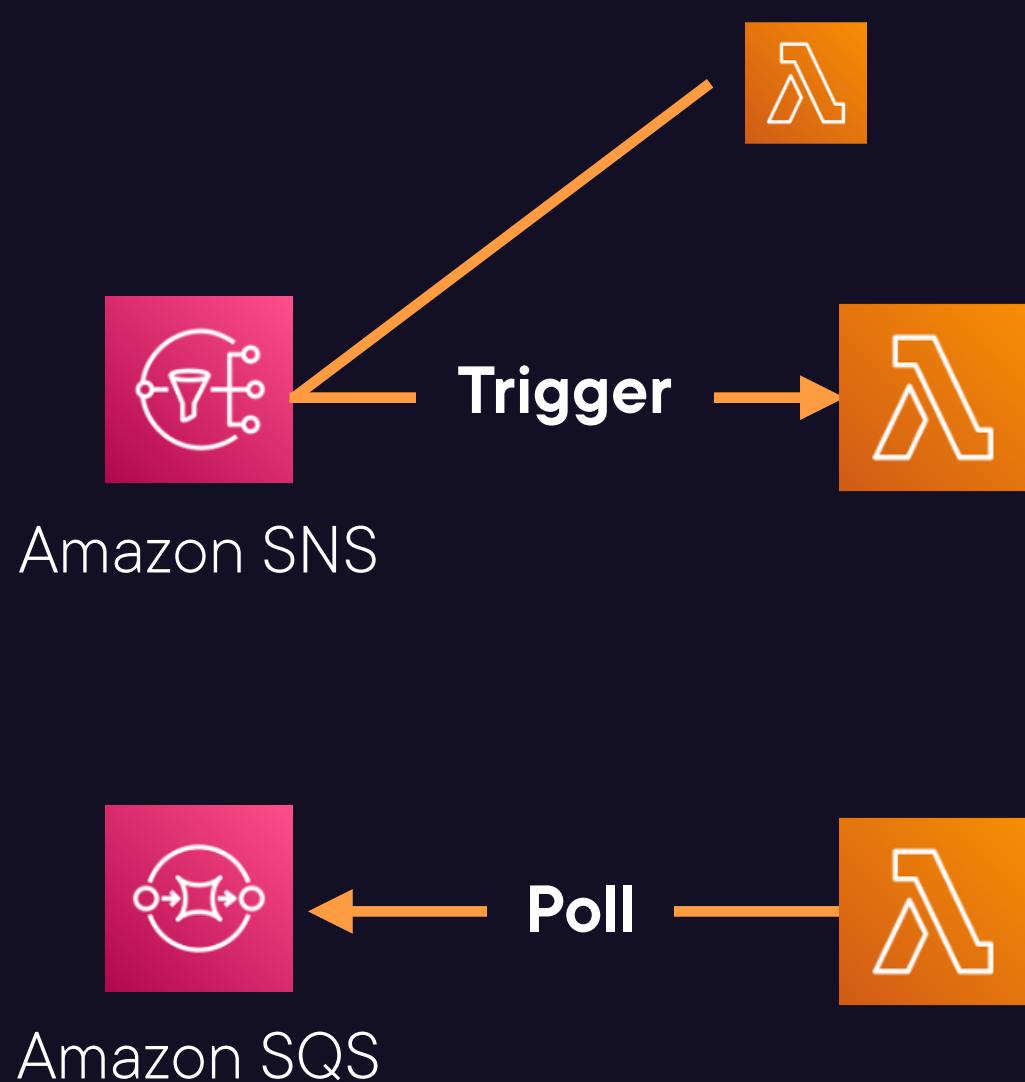


## Serverless Applications

- Integrates with API Gateway to expose the Lambda function to external applications (e.g., web or mobile app).



# Processing Messages



## Simple Notification Service (SNS)

**Trigger** Lambda to process SNS messages.

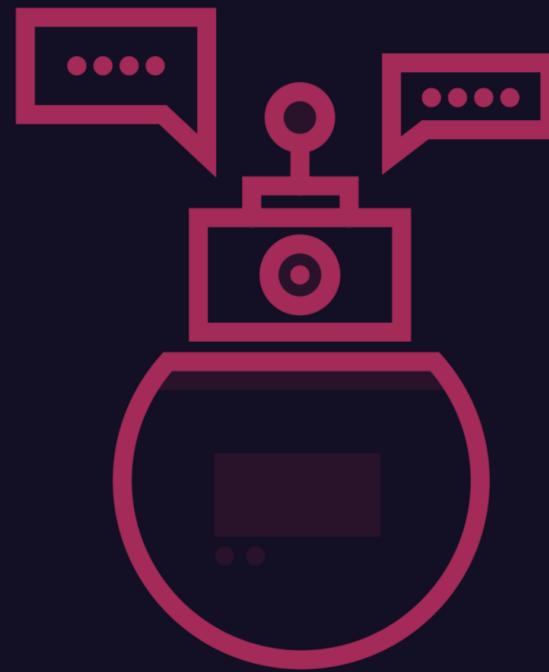
## Simple Queue Service (SQS)

Get Lambda to **poll** SQS to look for messages in the queue.

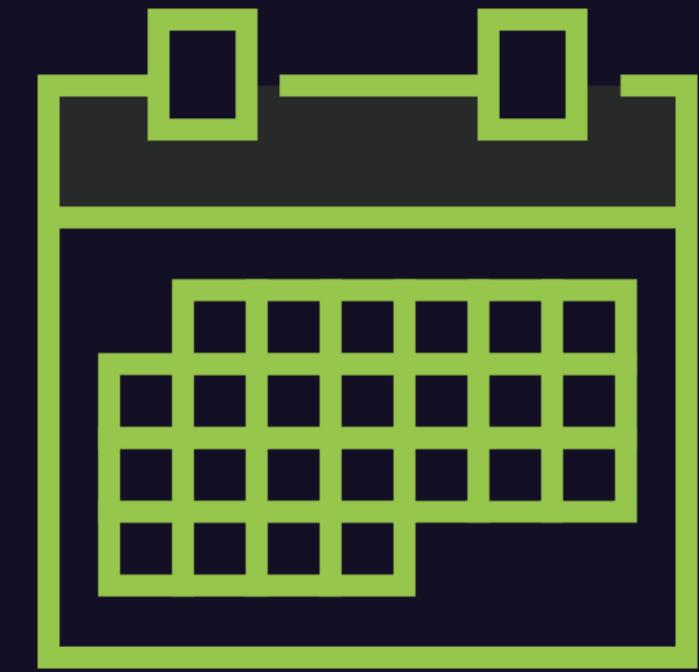


# Alexa

Integrate with Alexa through Alexa Skills.



Build real-time chatbots using Lambda, SQS, SNS, and AWS Lex.



Schedule jobs (i.e., starting/stopping EC2 instances).



# Lesson Summary



## Processing Data

Integrating **Lambda** with **S3** to:

- Convert documents rapidly
- Compress files
- Validate data
- Conform data

## Orchestrating Workflows

Integrating **Lambda** with **Step Functions** for branching logic.

## Processing Messages

- Getting triggered by **SNS**
- Polling from **SQS**

## Other Use Cases

- **Alexa**
- **Chatbots**
- Scheduling jobs



# When to Use Lambda



**Noreen Hasan**

Training Architect



# When to Use Lambda

## LESSON BREAKDOWN

Lambda Constraints

Alternative Services: Overview

EC2

Elastic Beanstalk

ECS

Deciding Factors

Lesson Summary



**Noreen Hasan**  
Training Architect



## When to Use Lambda

### Lambda Constraints

Lambda is no



*Execution Timeout:*



< 15 mins



## When to Use Lambda

### Lambda Constraints

*Disk Space:*

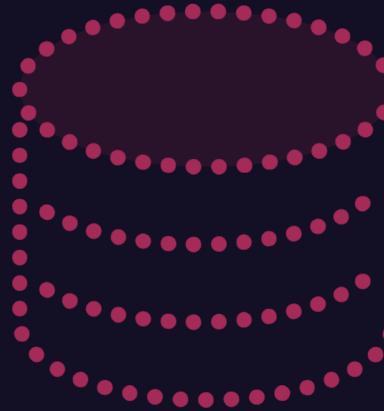
**512 MB >  < 10 GB**



## When to Use Lambda

### Lambda Constraints

*Memory:*

**128 MB >  < 10 GB**

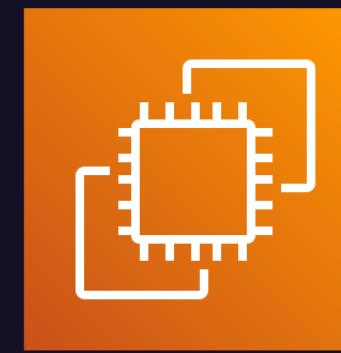


## When to Use Lambda

# Alternative Services: Overview



**AWS Elastic Beanstalk**



**Amazon Elastic  
Compute Cloud (EC2)**



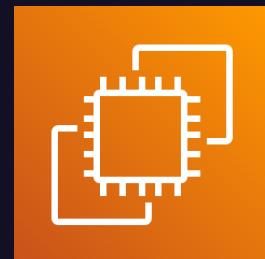
**Amazon Elastic  
Container Service (ECS)**



# When to Use Lambda Alternative Services



Lambda ..... → Function as a Service (FaaS)



Elastic Beanstalk ..... → Platform as a Service (PaaS)



EC2 ..... → Infrastructure as a Service (IaaS)



ECS ..... → Container as a Service (CaaS)



# AWS Elastic Beanstalk

What is it? When to use it?



# ELASTIC BEANSTALK

## AWS Elastic Beanstalk

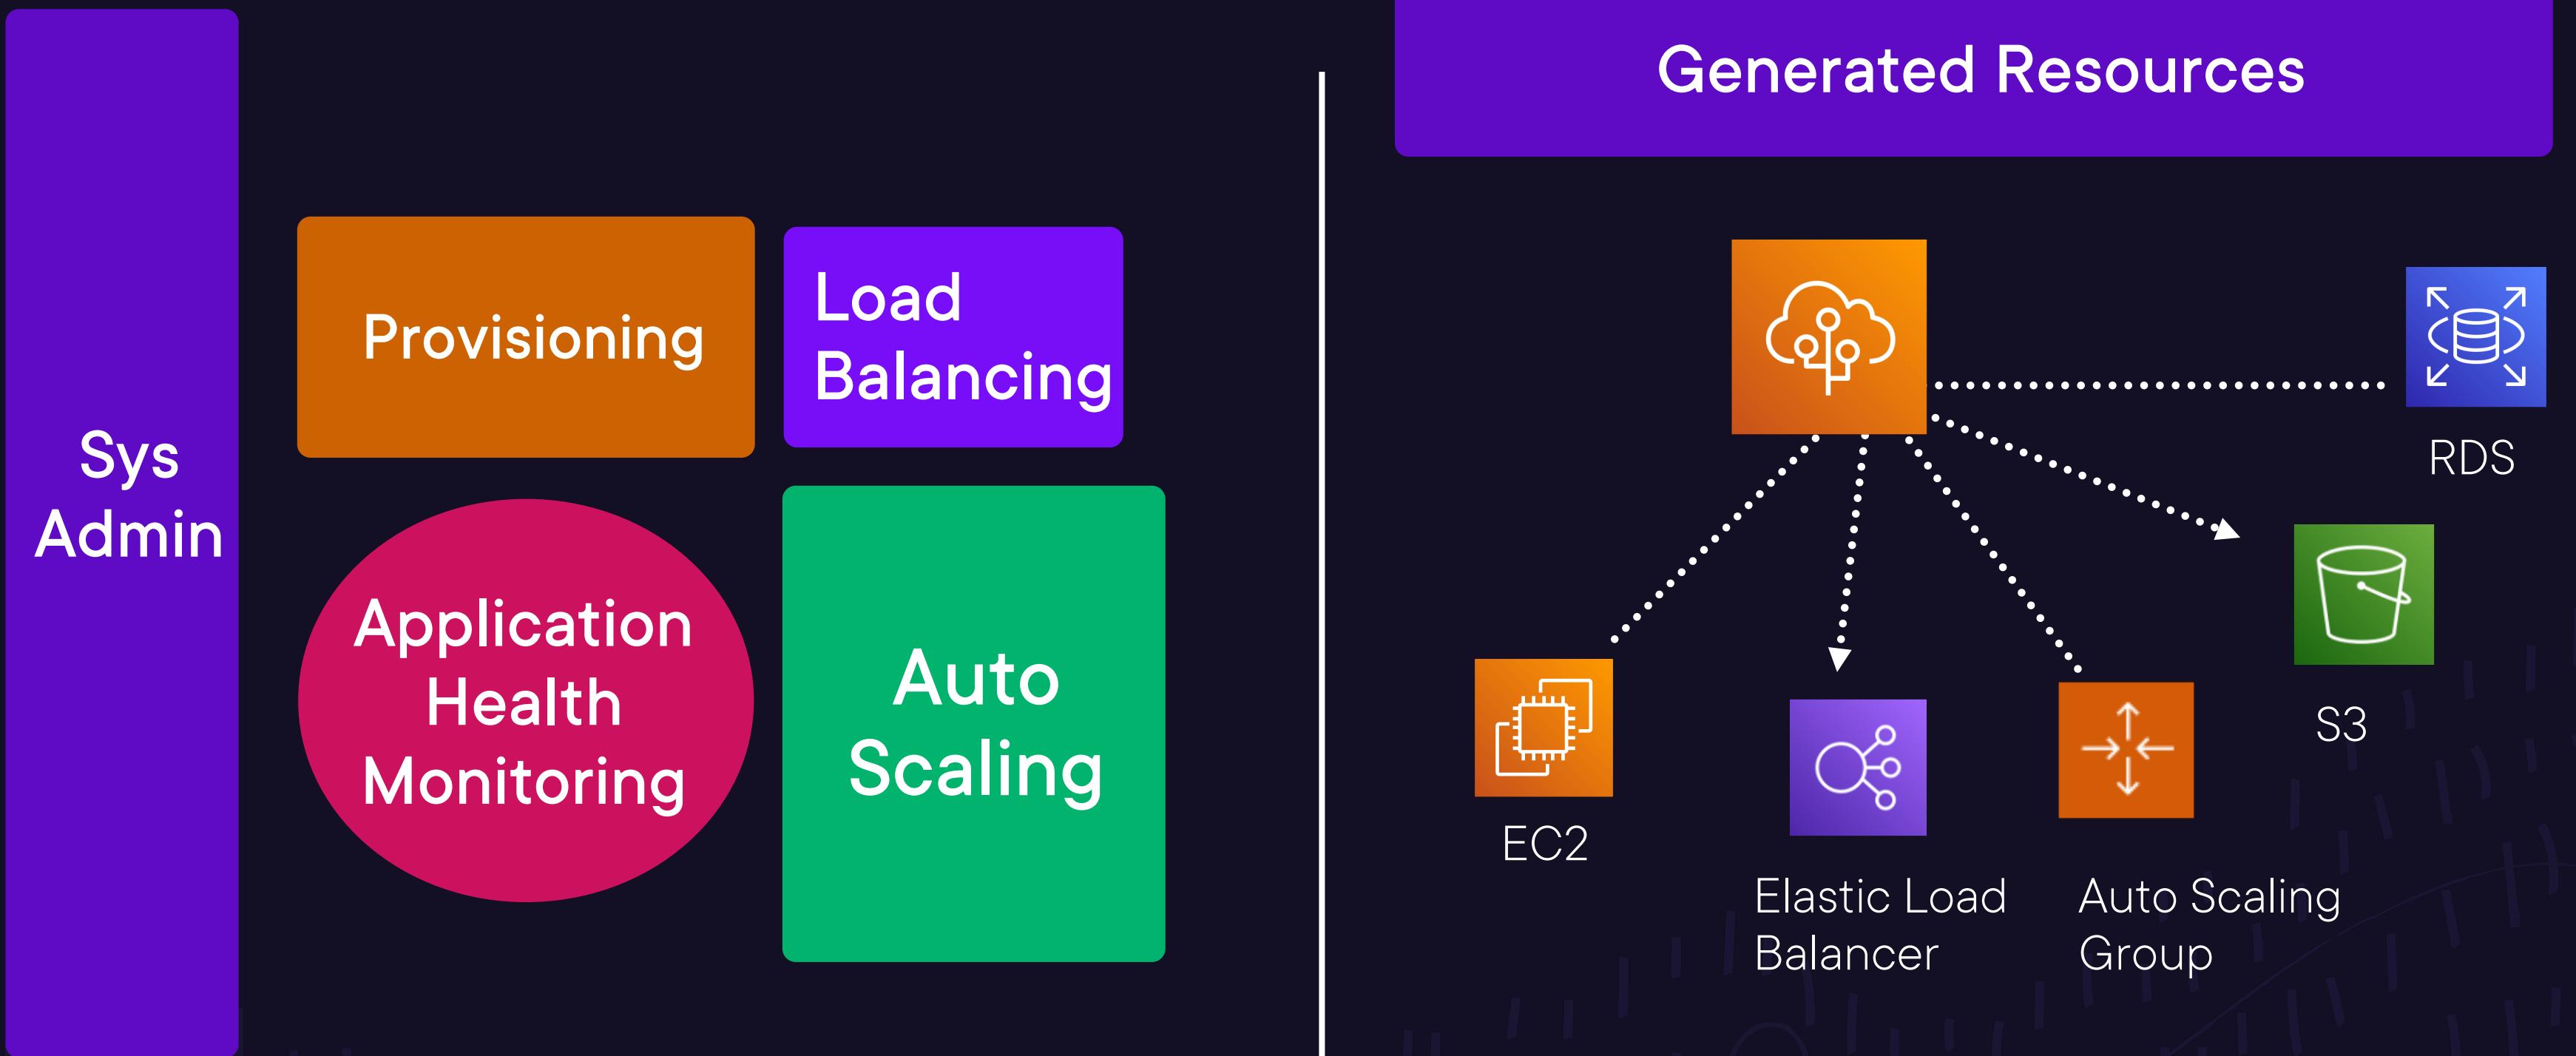
- Compute service for hosting applications
- Abstracts system administration

## Use Cases

- Automation
- Hosting and deploying **fully functional** applications



# When to Use Lambda Elastic Beanstalk



# AWS Elastic Beanstalk

How does it differ from Lambda?





Elastic Beanstalk reduces  
complexity without restricting your  
choices and control.



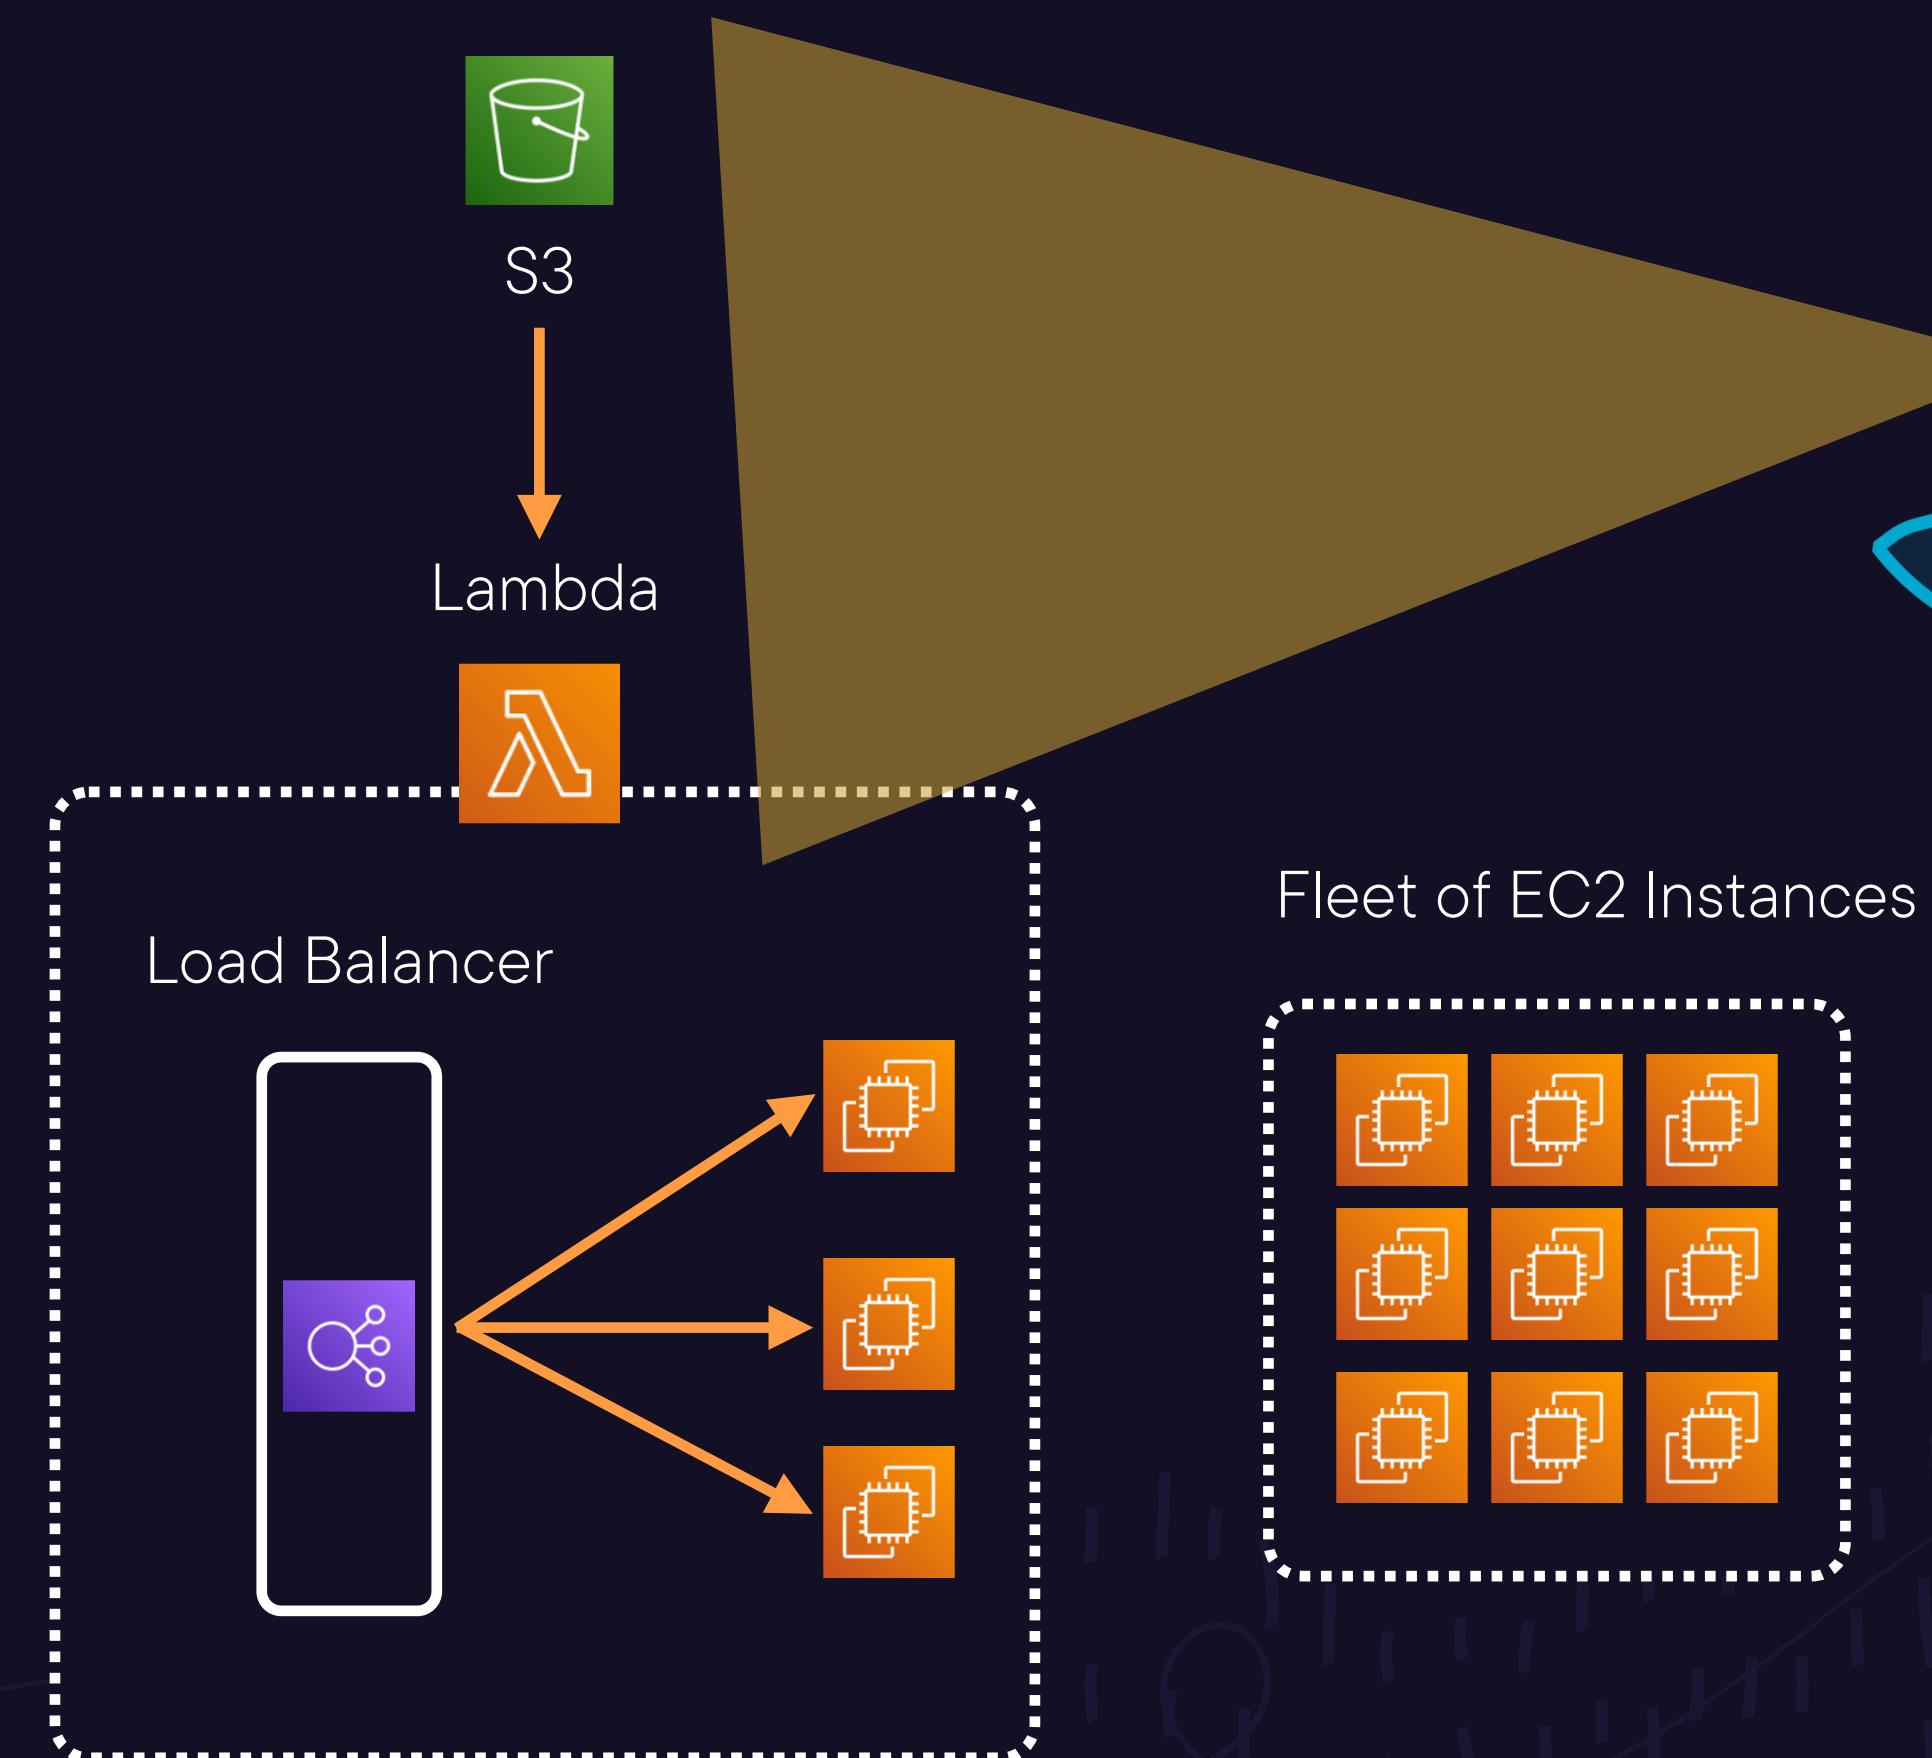


Lambda reduces complexity  
by abstracting provisioning.  
Due to limited choices, you  
can set it and forget it.



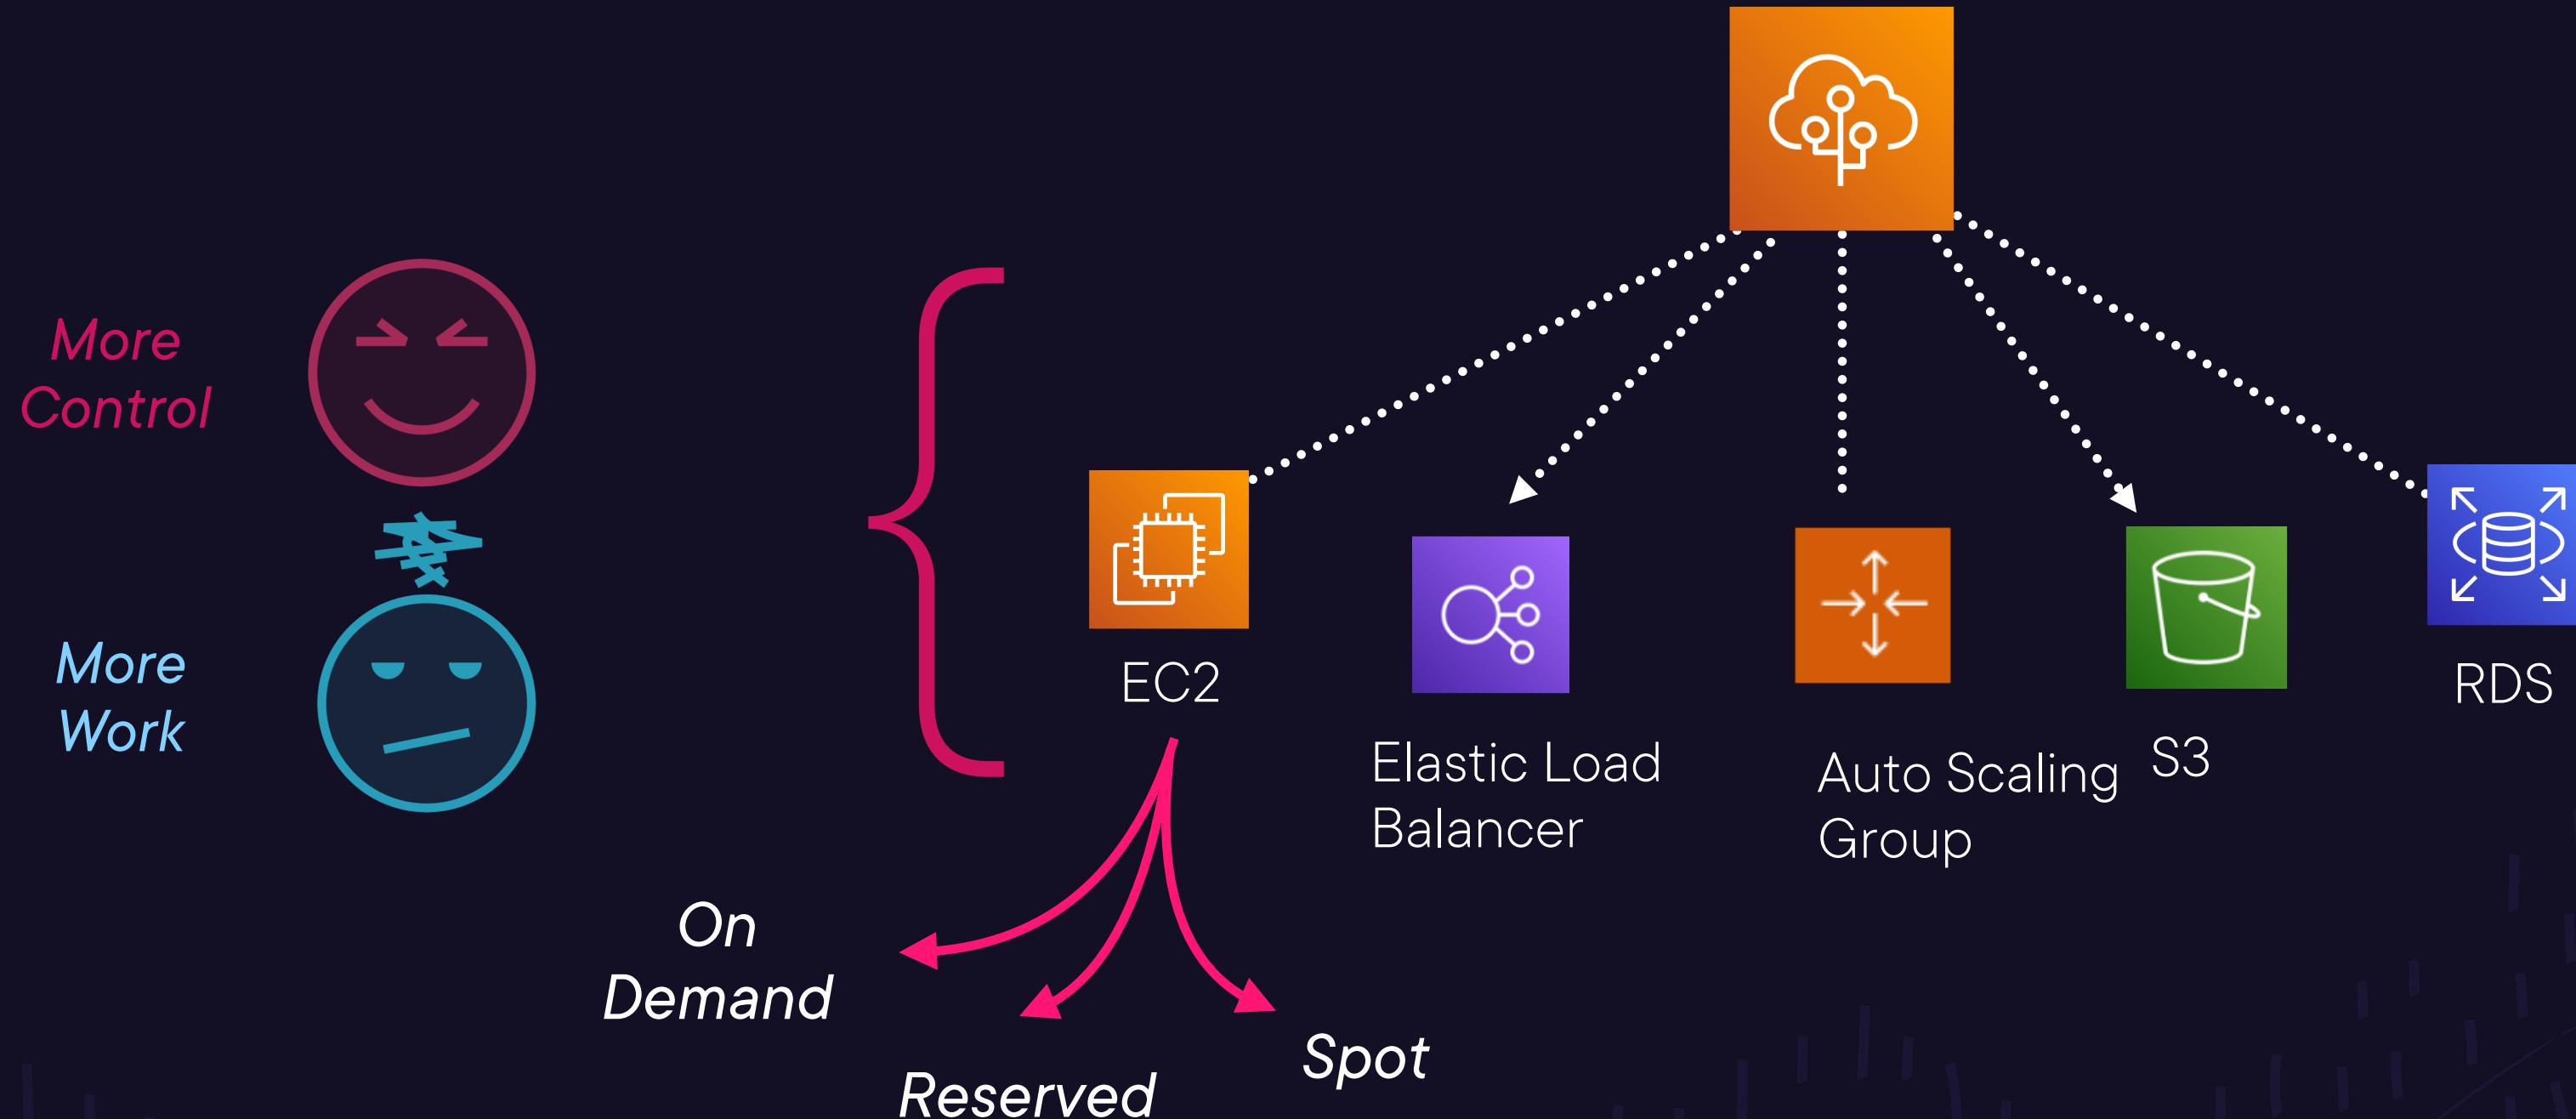
## When to Use Lambda

### Elastic Beanstalk vs. Lambda



## When to Use Lambda

# Elastic Beanstalk Offers Flexibility



## When to Use Lambda

### Elastic Beanstalk: Pricing

- Elastic Beanstalk is free.
- You only get charged for the used resources.



# EC2

What is it? When to use it?





## EC2

- Rentable virtual machines (**instances**)
- Allows you to spin up/down instances

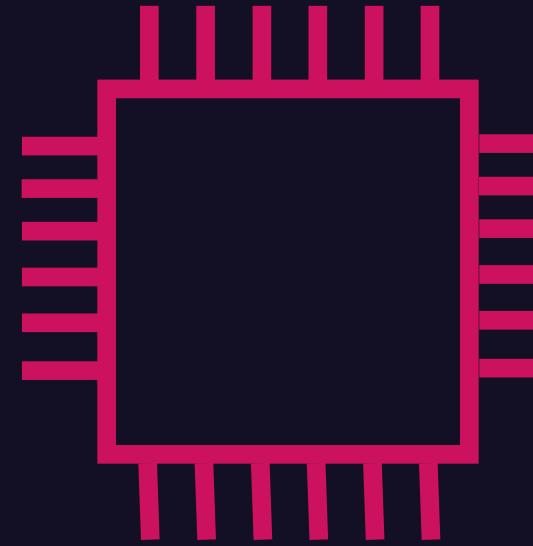
## Use Cases

- Hosting applications or databases (e.g., WordPress and MySQL)
- **High flexibility** needed and **high expertise** available



## When to Use Lambda

# EC2 Instance Components



Amazon Machine Image (AMI)  
(Windows or Linux)

Instance Type  
(Processing Power)

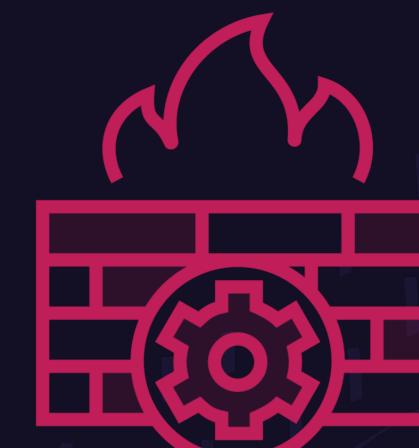
Memory  
(RAM)



EBS  
(Local Storage)



IP Addressing  
(Access)



Security Groups



# EC2

How does it differ from Lambda?



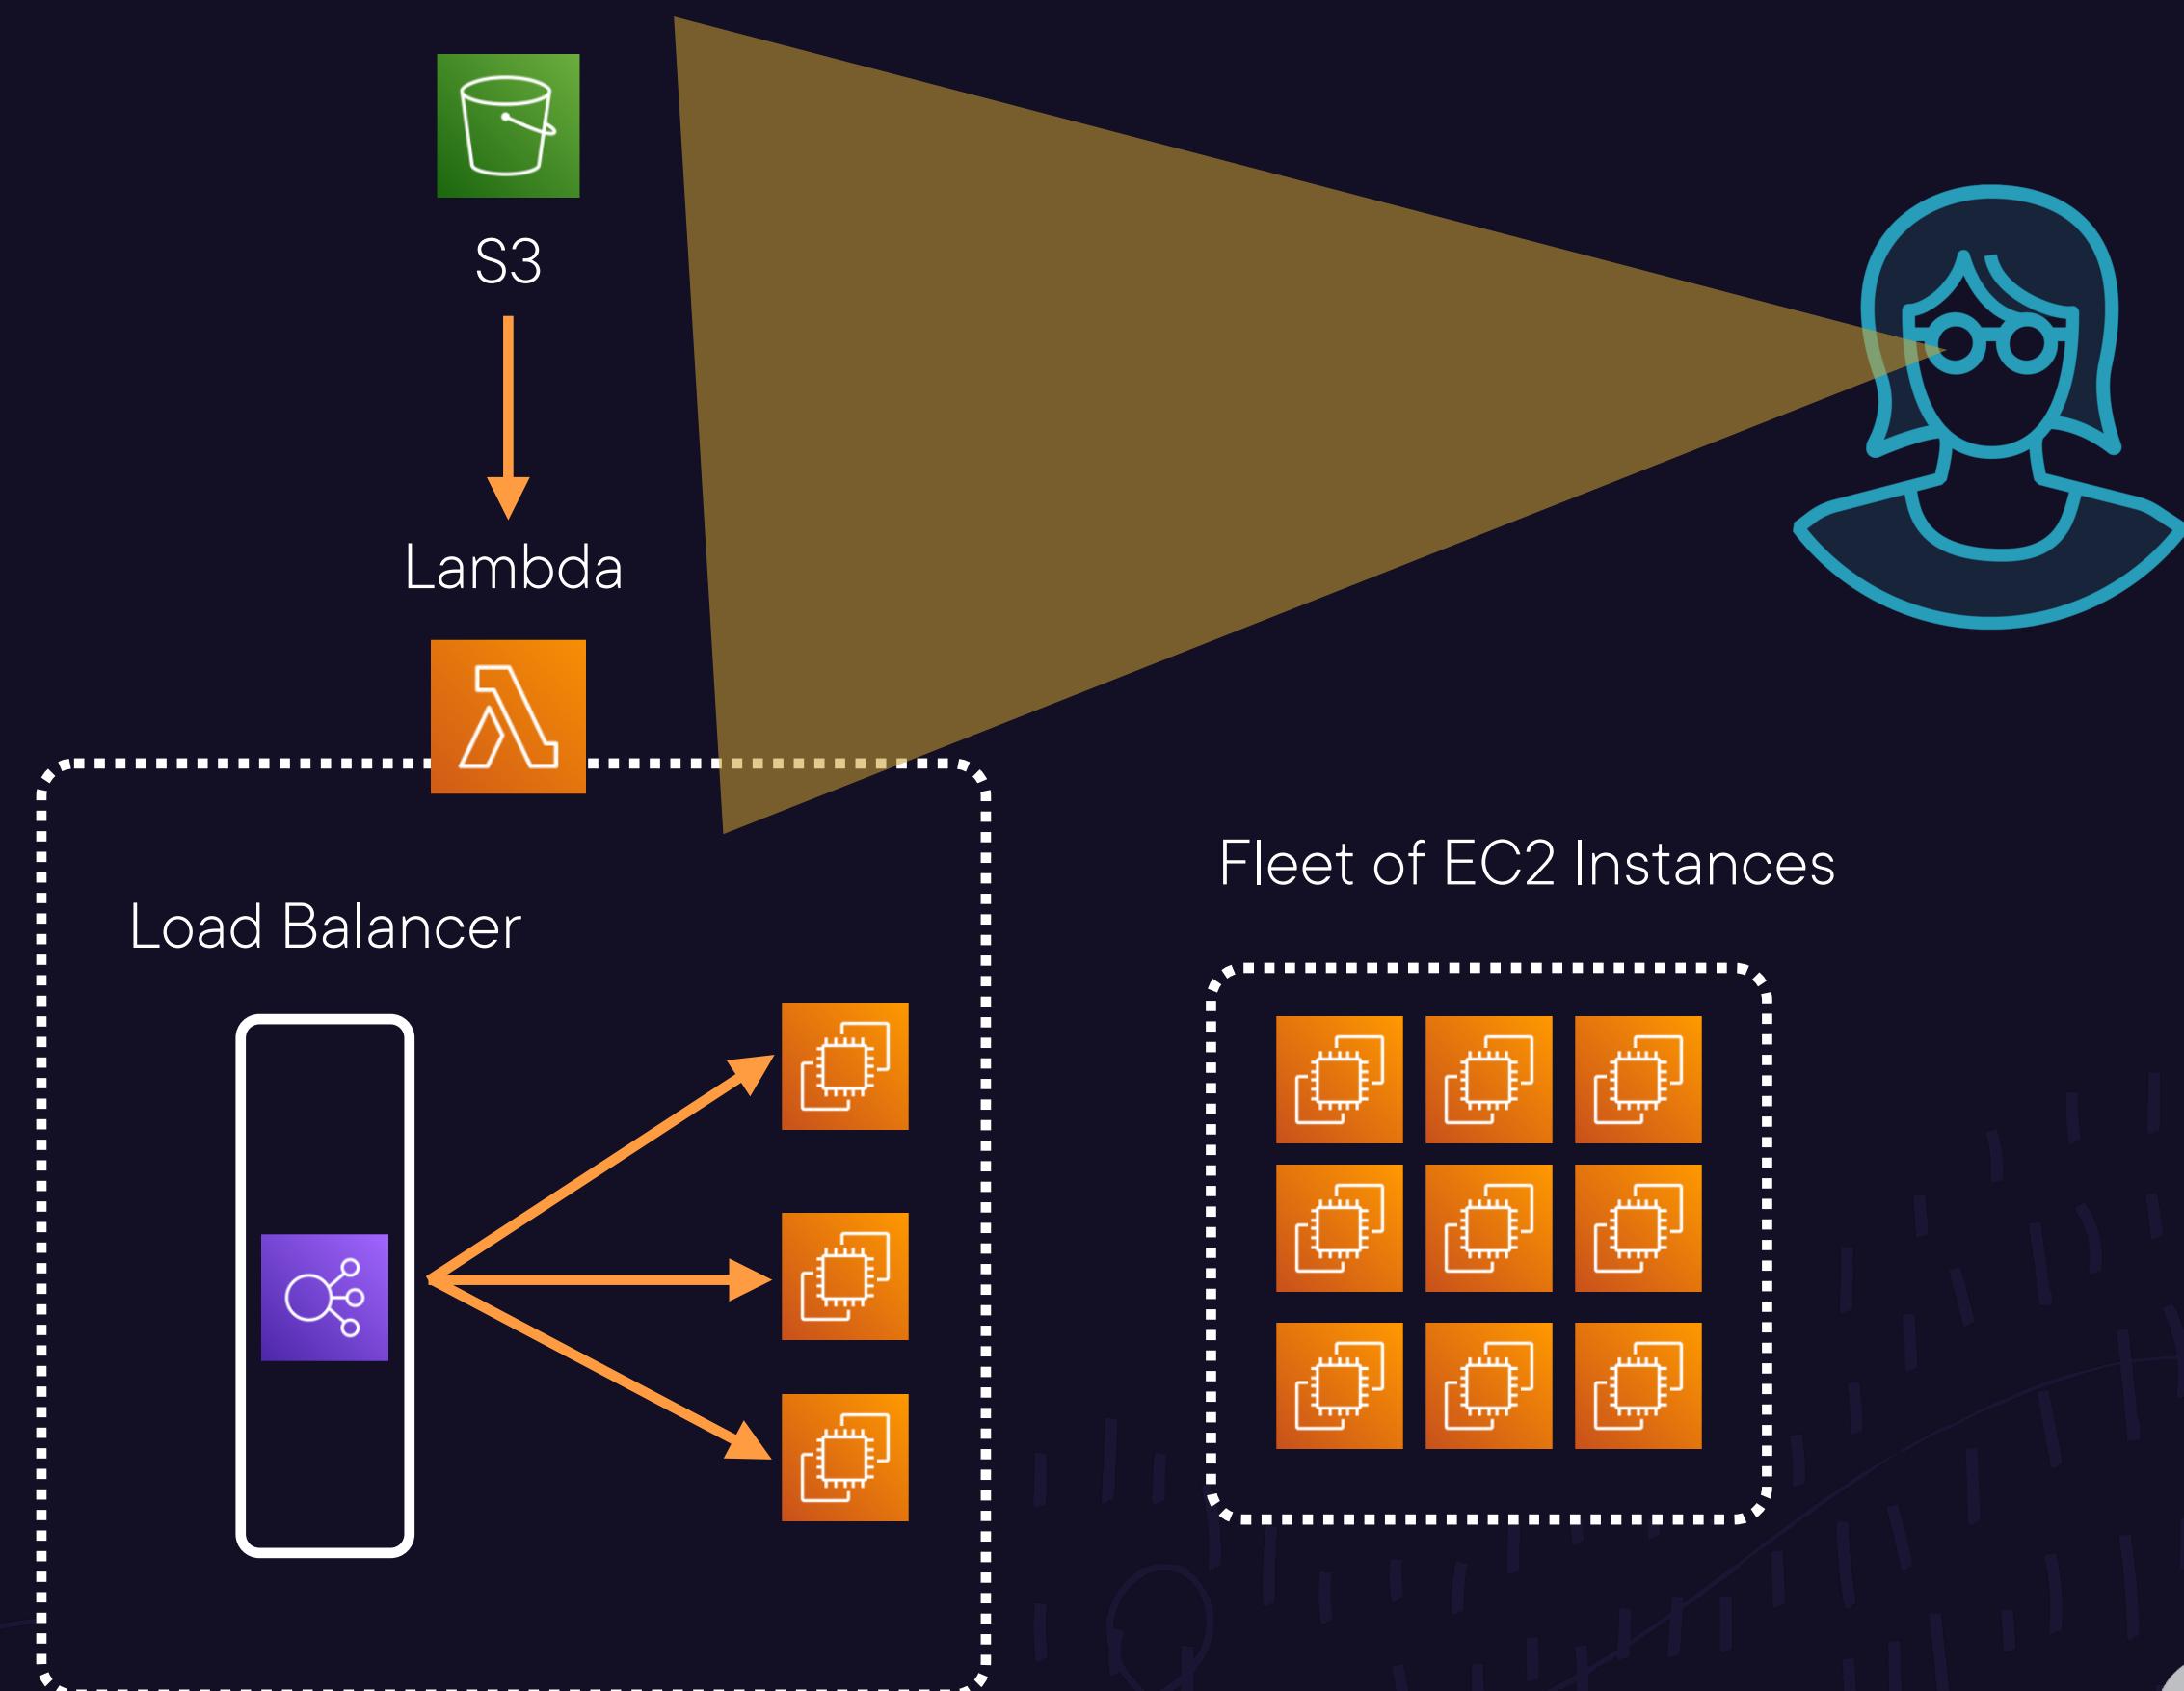
## When to Use Lambda

### EC2 vs. Lambda

Less Work

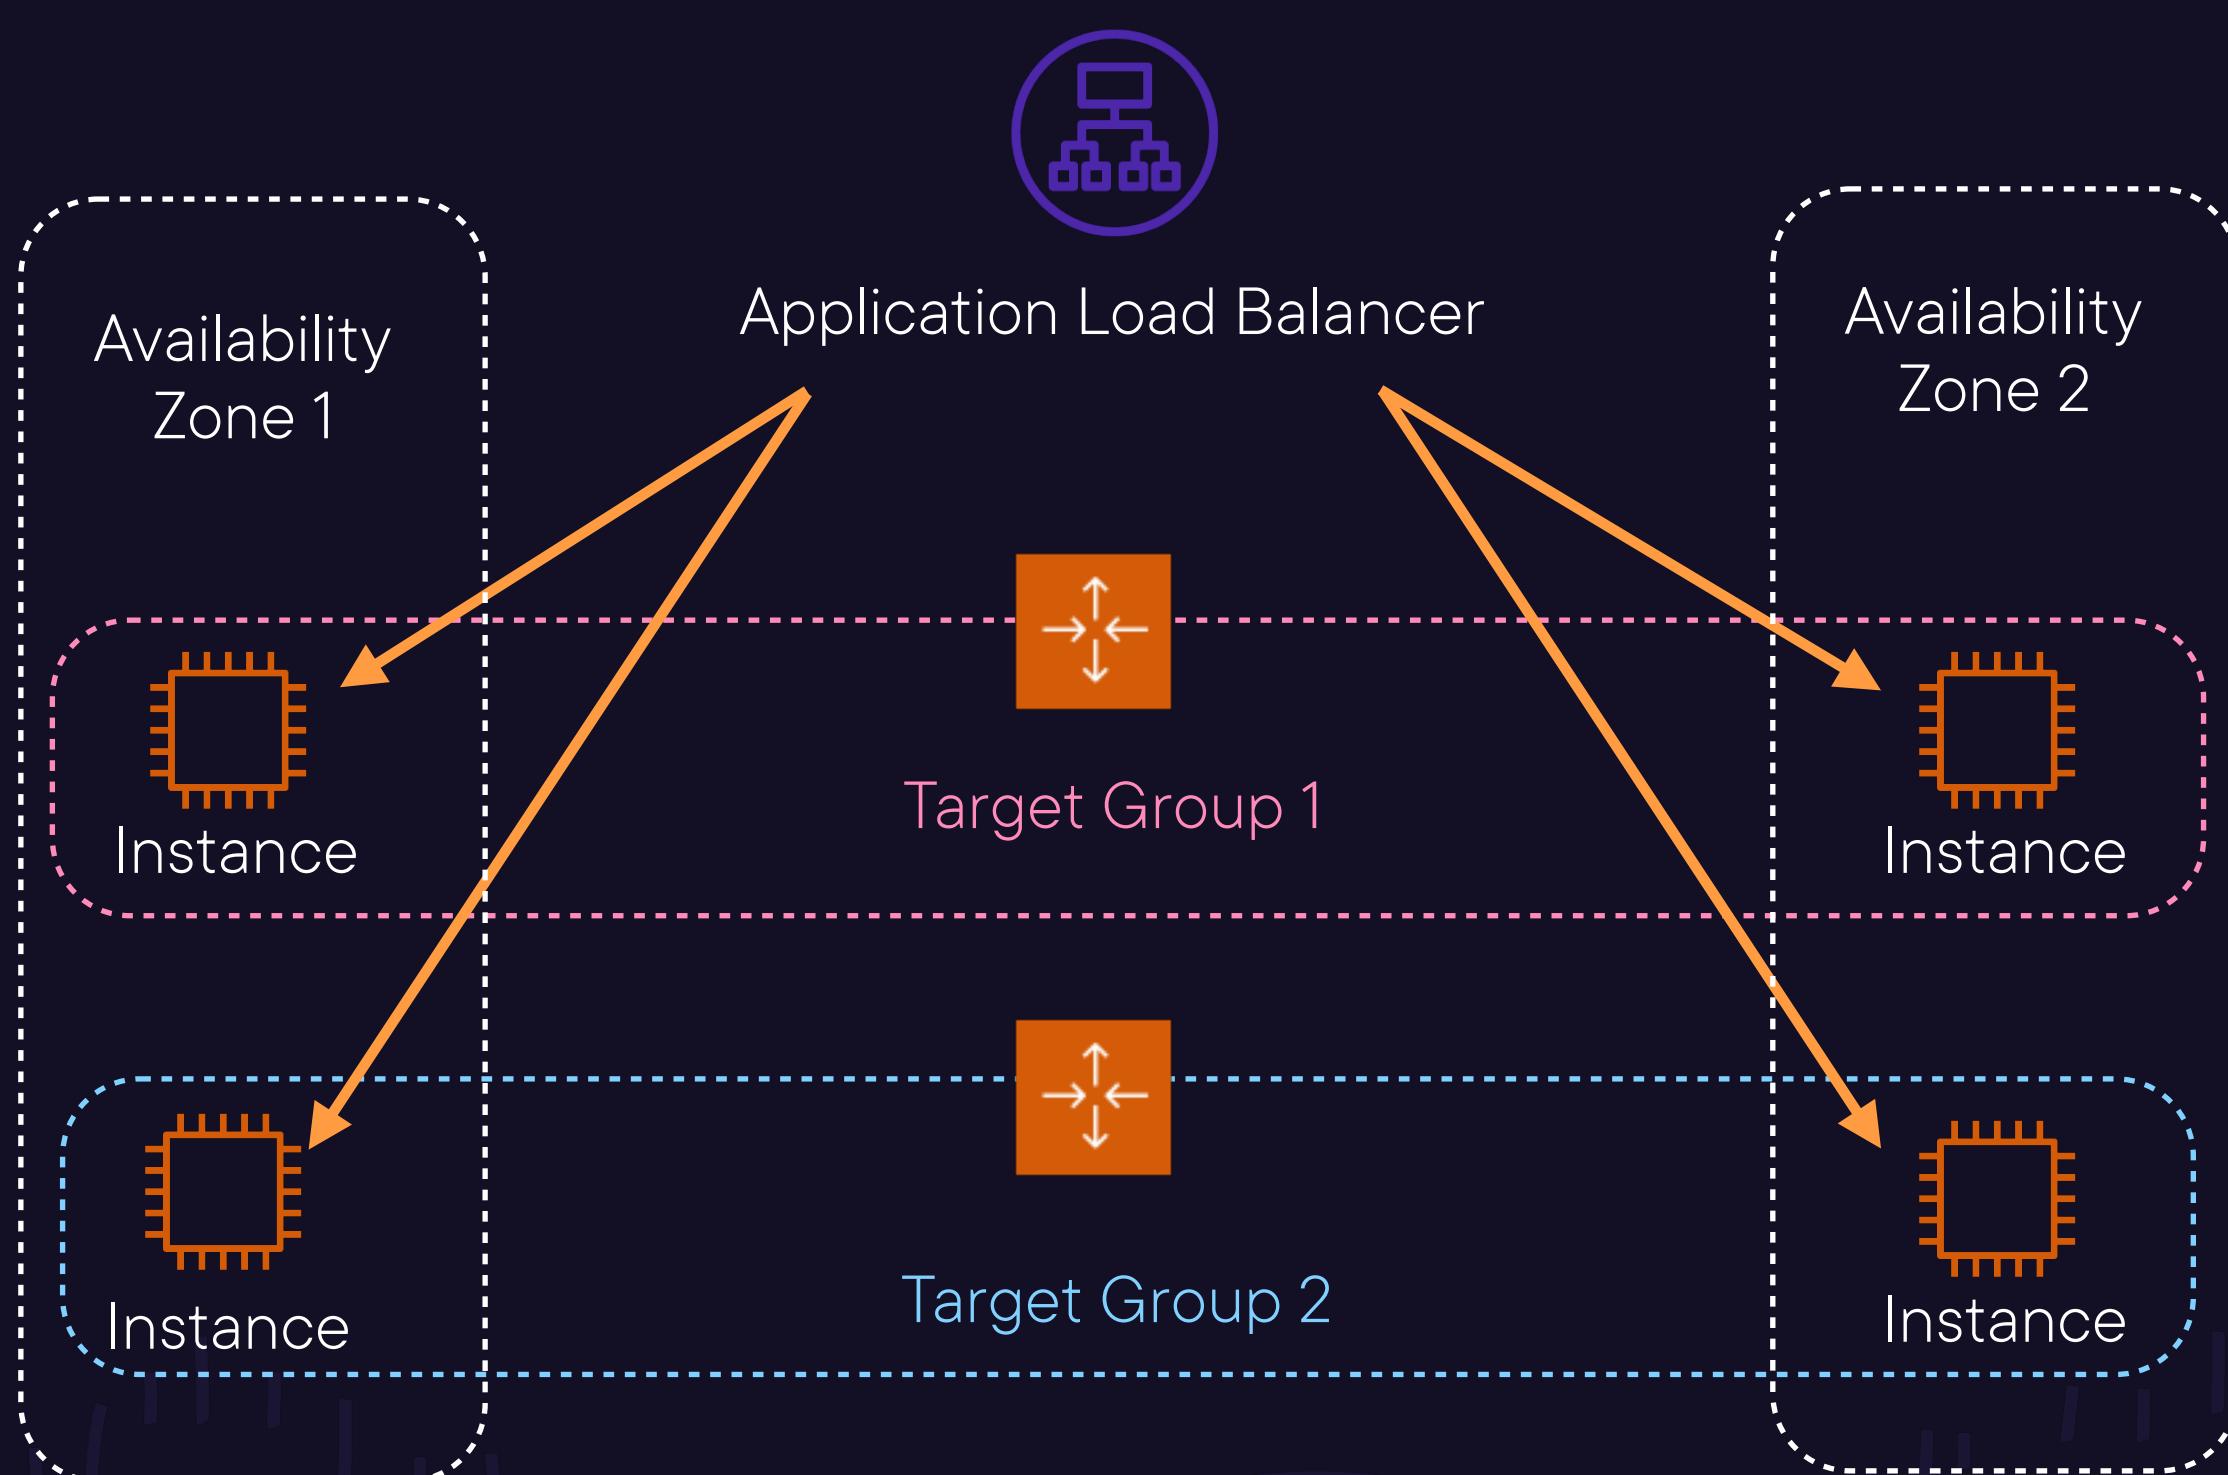


Less Control



## When to Use Lambda

# Deployment with EC2



More Control



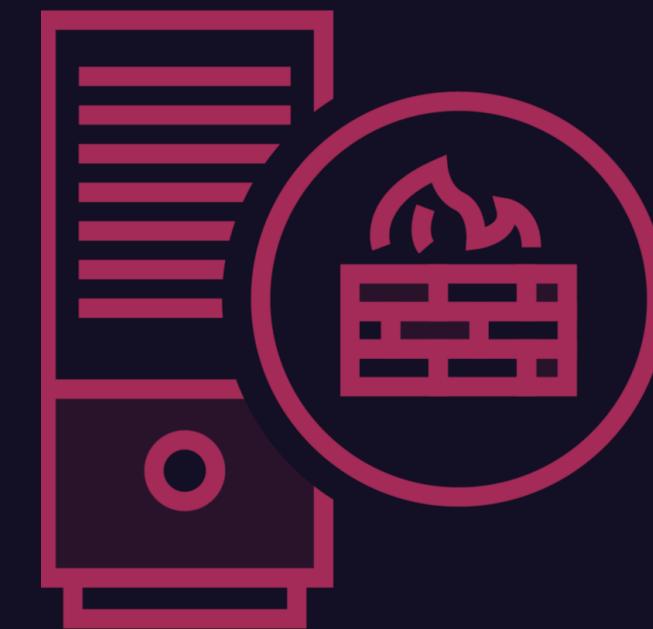
More Work



## EC2: Your Responsibilities with EC2



Pick an **instance type** for your **use case** (compute, memory, or general purpose).



Handle **security** through **security groups** and installing **security patches**.



Handle **scaling** and **traffic distribution** through setting up **Auto Scaling groups** and **target groups**.



Shut down the instance when no longer needed, either manually or via a **shutdown task**.



# ECS

What is it? When to use it?





## ECS

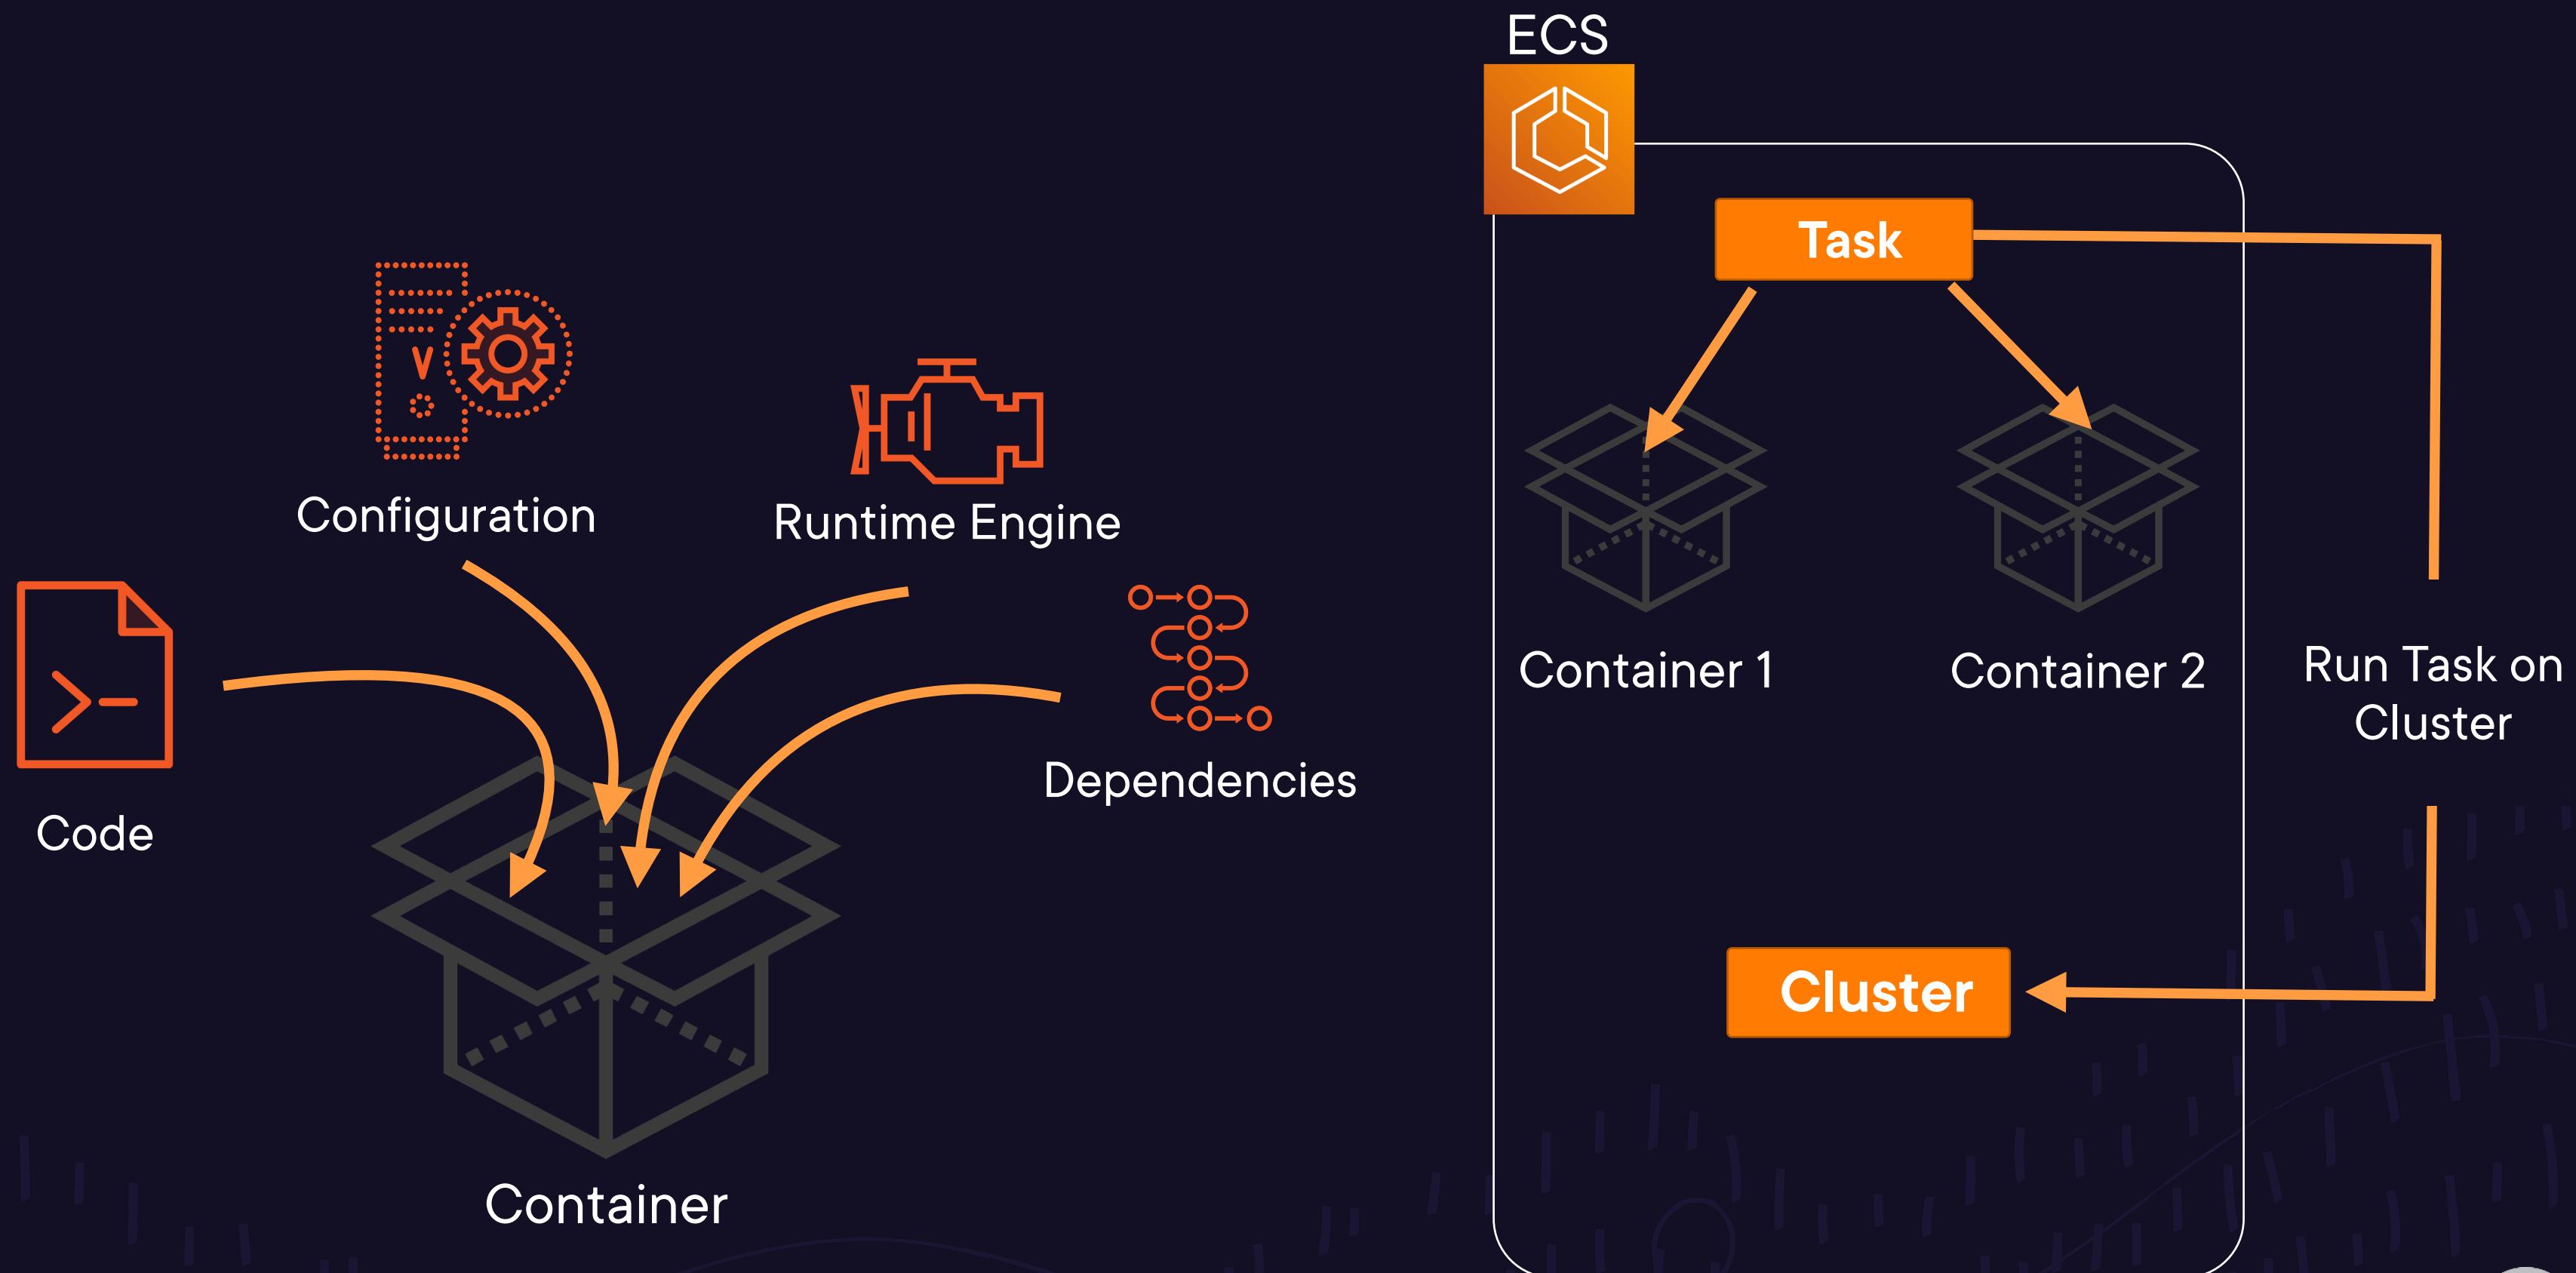
- Orchestrator **containers** (e.g., Docker)
- Operates on **clusters** (EC2 instances or Fargate) and **tasks**

## Use Cases

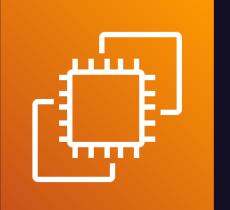
- Full-scale applications or ad-hoc jobs
- Running pre-packaged images or third-party software
- Predictable traffic



# When to Use Lambda ECS and Containers



So, if  doesn't meet your runtime needs...

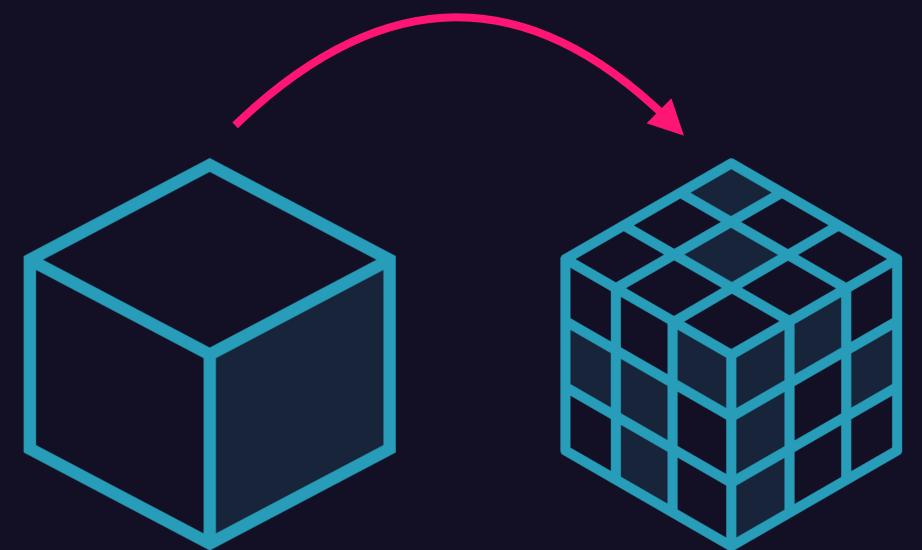
Consider  Elastic Beanstalk  
or  EC2



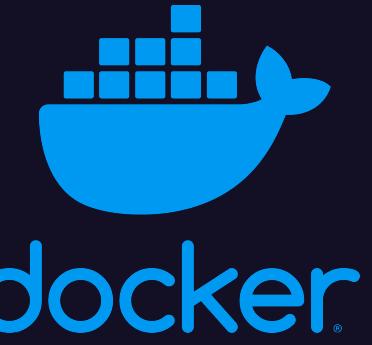
If your deployment package is

**> 512 MB**

**break up** your package  
into **smaller**  
packages and functions.



*If*

- You're working with 
- Your jobs take > 15 mins
- You need to schedule jobs

Consider  ECS.



*If*

- You want to focus **more** on the **code**
- You want to focus **less** on **servers**
- Your jobs take < **15 mins**
- Never pay **\$** for idle time

Consider  Lambda.



# Lesson Summary



## Lambda Constraints

Execution timeout, disk space, and memory.

## AWS Services

- **Lambda**: Smaller applications that take less than 15 mins to run
- **EC2**: Highly flexible, but high expertise needed
- **Elastic Beanstalk**: Fully managed customizable applications
- **ECS**: Container orchestration



# Accessing Lambda



**Noreen Hasan**

Training Architect



# Accessing Lambda

## LESSON BREAKDOWN

Ways to Access Lambda's Resources

Lambda Console: Ways to Create Functions

Demo

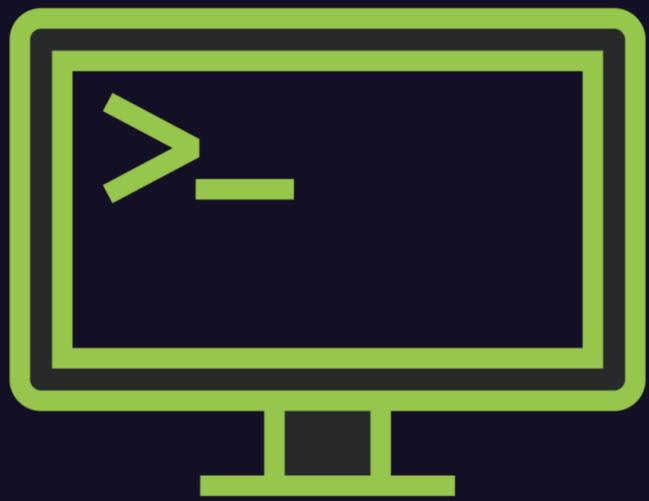
Lesson Summary



**Noreen Hasan**  
Training Architect



## Ways to Access Lambda's Resources



AWS Command Line  
Interface (CLI)

# SDK



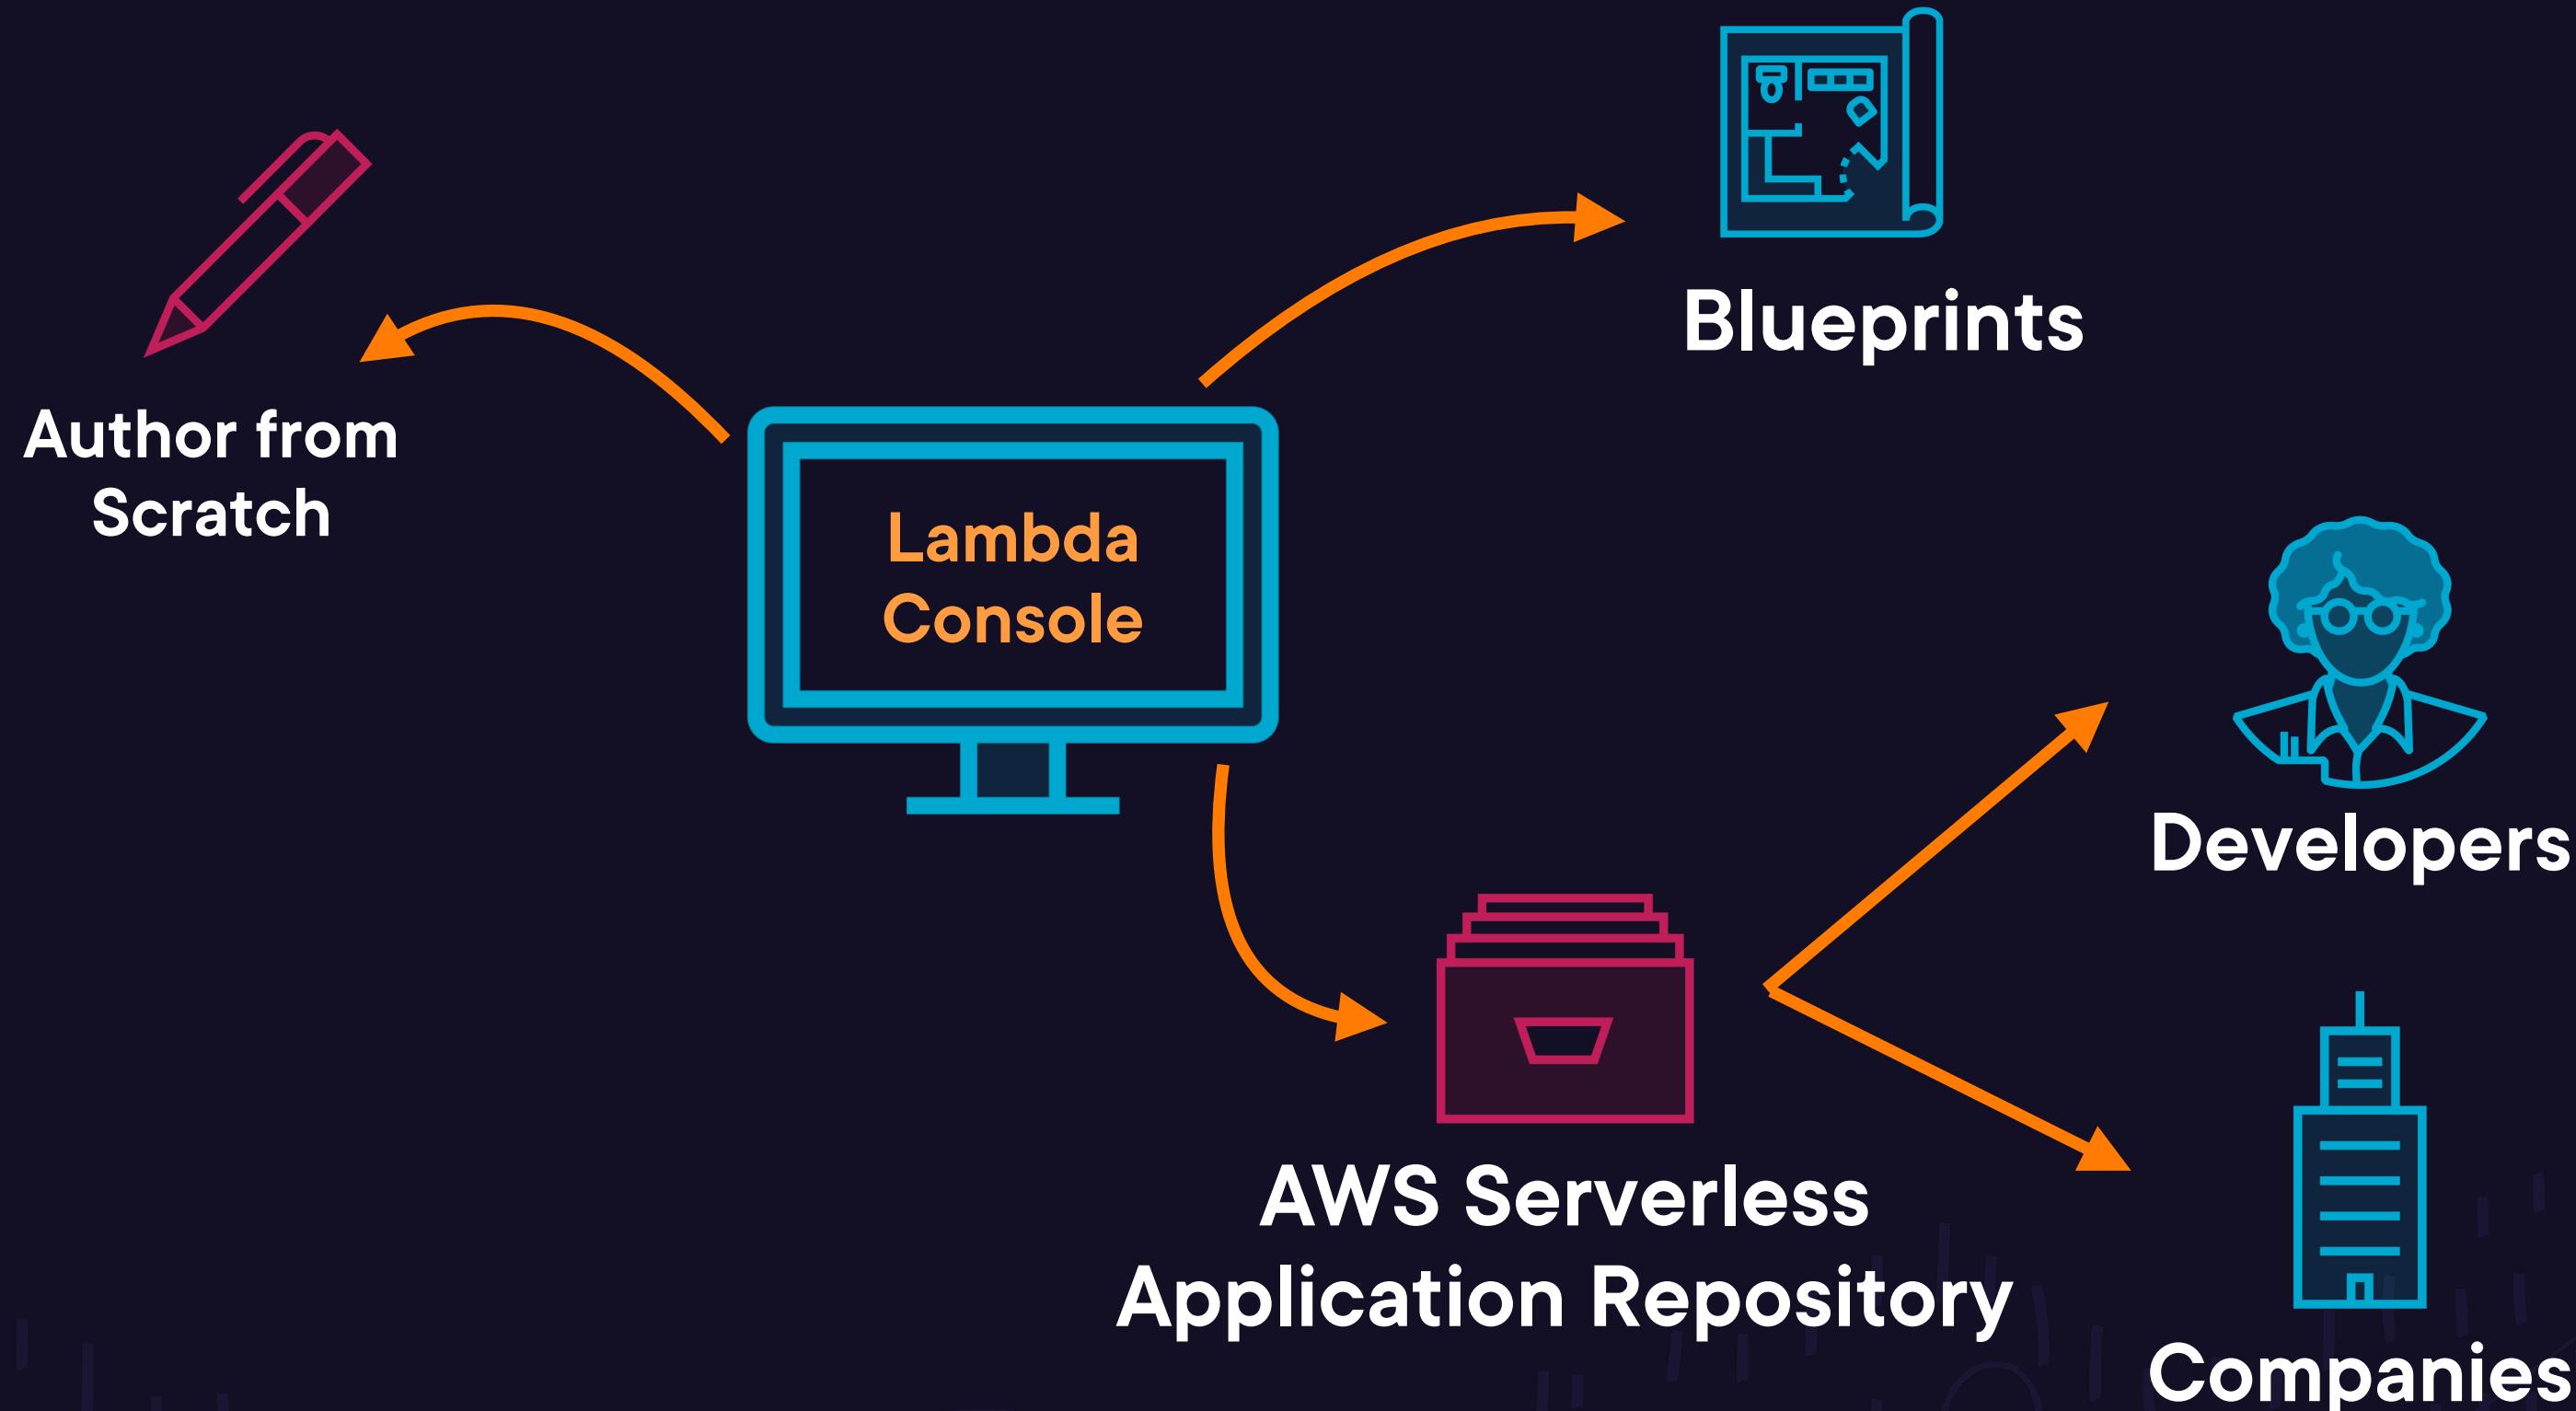
AWS Software  
Development Kit (SDK)

The AWS Console -  
Lambda Console



## Accessing Lambda

# Lambda Console: Ways to Create Functions



# Demo



## Lambda Console

- Walkthrough of the console
- Create a Hello World  application



# Lesson Summary



## Ways to Access Lambda's Resources

- AWS CLI
- SDK
- Lambda console

## Lambda Console: Ways to Create Functions

- Author from scratch
- Blueprints
- AWS Serverless Application Repository

## Demo

- Create a Hello World  application



# Introduction to Processing Data with Lambda



**Noreen Hasan**

Training Architect



## SECTION BREAKDOWN

# Introduction to Processing Data with Lambda

### Section Overview



**Noreen Hasan**  
Training Architect



# Introduction to Processing Data with Lambda

## Section Overview



Processing Data:  
Overview

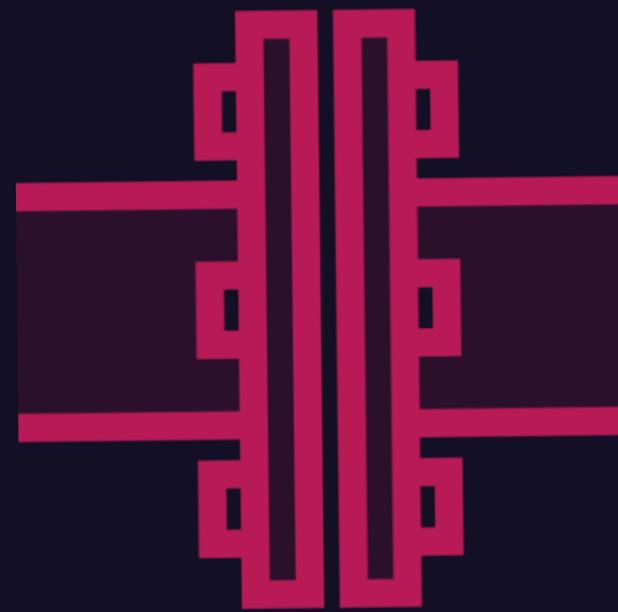
Understanding Lambda  
Layers

Demo: Converting JSON  
to CSV Using Lambda



# Introduction to Processing Data with Lambda

## Section Overview



**Demo:** Ingesting CSV Files from S3 to **DynamoDB**



**Processing Real-Time Data Using Kinesis and Lambda: Overview**



**Demo:** Processing Real-Time Data Using **Kinesis** and **Lambda**



**Testing and Debugging Lambda Failures**



# Let's get started!



# Processing Data: Overview



**Noreen Hasan**

Training Architect



# Processing Data: Overview

## LESSON BREAKDOWN

What Is Data Processing?

Data Lakes

Extract, Transform, and Load (ETL)

Extract, Load, and Transform (ELT)

Types of Data Transformations

Lesson Summary



**Noreen Hasan**  
Training Architect



# What Is Data Processing?



## Data Processing

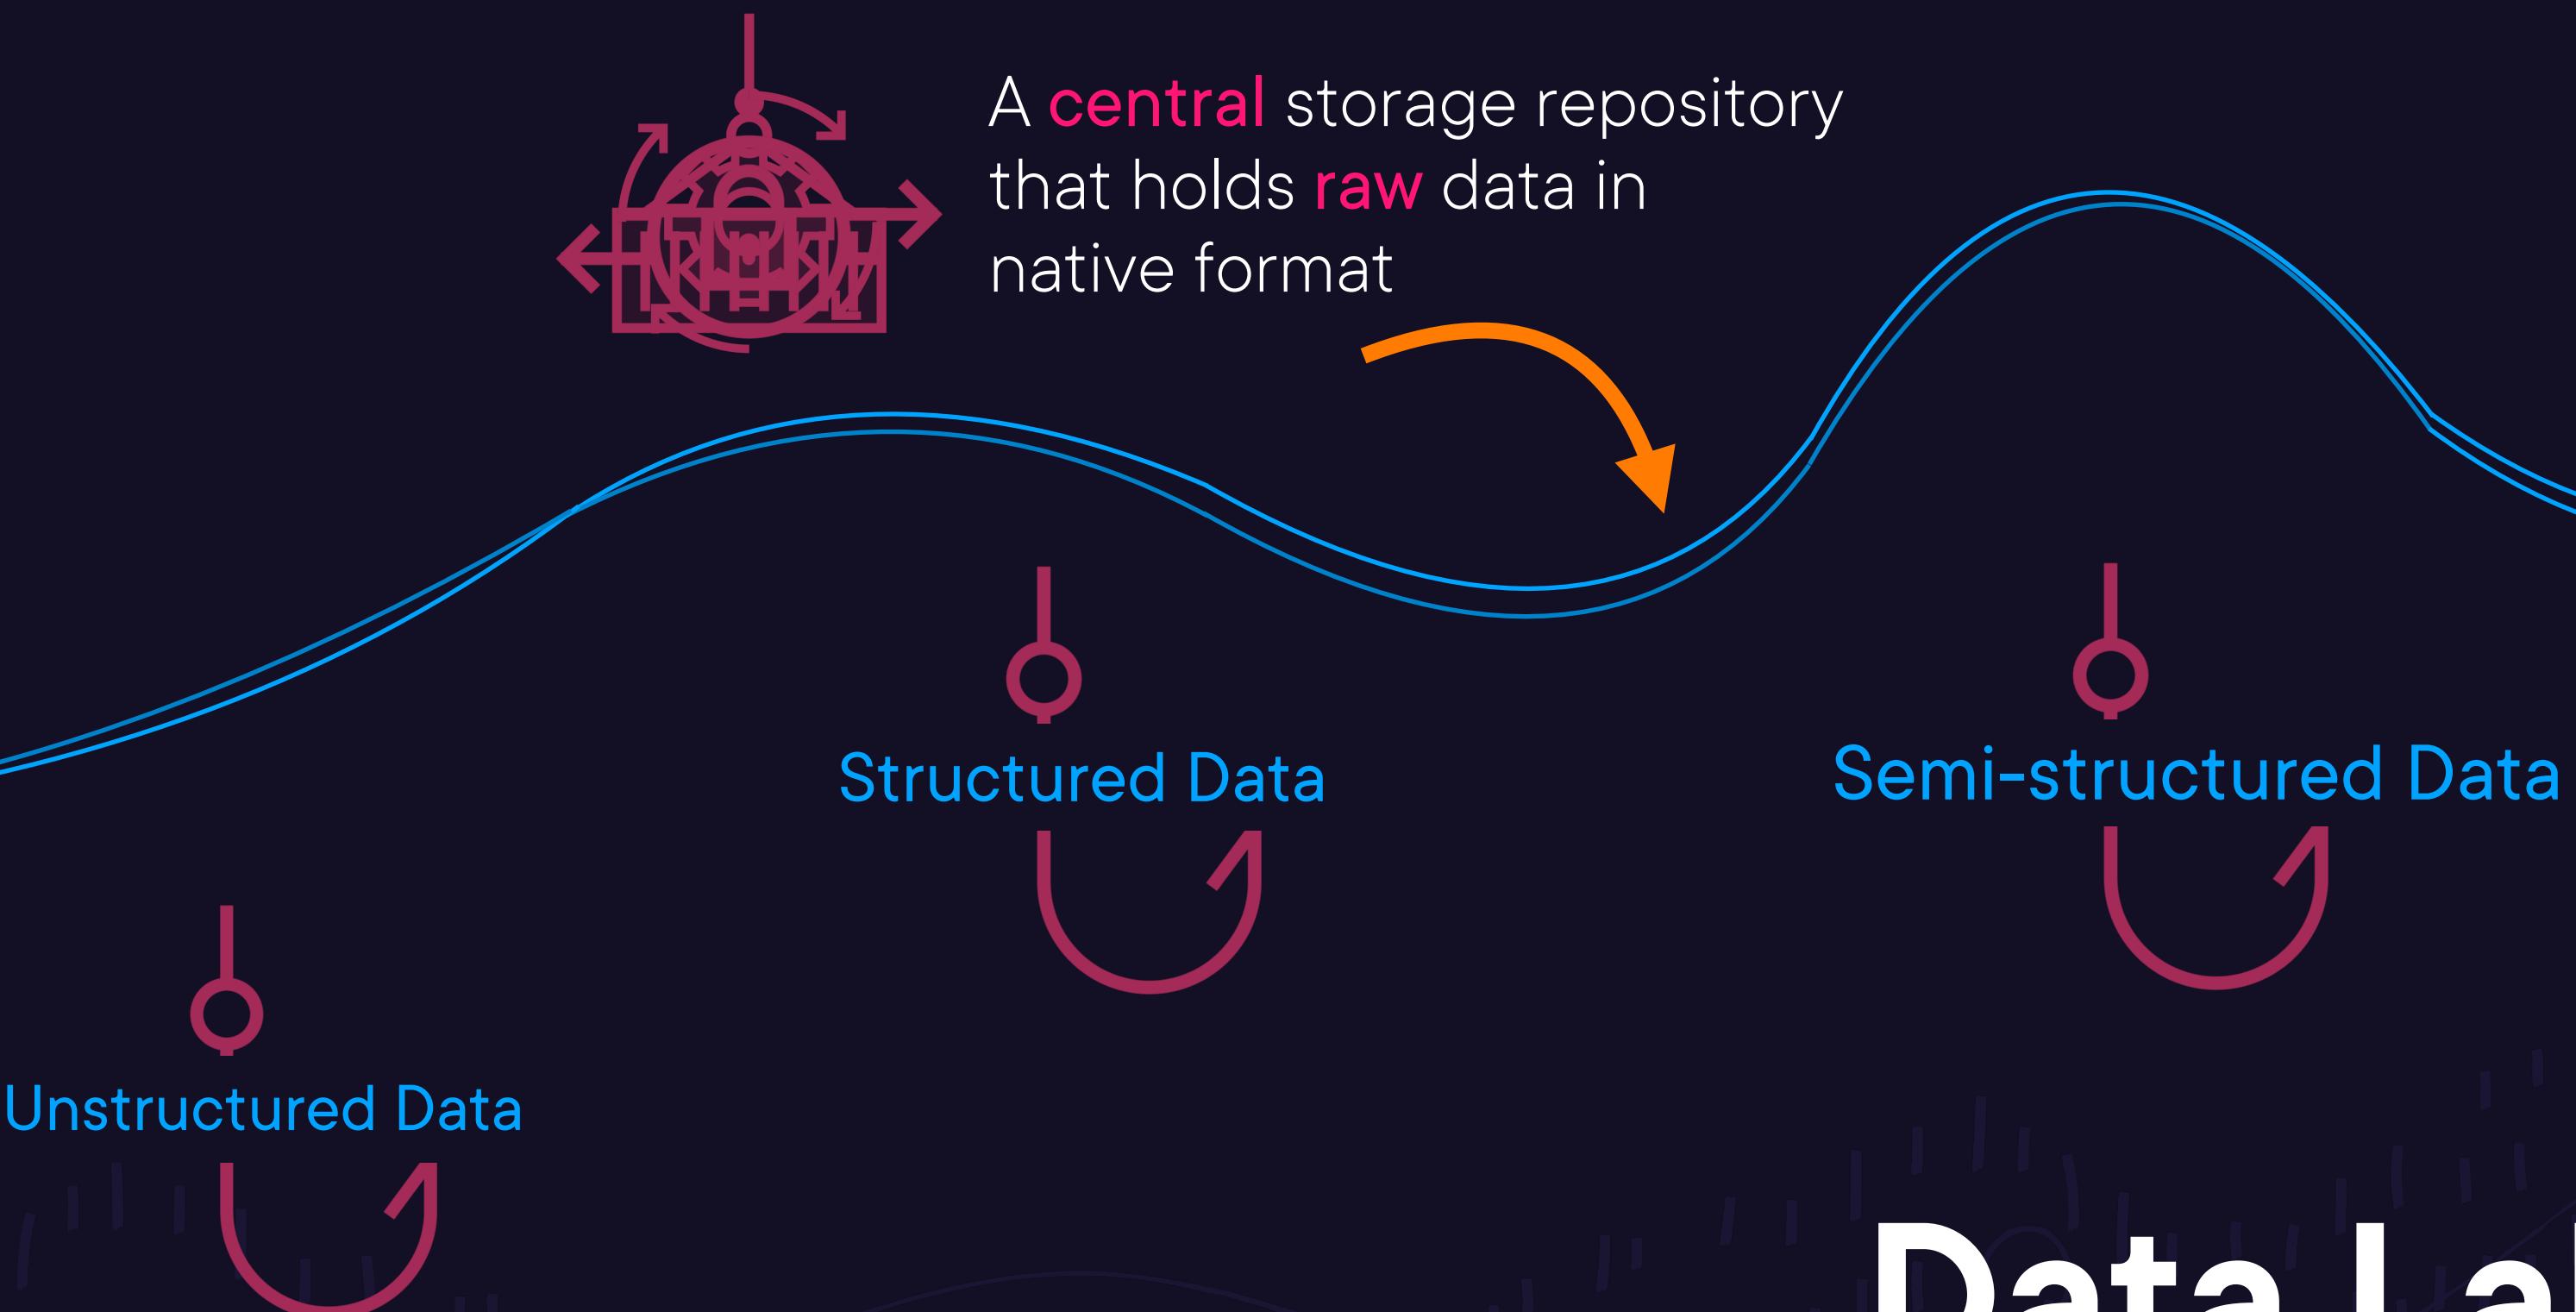
The process of converting raw data to meaningful information

## How Is It Done?

- Manual: pen and paper
- Mechanical: typewriters
- Electronic: computers programs



# Data Lakes



## Processing Data: Overview

# Extract, Transform, and Load (ETL)



Extract



Transform



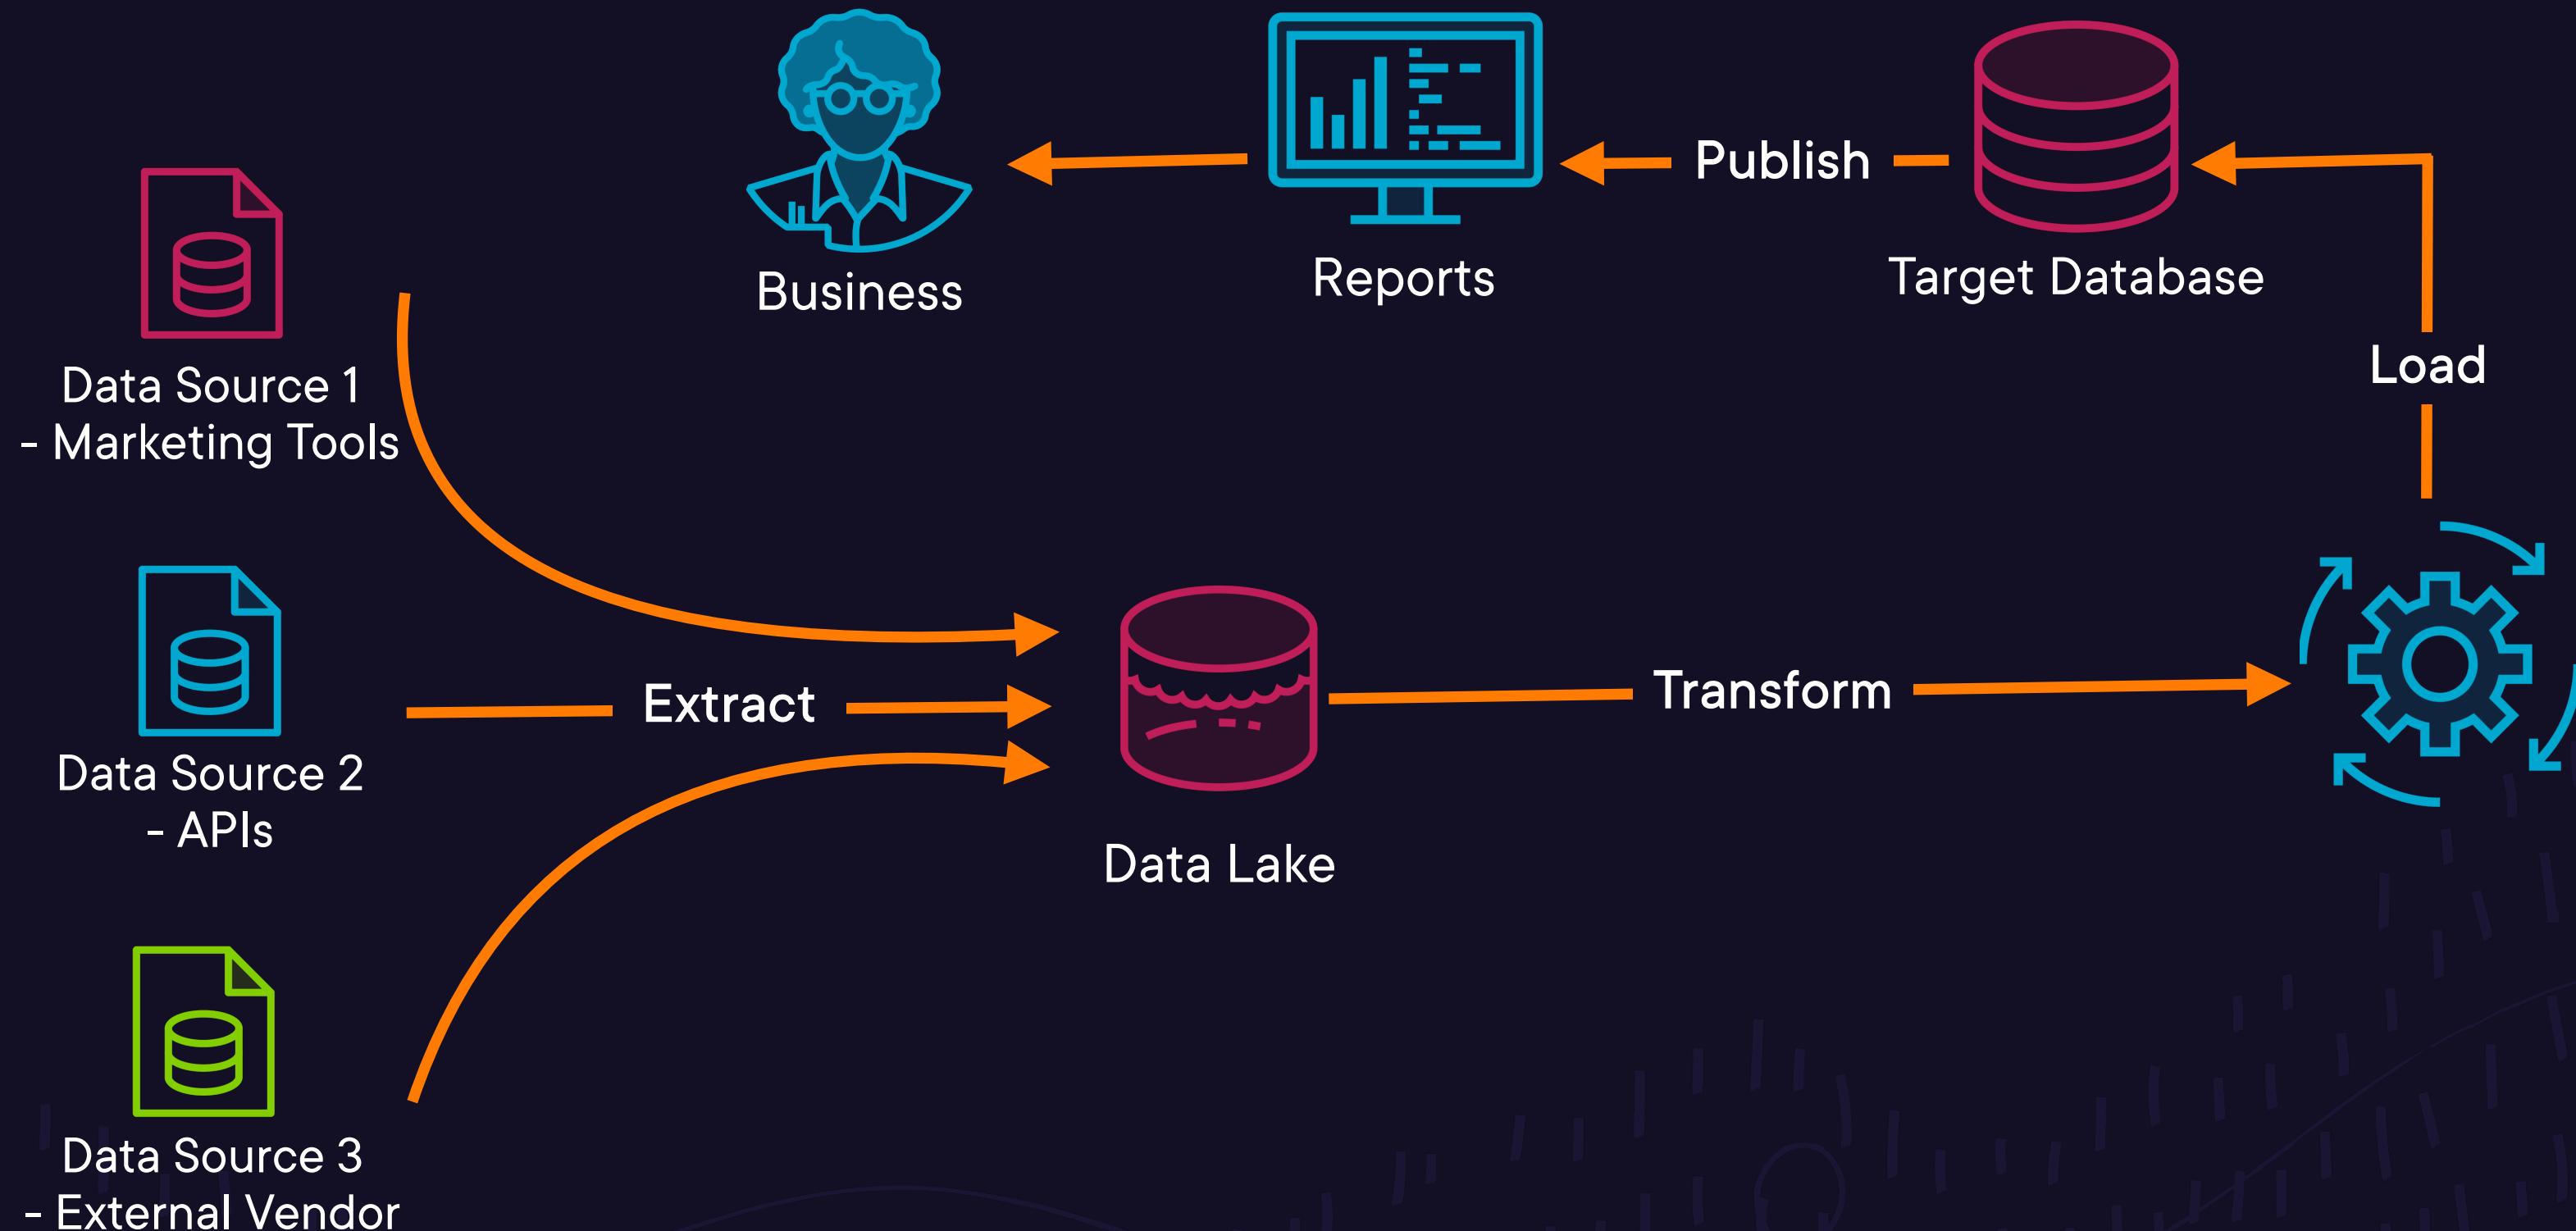
Load

Data Lake



## Processing Data: Overview

### Data Lakes and ETL



## Processing Data: Overview

# Extract, Load, and Transform (ELT)



Extract



Load



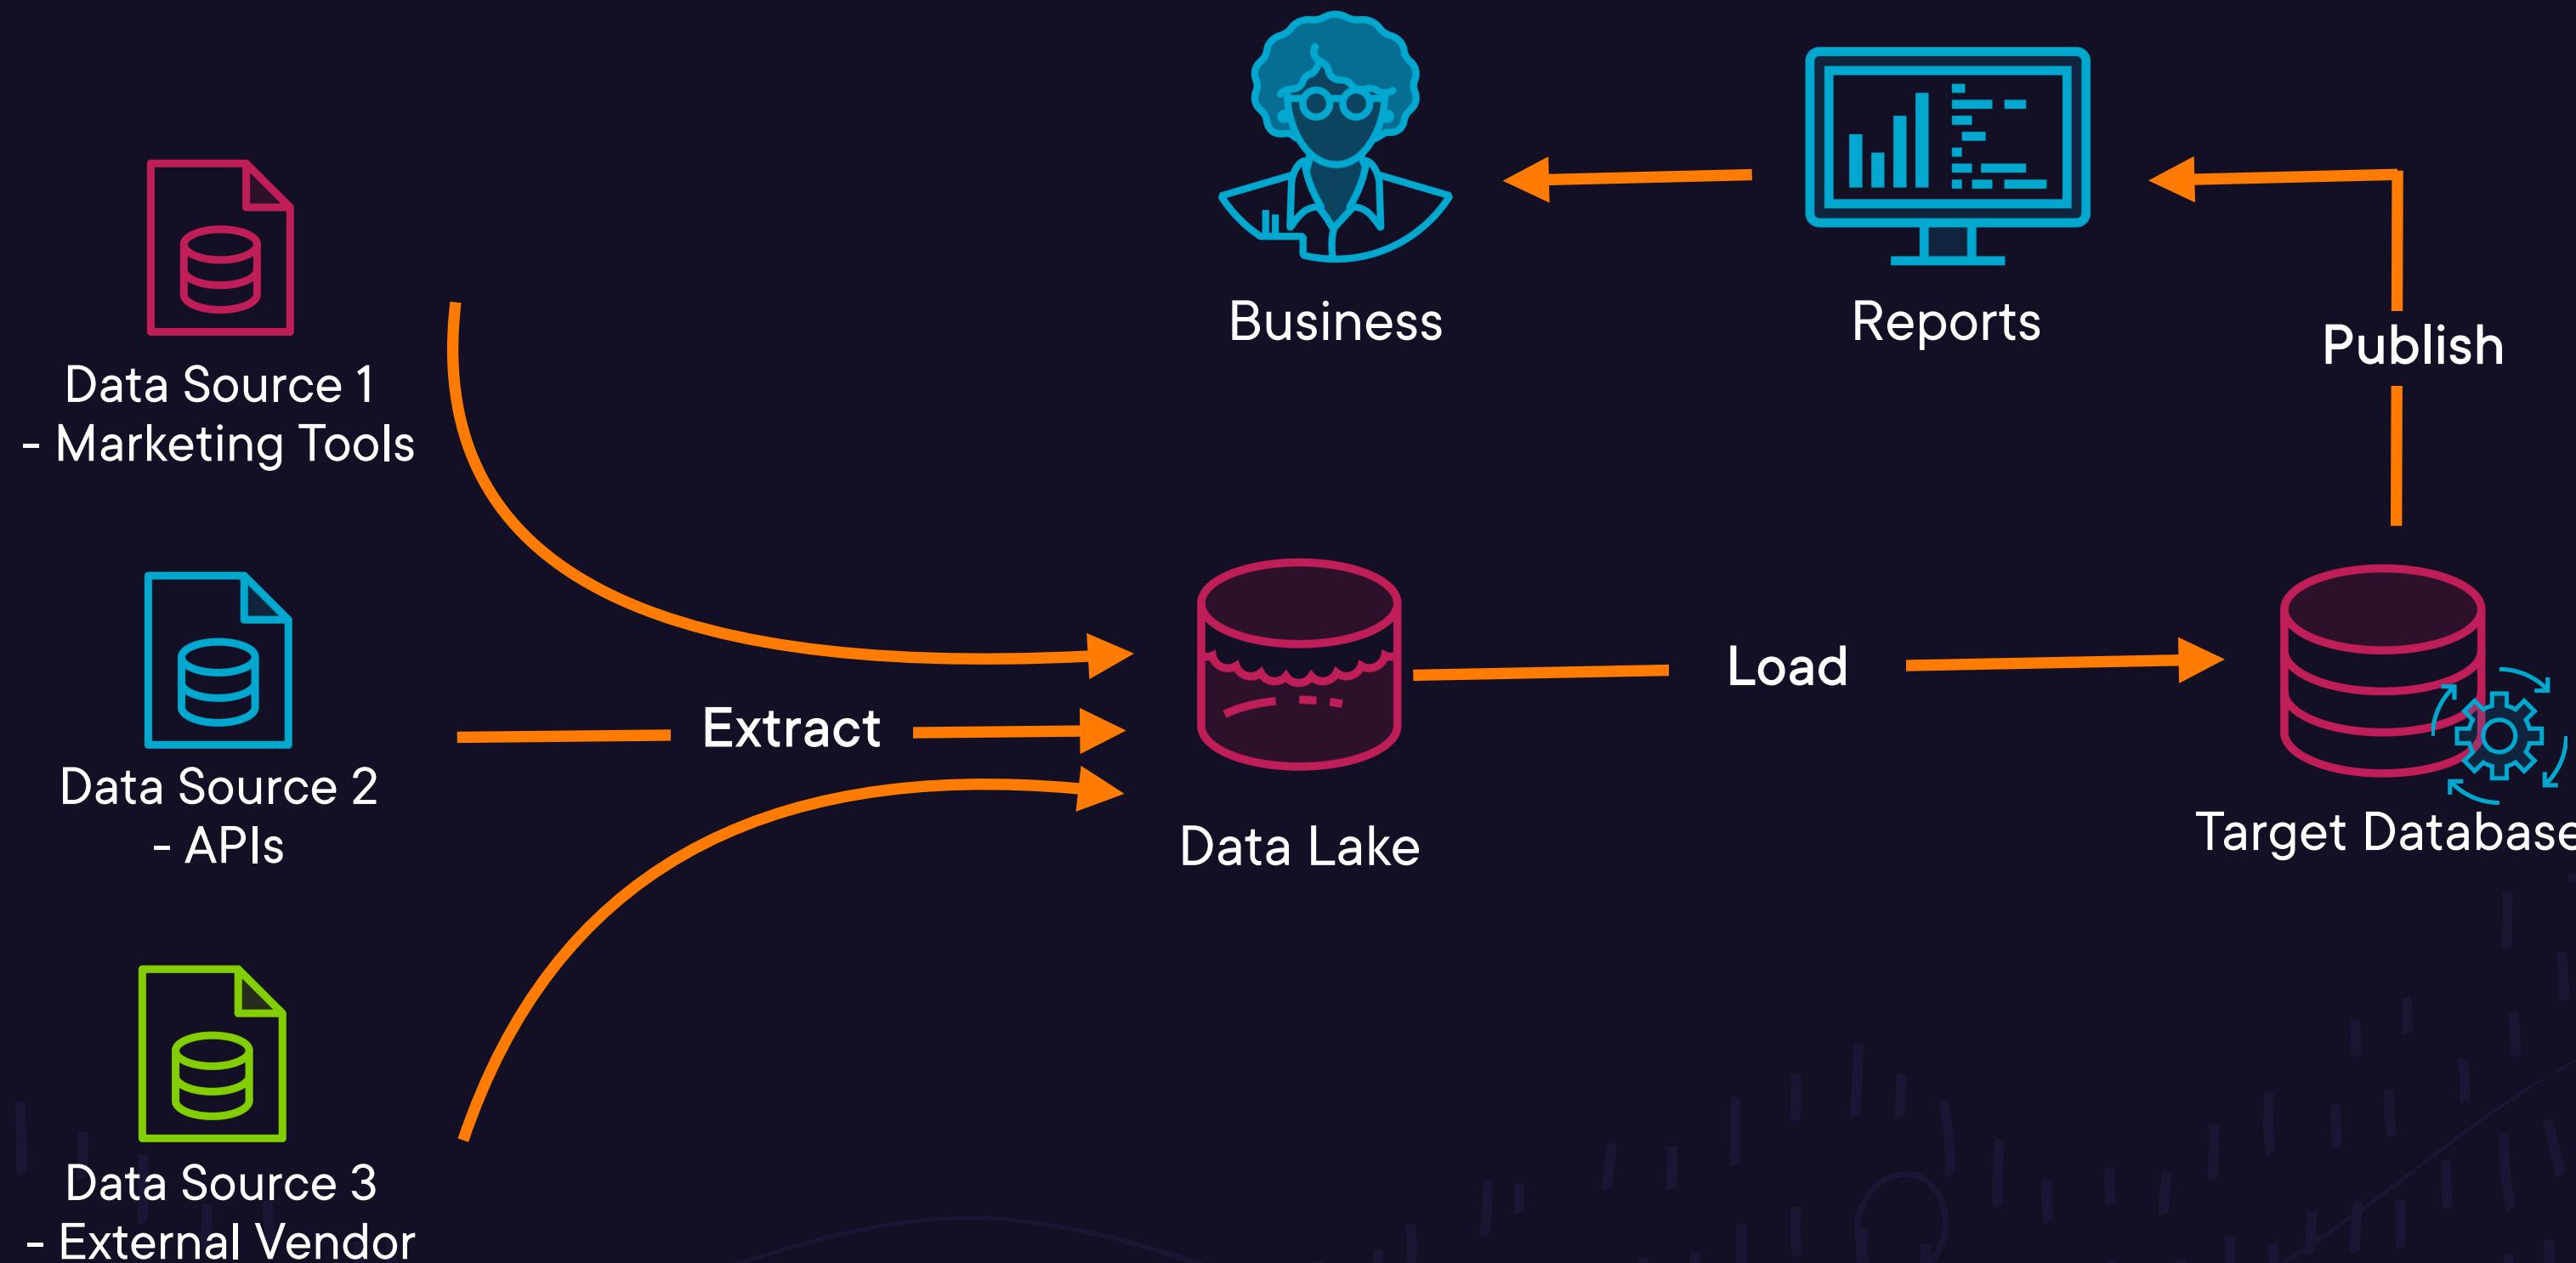
Transform

Data Lake



## Processing Data: Overview

### Data Lakes and ELT

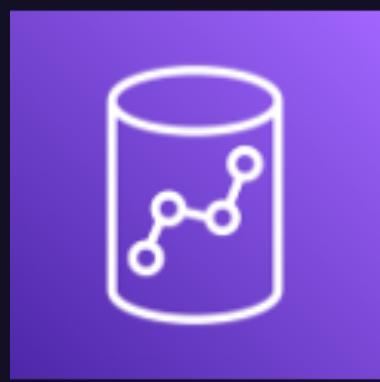


## Processing Data: Overview

# How Is ELT Made Possible?



# How Is ELT Made Possible?



**Amazon Redshift**



**Snowflake**



**Azure**

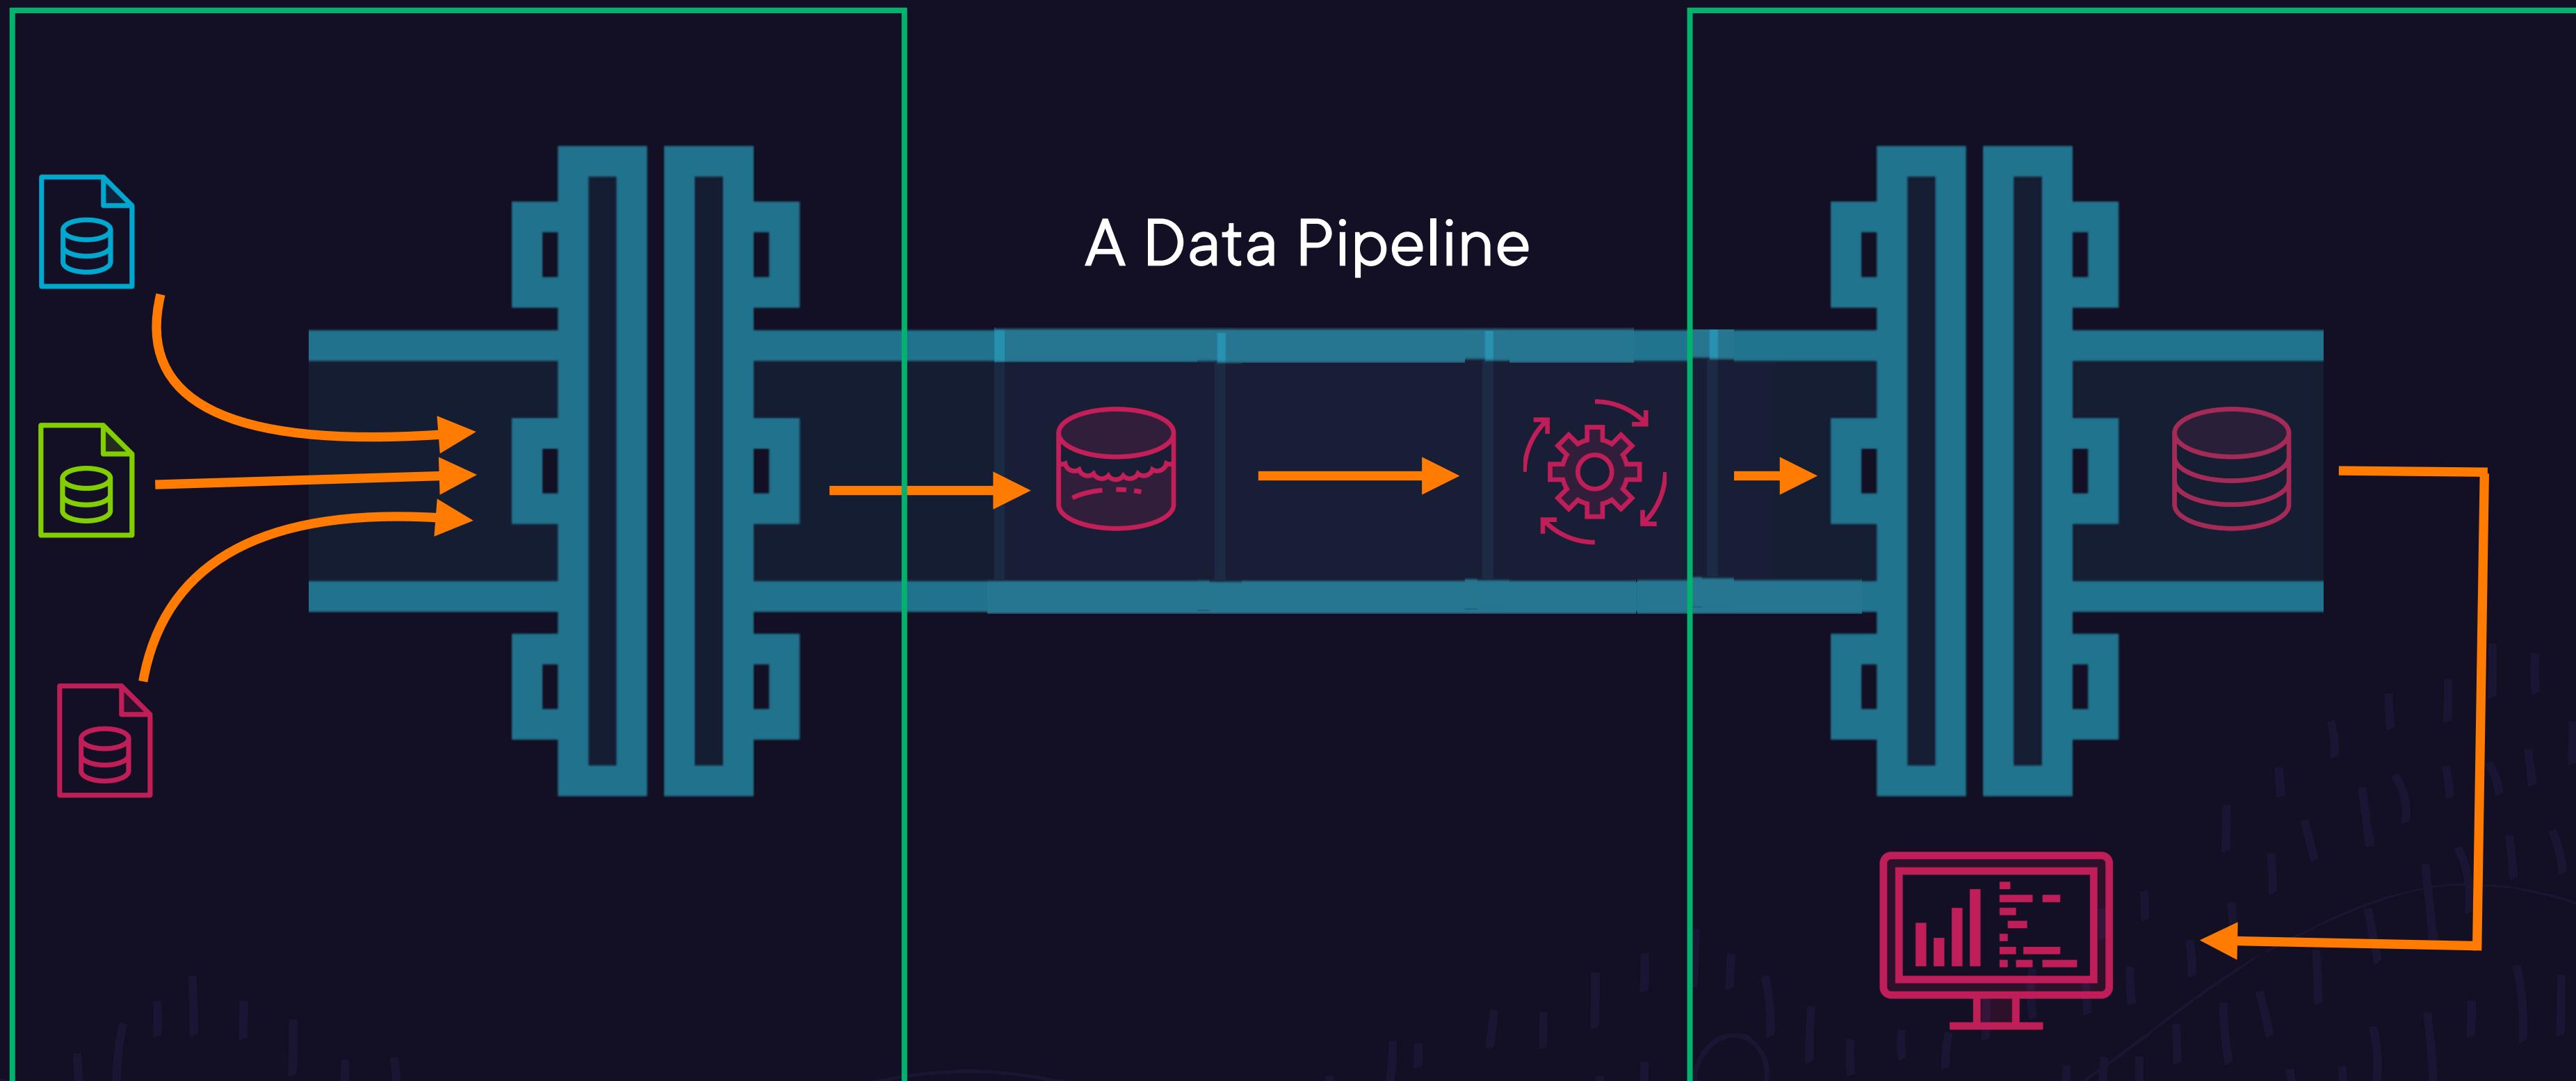


**Google BigQuery**



## Processing Data: Overview

### Data Pipeline



### Types of Data Transformations



#### Filtering Data

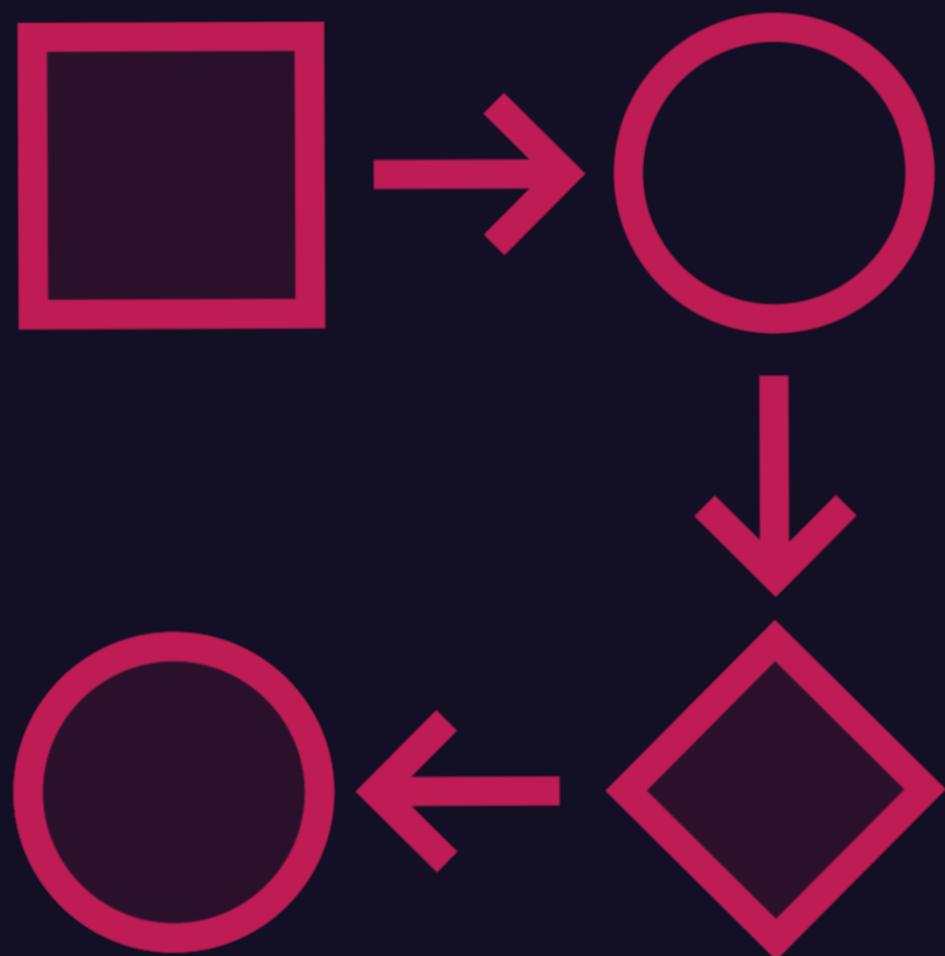
Removing certain columns or rows before inserting the data into the final destination

#### Example

- Filtering out historical records to create the table `certification_exams_2023`



### Types of Data Transformations



#### Mapping Data

Taking an input from a data source and changing **(conforming)** it into an equivalent in another data format

#### Example

- Transforming date fields from local time zones to a standardized timestamp
- Transforming textual fields to numerical fields (i.e., surveys)
  - Like= 1
  - Dislike= 0



### Types of Data Transformations



#### Deriving Variables

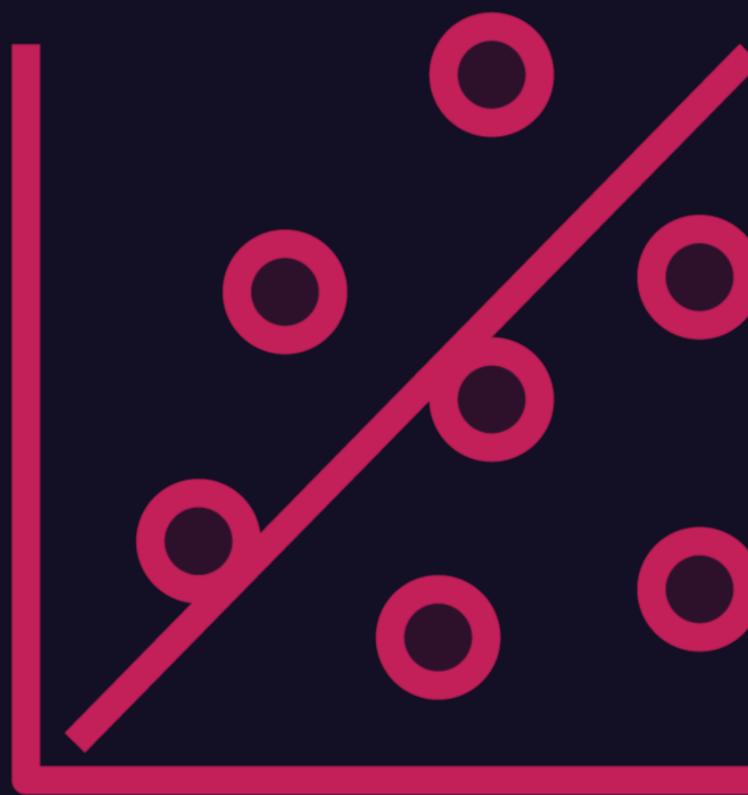
Variables that are computed based on other data variables

#### Example

- Creating a column, **Final Price**, that computes the final price of a course by subtracting the **Discount** from the **Original Price**



### Types of Data Transformations



#### Aggregating Data

Creating a **summarizing field** that groups and **aggregates** data across a particular **dimension**

#### Example

- Calculating the sum, average, mean, or median



## Processing Data: Overview

### Types of Data Transformations

CourseID	CourseTitle	Author
837	Intro to Lambda	Tia Williams
837	Intro to Lambda	Tia Williams

#### Deduplicating Data

Removing duplicated data to decrease needed storage

It can be...

- part of the transform process **before** the data is inserted to the final database.
- a process that takes place **after insertion**.

Example

- Removing course records that are identical across all fields



### Types of Data Transformations



#### Splitting Data

Refining unstructured data by splitting certain fields before inserting them into the final table

#### Example

- **Splitting** an **address** field into smaller fields, such as **street**, **city**, and **zip\_code**



# Lesson Summary



## What Is Data Processing?

- Converting raw data to information

## Data Lakes

- A central repo  with raw data
- ETL vs. ELT

## Types of Data Transformation

- Data filtering
- Data mapping
- Deriving variables
- Aggregating data
- Deduplicating data
- Splitting data



# Understanding Lambda Layers



**Noreen Hasan**

Training Architect



# Understanding Lambda Layers

## LESSON BREAKDOWN

What Are Lambda Layers?

Why Layer Up?

Challenges

Demo

Lesson Summary



**Noreen Hasan**  
Training Architect





## Lambda Layers

- A **layer** is a **ZIP file** with **additional code or data**.
- Layers allow you to share common code between **functions** and **across accounts**.
- To use a layer, you need to **import** it to your function.





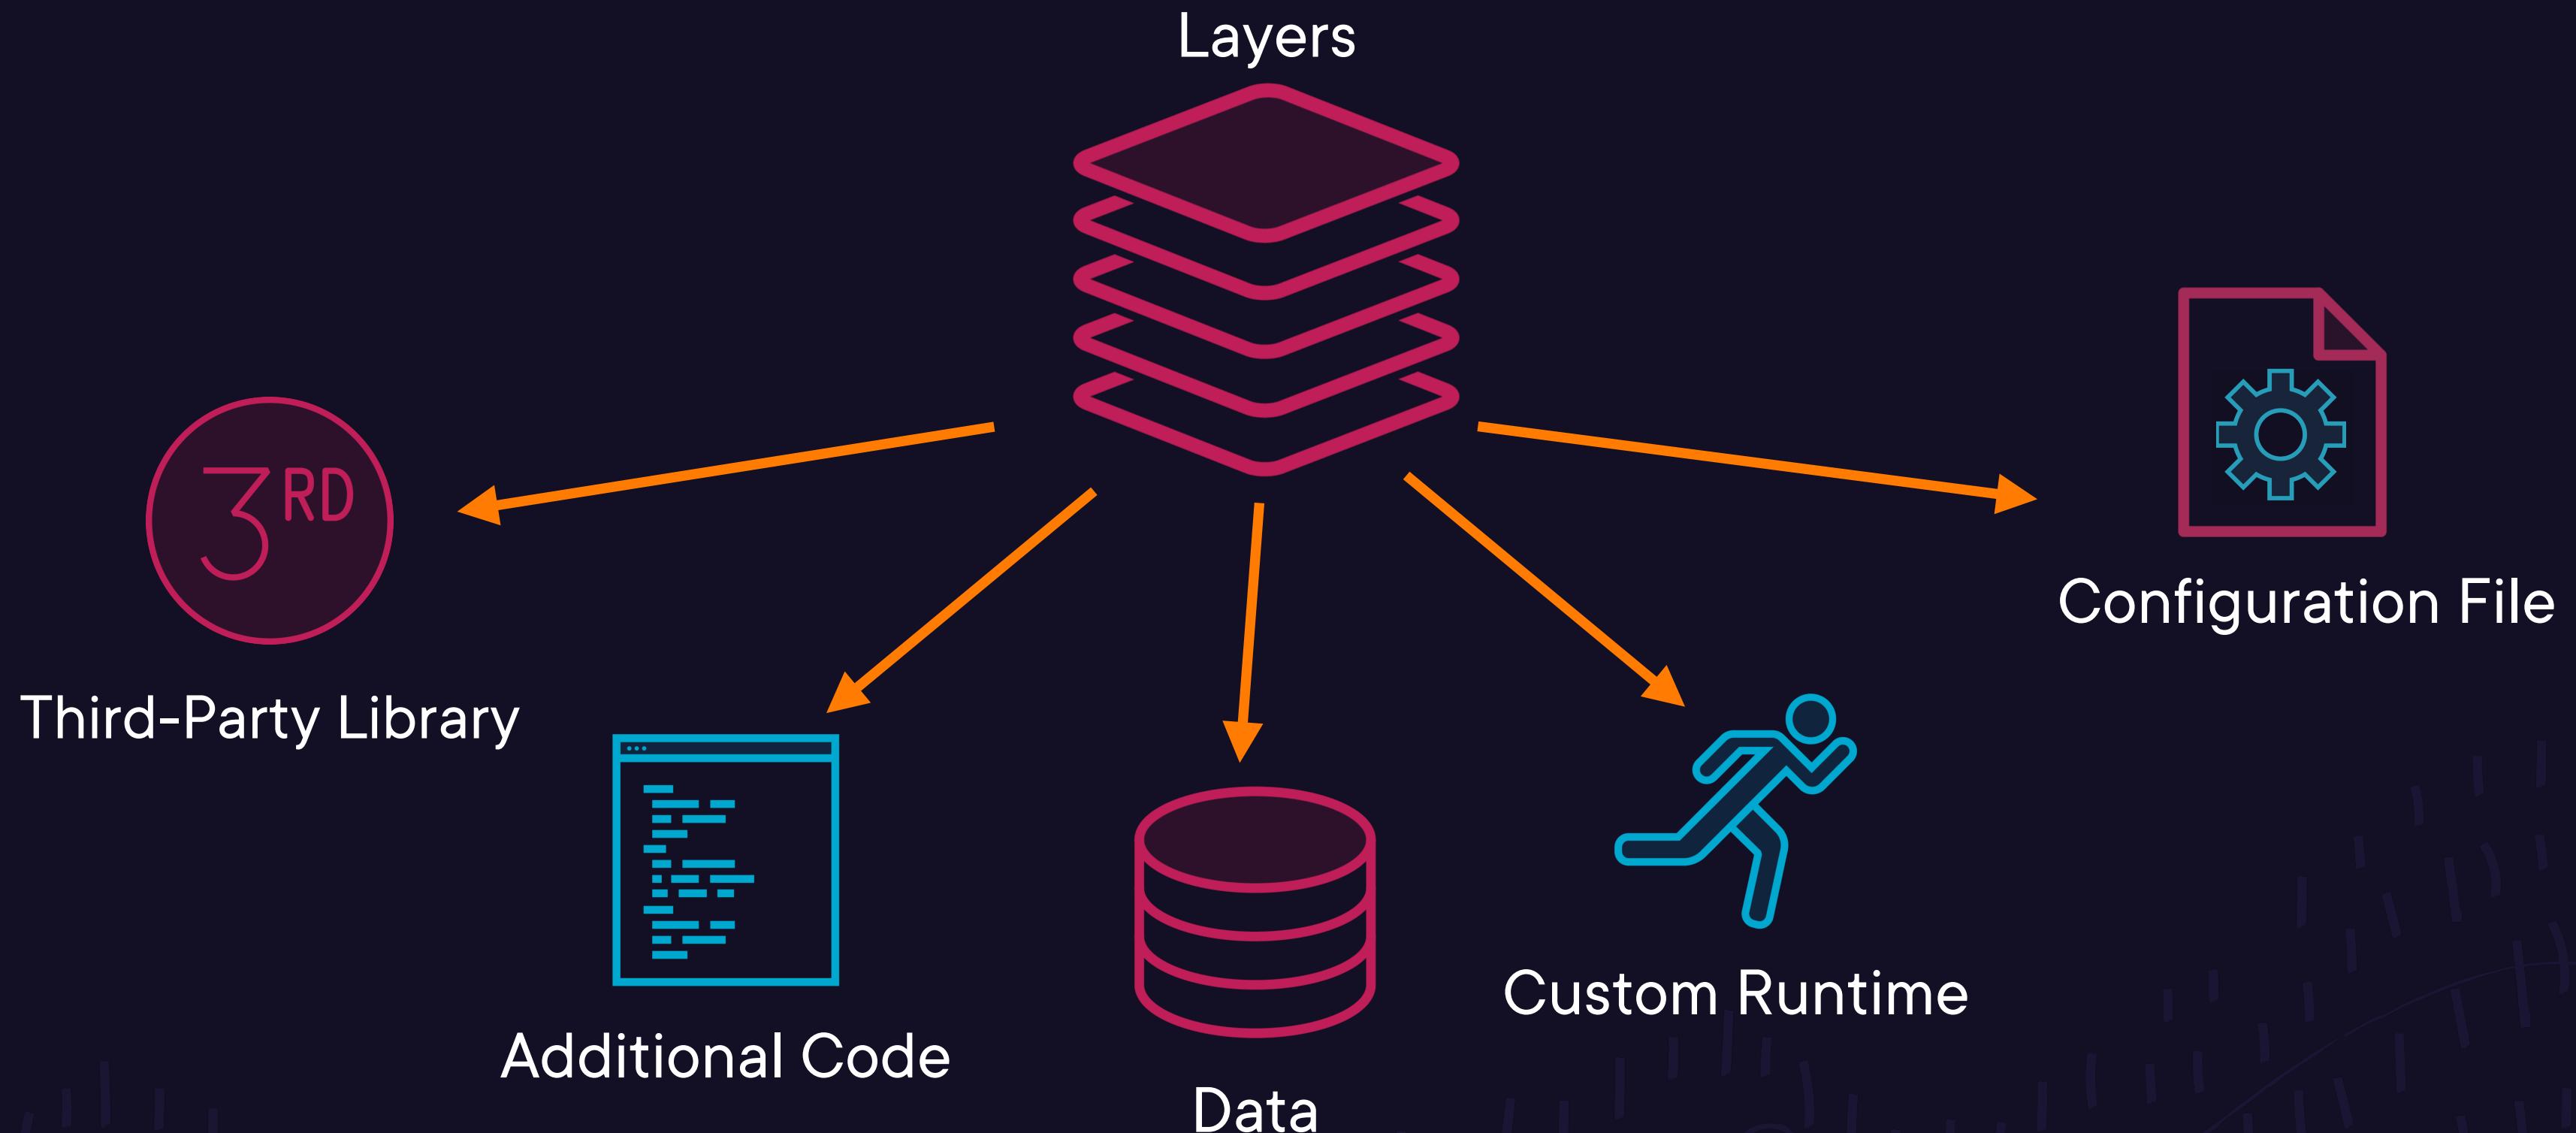
## Lambda Layers

- You can add up to **five layers** to a function.
- Each **layer** can potentially **overwrite** the previous layer; thus, the **order** is **important**.
  - **First layer**= custom runtime
  - **Second layer**= third-party library
  - **Third layer**= config file



# Understanding Lambda Layers

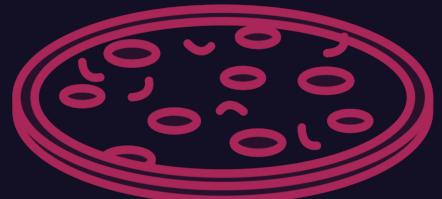
## What's in a Layer?



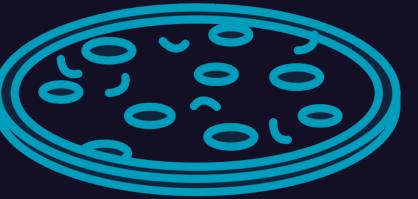
# Understanding Lambda Layers

## Why Layer Up? Story Time

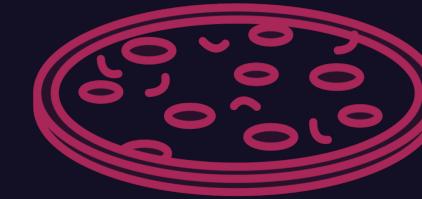
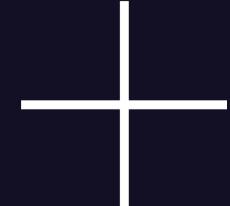
Grandma's Pizza



My Pizza



Grandma's Sauce

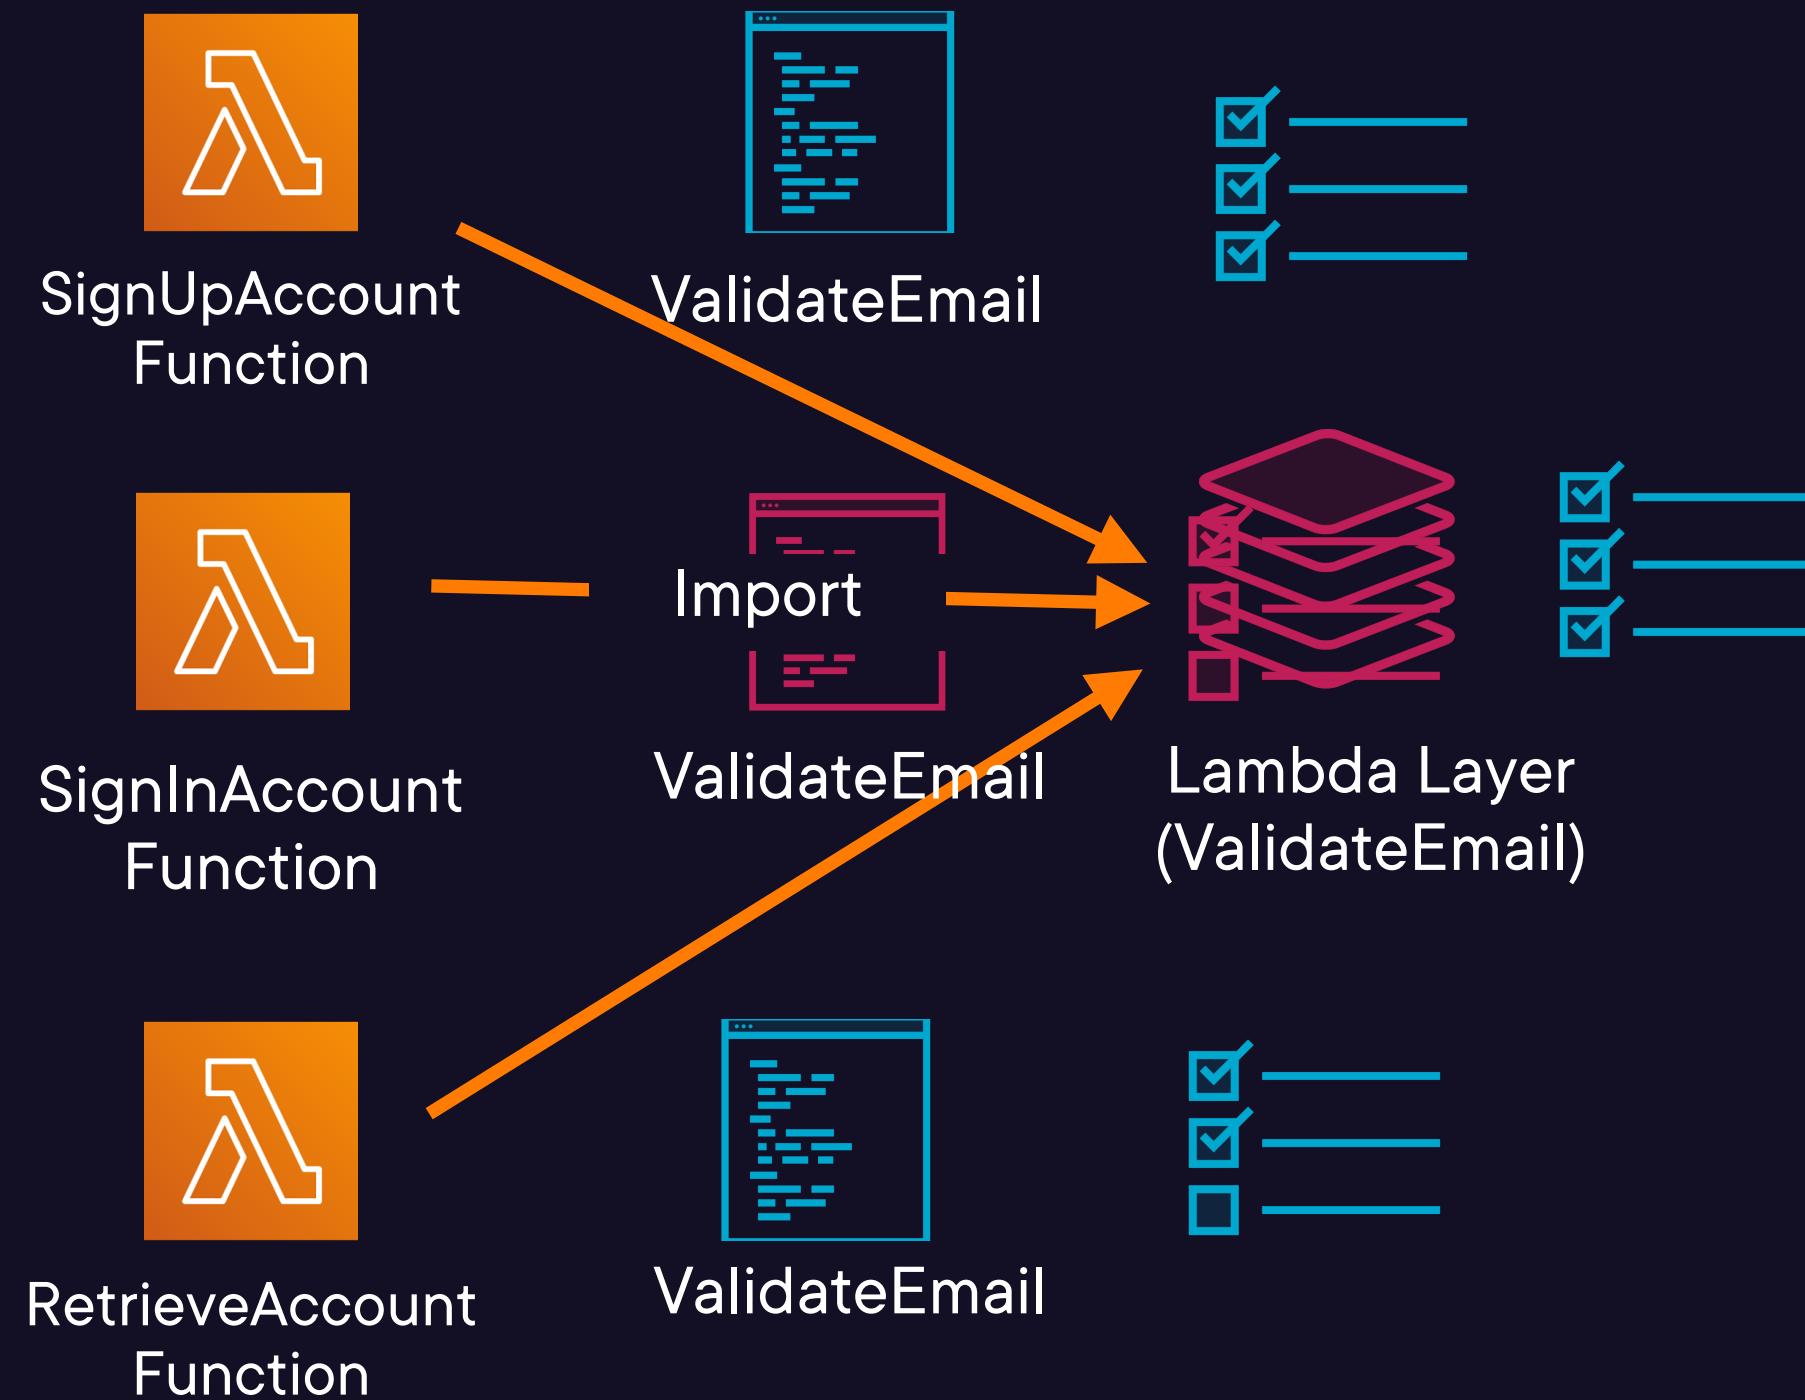


Any Pizza

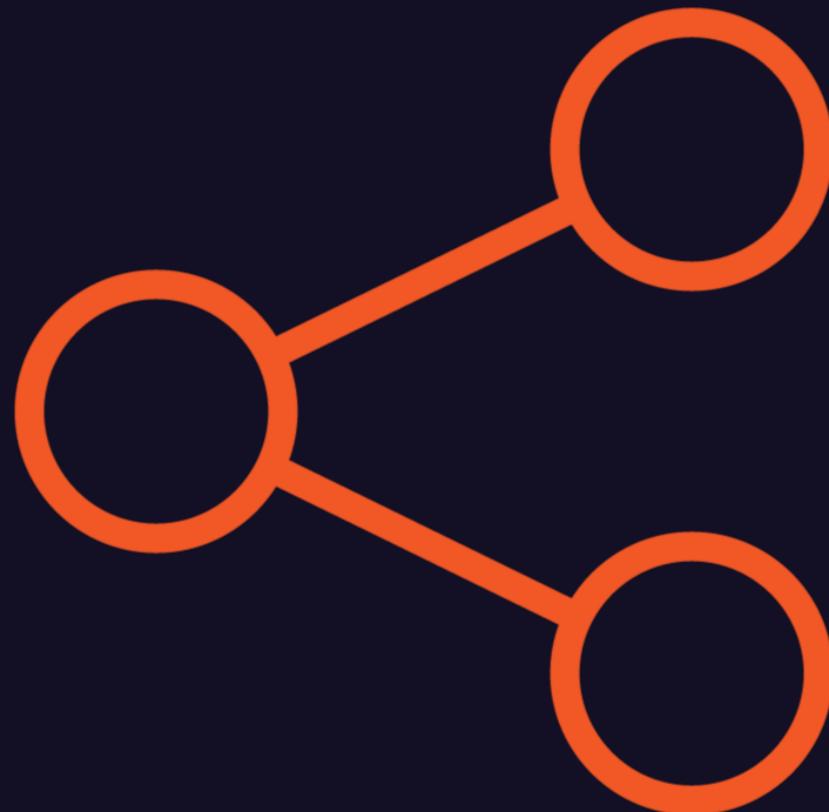


## Understanding Lambda Layers

# Why Layer Up? Code Reusability



# Why Layer Up? Code Reusability



**Layers can be shared...**

- Within an AWS account
- Between different AWS accounts
- Publicly with the developer community





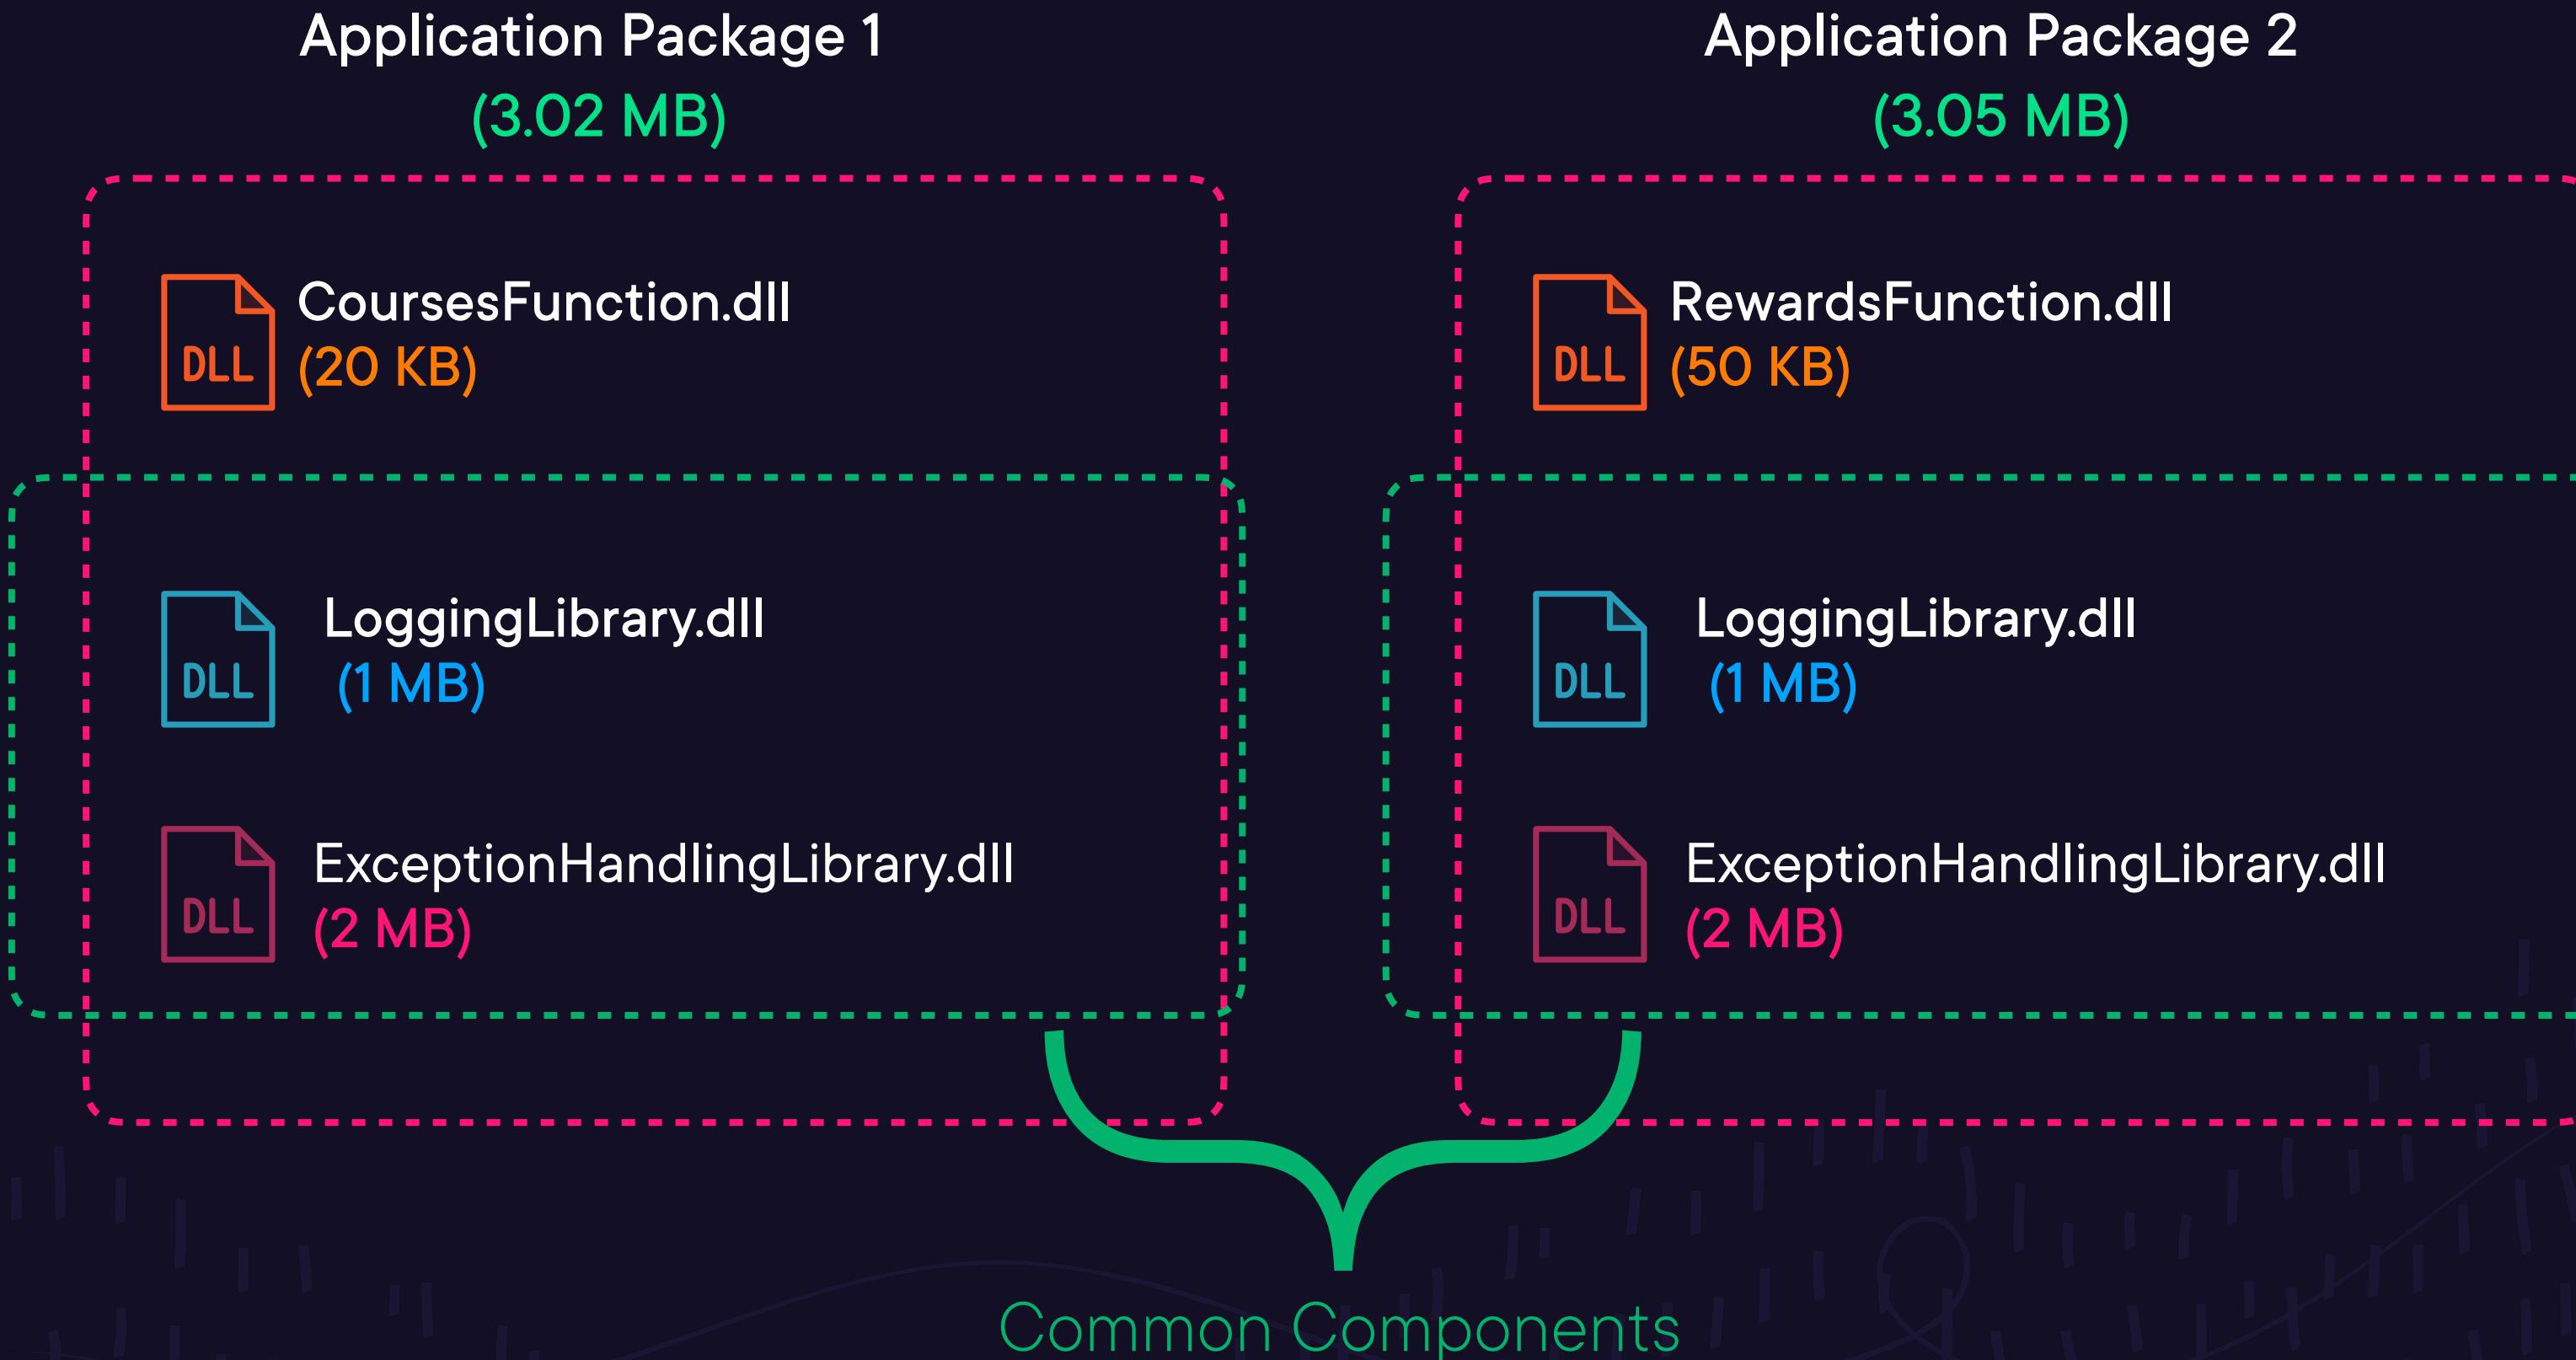
## Code Reusability

- Don't waste time rebuilding existing functionality.
- Simplify code updates.
- Bugs 🐞 are contained in layers.
- Smaller deployment packages result in faster deployments.



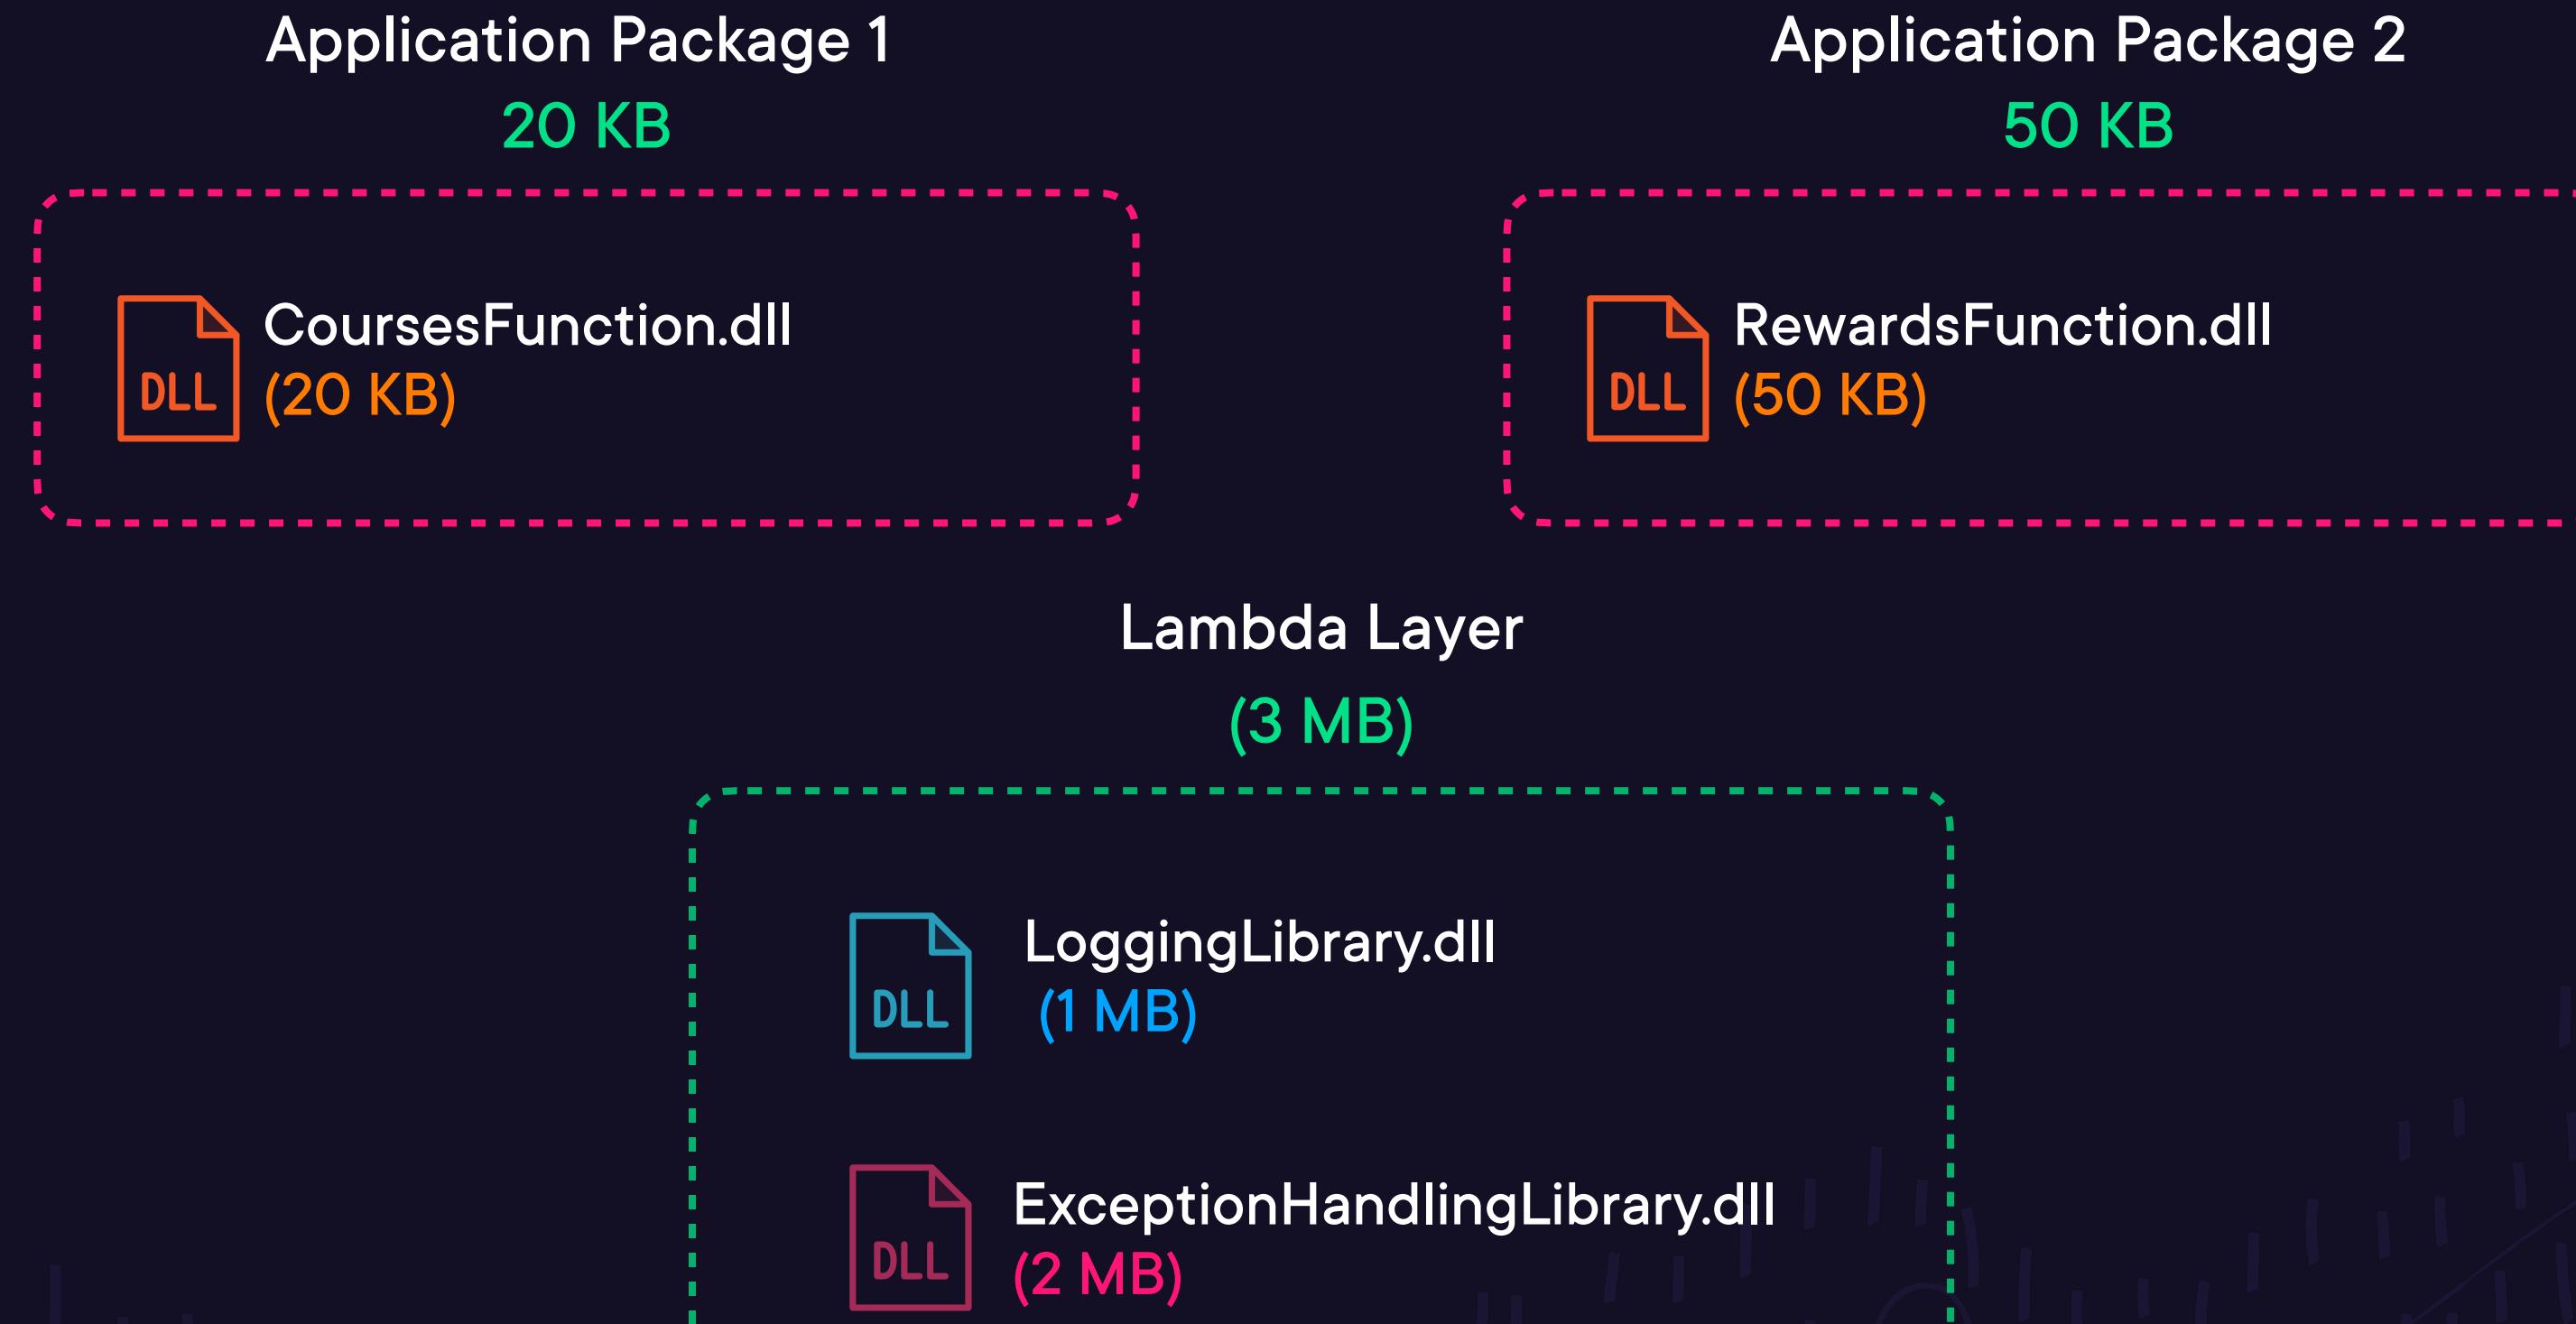
## Understanding Lambda Layers

# Why Layer Up? Smaller Deployment Packages



## Understanding Lambda Layers

# Why Layer Up? Smaller Deployment Packages





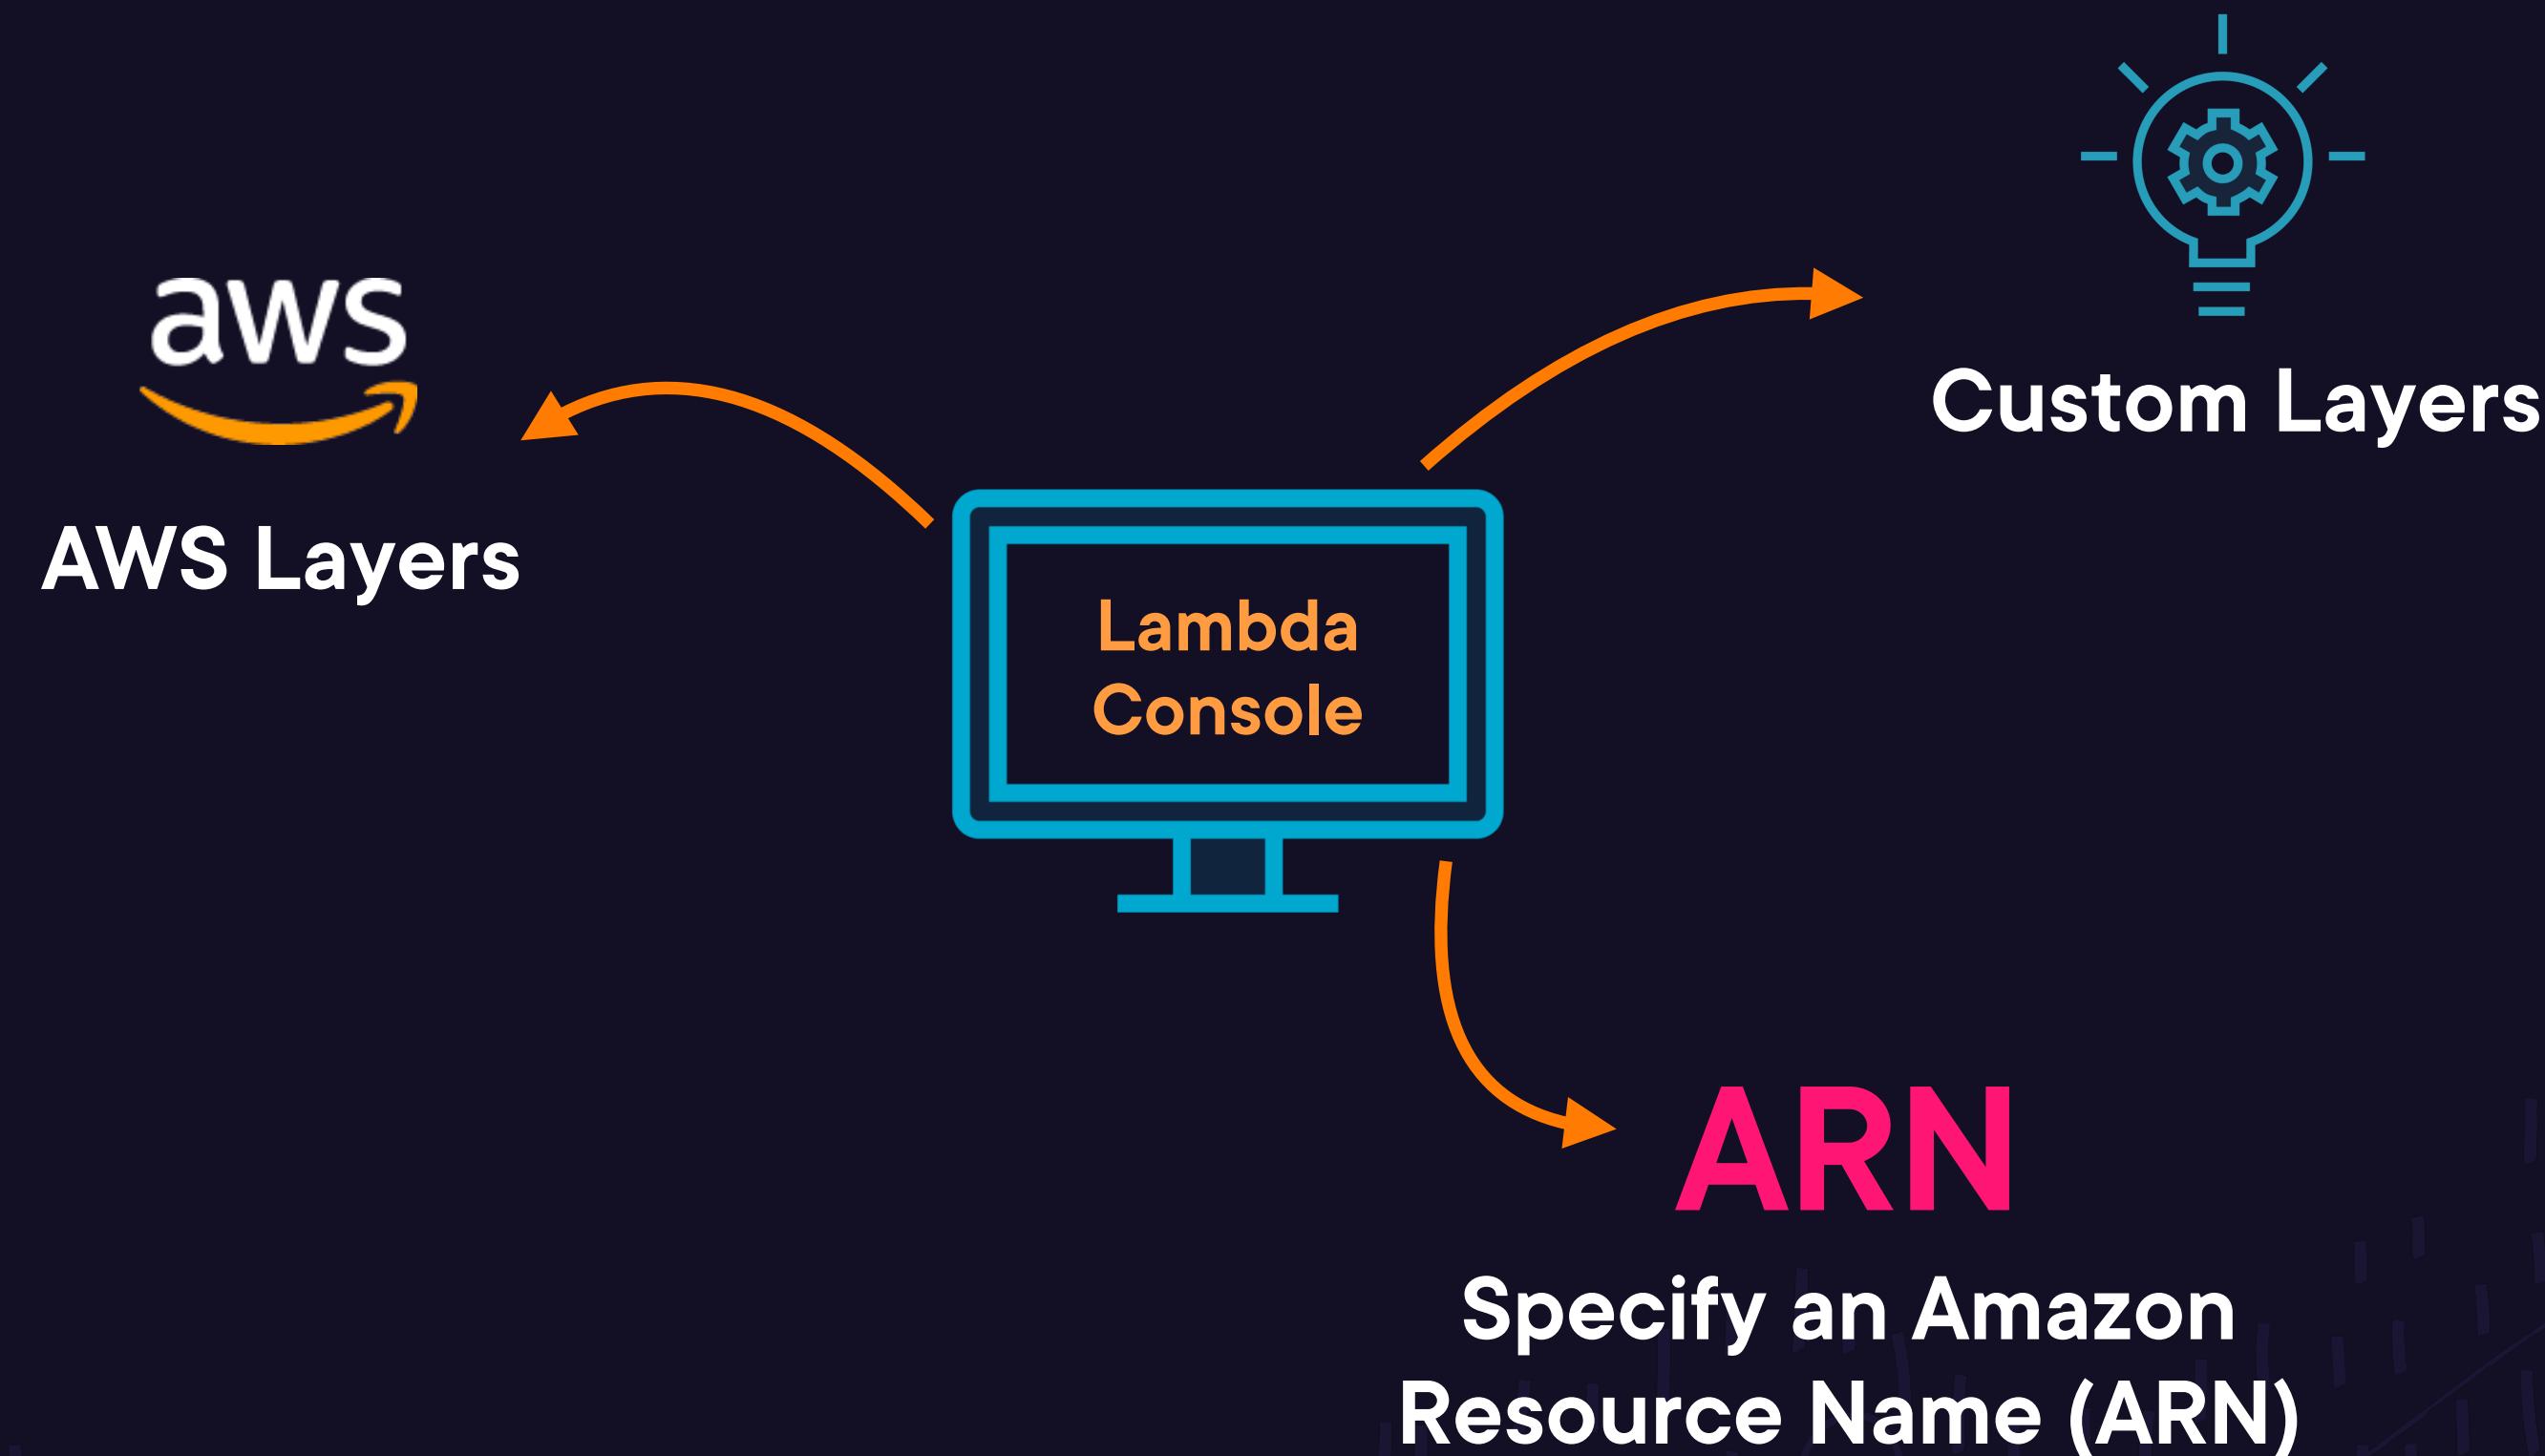
## Challenges

- Invoking and **testing** functions locally becomes **harder**.
- You **won't be able to update** the **function** if:
  - The layer is **deleted**
  - **Permissions** to use the layer are **revoked**
  - **Third party retires** an older version of their library
- The total **unzipped size** of the **function + all layers < 250 MB**.

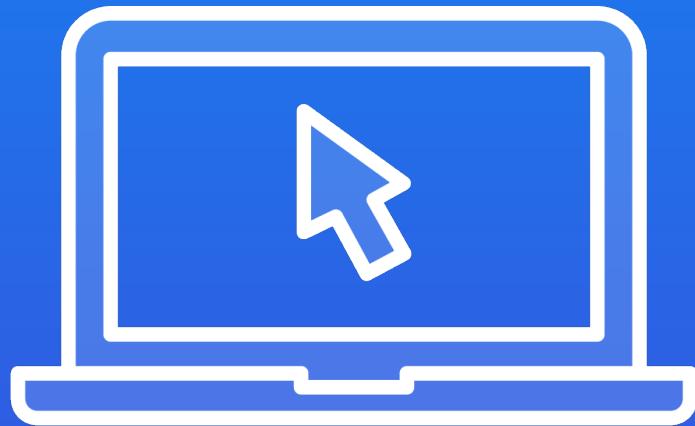


## Understanding Lambda Layers

### Lambda Console: Ways to Import Lambda Layers



# Demo



## Lambda Layers

- Download the **requests library** and create a ZIP folder.



- Use the package to create a new **custom layer** in Lambda.



# Lesson Summary



## What Are Lambda Layers?

- Enables **importing packages** into functions
- Packages can contain **additional code, third-party libraries**, config files, or **custom runtime environments**

## Why Layer Up?

- Code reusability
- Smaller packages= faster deployments

## Challenges

- Testing is harder
- Updating functions can get tricky
- Unzipped function + layers < 250 MB



# Lesson Summary



## Demo

- Created a **custom layer** for the **requests** library



# Demo: Converting JSON to CSV Using Lambda



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# Demo: Converting JSON to CSV Using Lambda



**Noreen Hasan**  
Training Architect

---

What's in a File Type?

The Workflow

Demo

Lesson Summary

---



## Demo: Converting JSON to CSV Using Lambda

### What's in a File Type?

A **central** storage repository  
that holds **raw** data in  
native format

Unstructured Data



Structured Data



Semi-structured Data

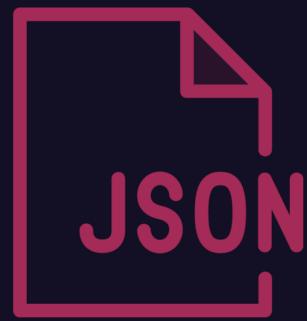


# Data Lake



## Demo: Converting JSON to CSV Using Lambda

### JSON vs. CSV



vs.



Semi-structured data

Larger file size

Less secure

Supports hierarchical relational data

Semi-structured data

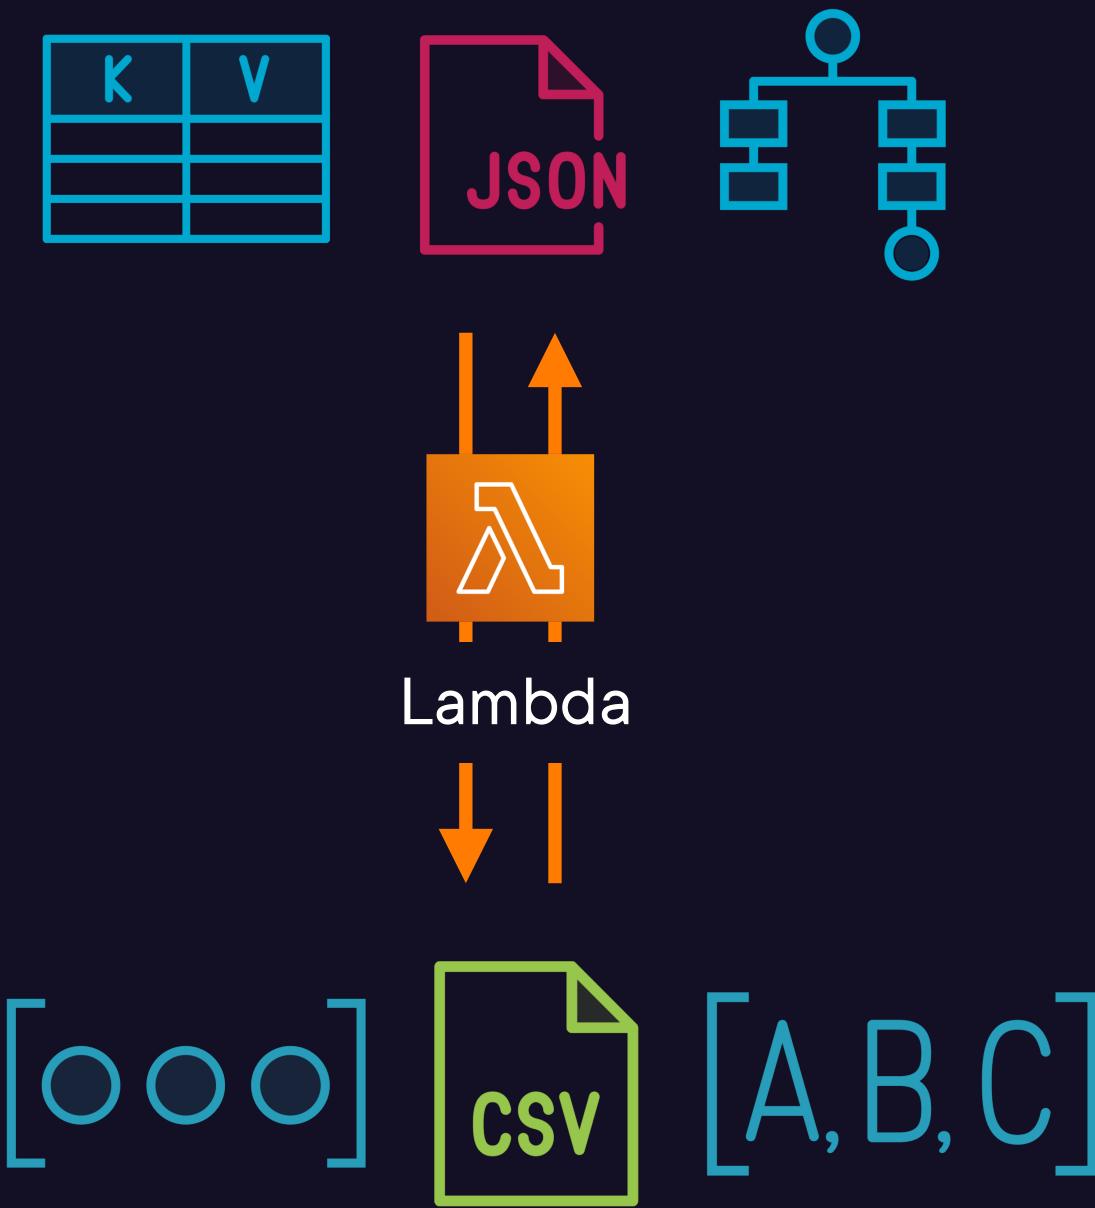
More compact

More secure

Doesn't support hierarchical structure

## Demo: Converting JSON to CSV Using Lambda

### JSON vs. CSV Use Cases



#### JSON

- Ideal for exchanging structured data
- Easily scalable and suitable for **large datasets**
- The compact size and readable nested structure makes it ideal for **APIs**

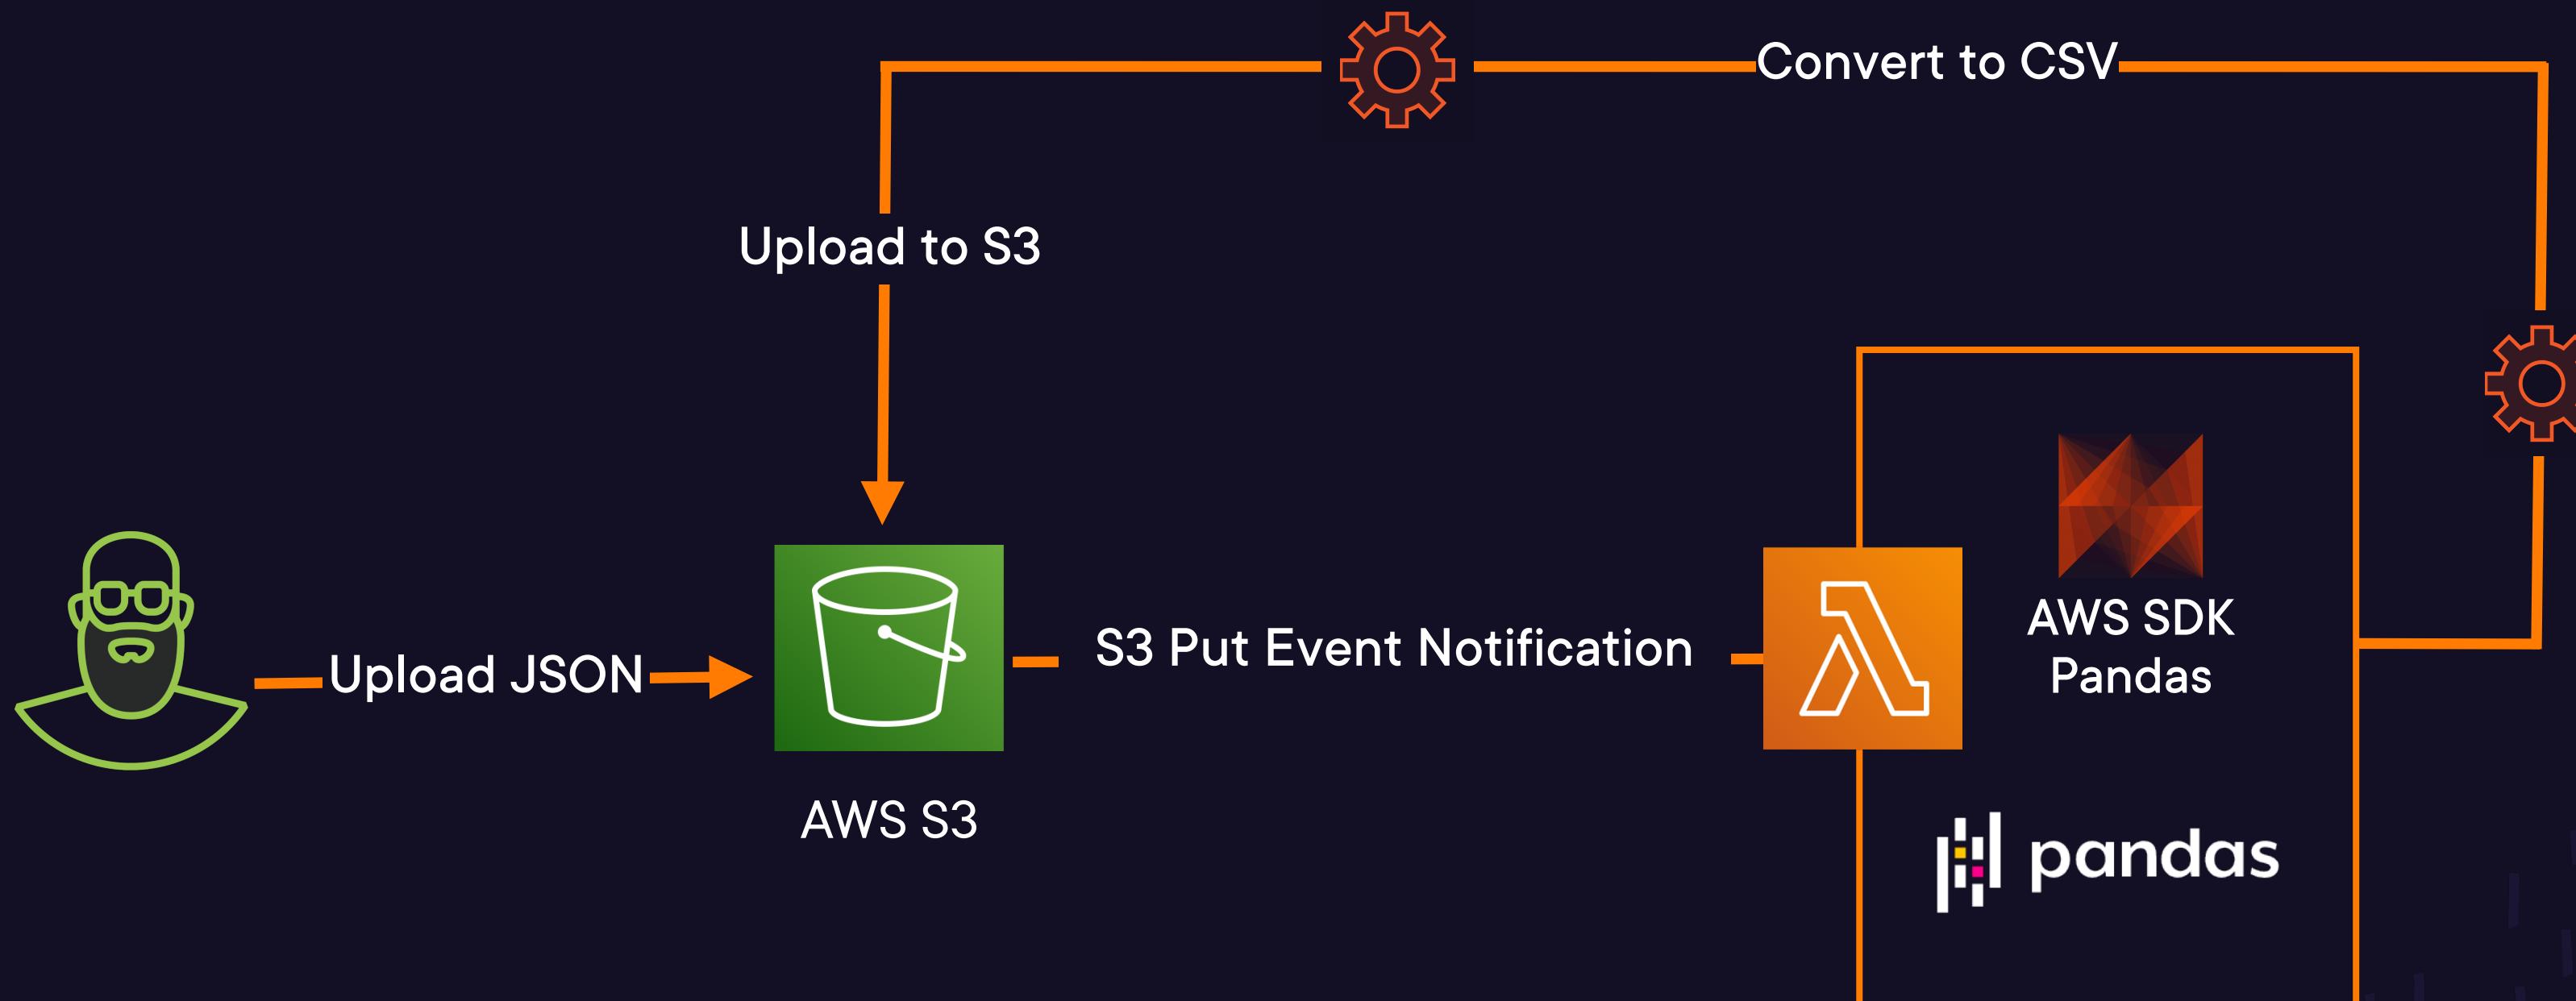
#### CSV

- Ideal for storing tabular data in a delimited text file
- Convenient for **small datasets**

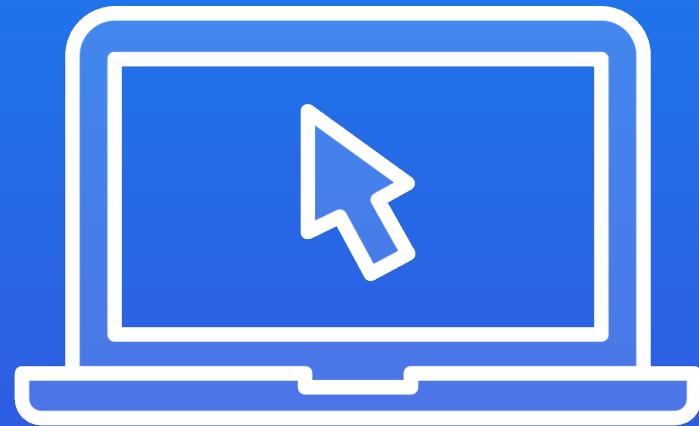


## Demo: Converting JSON to CSV Using Lambda

### The Workflow



# Demo



## Lambda Function + Layer

The function code will use the **AWS SDK for Pandas** library to convert JSON files to CSV files.

## IAM Policy

Allow **Lambda** to **list** the S3 buckets, **read**, and **upload** S3 objects to the bucket.

## S3 Bucket

Add **S3 Put Event Notification** to **trigger Lambda** each time a JSON file is uploaded to S3.



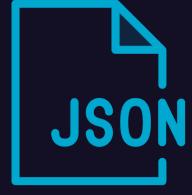
# Lesson Summary



## What's in a File Type?

- How structured, unstructured, and semi-structured data fit in a data lake 
- JSON versus CSV

## The Workflow and Demo

- Trigger Lambda when **{JSON}** files are uploaded to S3
- Function: Convert  files to 



# Demo: Ingesting CSV Files from S3 to DynamoDB



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# Demo: Ingesting CSV Files from S3 to DynamoDB



**Noreen Hasan**  
Training Architect

---

DynamoDB: Overview

The Workflow

Demo

Lesson Summary

---



## Demo: Ingesting CSV Files from S3 to DynamoDB

### DynamoDB: Overview

#### NoSQL Database

K	V

Table 1

K	V

Table 2

#### SQL Database


Table 1


Table 2

#### NoSQL Database

- **Schema-less**: it doesn't enforce relationships between the tables.
- It supports **key-value** and **document** data models, so every table must have a **unique key**.

- DynamoDB can support **semi-structured** data such as **CSV** and **JSON**.

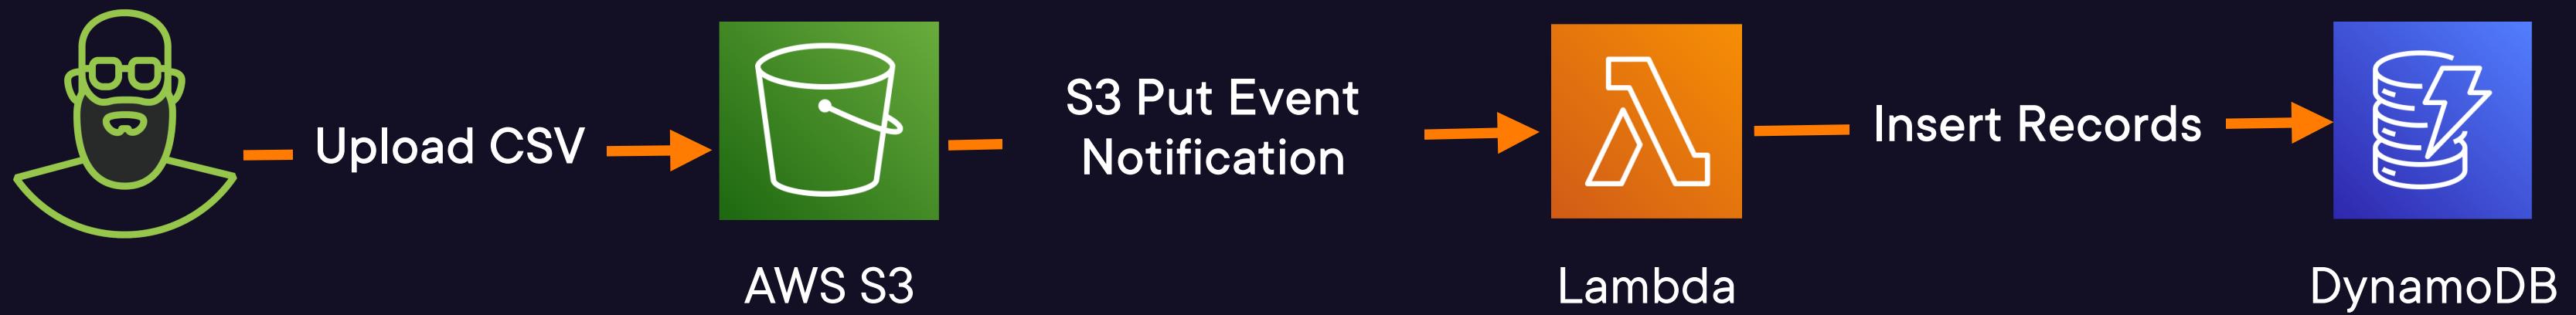
#### Fully Managed

- AWS takes care of operating and scaling the database behind the scenes.

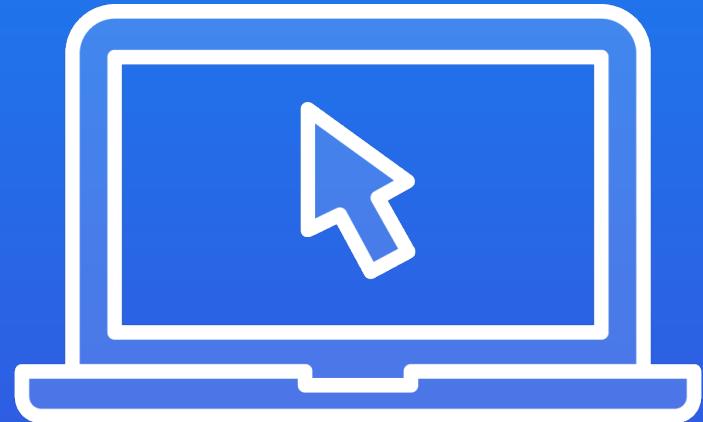


## Demo: Ingesting CSV Files from S3 to DynamoDB

### The Workflow



## Demo



### DynamoDB

Create a DynamoDB table to store the records.

### S3 Bucket

Create an S3 bucket and add an **S3 Put event notification** to **trigger Lambda** each time a CSV file is uploaded to S3.

### Lambda Function

Create a function that **inserts** the records from the CSV file to the **DynamoDB** table.

### IAM Policy

Allow **Lambda** to access **S3** and **DynamoDB**.



# Lesson Summary



## DynamoDB: Overview

- NoSQL DB
- Fully managed



## The Workflow and Demo

- **Function:** insert records to a DynamoDB table
- **Trigger:** a CSV file is uploaded to S3



# Processing Real-Time Data Using Kinesis and Lambda: Overview



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# Processing Real-Time Data Using Kinesis and Lambda: Overview



**Noreen Hasan**  
Training Architect

---

Streaming Data: Use Cases

The Kinesis Umbrella

Kinesis Data Streams

---

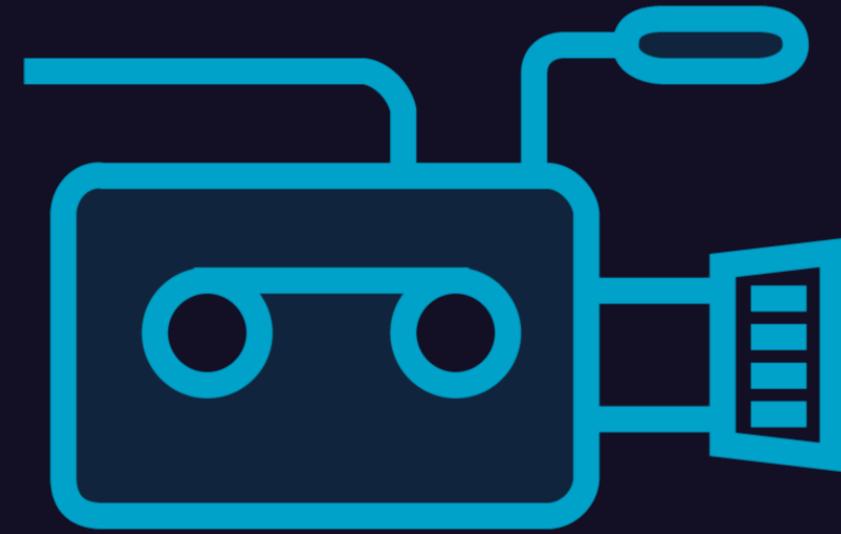
Lesson Summary

---

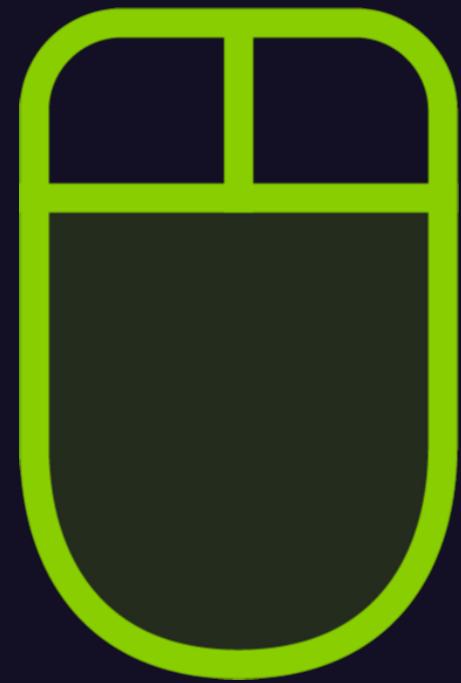


# Processing Real-Time Data Using Kinesis and Lambda: Overview

## Streaming Data: Use Cases



Video Processing



Clickstream Events



Data Lakes



## The Kinesis Umbrella

### Amazon Kinesis Data Streams

Store, process, and aggregate data streams.

### Amazon Kinesis Video Streams

Store and process video streams.

### Amazon Kinesis Data Analytics

Use SQL to analyze data streams.

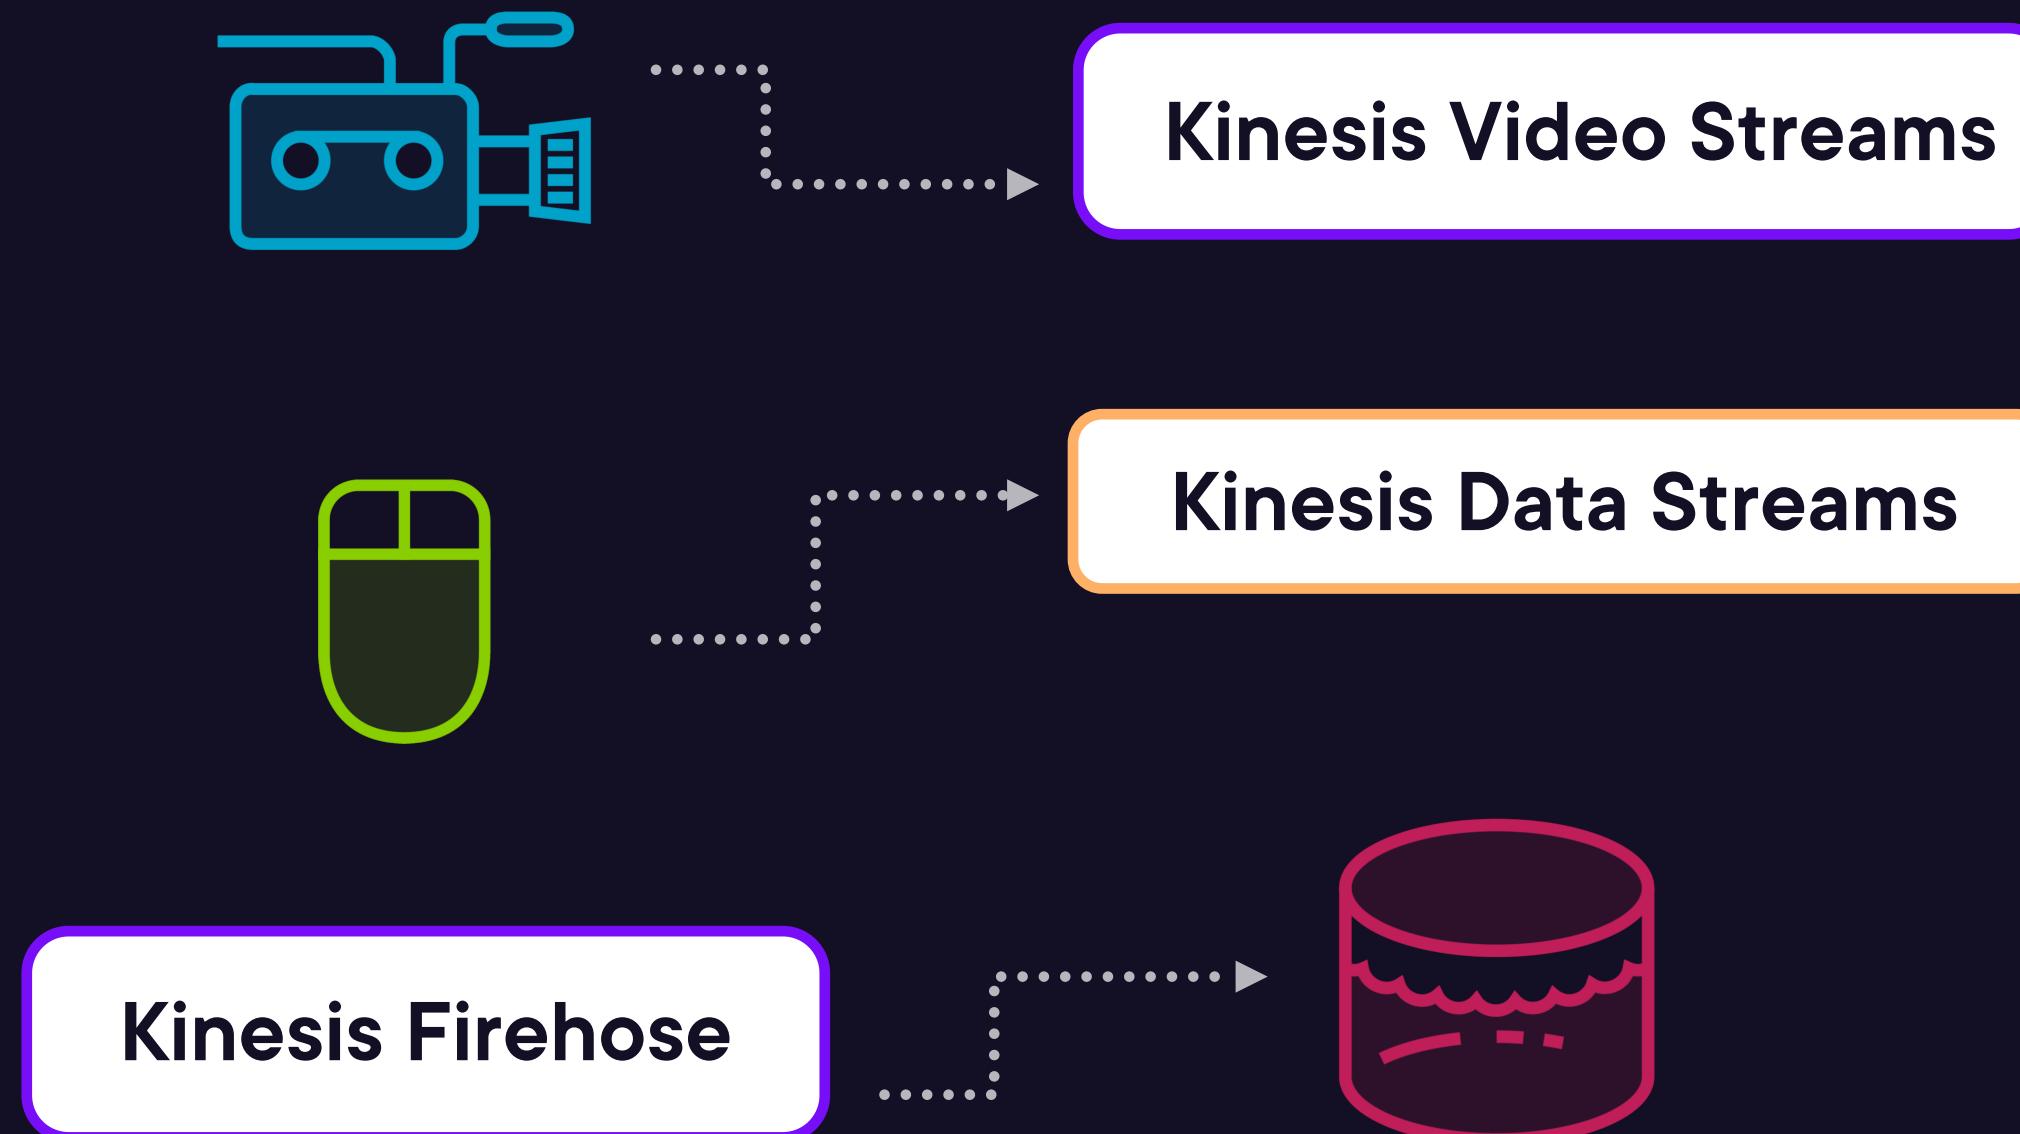
### Amazon Kinesis Data Firehose

Load data streams into AWS data stores.



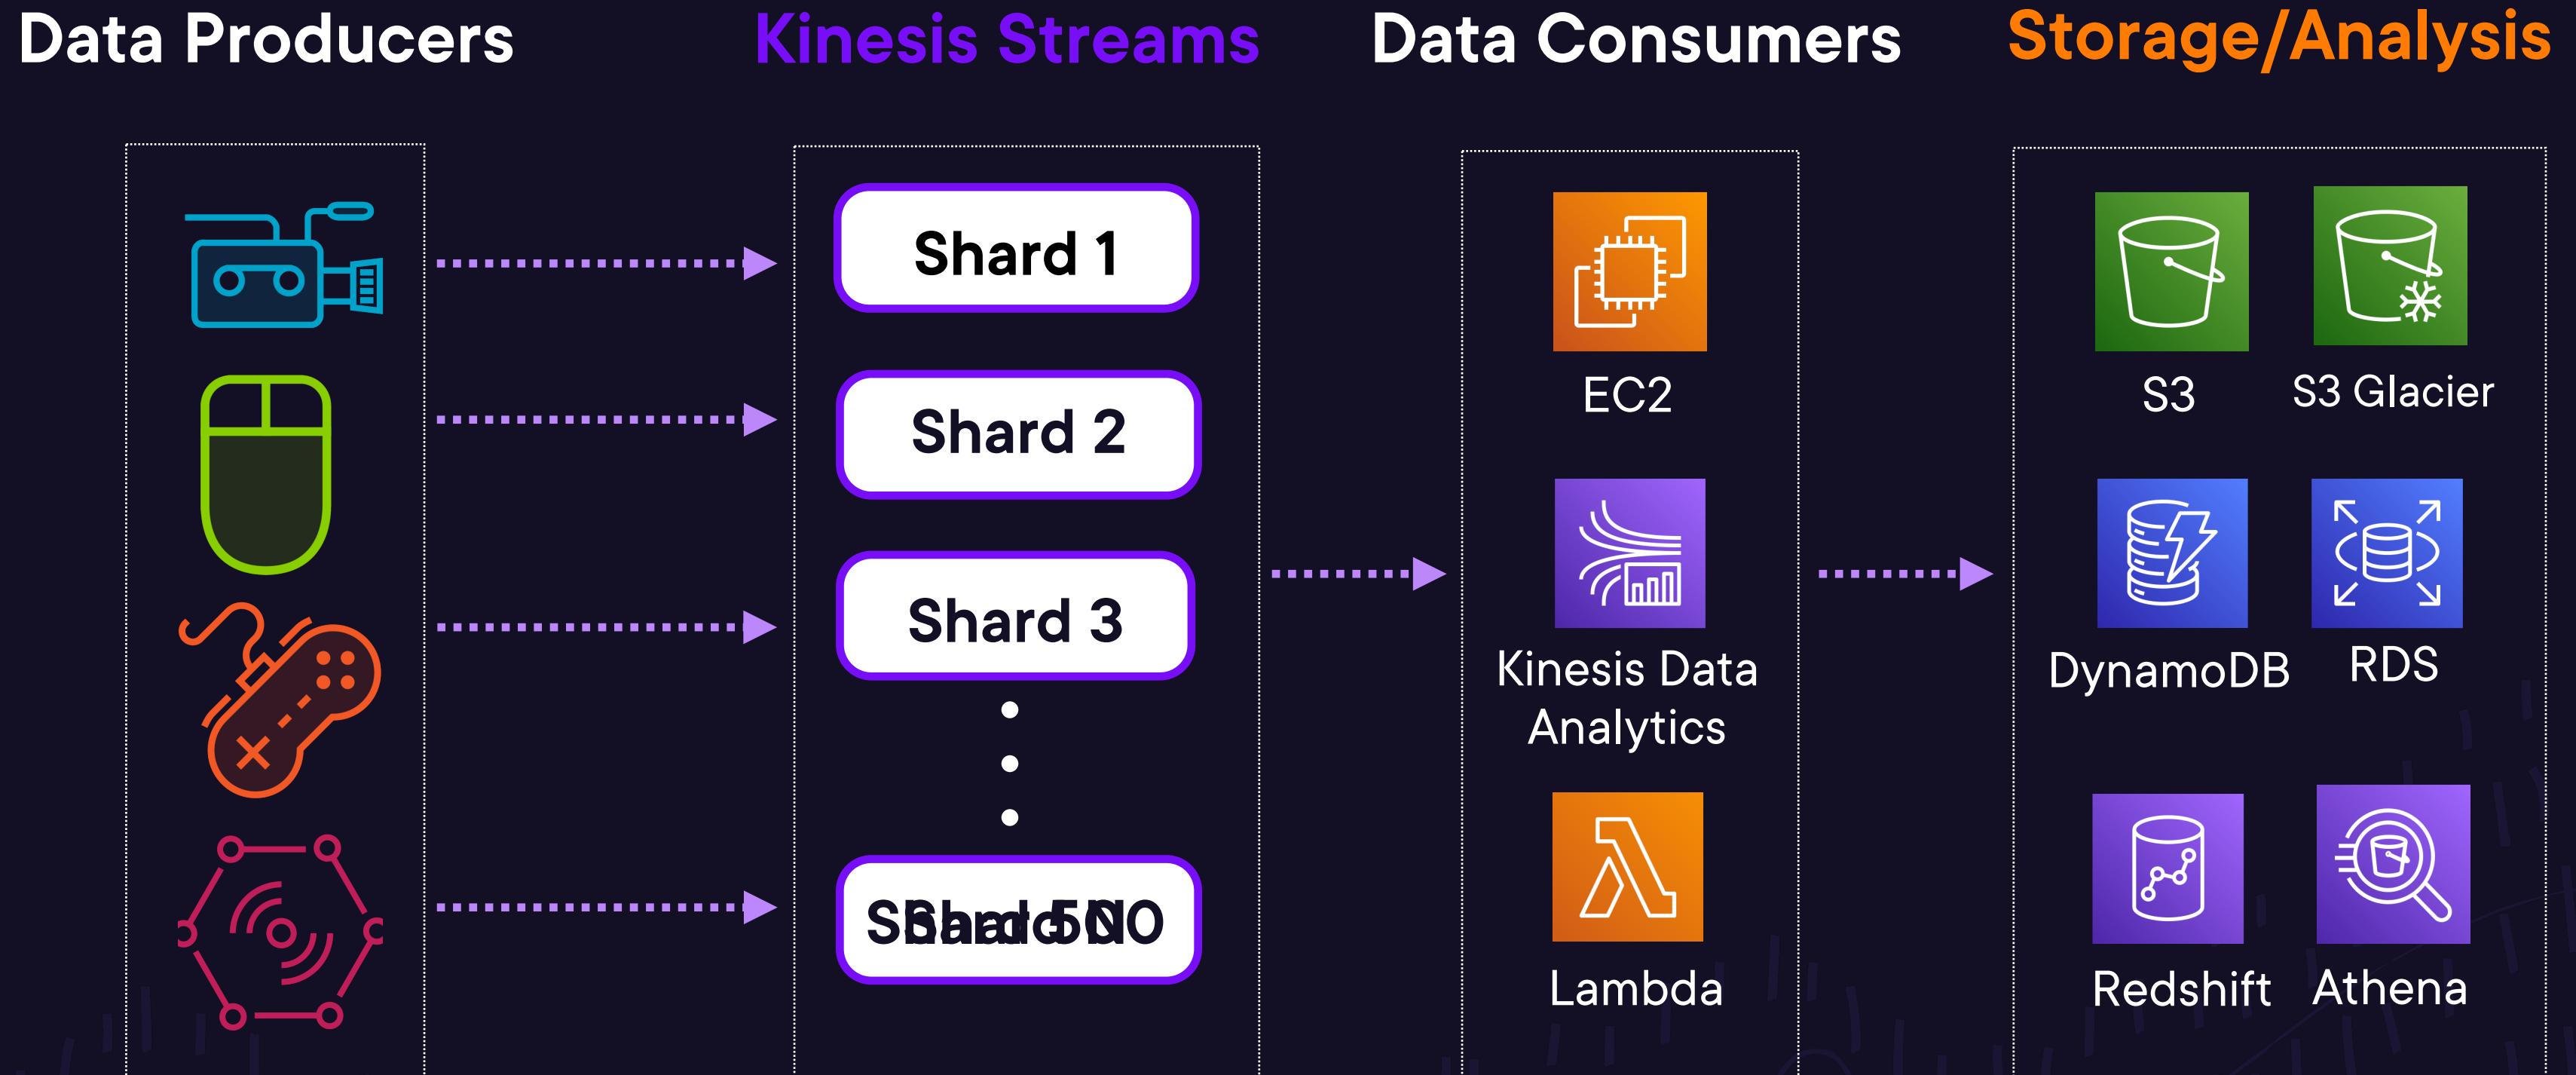
# Processing Real-Time Data Using Kinesis and Lambda: Overview

## Use Cases and the Kinesis Umbrella



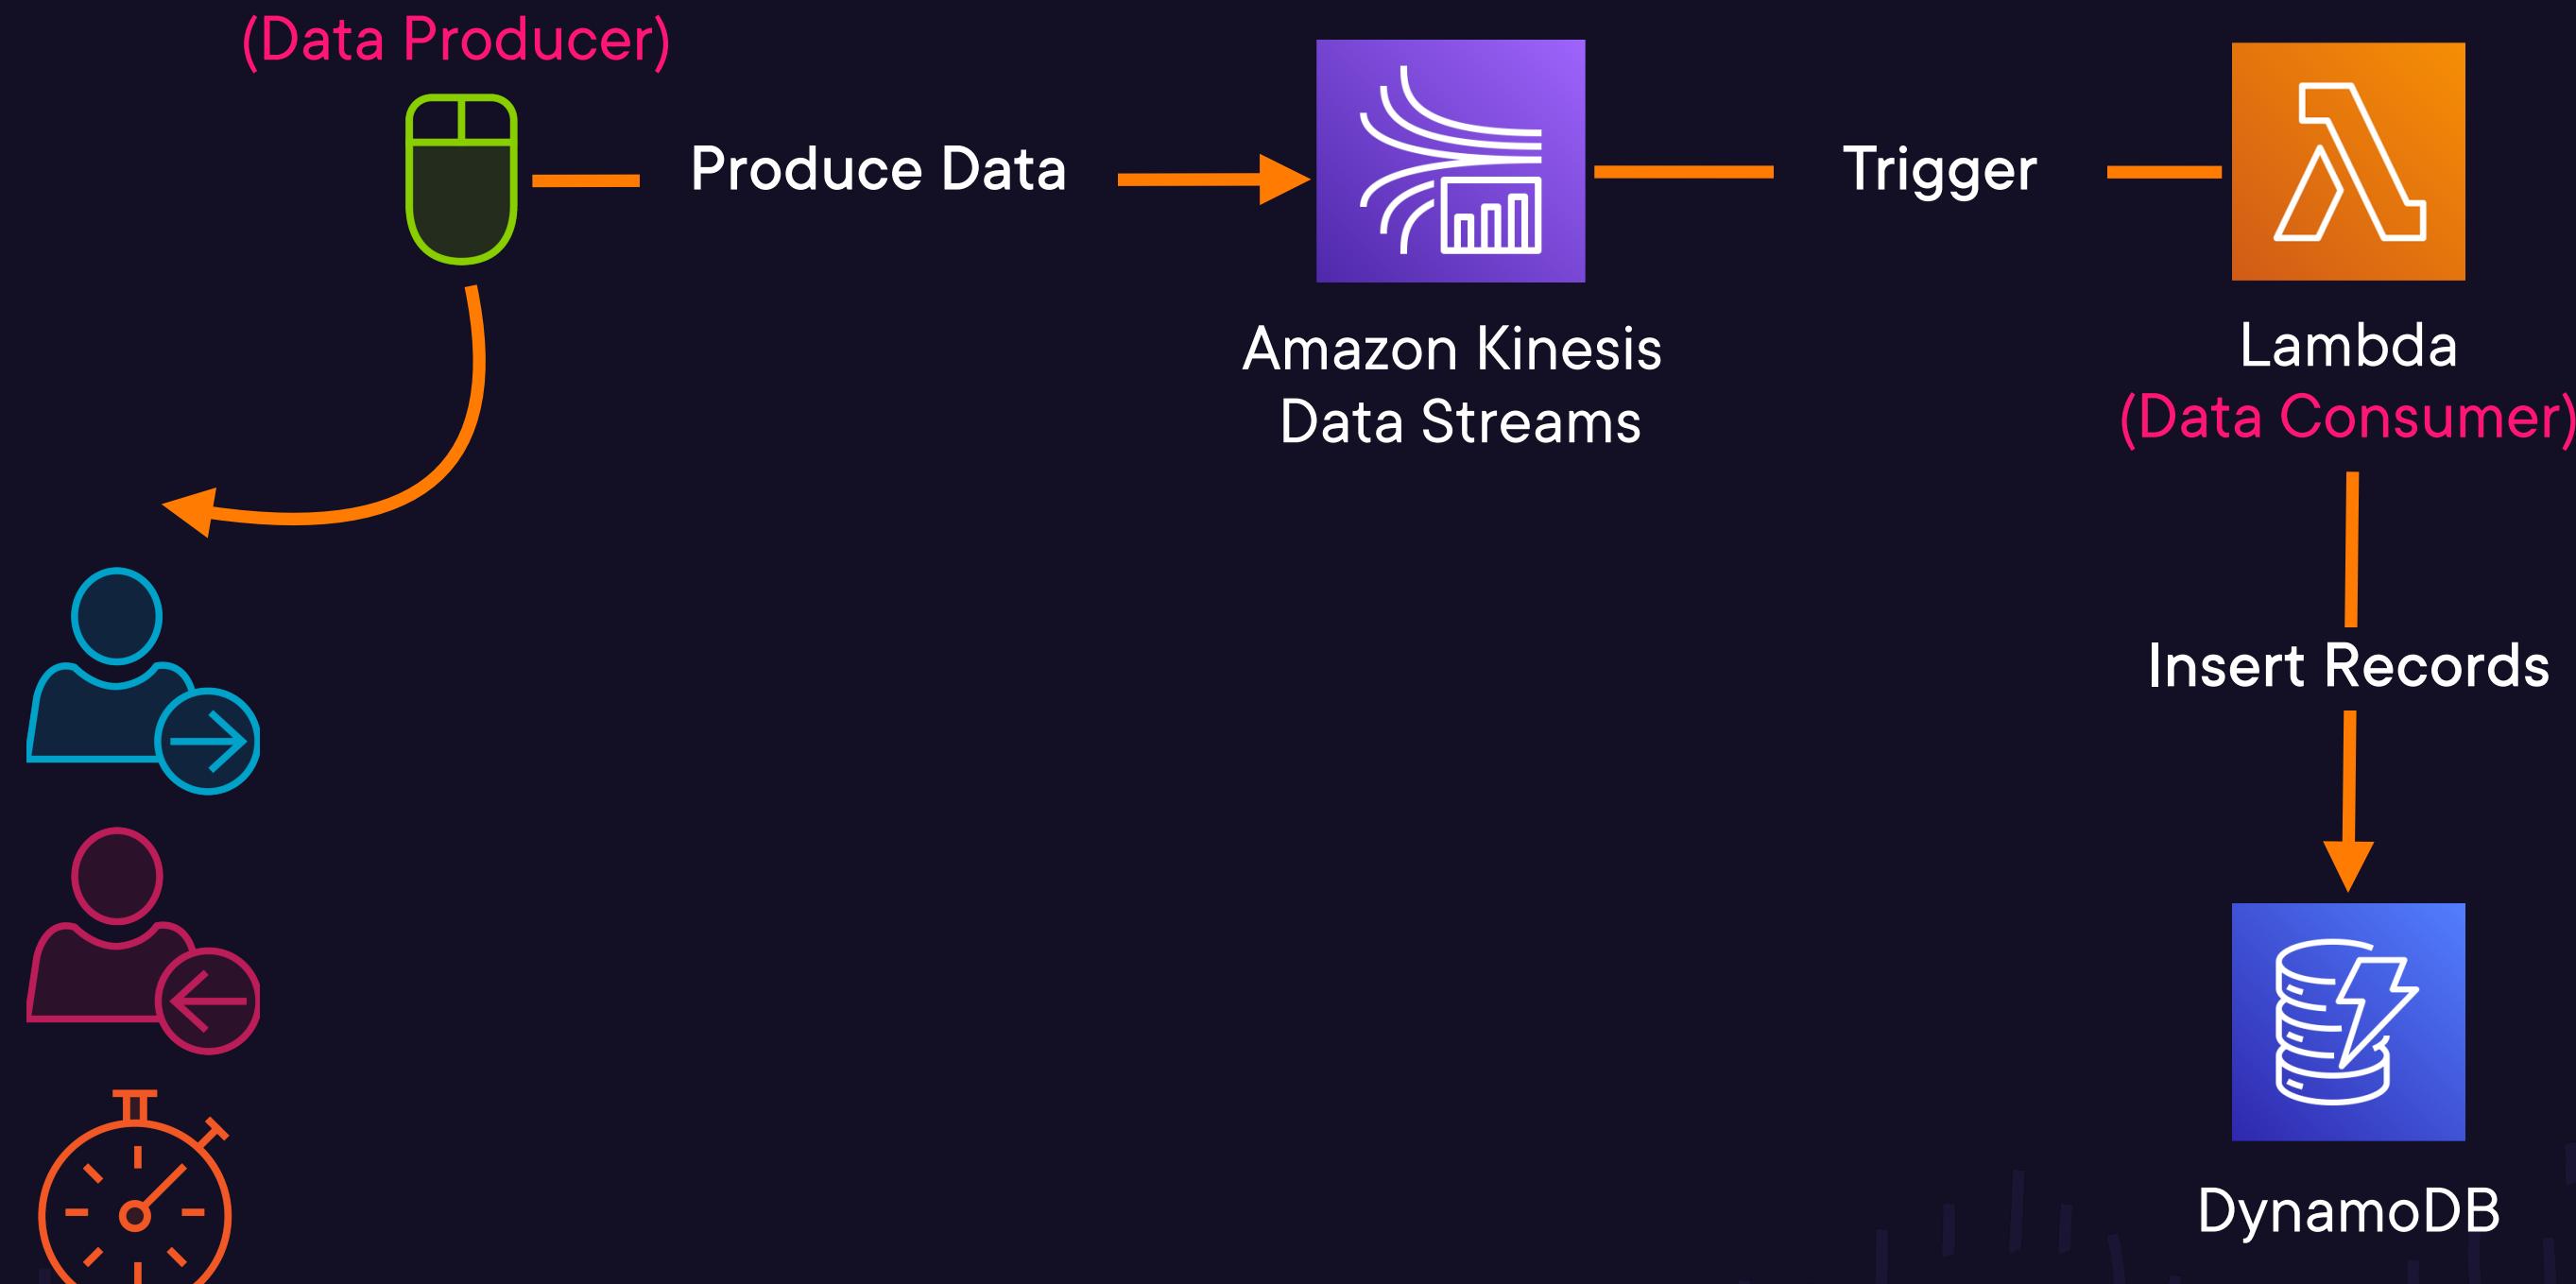
# Processing Real-Time Data Using Kinesis and Lambda: Overview

## Kinesis Data Streams



# Processing Real-Time Data Using Kinesis and Lambda: Overview

## Kinesis Data Streams: Sample Workflow



# Lesson Summary



## Streaming Data: Use Cases

- Video processing 
- Clickstream events 
- Data lakes 

## The Kinesis Umbrella

- Kinesis Data Streams
- Kinesis Video Streams
- Kinesis Data Analytics
- Kinesis Data Firehose



## Main Components of Kinesis Data Streams

- Data producers
- Kinesis streams
- Data consumers



# Demo: Processing Real-Time Data Using Kinesis and Lambda



**Noreen Hasan**

Training Architect



## LESSON BREAKDOWN

# Demo: Processing Real-Time Data Using Kinesis and Lambda



**Noreen Hasan**  
Training Architect

---

Kinesis, Lambda, and Shards

Demo

Lesson Summary

---



## Demo: Processing Real-Time Data Using Kinesis and Lambda

# Kinesis, Lambda, and Shards

### Data Producers



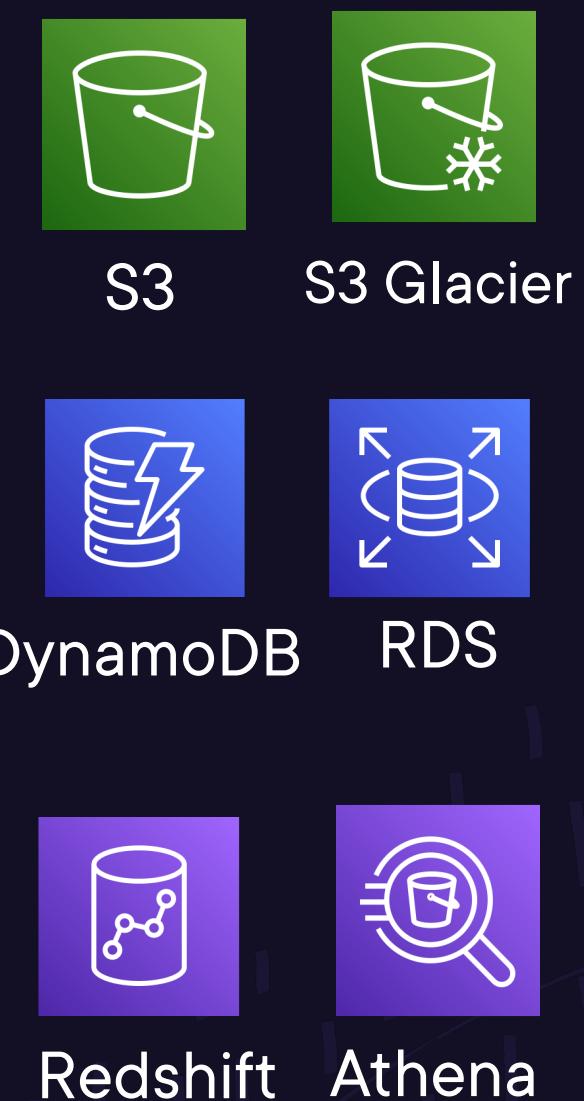
### Kinesis Streams



### Data Consumers



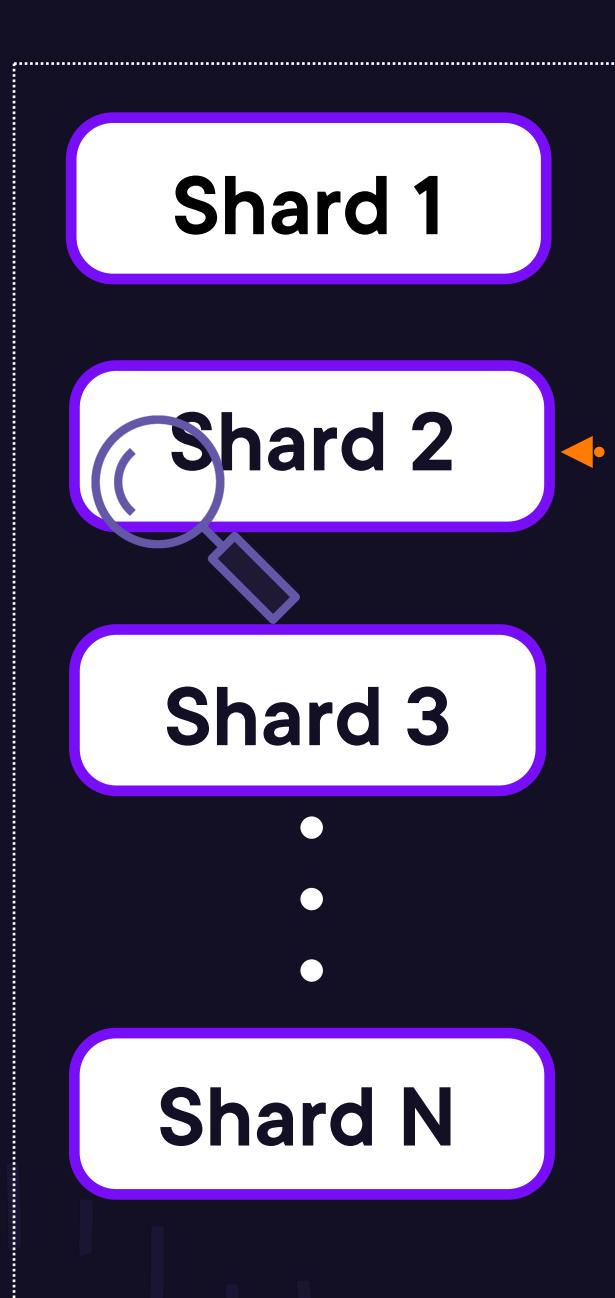
### Storage/Analysis



## Demo: Processing Real-Time Data Using Kinesis and Lambda

### Kinesis Data Streams: Shards and Lambda

#### Kinesis Streams



#### A Sequence of Records

Record 1

Record 2

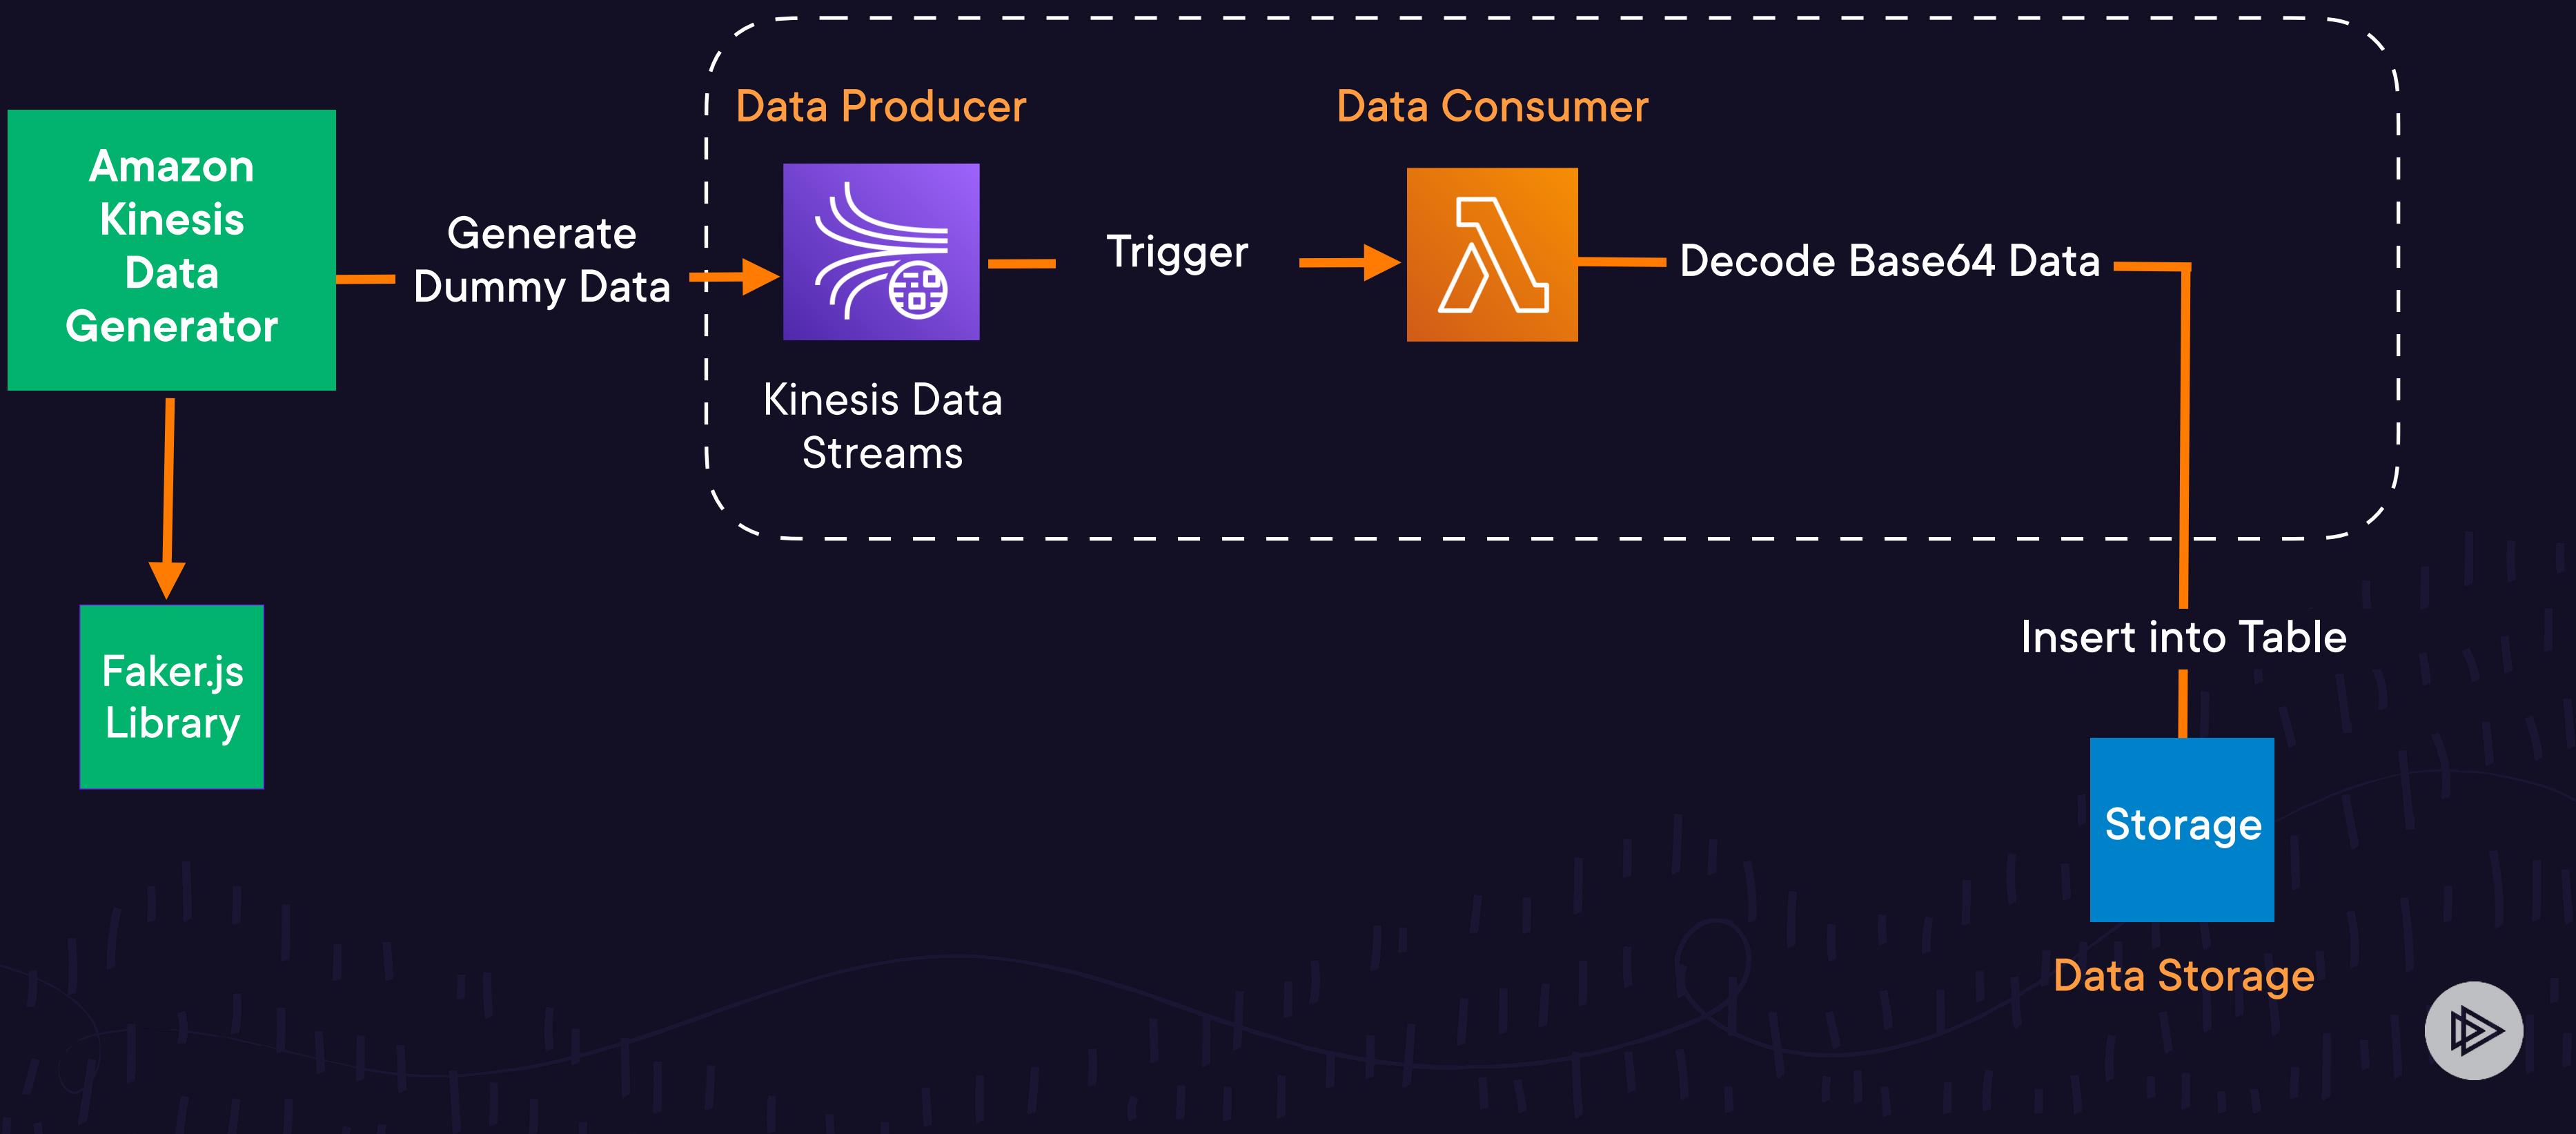
Record 3

```
"kinesis": {  
  "kinesisSchemaVersion": "1.0",  
  "partitionKey": "2",  
  "sequenceNumber": "496358289288593949635961132  
21898552279281960528115662850",  
  "data": "SGVsbG8sIHRoaXMgaXMgYSB0ZXN0Lg==",  
  "approximateArrivalTimestamp": 1545084650.987}
```

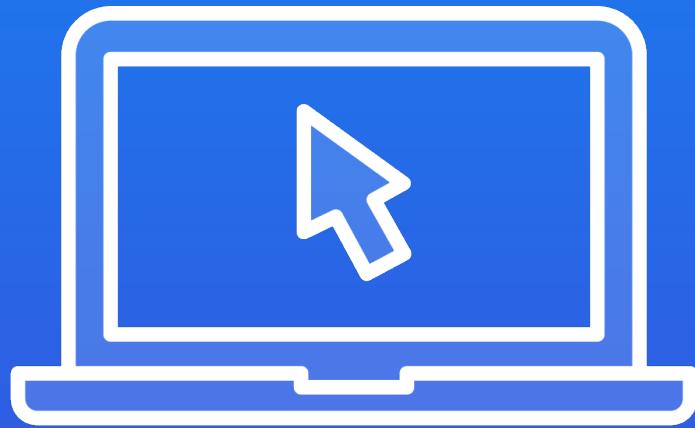


## Demo: Processing Real-Time Data Using Kinesis and Lambda

### The Workflow: Decode Shards



# Demo



## Kinesis + Lambda

Create a **Kinesis Data Stream** and **Lambda function** to decode the incoming records.

## IAM Policy

Allow Lambda access to Kinesis Data Streams.

## CloudWatch Logs

Utilize CloudWatch to view the data stream records.



# Lesson Summary



## Kinesis, Lambda, and Shards

A Kinesis data stream consists of **shards**, which consist of a set of decodable records.

## The Workflow and Demo

Created a **Lambda** function to decode Base64 data coming from **Kinesis data streams**.



# Testing and Debugging Lambda Failures



**Noreen Hasan**

Training Architect



# Testing and Debugging Lambda Failures

## LESSON BREAKDOWN

Invocation Types

Types of Failures

Testing

Monitoring

Lesson Summary



**Noreen Hasan**  
Training Architect





## Synchronous Invocation

- The invoking service **waits for a response** from Lambda.
- Lambda **returns a response** to the invoking service.

### Lambda's Retry Behavior: None

The client (i.e., invoking service) is in charge of retries on failure.





## Asynchronous Invocation

- The invoking service **doesn't wait** for **response** from Lambda.
- The invoking service is **only** in charge of **delivering the event**, Lambda handles the rest.

## Lambda's Retry Behavior: 0-2 Attempts

- Lambda is in charge of retries.
- Extra charges.
- You can configure

**MaximumRetryAttempts.**



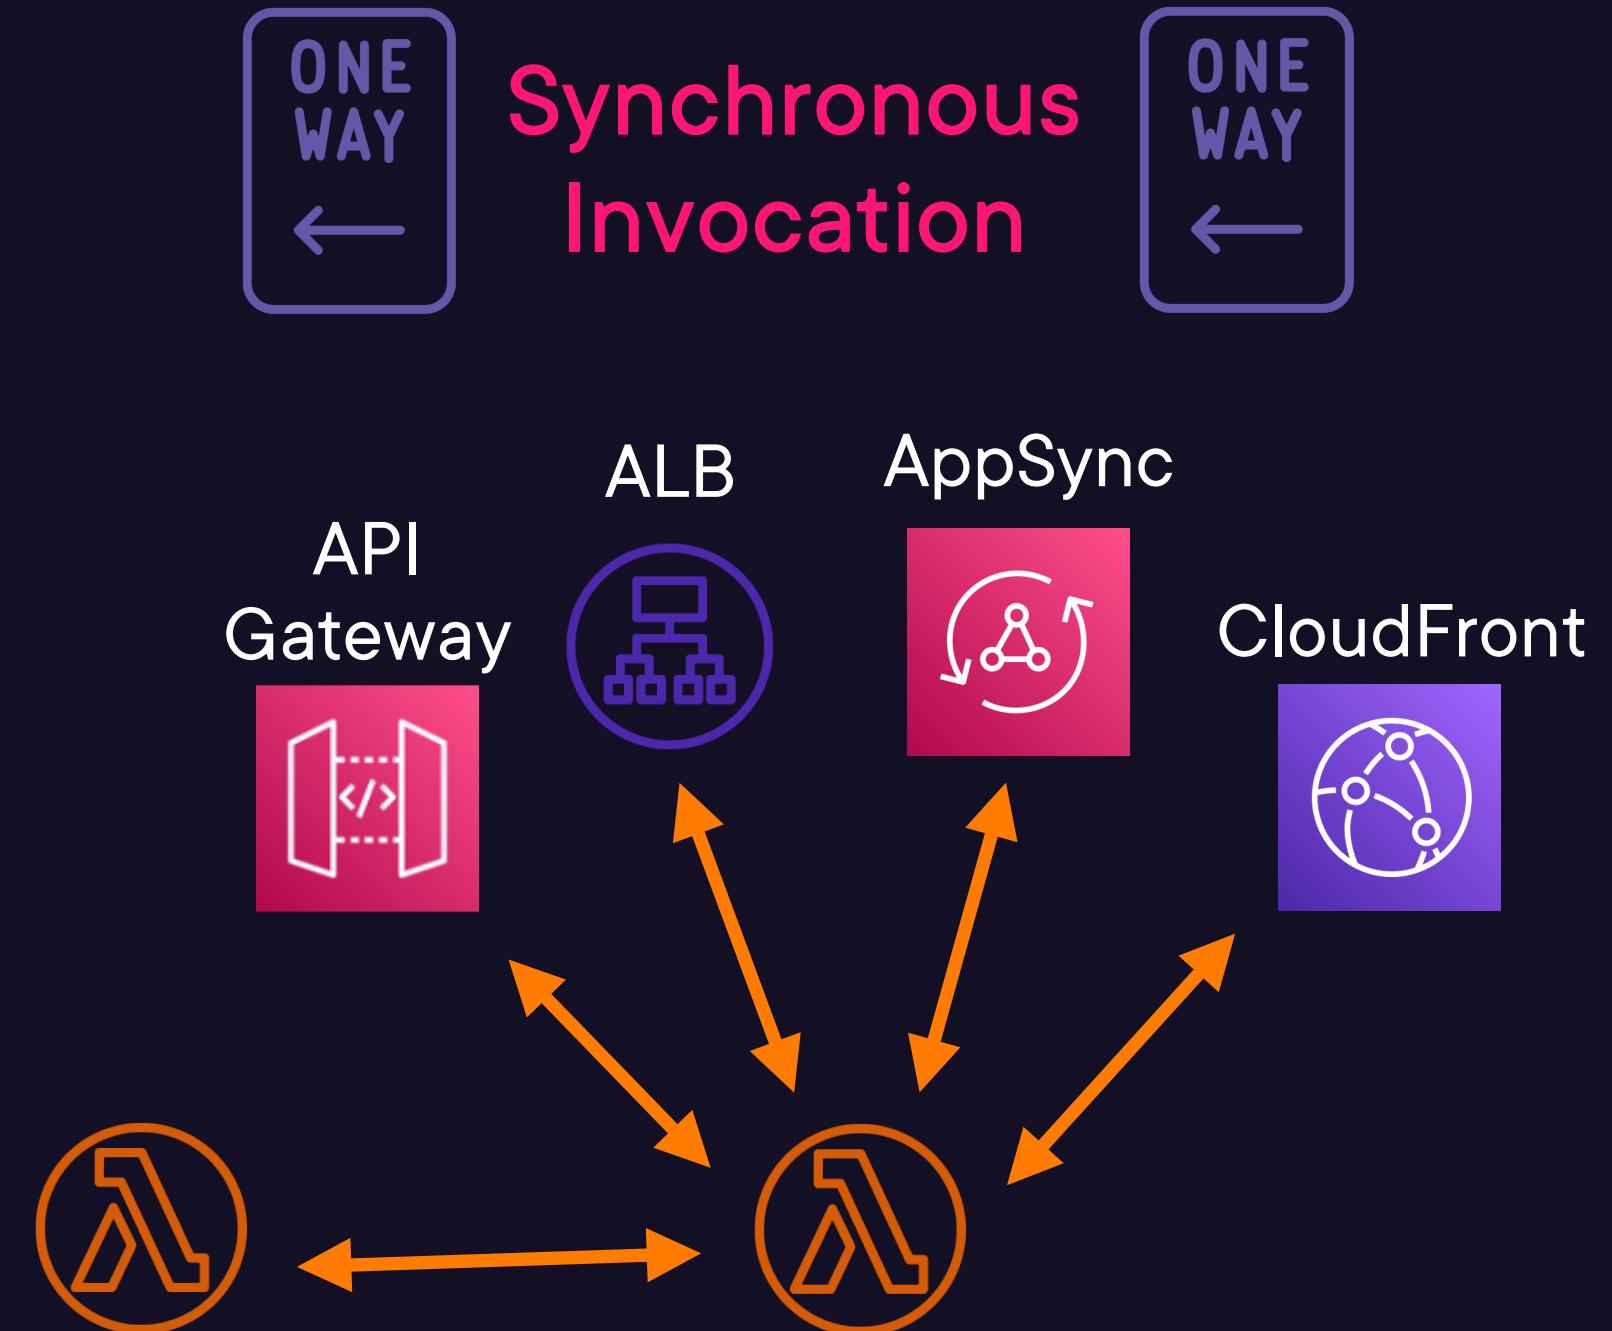
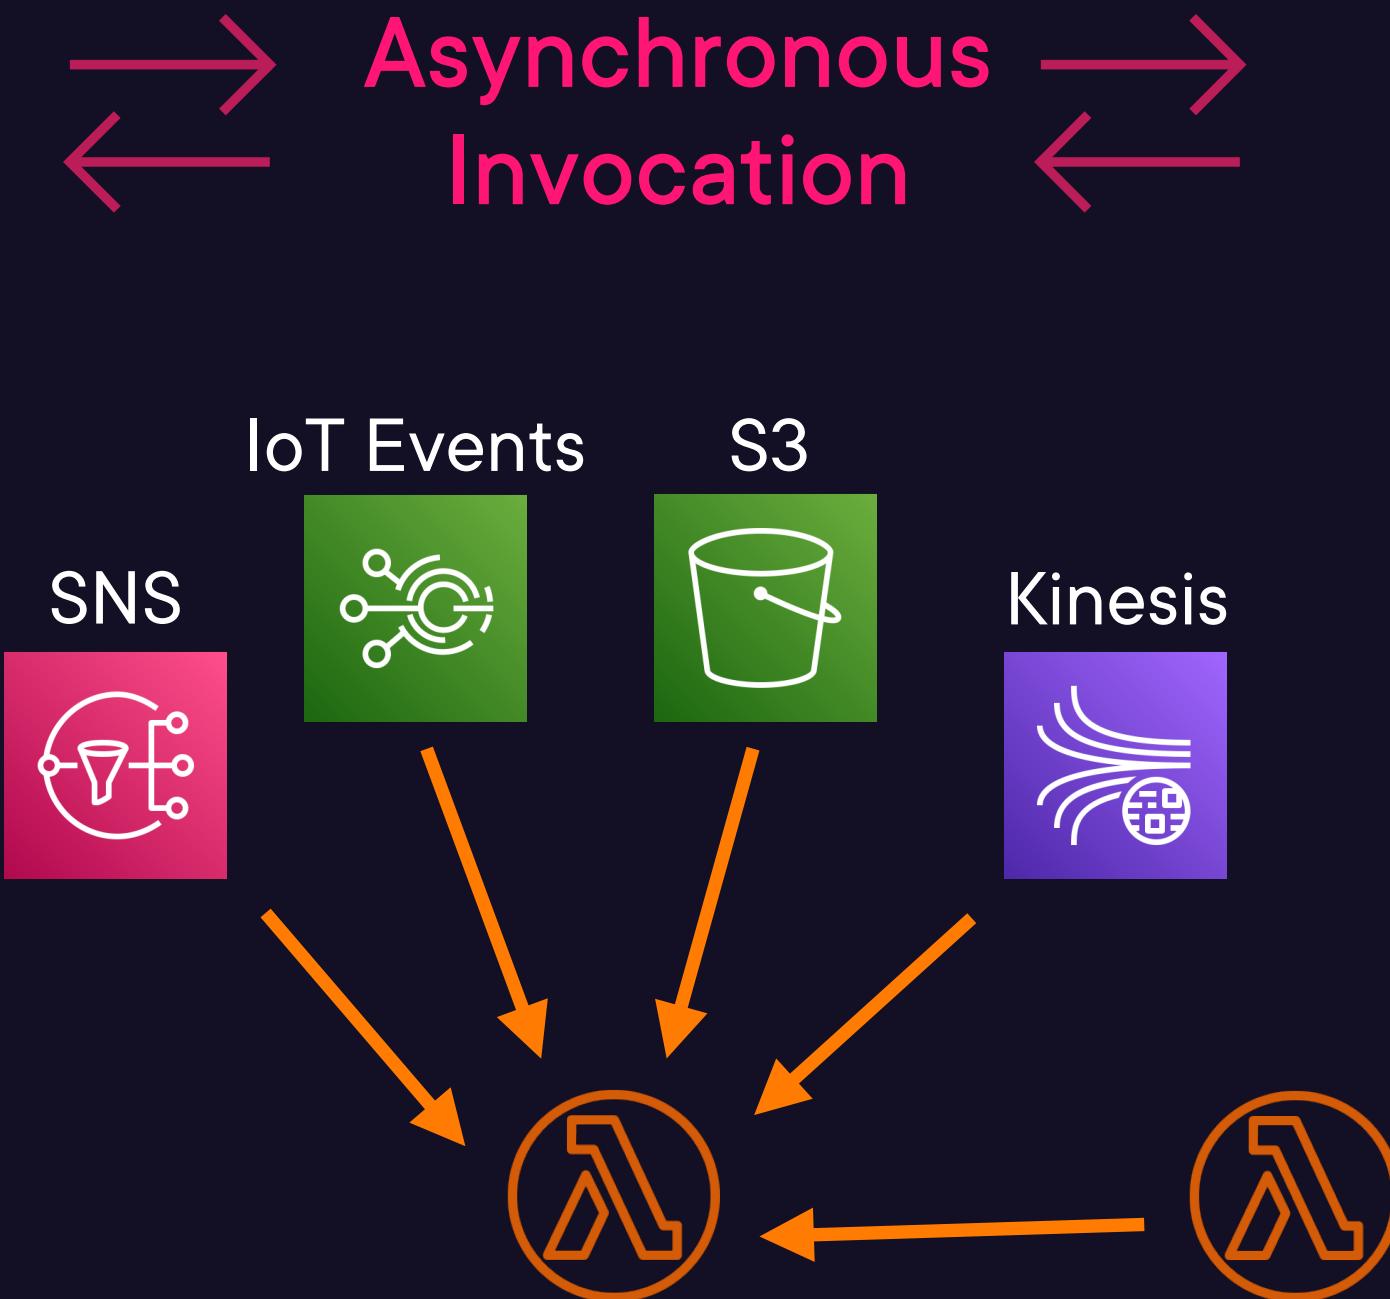
## Asynchronous Invocation

- Offers the option to route results to destinations
  - ▶ OnSuccess
  - ▶ OnFailure
- No additional code needed
- Possible Destinations:
  - ▶ Amazon SNS
  - ▶ Amazon SQS
  - ▶ AWS Lambda



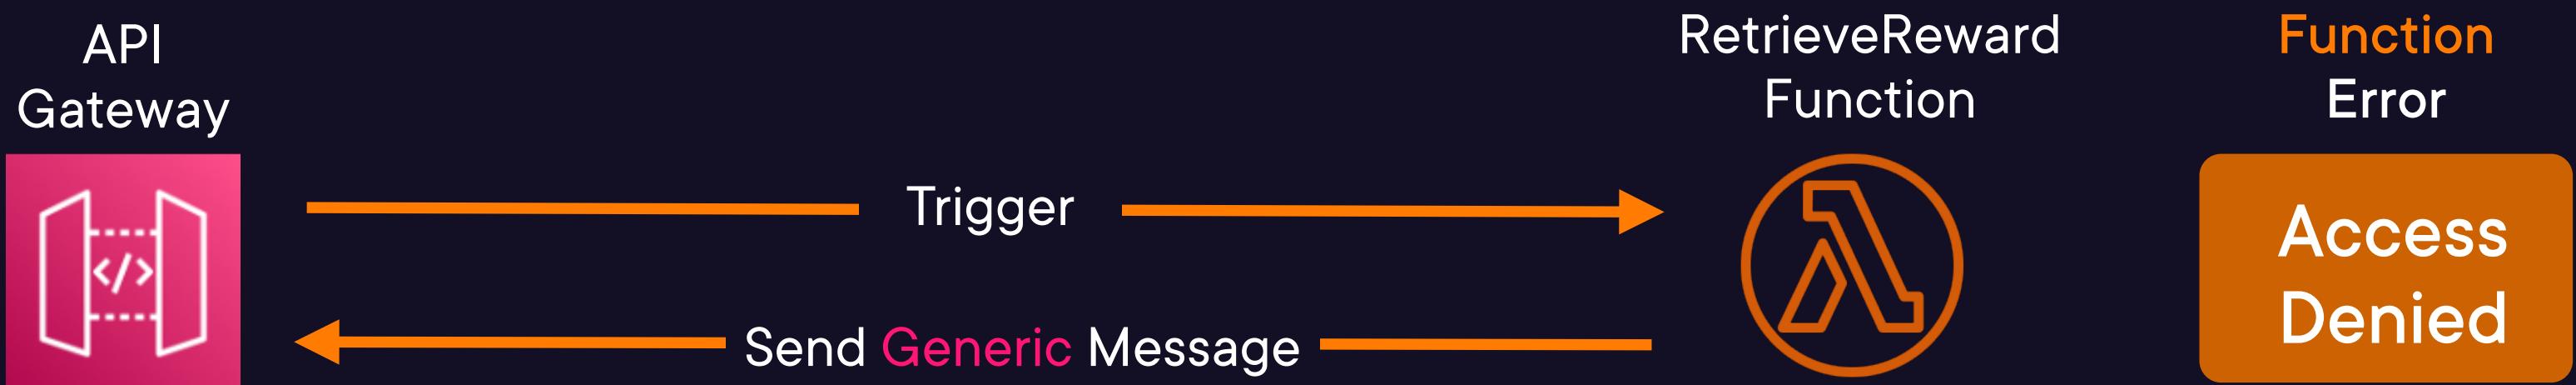
# Testing and Debugging Lambda Failures

## Invocation Types: Sample Services



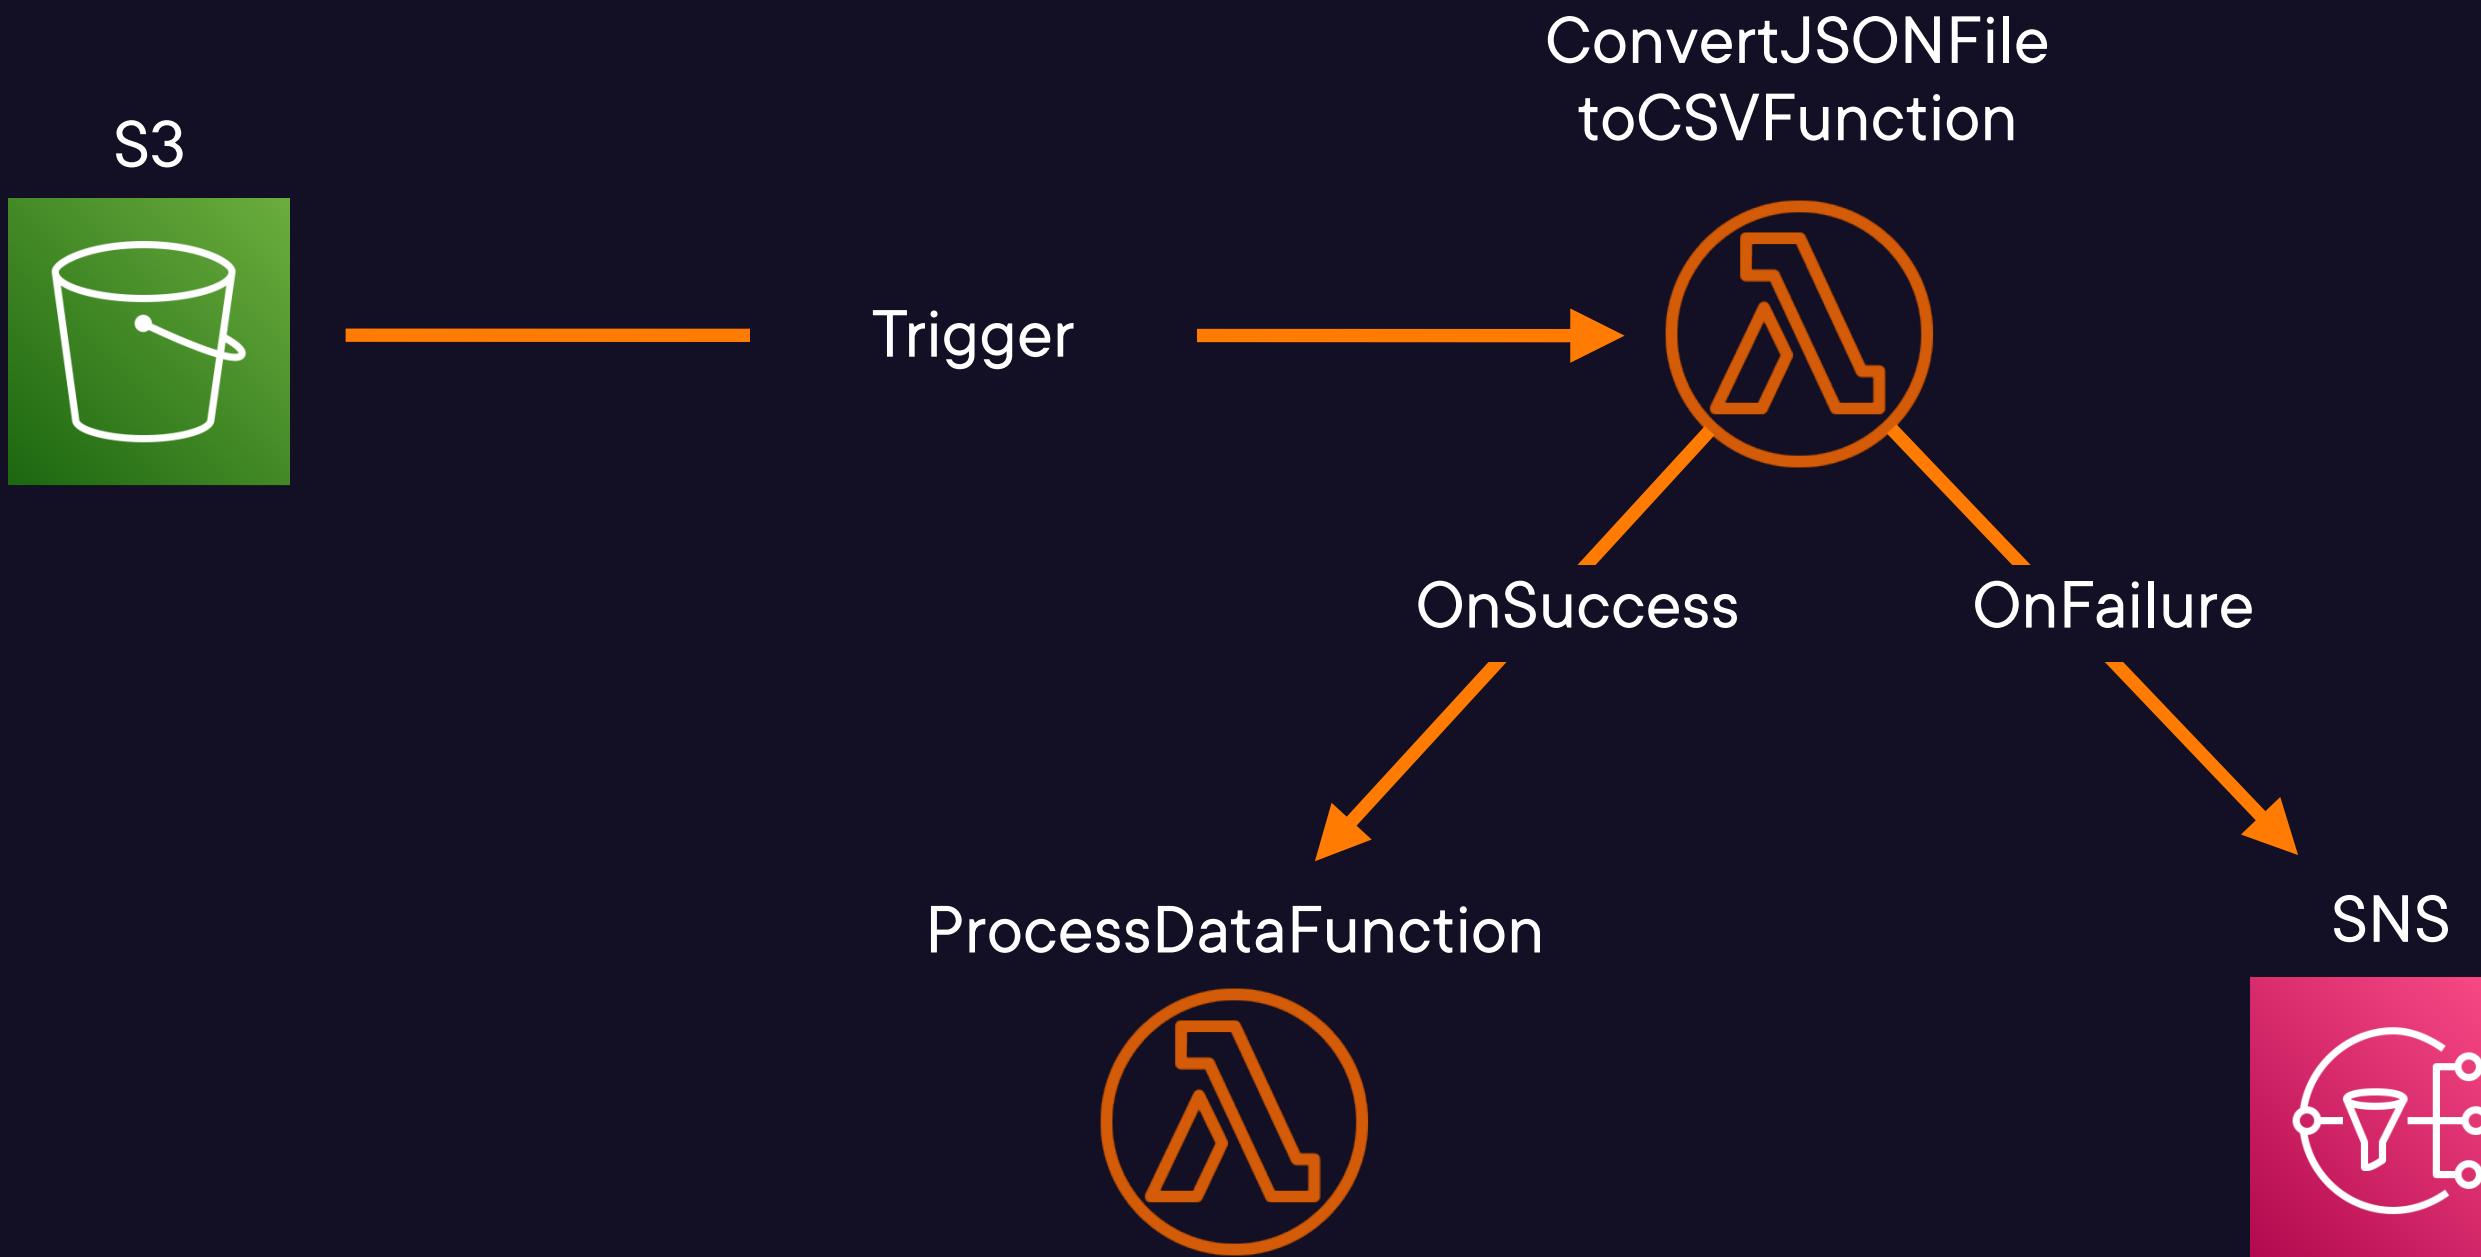
## Testing and Debugging Lambda Failures

# Synchronous Invocation: Handling Errors



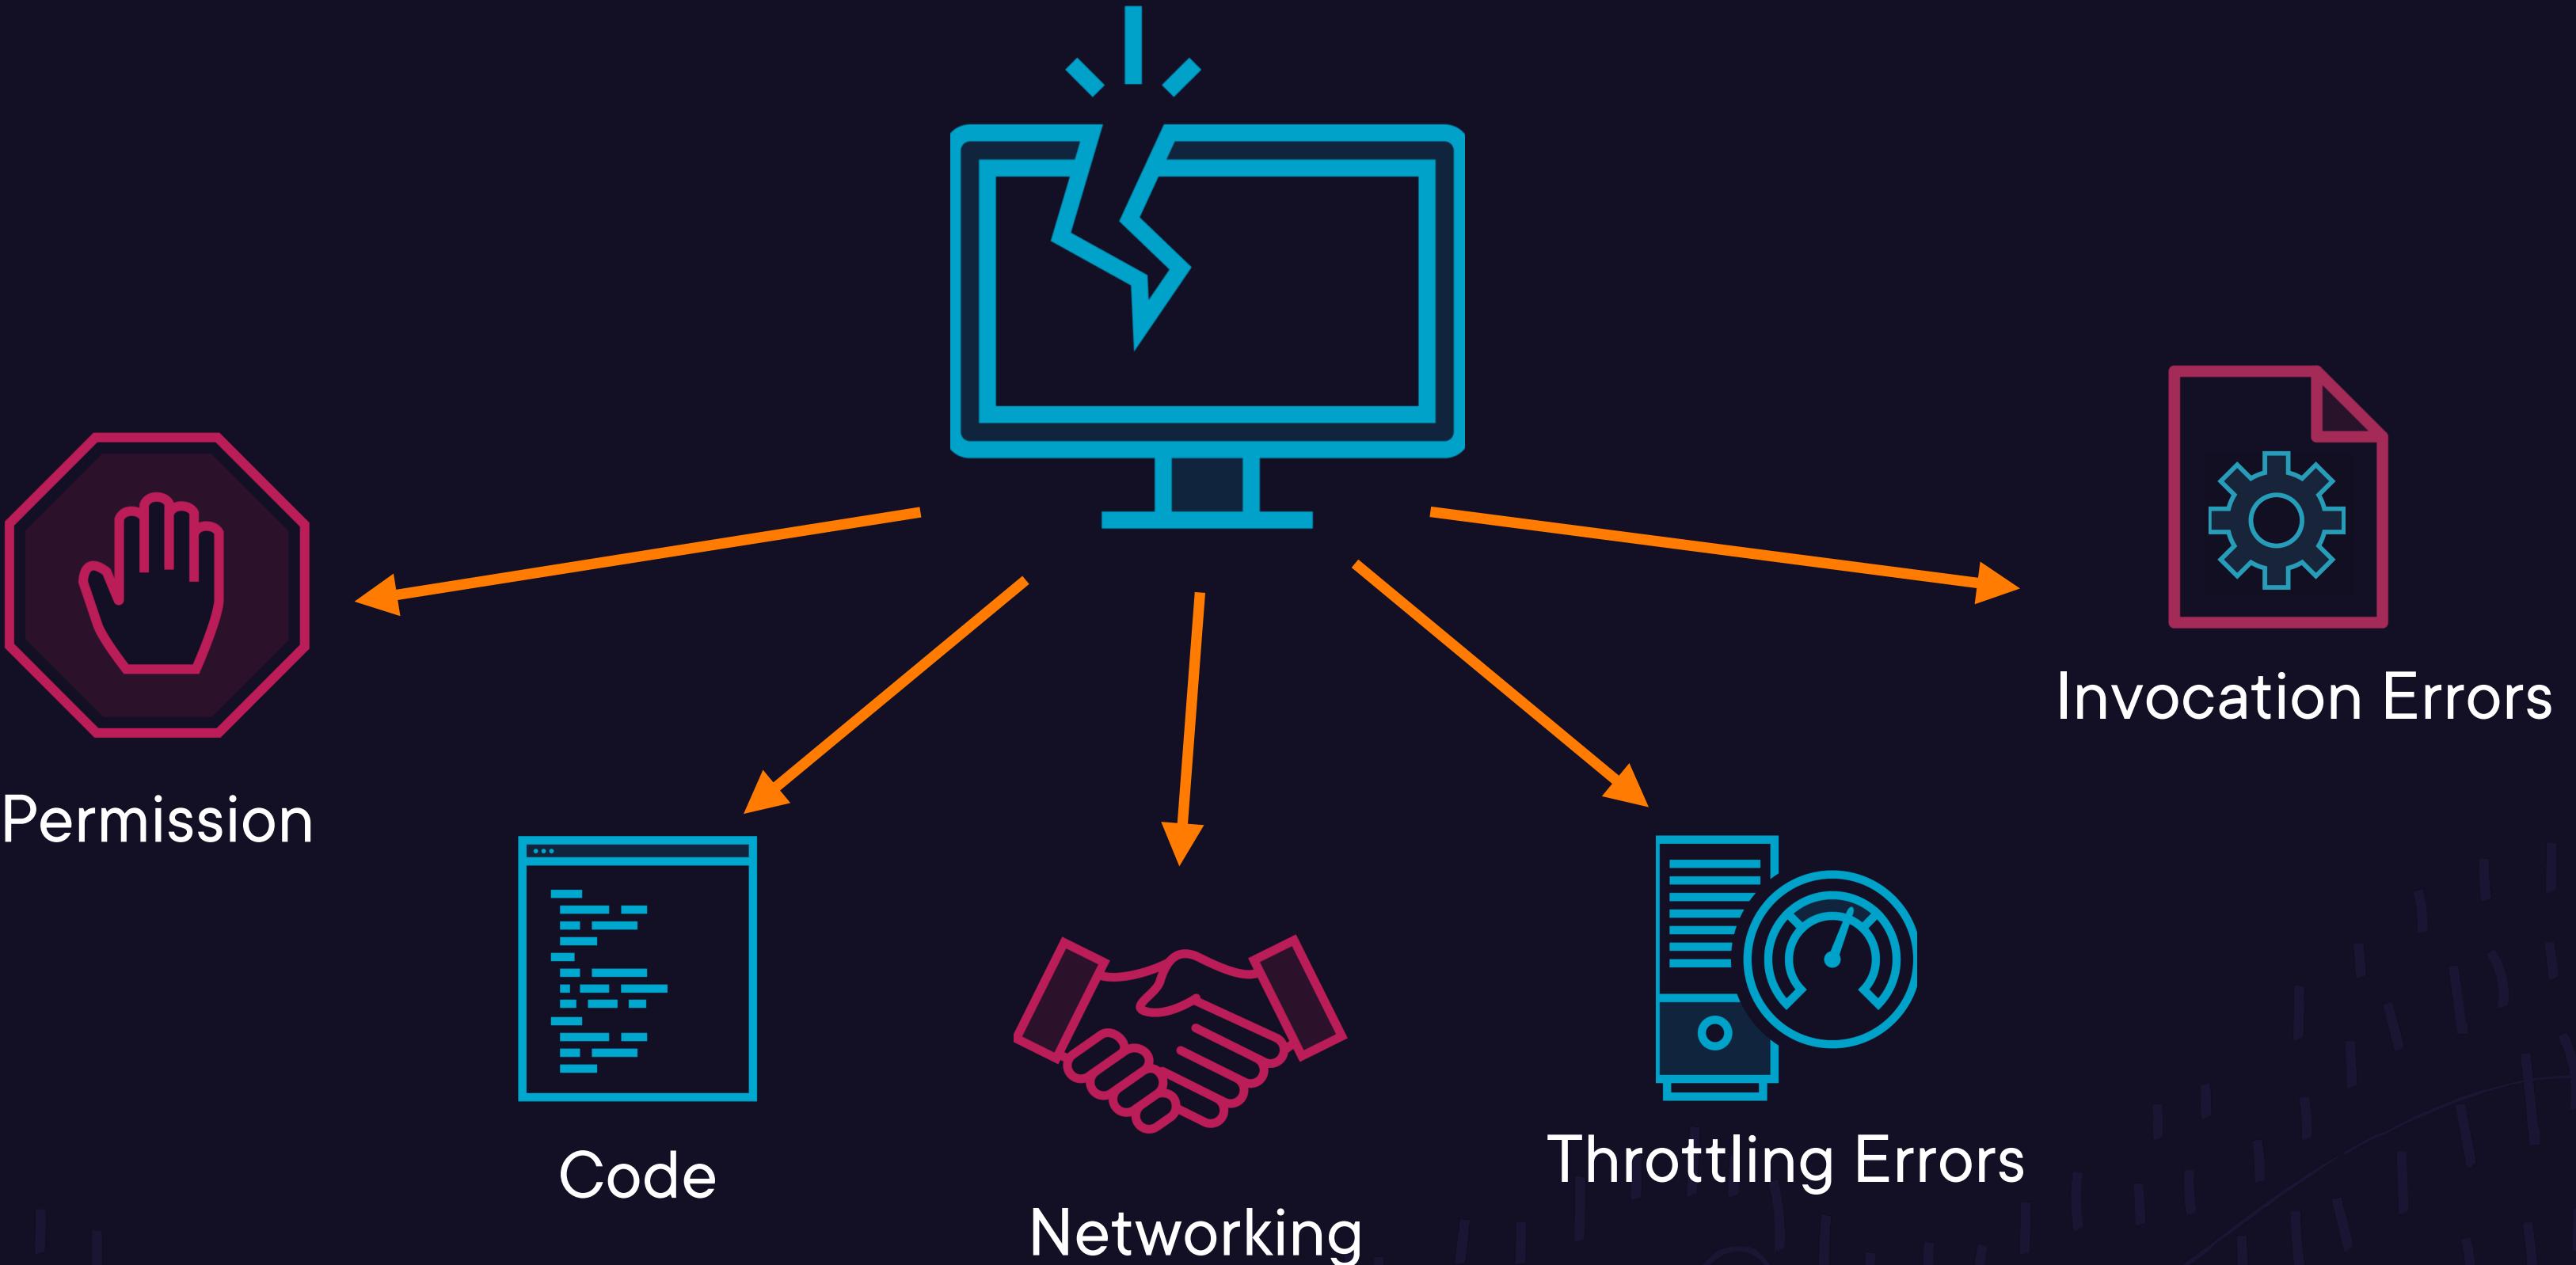
## Testing and Debugging Lambda Failures

# Asynchronous Invocation: Handling Errors



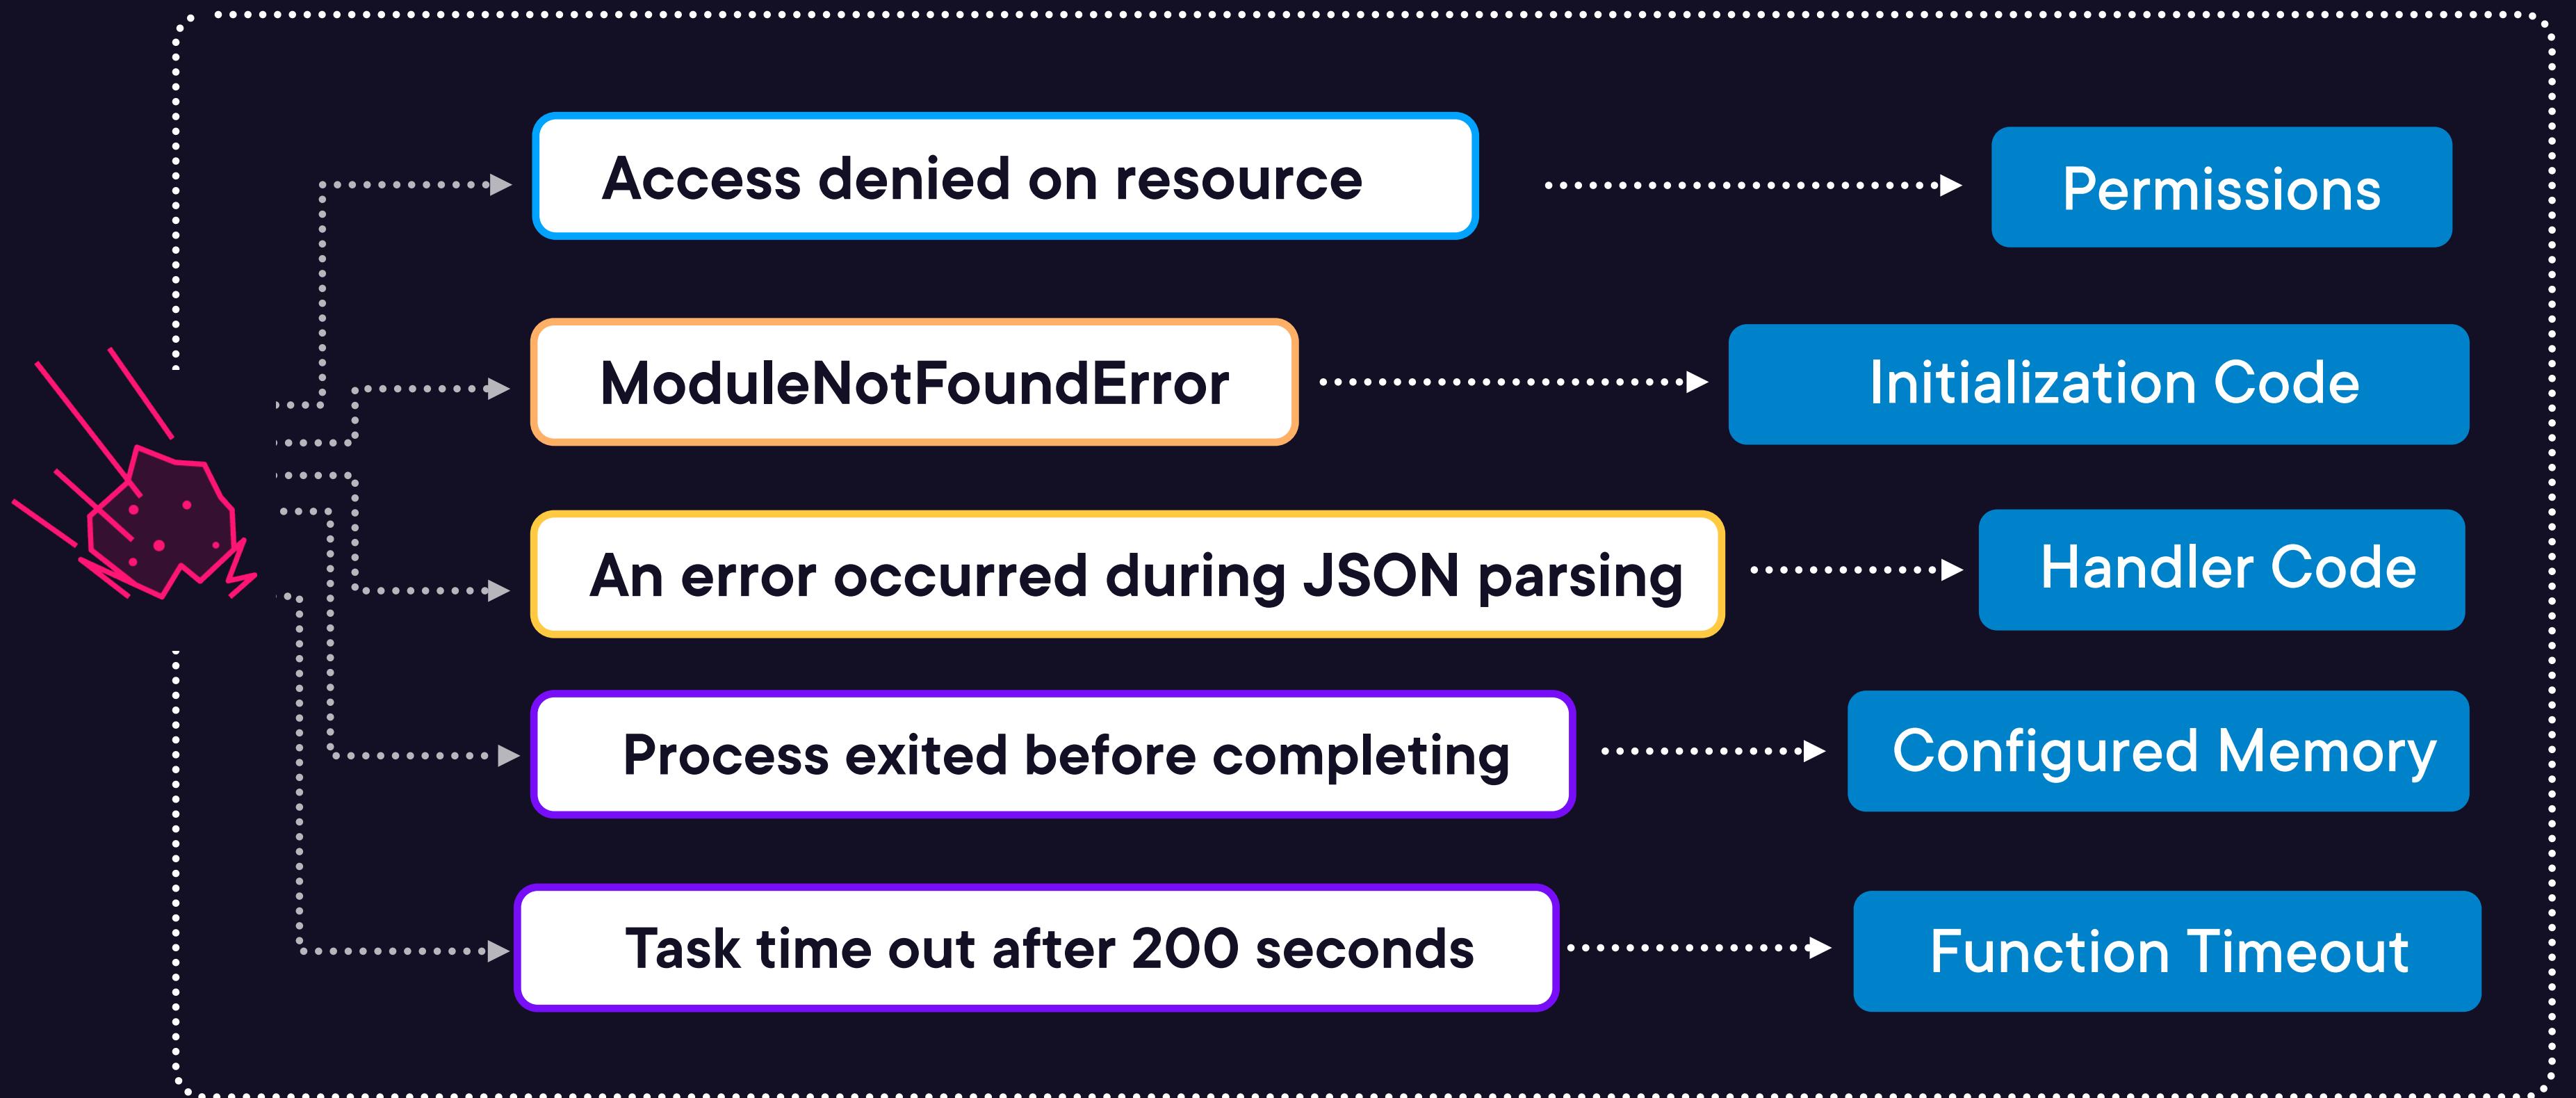
# Testing and Debugging Lambda Failures

## Types of Failures



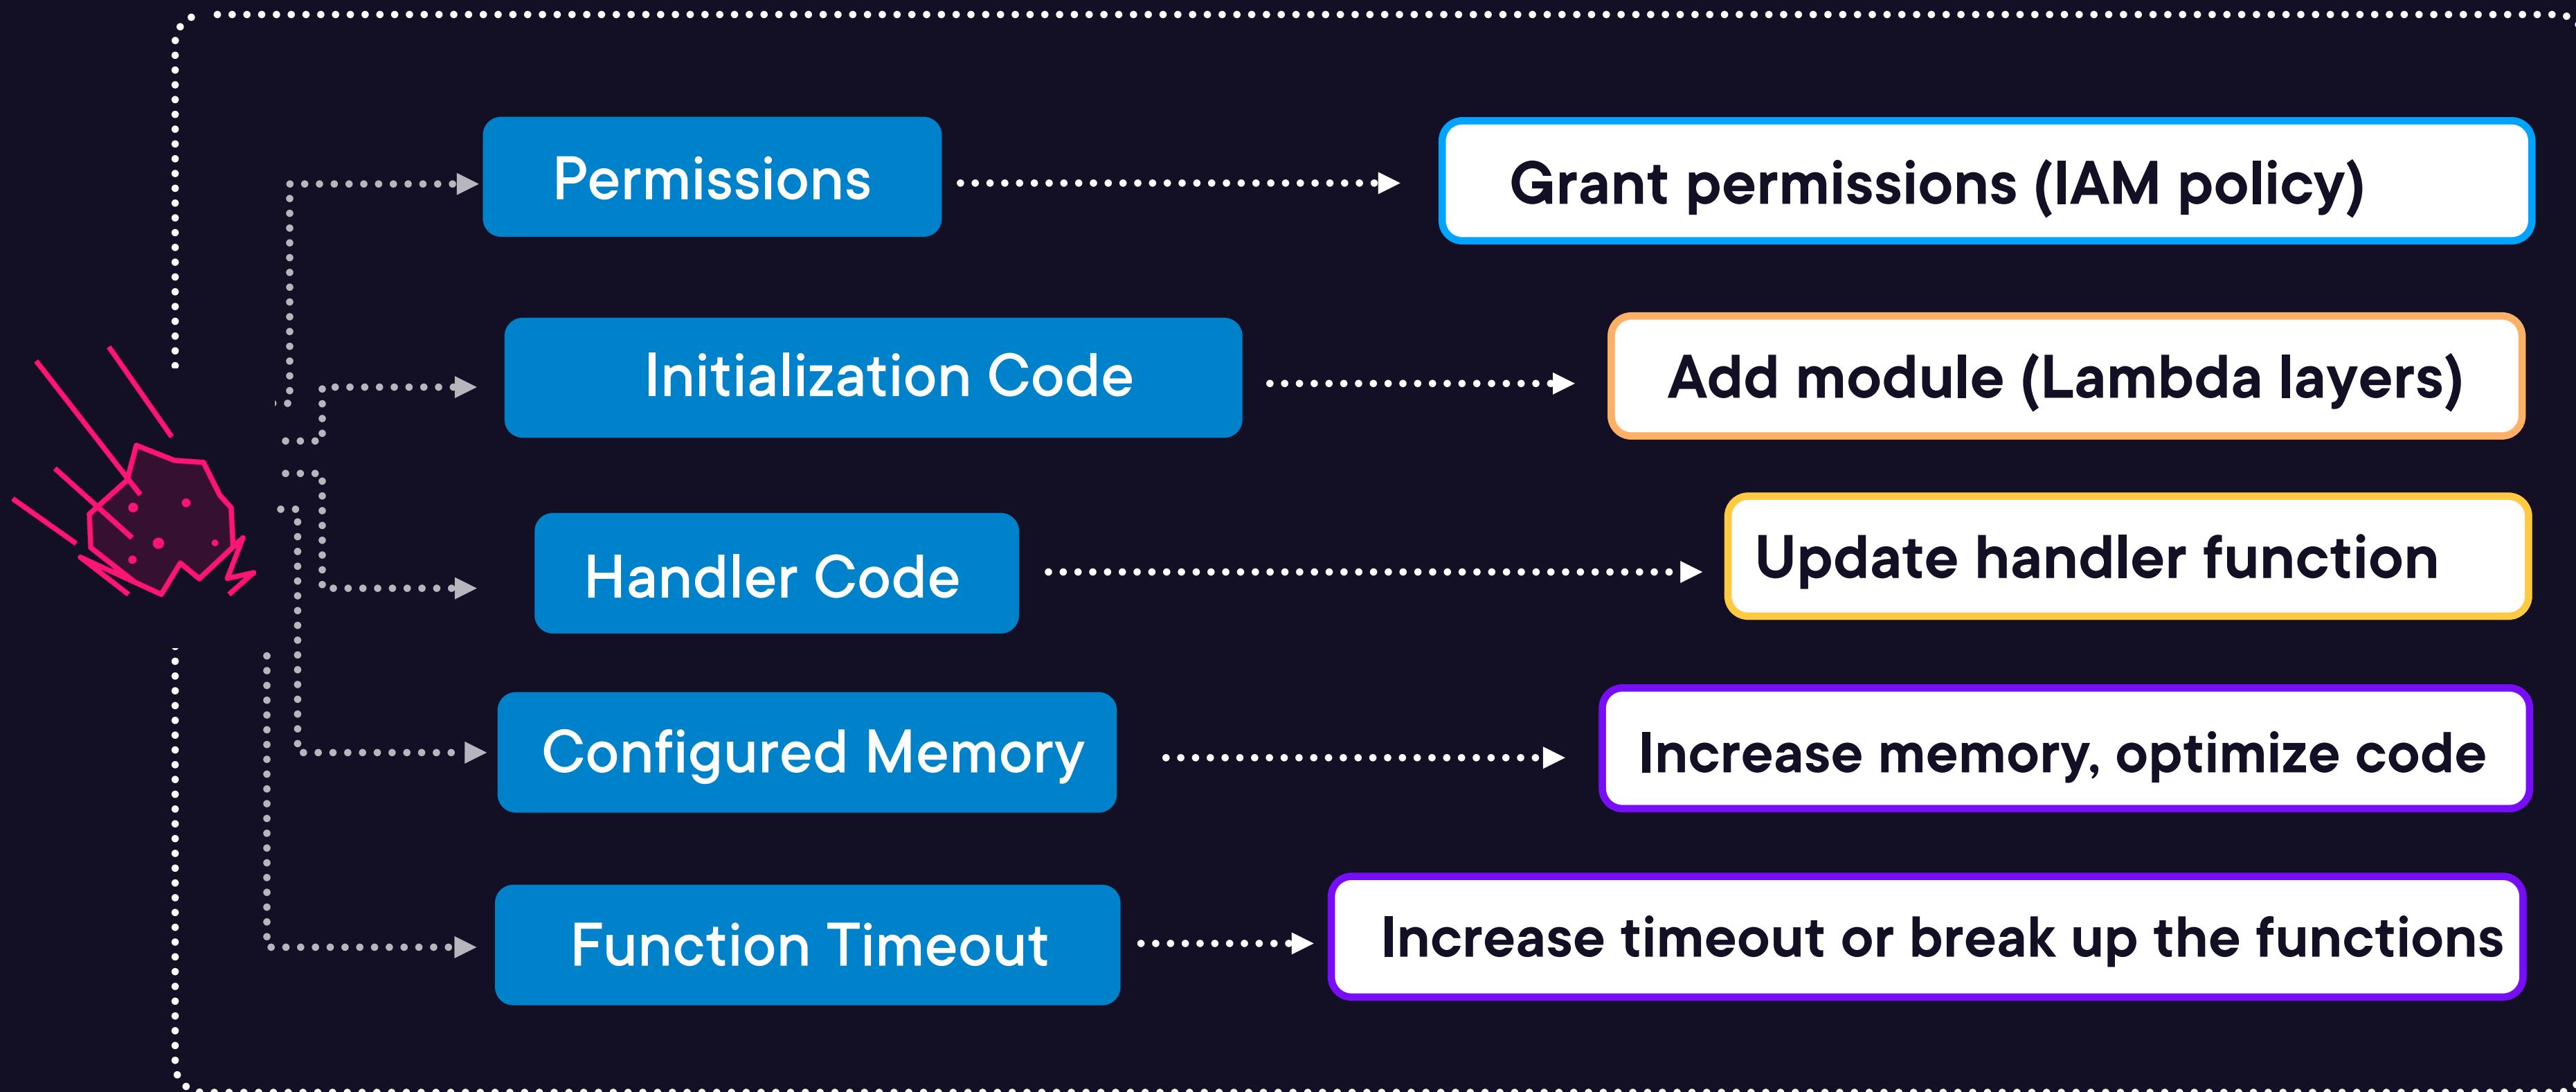
# Testing and Debugging Lambda Failures

## Common Lambda Errors



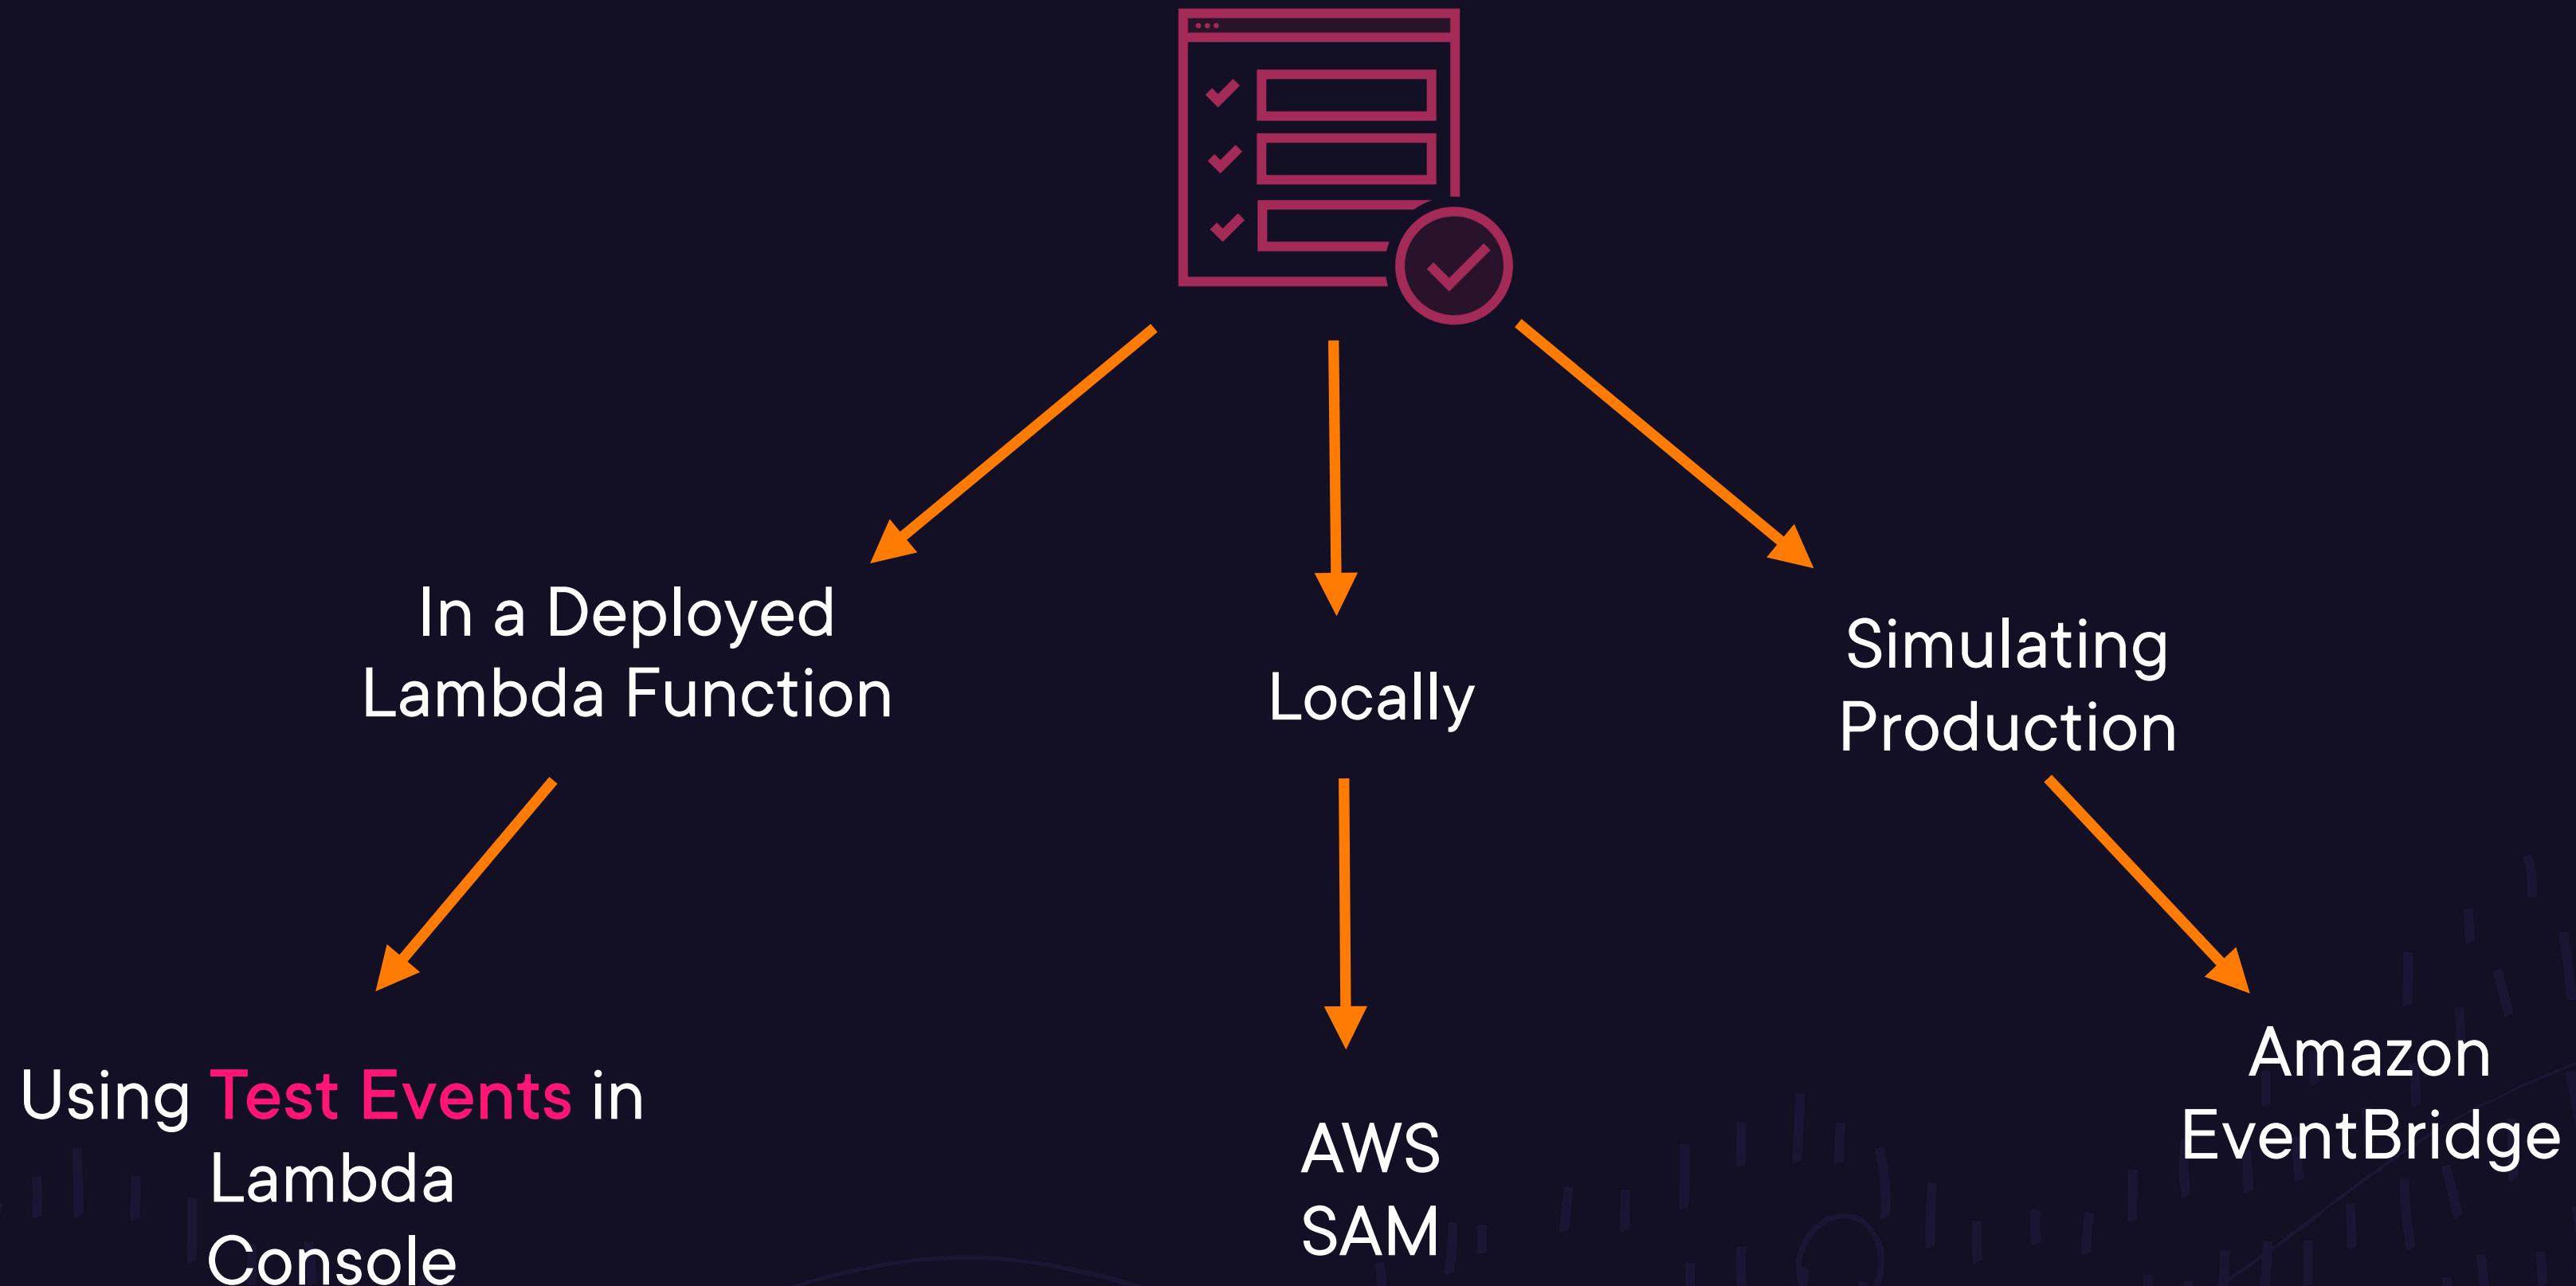
# Testing and Debugging Lambda Failures

## Common Lambda Errors



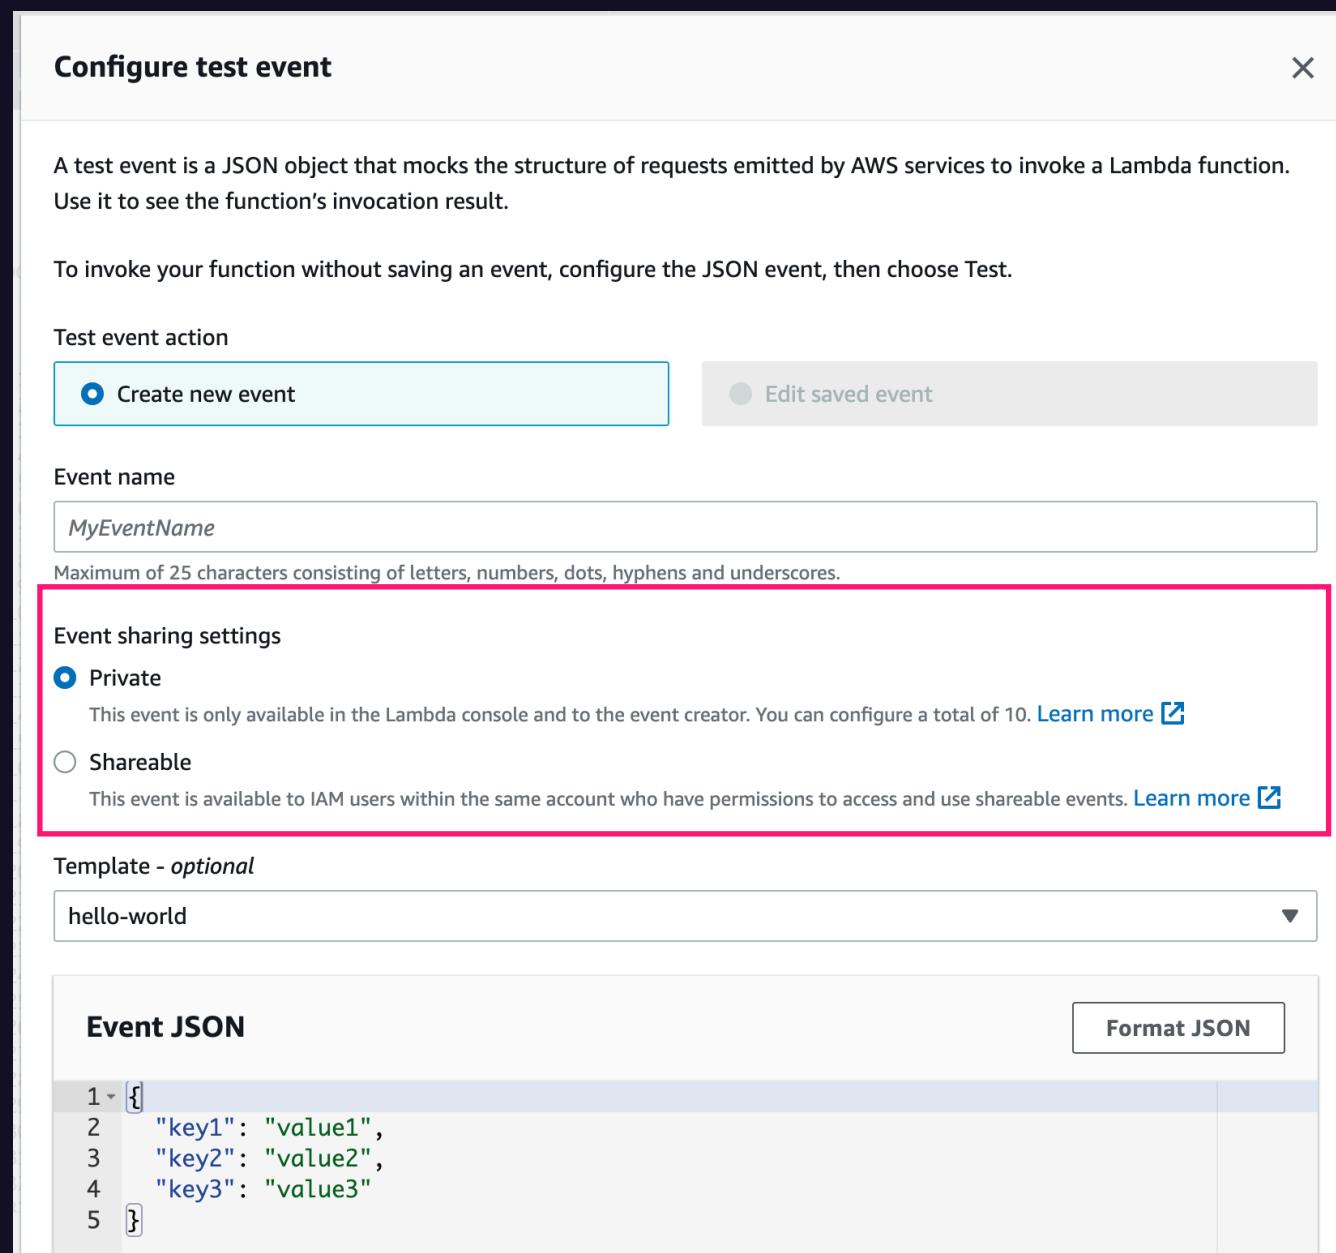
## Testing and Debugging Lambda Failures

### Testing Environments



# Testing and Debugging Lambda Failures

## Lambda Console: Testing



- from the AWS Lambda console

## Test Events

### Private

- Can only be accessed by the event creator
- You can create up to **10 private test events/function**

### Shareable

- Shared with other IAM users in the same account
- Lambda saves shareable test events as schema in an **Amazon EventBridge schema registry**



## Testing and Logging Goals

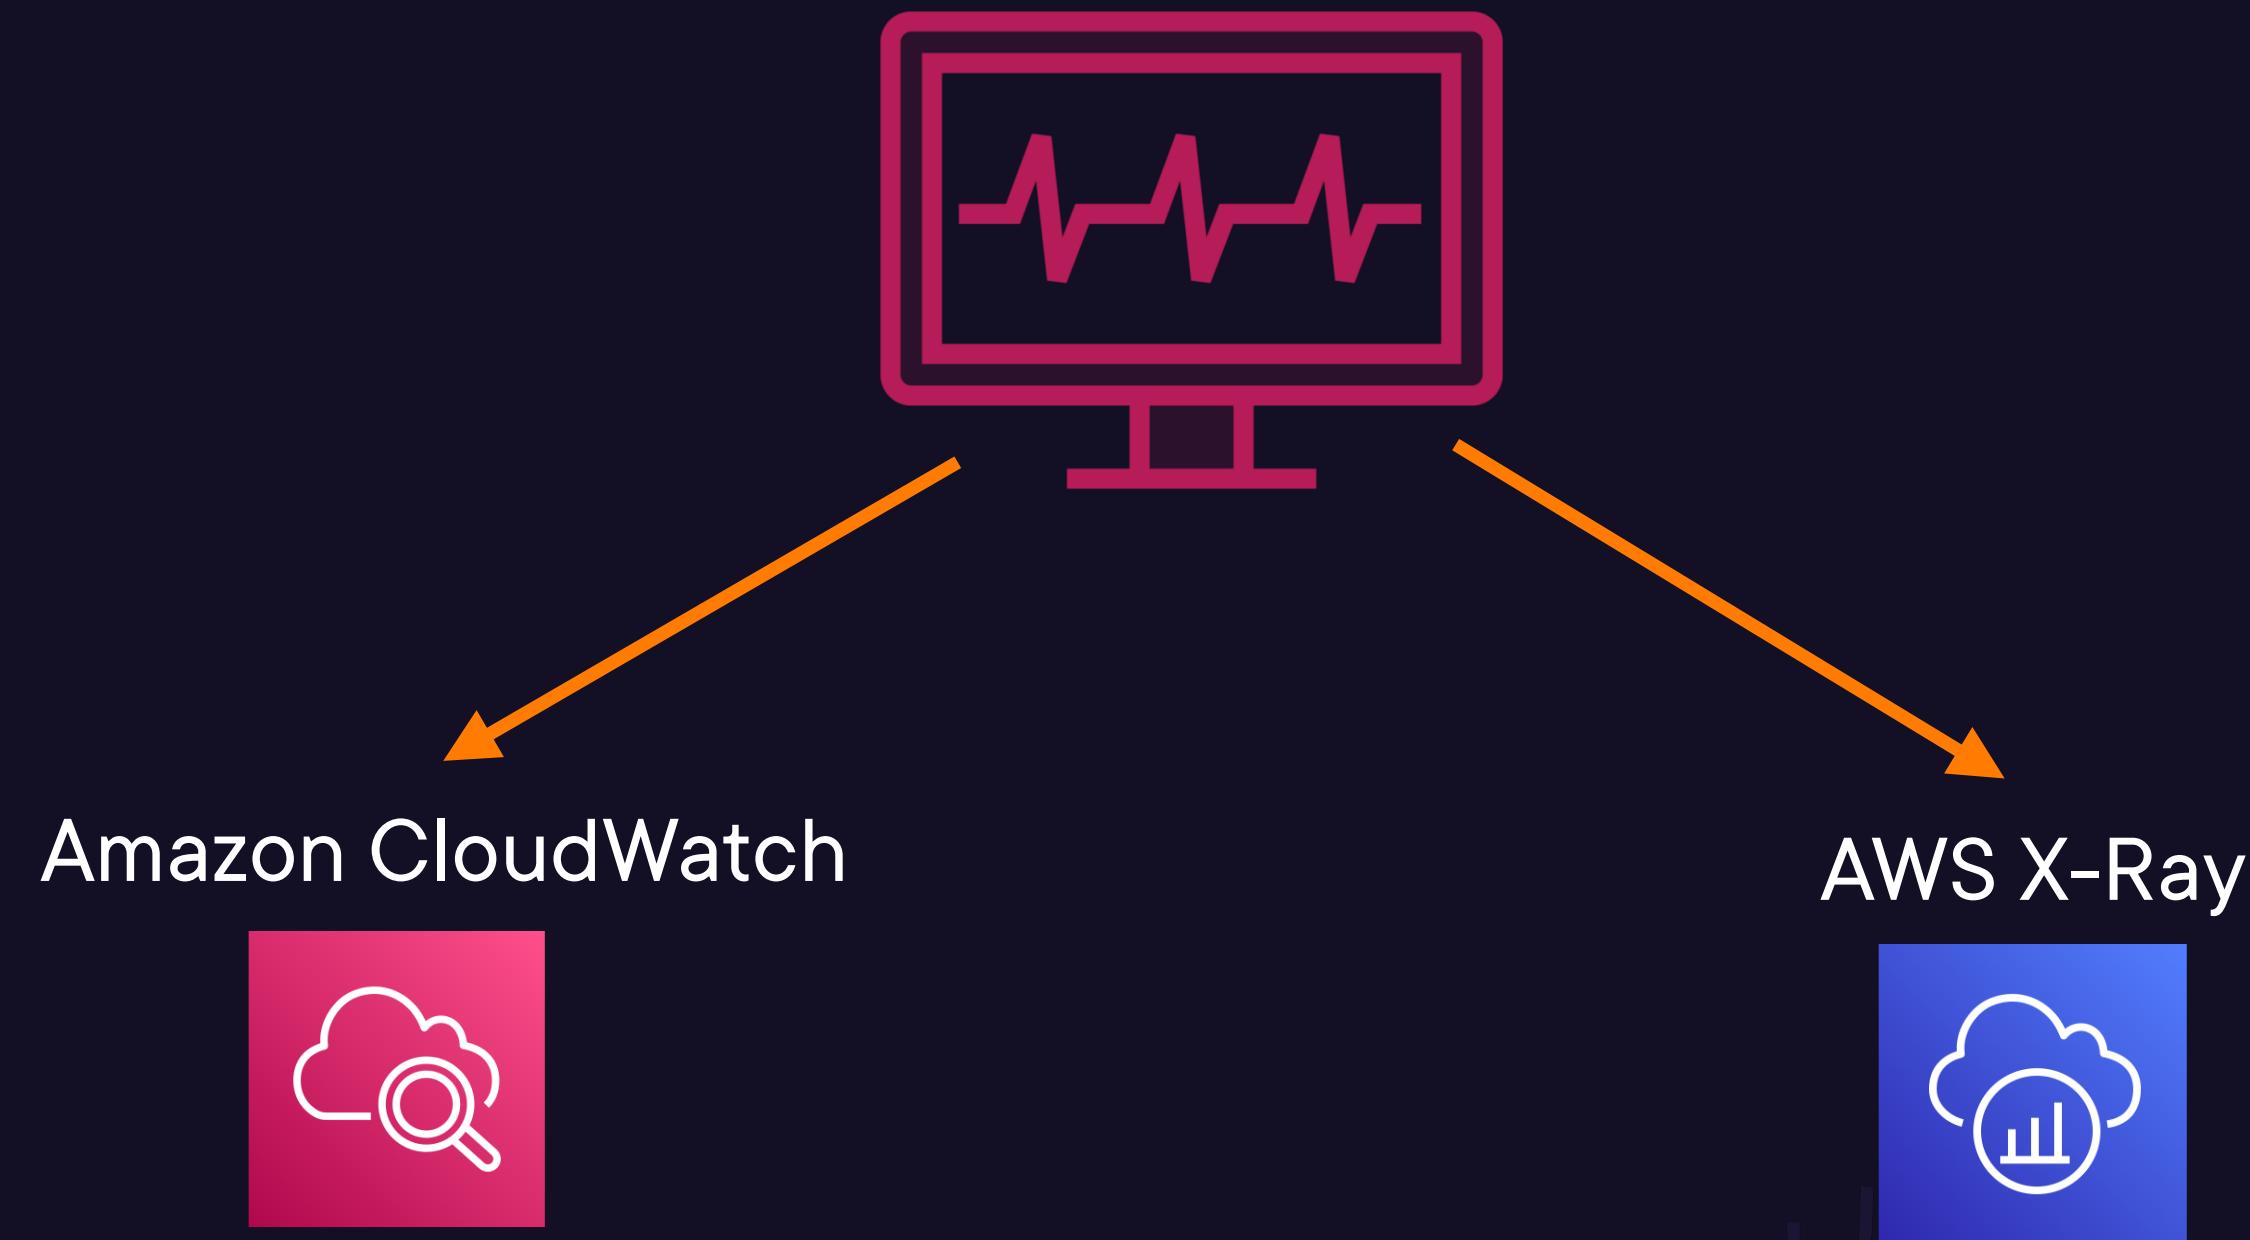


- Log errors in a consistent way.
- Enable alarms on necessary metrics.
- Synchronous Invocation: Send back sensible error messages to the invoking service.
  - ▶ Lambda sends back a meaningful error message when the client benefits from being informed about the issues.
  - ▶ Lambda sends back a generic error message when the issue can't be resolved by the client (i.e., throttles).



## Testing and Debugging Lambda Failures

### Monitoring Tools



# Things Worth Tracking



## Invocation Metrics

The **sum** of the below metrics:

- Errors
- Throttles
- (i.e., TooManyRequestsException)
- DestinationDeliveryFailures

## Performance Metrics

The **average** or **max** of the below metrics:

- Duration
- IteratorAge



# Lesson Summary



## Invocation Types

- Synchronous
- Asynchronous

## Types of Failures

- Different types of failures (permission, code, networking, etc.)
- Common Lambda errors

## Testing

- Environments (Lambda console, AWS SAM, and EventBridge)
- Testing goals and types of test events

## Monitoring

- Monitoring tools
- Metrics



# Processing Data with Lambda: Recap



**Noreen Hasan**

Training Architect



## SECTION BREAKDOWN

# Processing Data with Lambda: Recap



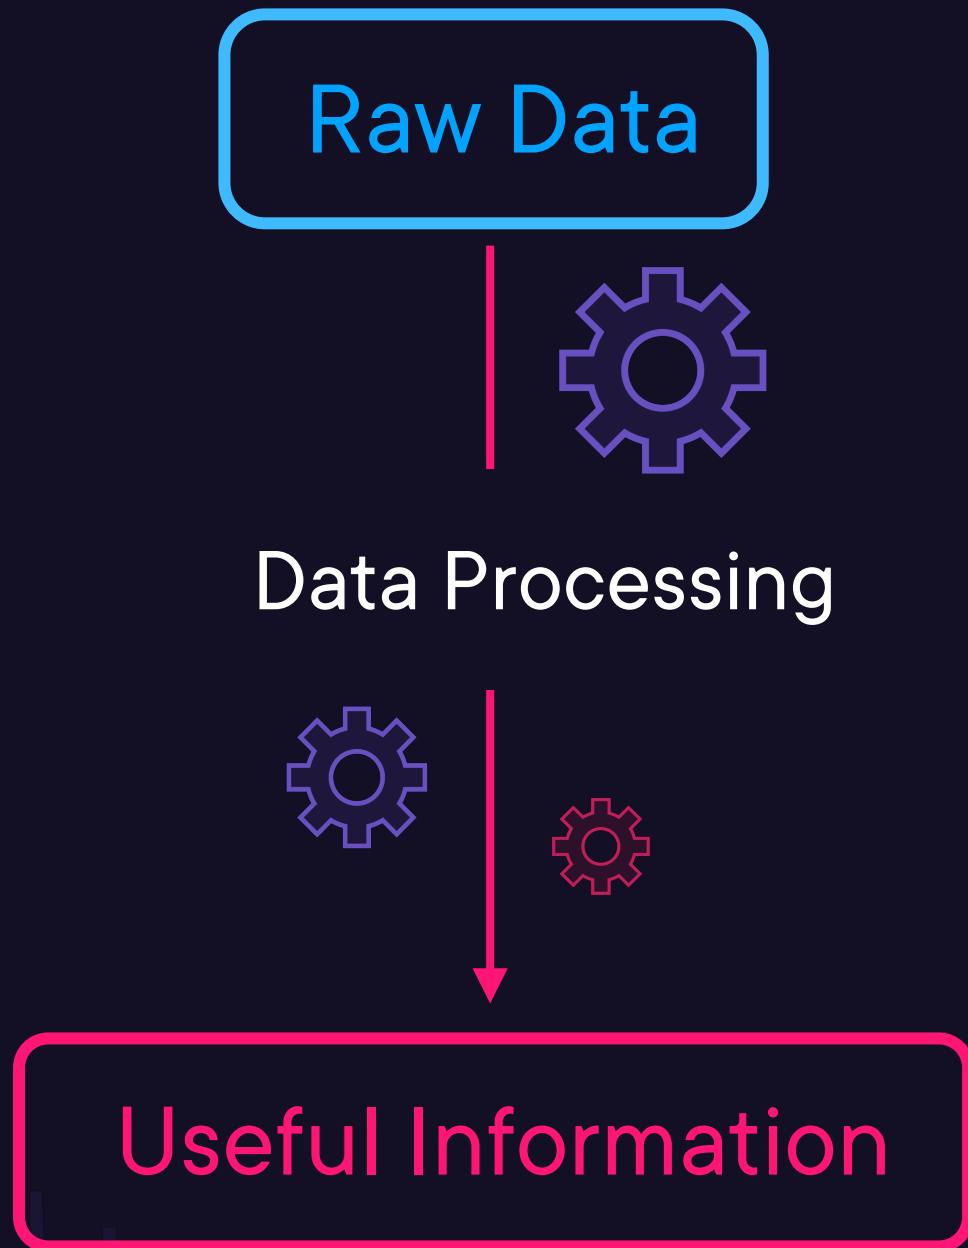
**Noreen Hasan**  
Training Architect

## Lessons Summary



## Processing Data with Lambda: Recap

### Processing Data: Overview



#### What Is Data Processing?

- Converting raw data to information

#### Data Lakes

- A central repo  with raw data
- ETL versus ELT

#### Types of Data Transformation

- Data filtering
- Data mapping
- Deriving variables
- Aggregating data
- Deduplicating data
- Splitting data



# Processing Data with Lambda: Recap

## Understanding Lambda Layers



### What Are Lambda Layers?

- Enables **importing packages** into functions
- Packages can contain **additional code, third-party libraries**, config files, or **custom runtime environments**

### Why Layer Up?

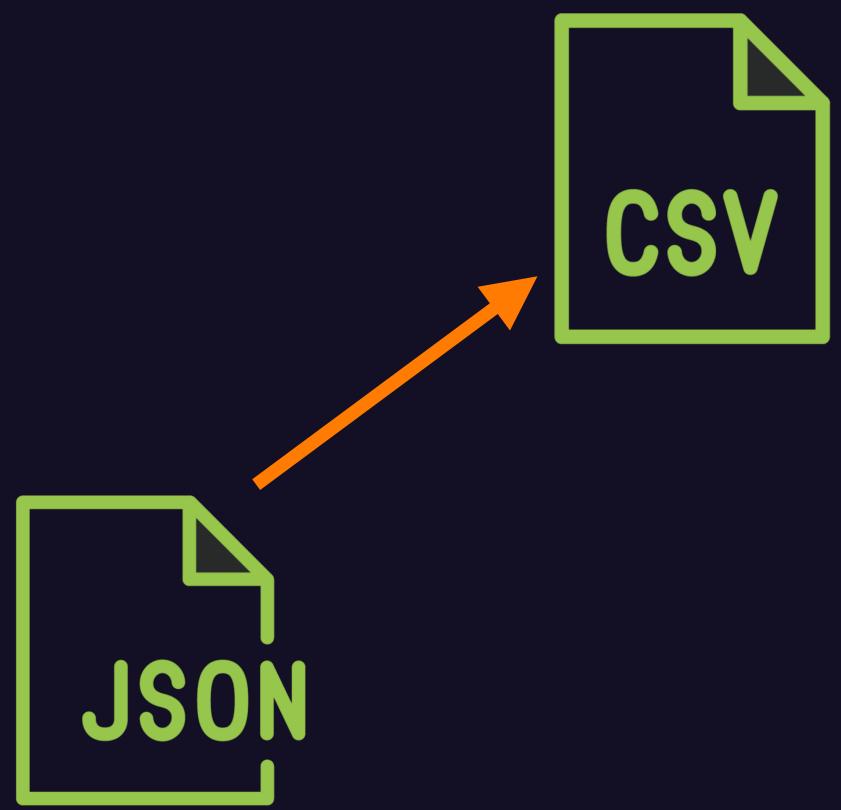
- Code reusability
- Smaller packages= faster deployments

### Challenges

- Testing is harder
- Updating functions can get tricky
- Unzipped function + layers < 250 MB



## Demo: Converting JSON to CSV Using Lambda



### What's in a File Type?

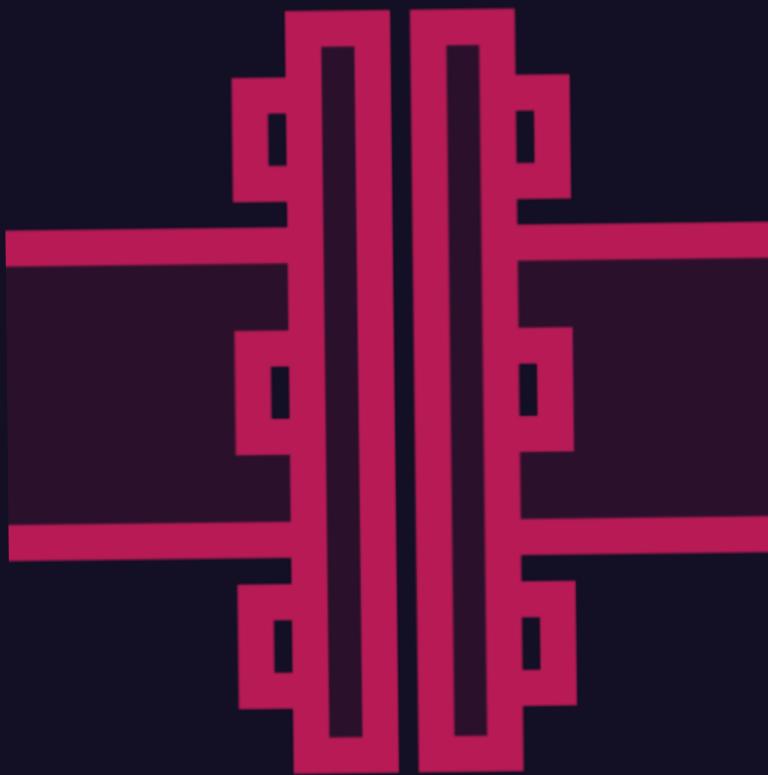
- How structured, unstructured, and semi-structured data fit in a data lake
- JSON versus CSV

### The Workflow and Demo

- Trigger **Lambda** when **{JSON}** files are uploaded to **S3**
- Function: Convert files to



# Demo: Ingesting CSV Files from S3 to DynamoDB



## DynamoDB: Overview

- NoSQL DB
- Fully managed



## The Workflow and Demo

- **Function:** insert records to a **DynamoDB** table
- **Trigger:** a CSV file is uploaded to S3



# Processing Real-Time Data Using Kinesis and Lambda: Overview



## Streaming Data: Use Cases

- Video processing 
- Clickstream events 
- Data lakes 

## The Kinesis Umbrella

- Kinesis Data Streams
- Kinesis Video Streams
- Kinesis Data Analytics
- Kinesis Data Firehose



## Main Components of Kinesis Data Streams

- Data producers
- Kinesis streams
- Data consumers



## Processing Data with Lambda: Recap

# Demo: Processing Real-Time Data Using Kinesis and Lambda



## Kinesis, Lambda, and Shards

A Kinesis data stream consists of **shards**, which consist of a set of decodable records.

## The Workflow and Demo

Created a **Lambda** function to decode Base64 data coming from **Kinesis data streams**.

# Testing and Debugging Lambda Failures



## Invocation Types

- Synchronous
- Asynchronous

## Types of Failures

- Different types of failures (permission, code, networking, etc.)
- Common Lambda errors

## Testing

- Environments (Lambda console, AWS SAM, and EventBridge)
- Testing goals and types of test events

## Monitoring

- Monitoring tools
- Metrics



**See you in the  
next section!**

